



Event sourcing and CQRS from the trenches



SIDNEY SHEK • ARCHITECT • ATLASSIAN • @SIDNEYSHEK



EVENTS

EVENTS EVERYWHERE

memegenerator.net



Universe of Users

users				
Id	Name	Username	APIKey	
1	Homer	homer	d0a	
2	Bart	bart	f00	
3	Maggie	maggie	baa	

Universe of Users

users				
Id	Name	Username	APIKey	
1	Homer	homer	d0a	
2	Bart	bart	f00	
3	Lisa Jr	maggie	baa	

Universe of Users

users			
Id	Name	Username	APIKey
1	Homer	homers	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa

Universe of Users

users			
Id	Name	Username	APIKey
1	Homer	homer	d0a
2	Bart	bart	f00
3	Maggie	maggie	baa

events

Seq	Event	Time
123	SetUsername(3, Maggie)	0

Universe of Users


users			
Id	Name	Username	APIKey
1	Homer	homer	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa

events

Seq	Event	Time
123	SetUsername(3, Maggie)	0
124	SetName(3, Lisa Jr)	10

Universe of Users

users			
Id	Name	Username	APIKey
1	Homer	homers	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa



Seq	Event	Time
123	SetUsername(3, Maggie)	0
124	SetName(3, Lisa Jr)	10
125	SetUsername(1, homers)	15

Universe of Users

users_new

Id	Name	Derived
1	Homer	Homer1
2	Bart	Bart2
3	Lisa Jr	Lisa Jr3

users

Id	Name	Username	APIKey
1	Homer	homers	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa

events

Seq	Event	Time
123	SetUsername(3, Maggie)	0
124	SetName(3, Lisa Jr)	10
125	SetUsername(1, homers)	15

Universe of Users

users_new

Id	Name	Derived
1	Homer	Homer1
2	Bart	Bart2
3	Lisa Jr	Lisa Jr3

users

Id	Name	Username	APIKey
1	Homer	homers	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa

events

Seq	Event	Time
123	SetUsername(3, Maggie)	0
124	SetName(3, Lisa Jr)	10
125	SetUsername(1, homers)	15



elastic

Our Identity System requirements



Our Identity System requirements

- Users, groups and memberships



Our Identity System requirements

- Users, groups and memberships
 - Searching for users



Our Identity System requirements

- Users, groups and memberships
 - Searching for users
 - Retrieve by email



Our Identity System requirements

- Users, groups and memberships
 - Searching for users
 - Retrieve by email
- High volume low latency reads



Our Identity System requirements

- Users, groups and memberships
 - Searching for users
 - Retrieve by email
 - Incremental synchronisation
- High volume low latency reads



Our Identity System requirements

- Users, groups and memberships
 - Searching for users
 - Retrieve by email
 - Incremental synchronisation
- High volume low latency reads
- Audit trails for changes



Our Identity System requirements

- Users, groups and memberships
 - Searching for users
 - Retrieve by email
 - Incremental synchronisation
- Audit trails for changes
- High volume low latency reads
- Highly available
 - Disaster recovery
 - Zero-downtime upgrades



Our Identity System requirements

- Users, groups and memberships
 - Searching for users
 - Retrieve by email
 - Incremental synchronisation
- Audit trails for changes
- High volume low latency reads
- Highly available
 - Disaster recovery
 - Zero-downtime upgrades
- Testing with production-like data

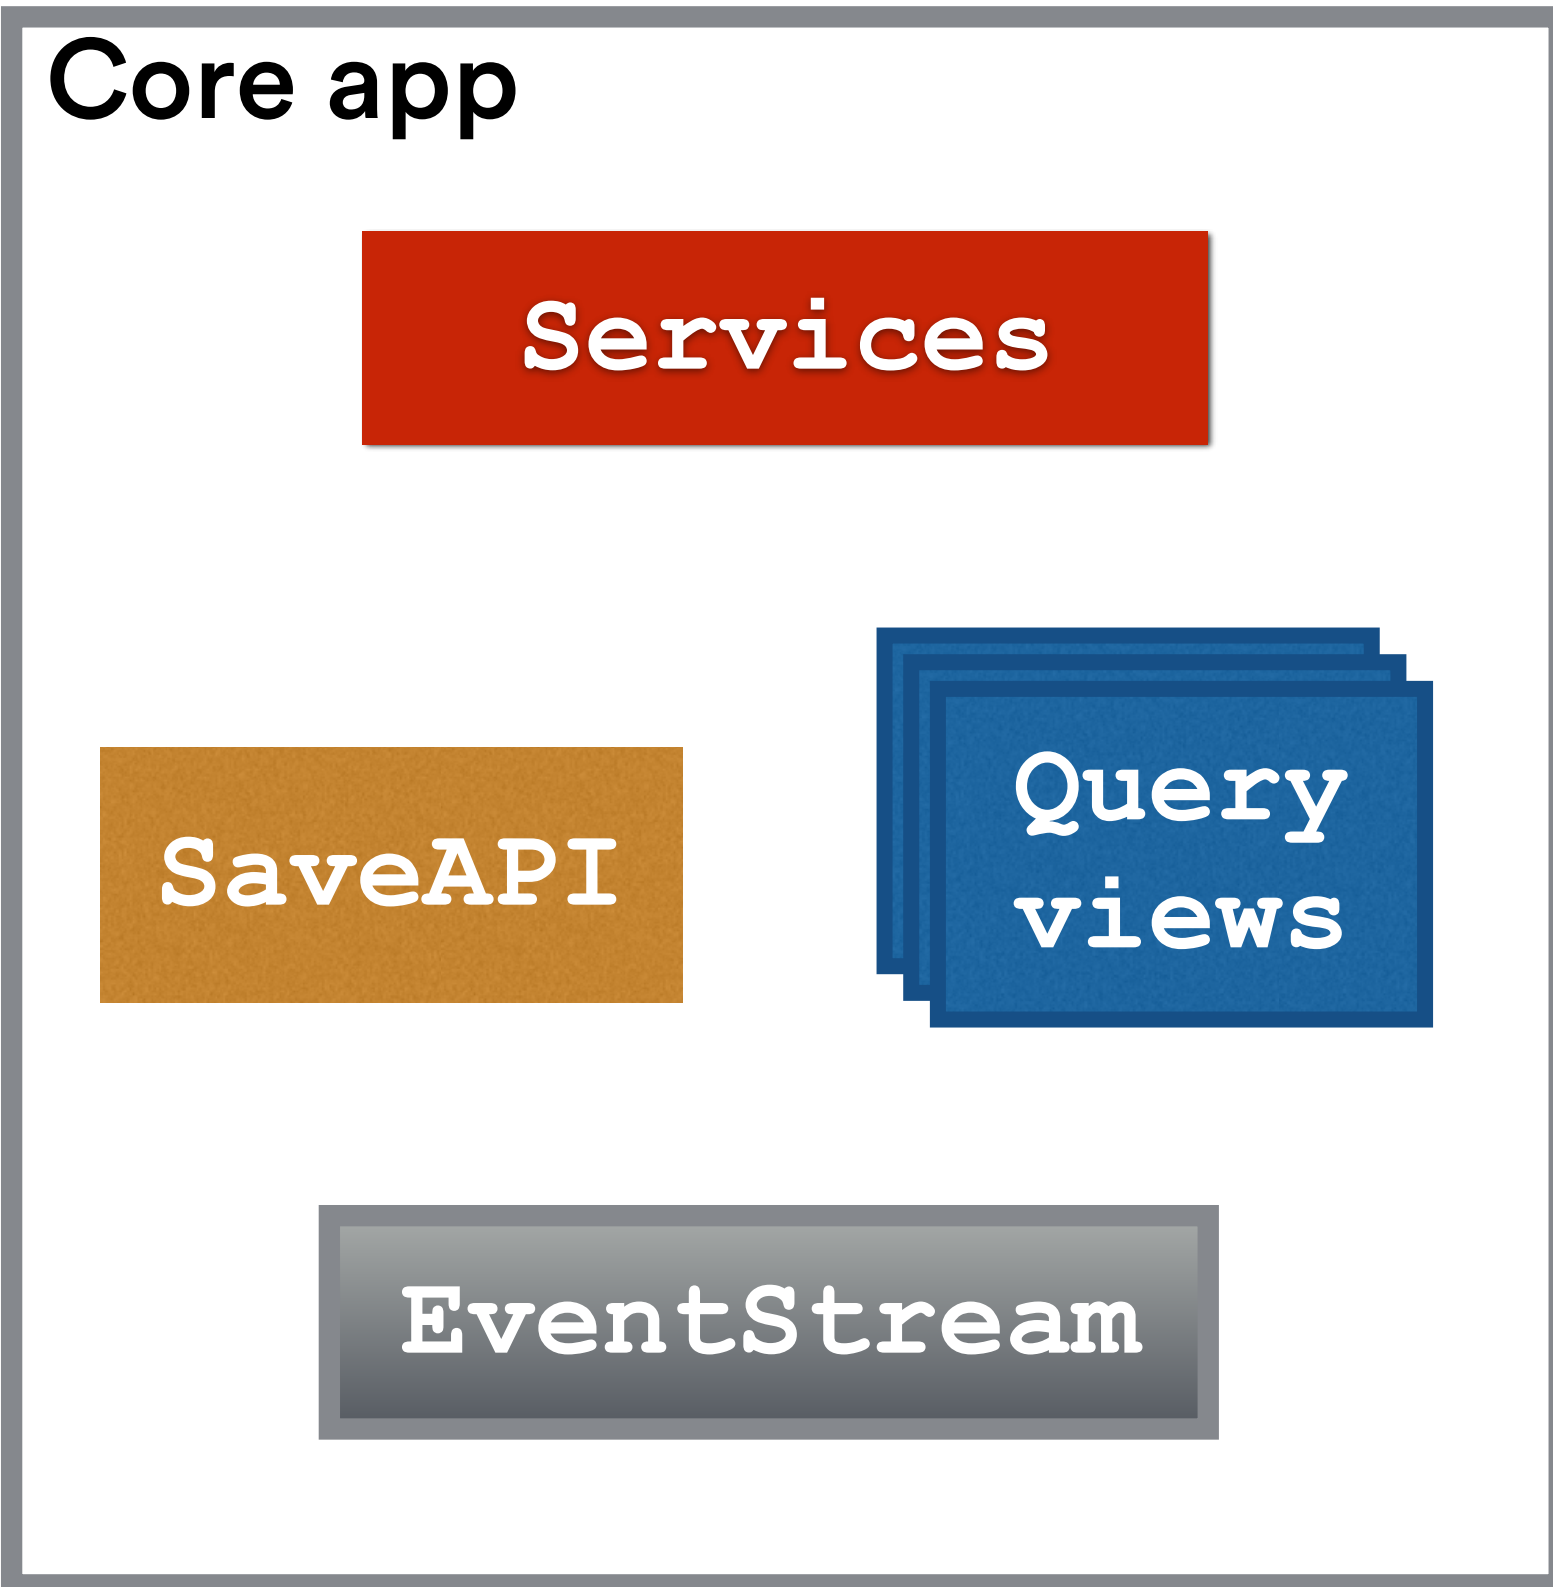




Evolving the architecture

REST calls

e.g. Add User



DynamoDB



REST calls

e.g. Add User



Core app

Services

SaveAPI

Query
views

EventStream

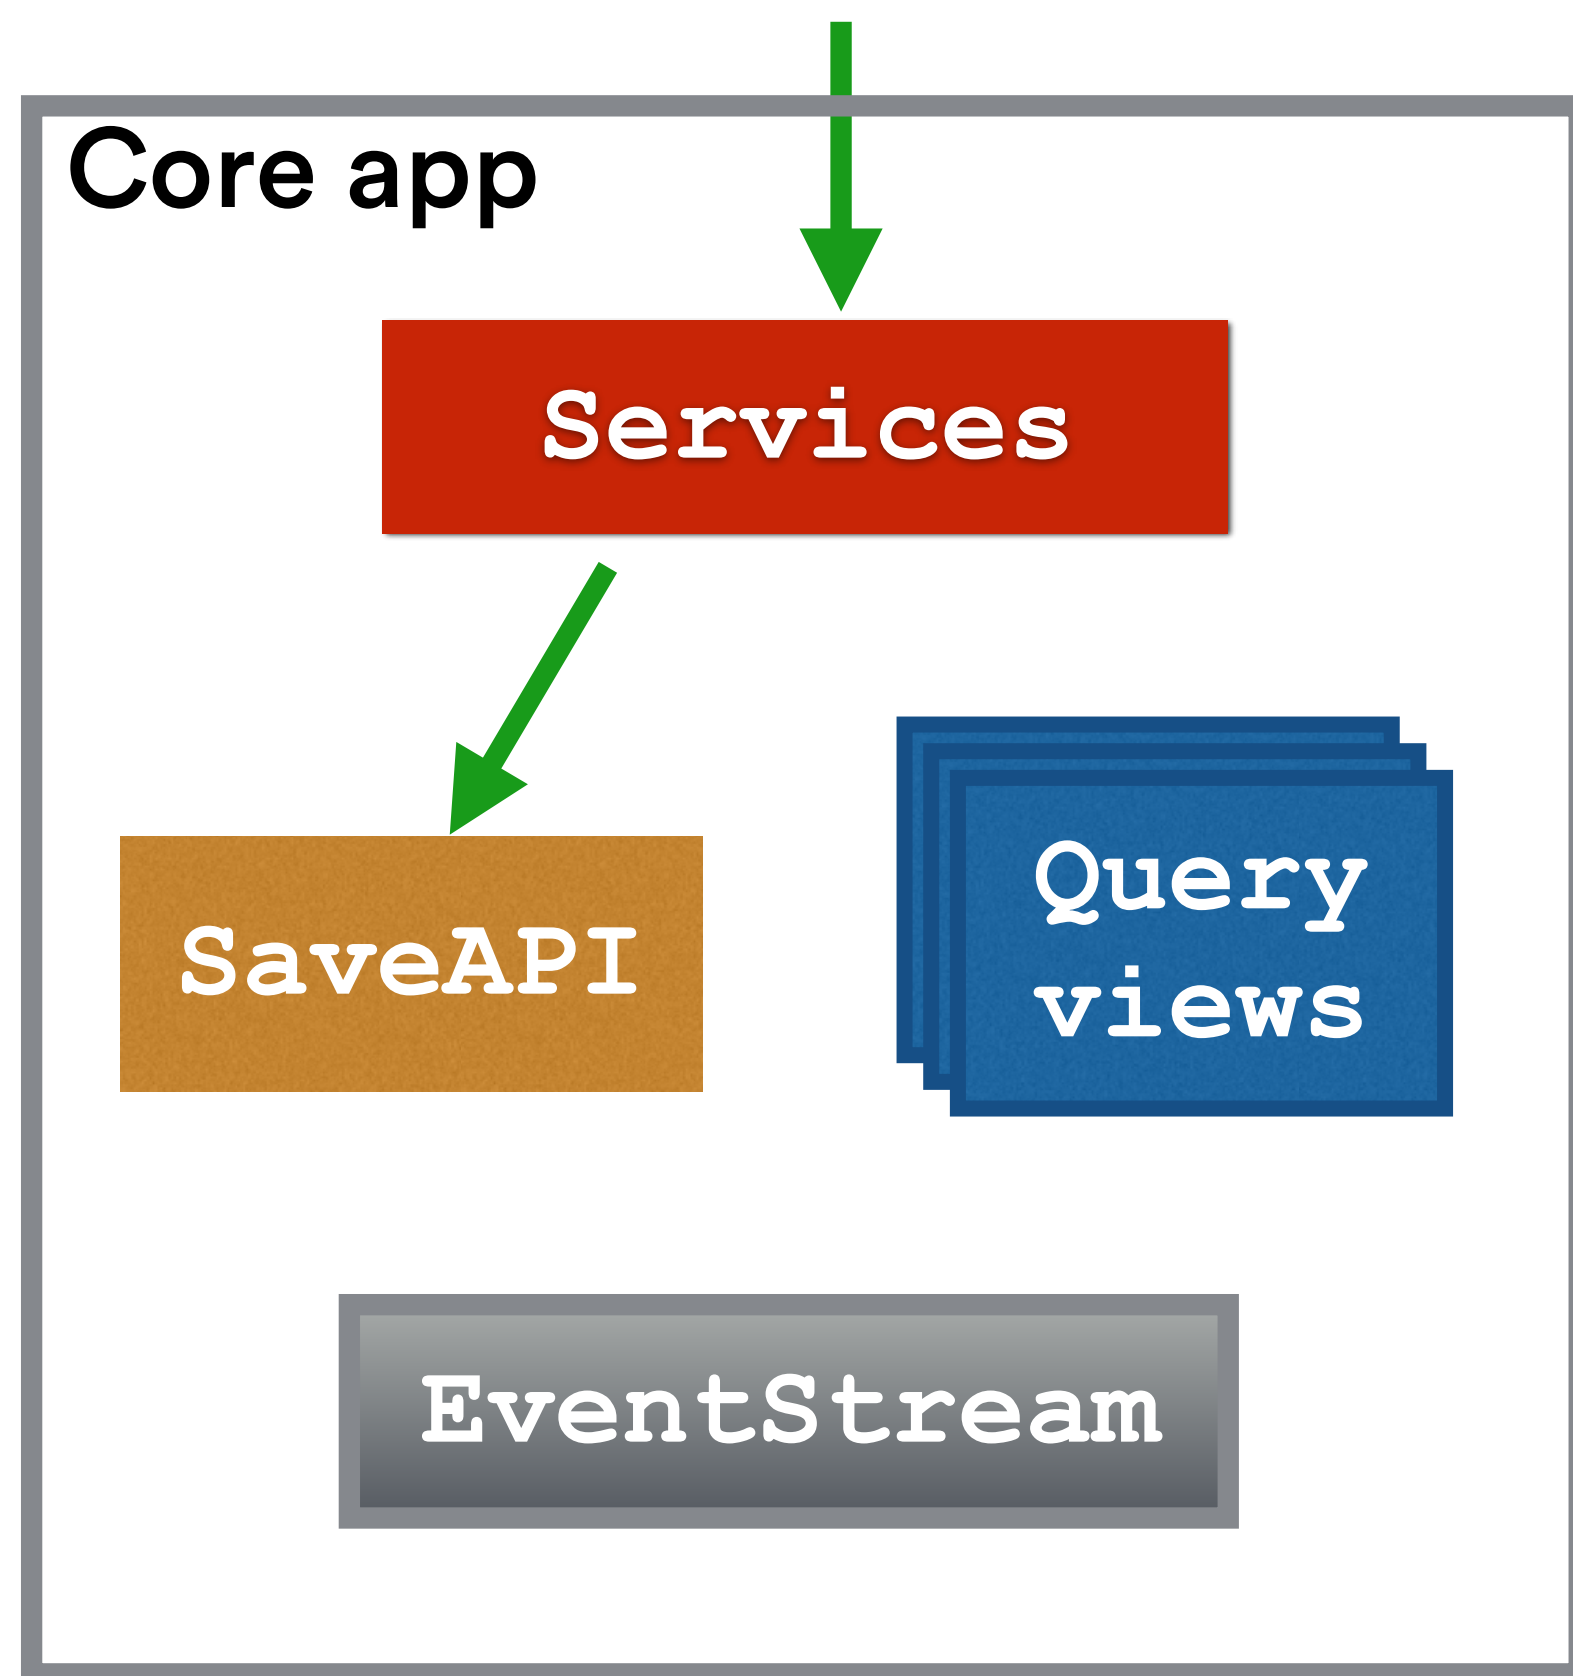


DynamoDB



REST calls

e.g. Add User

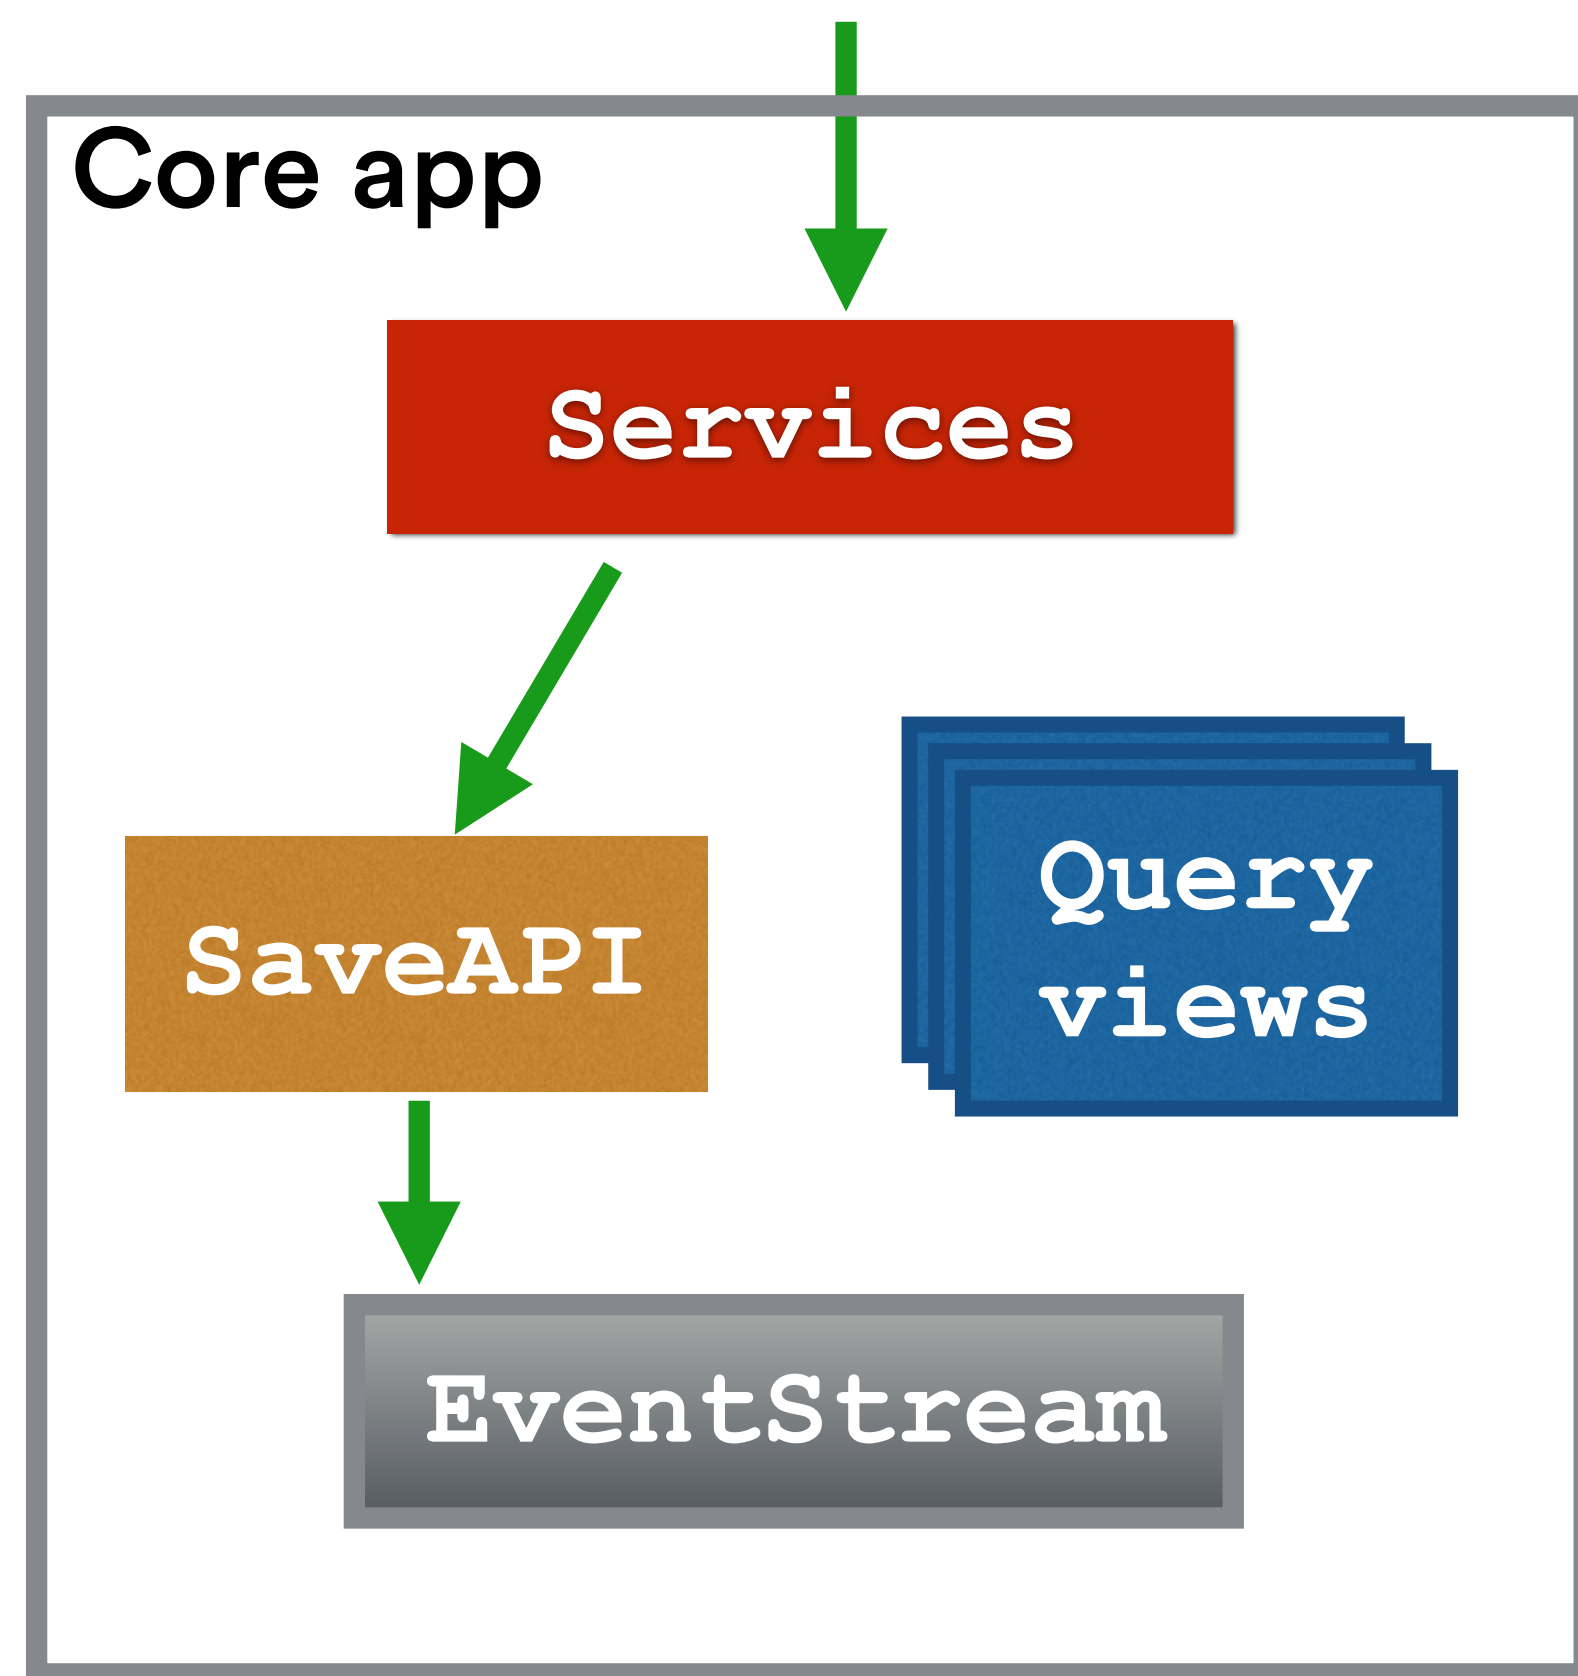


DynamoDB



REST calls

e.g. Add User

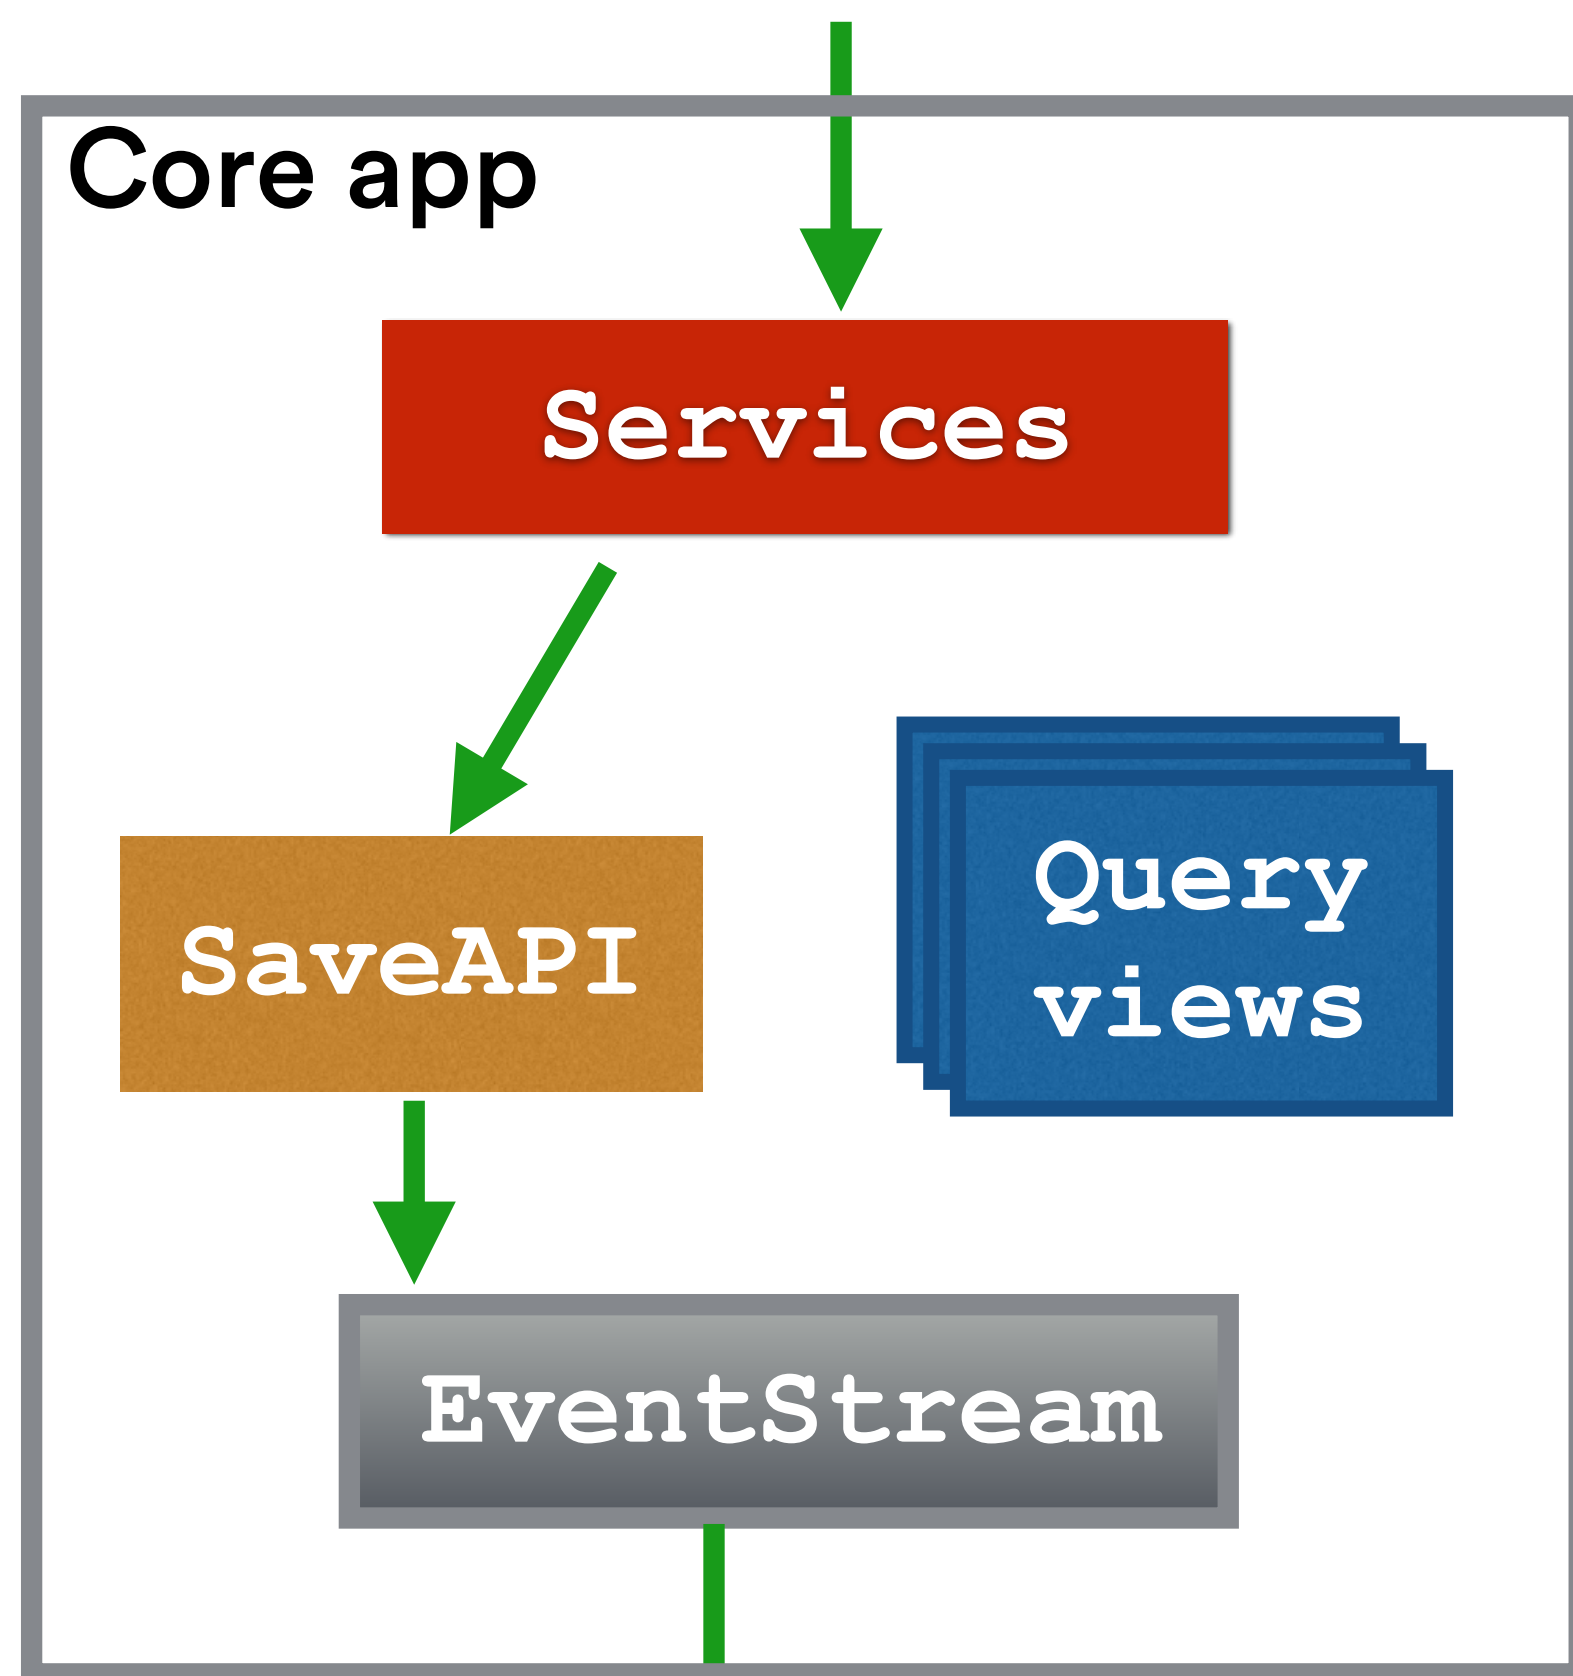


DynamoDB



REST calls

e.g. Add User

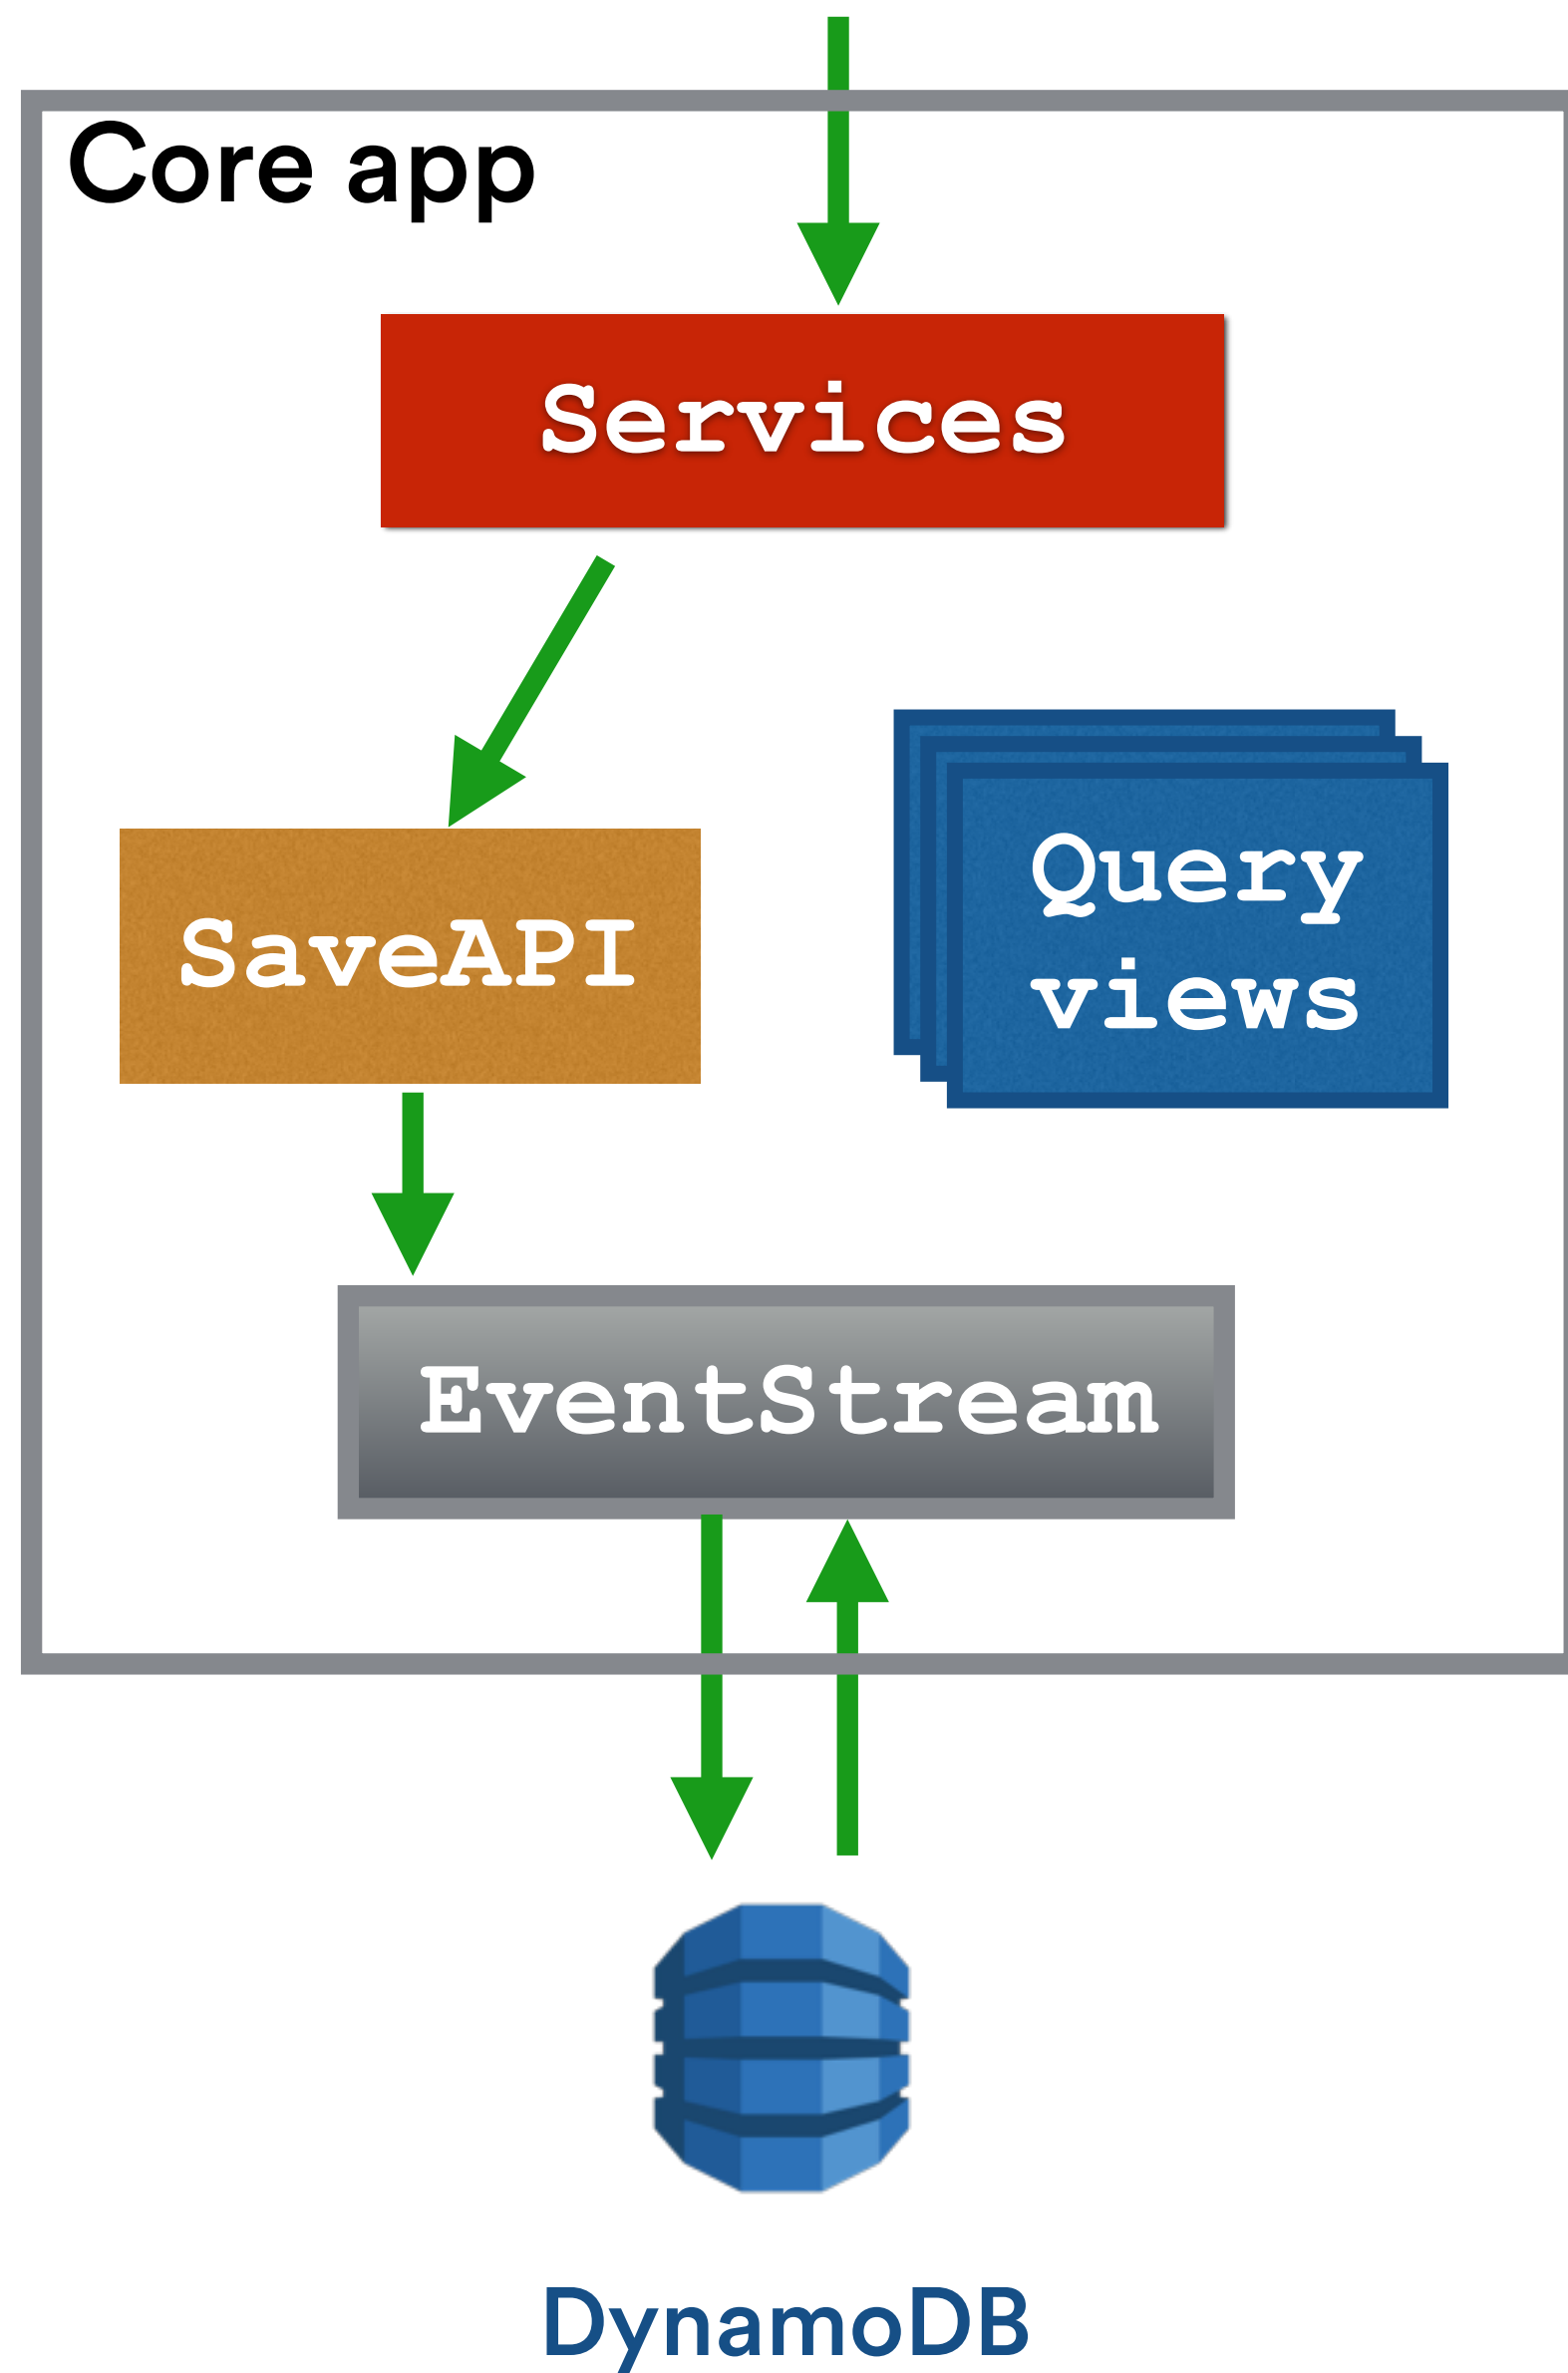


DynamoDB



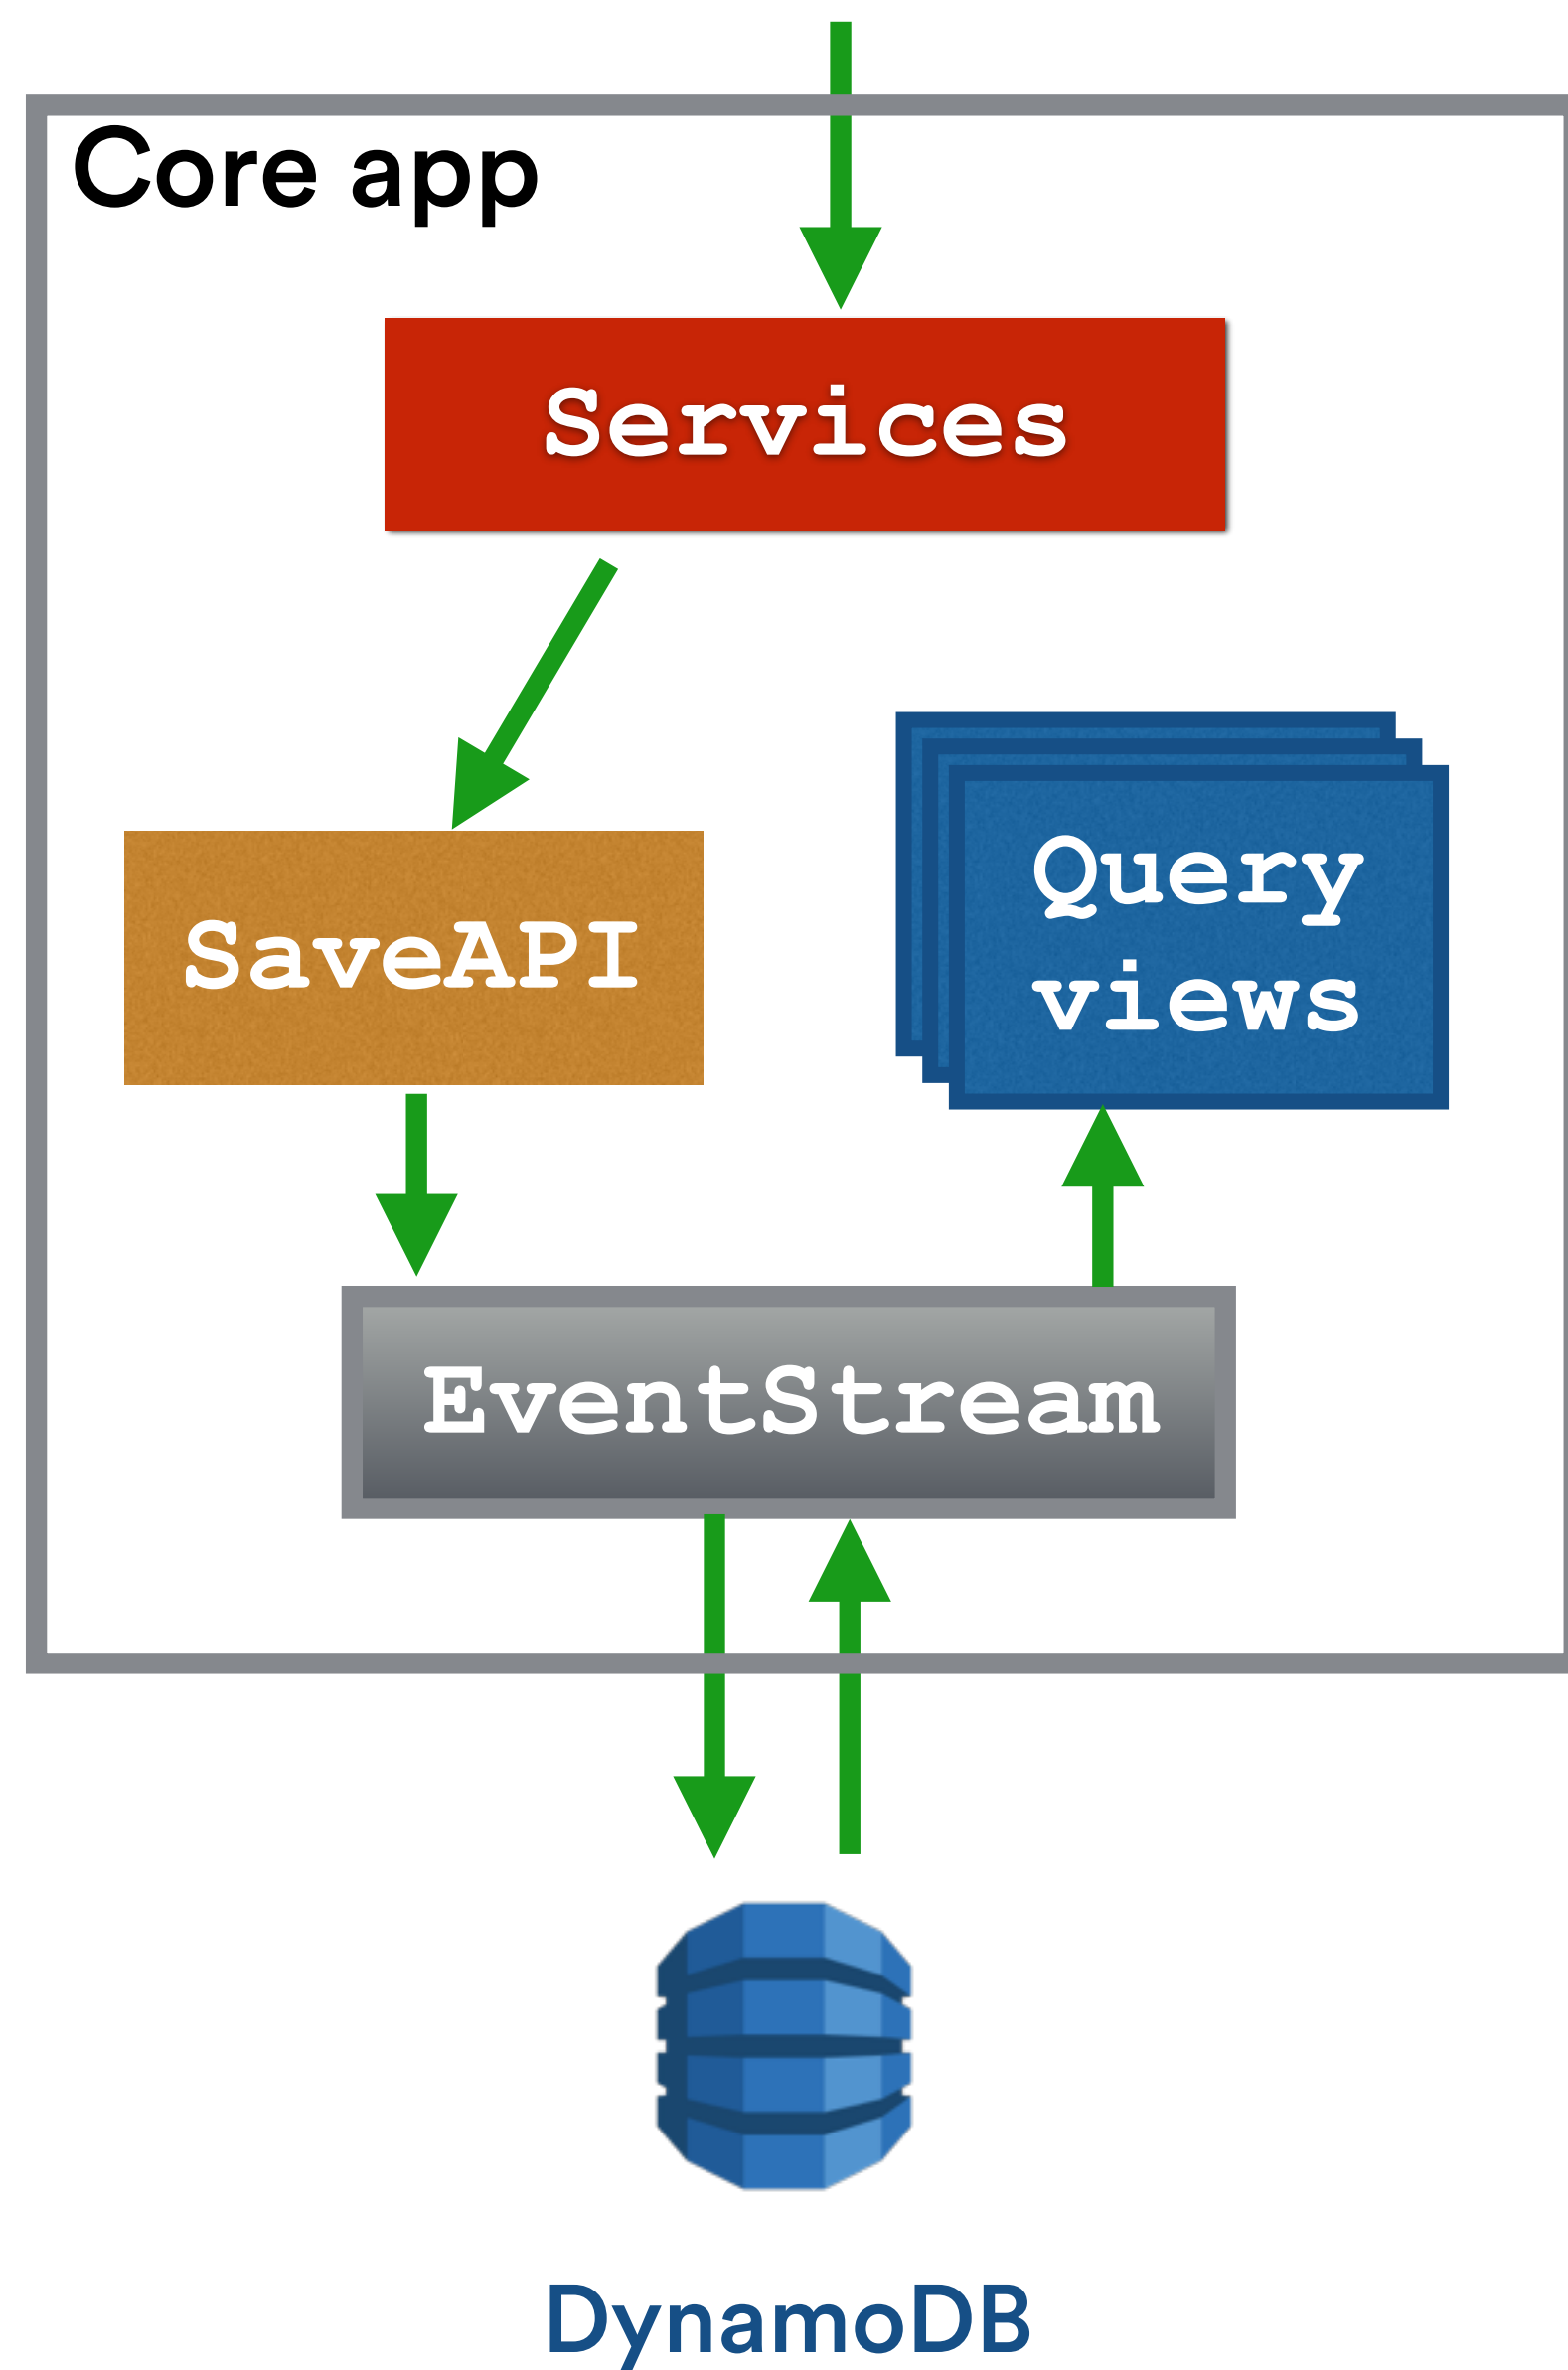
REST calls

e.g. Add User



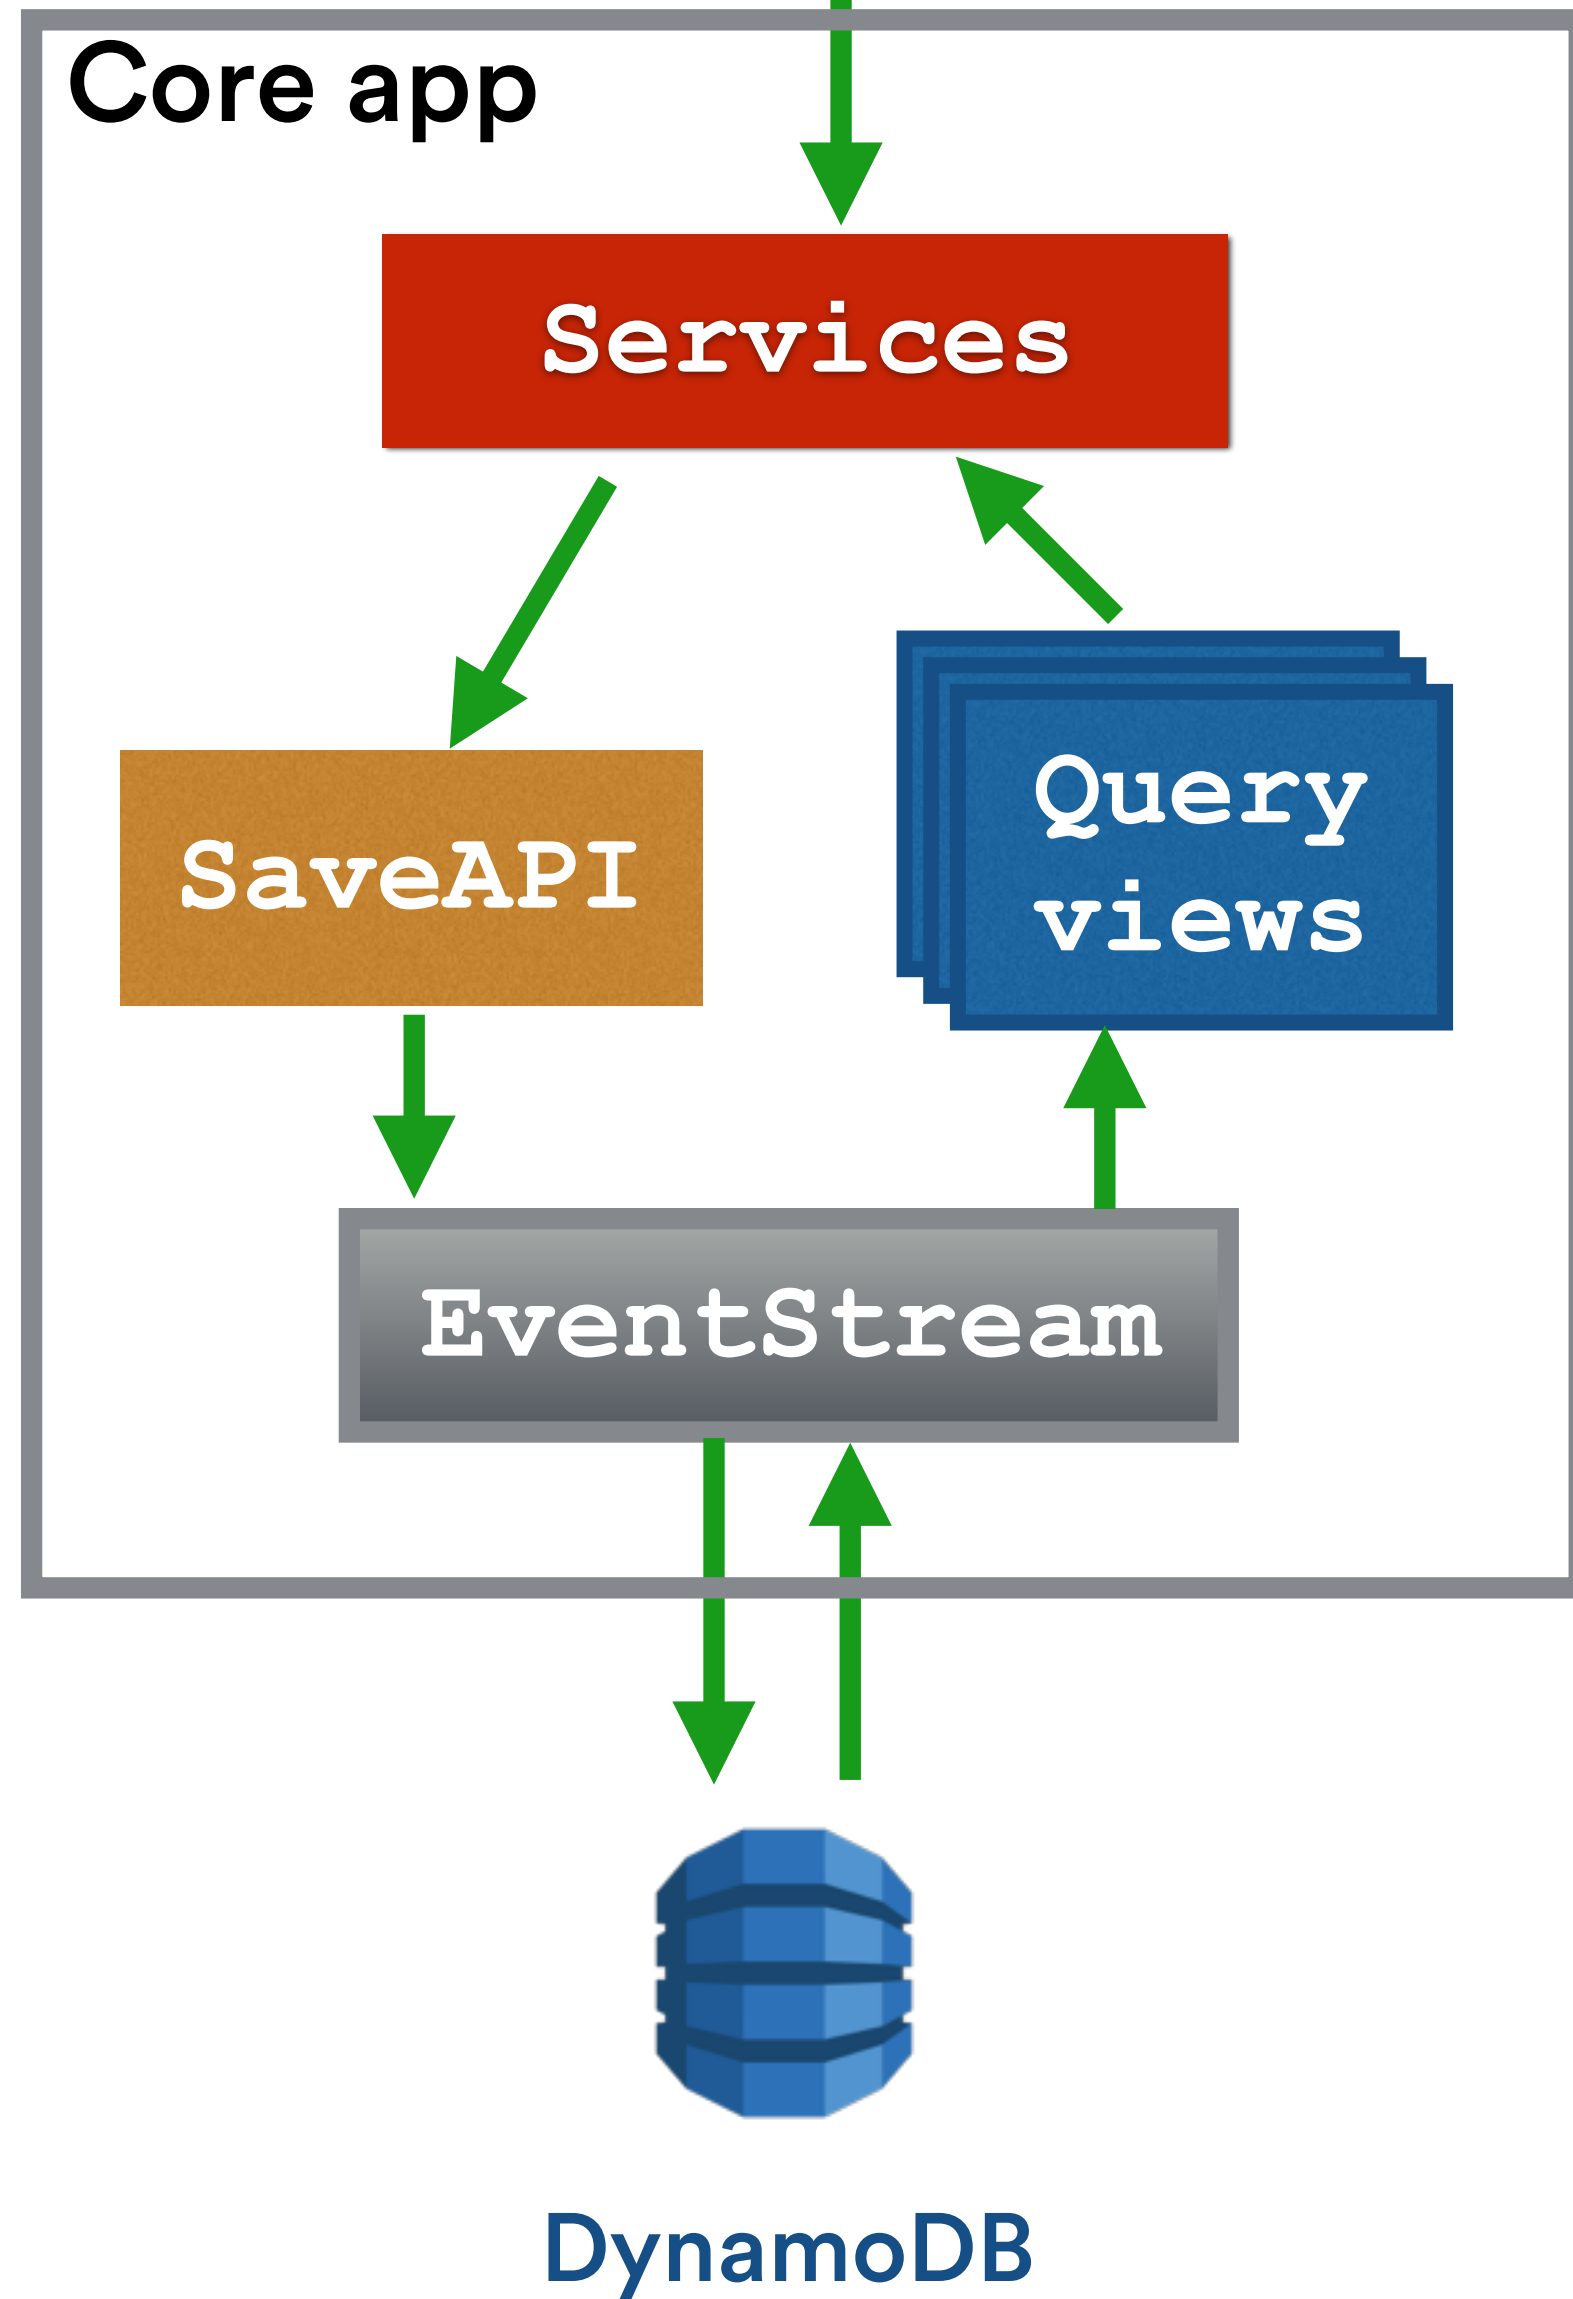
REST calls

e.g. Add User



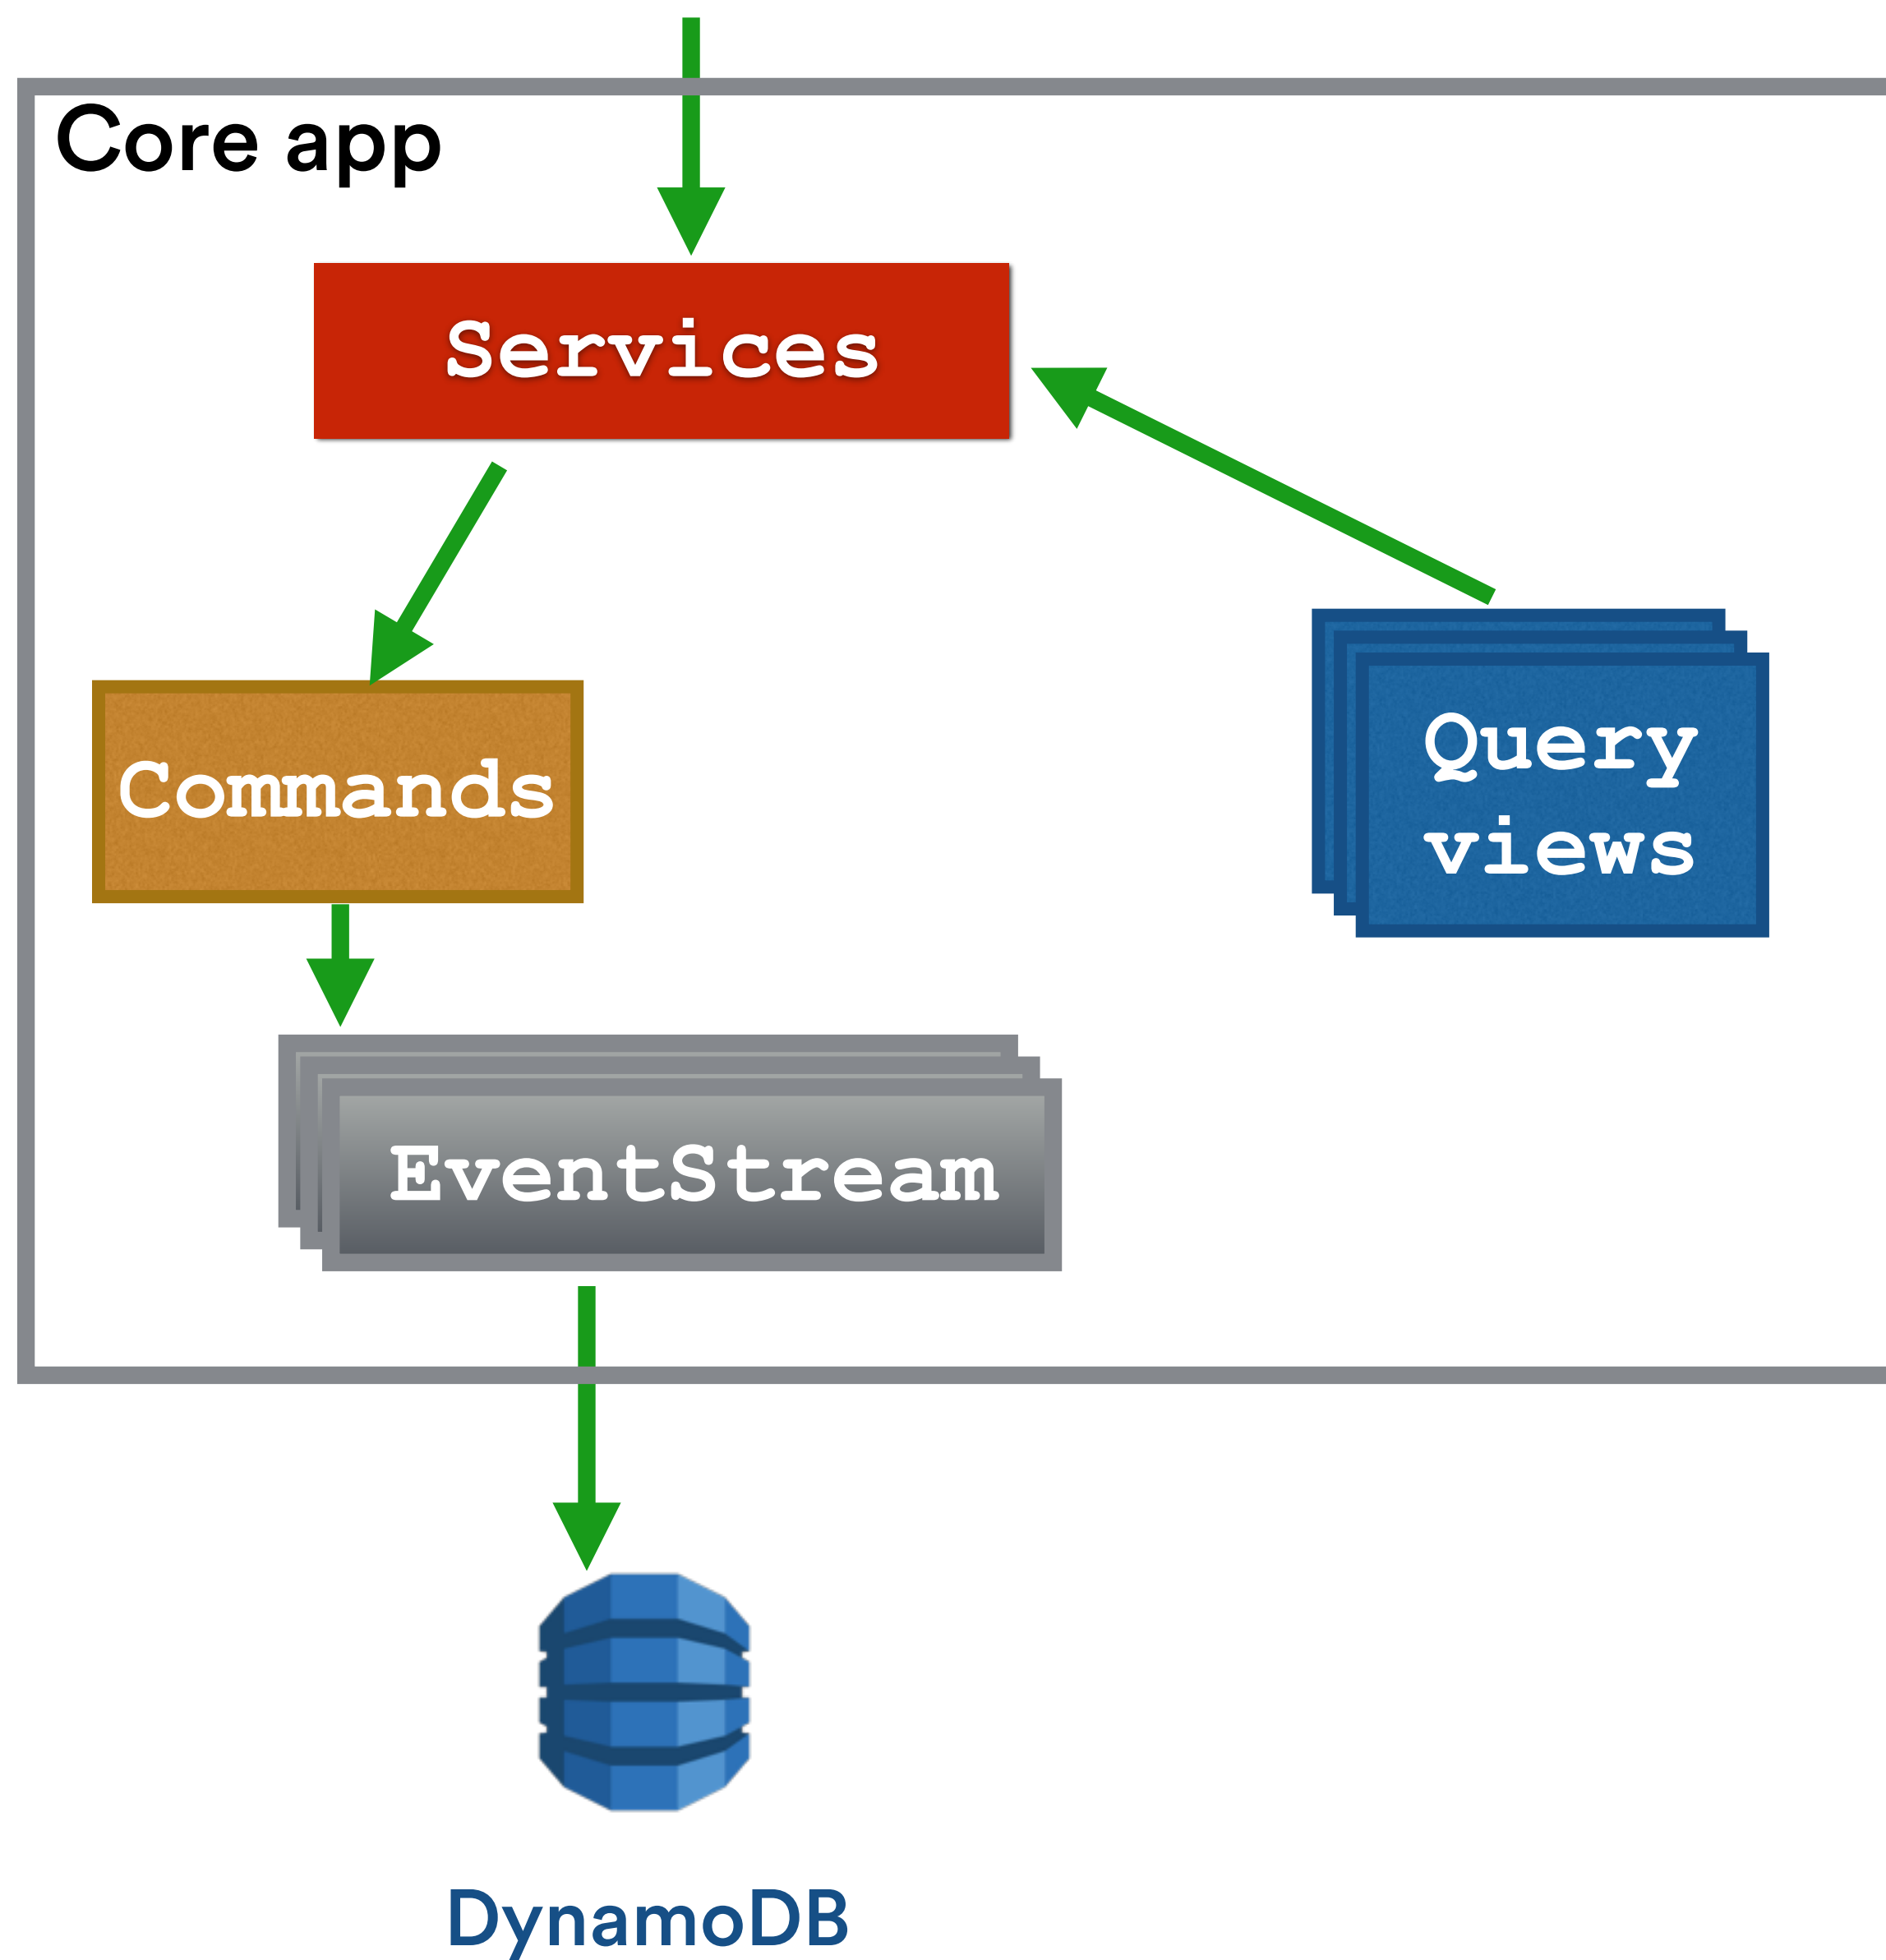
REST calls

e.g. Add User

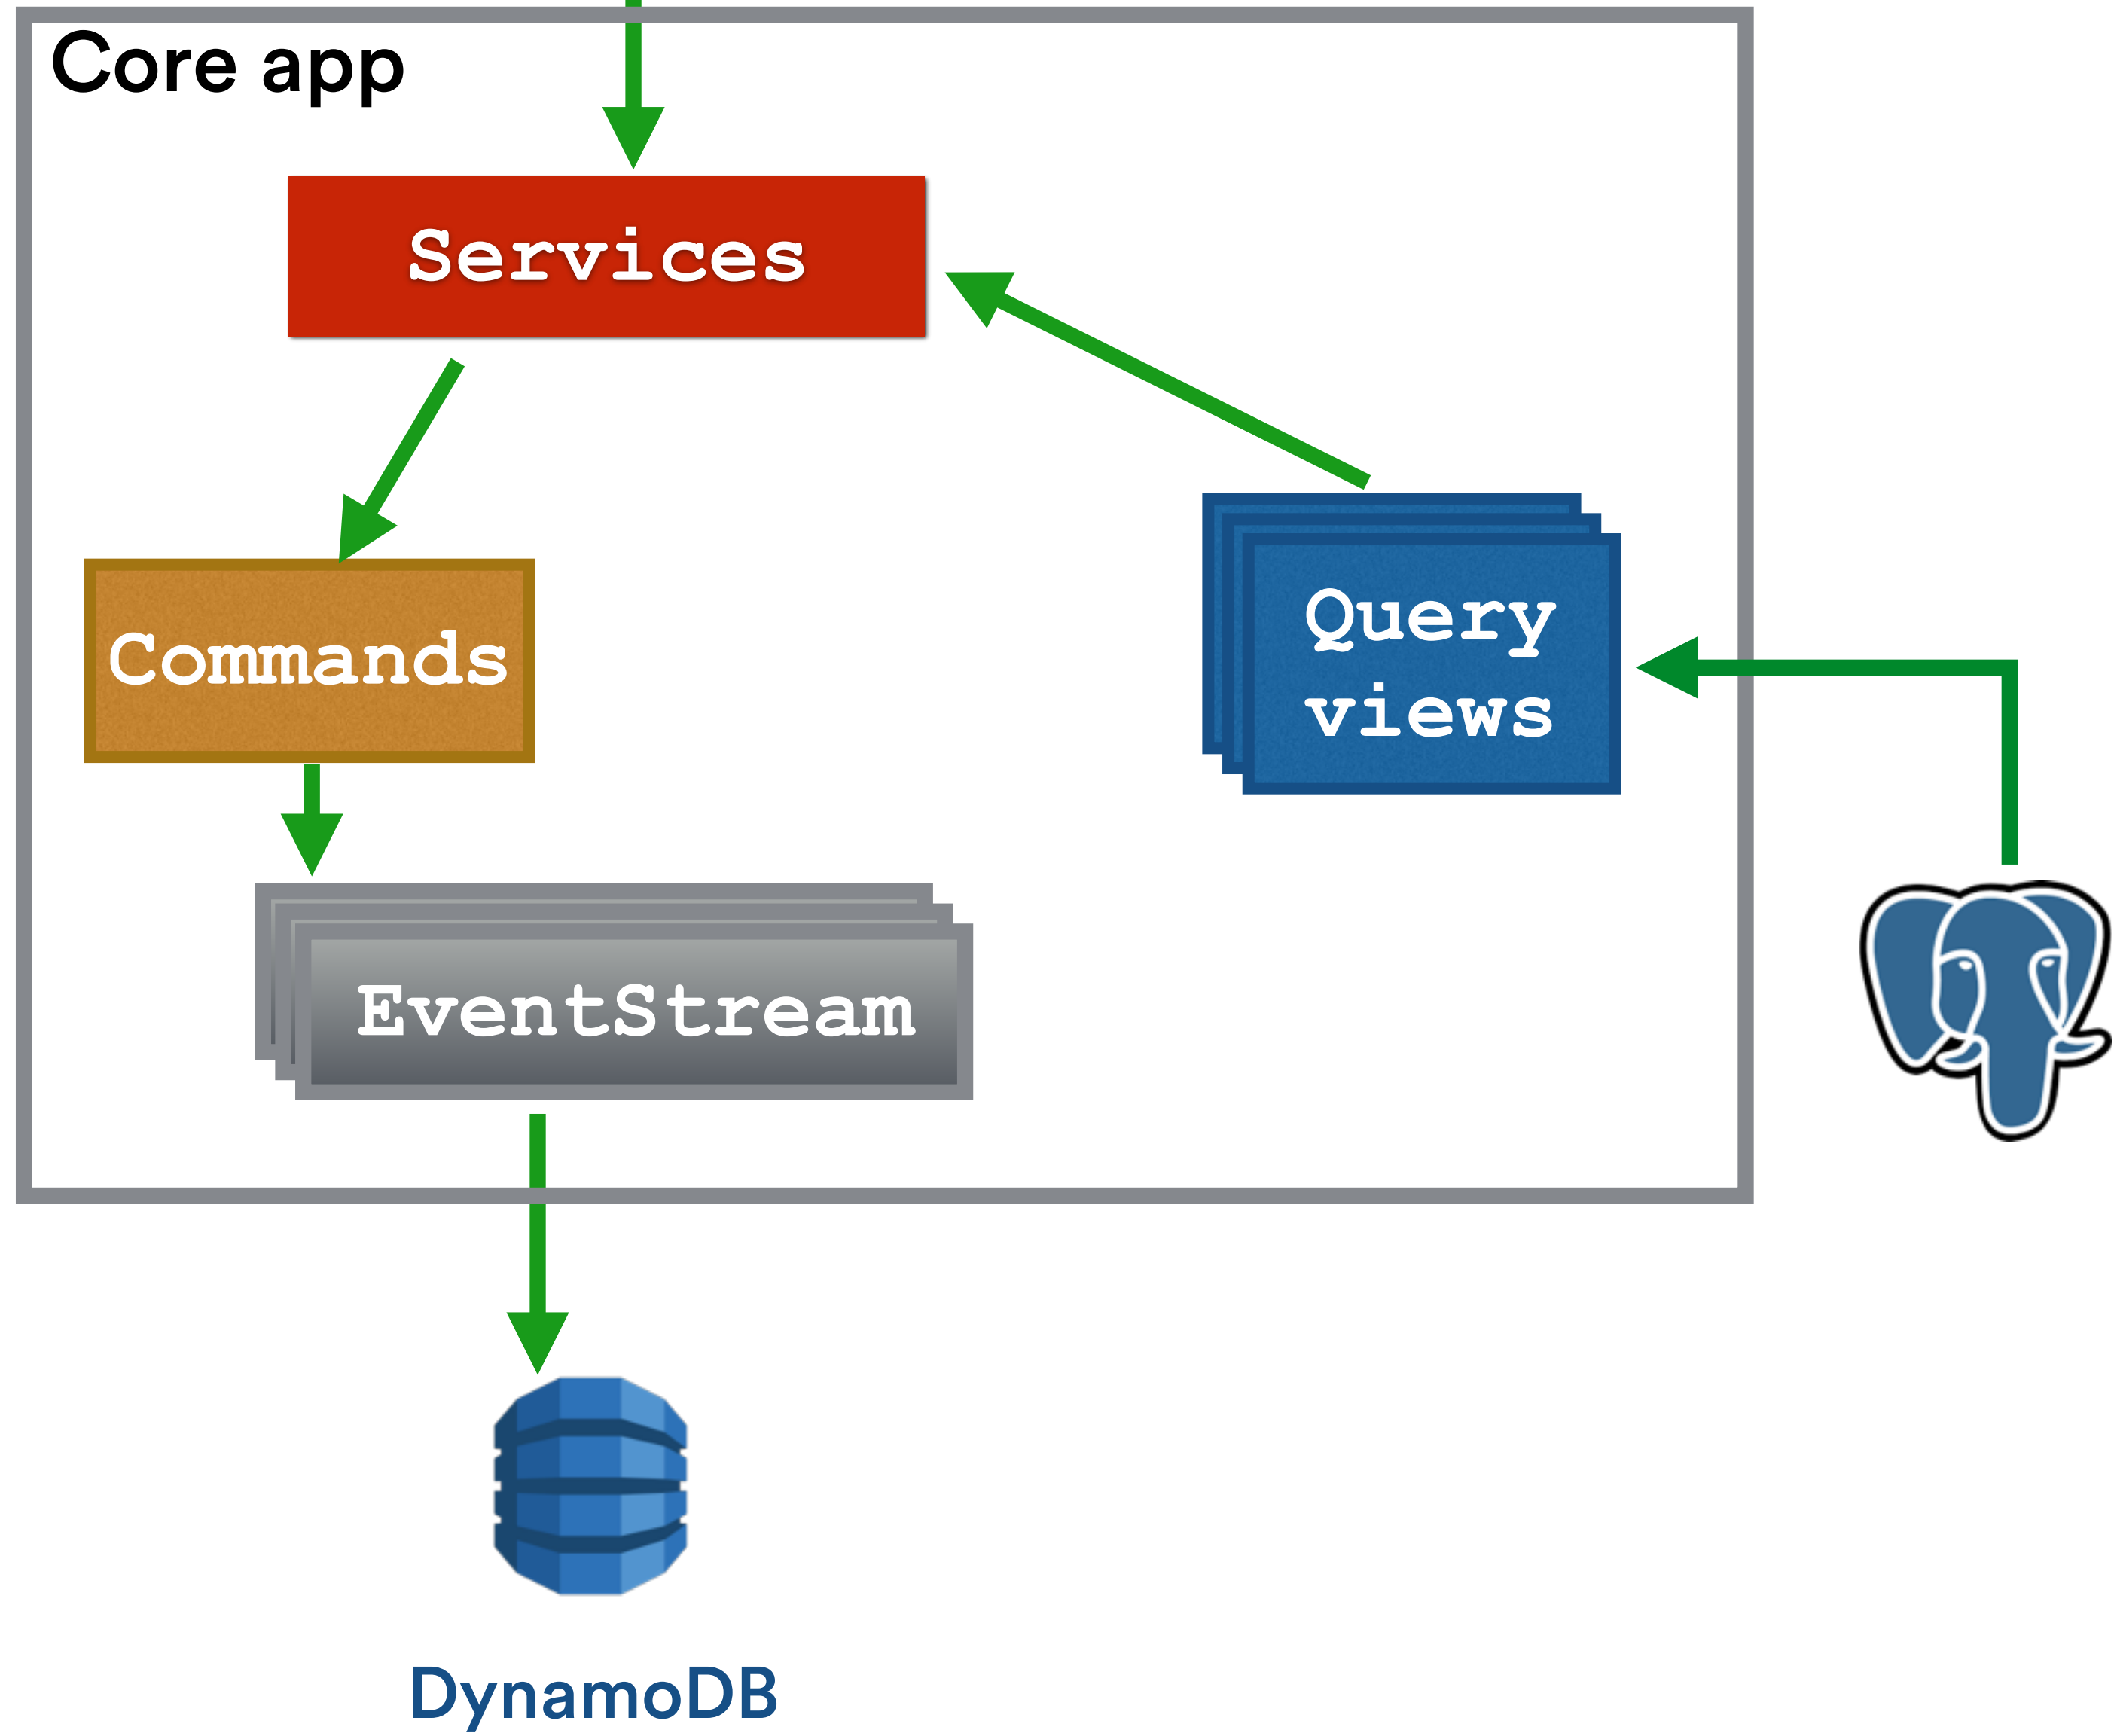


REST calls

e.g. Add User

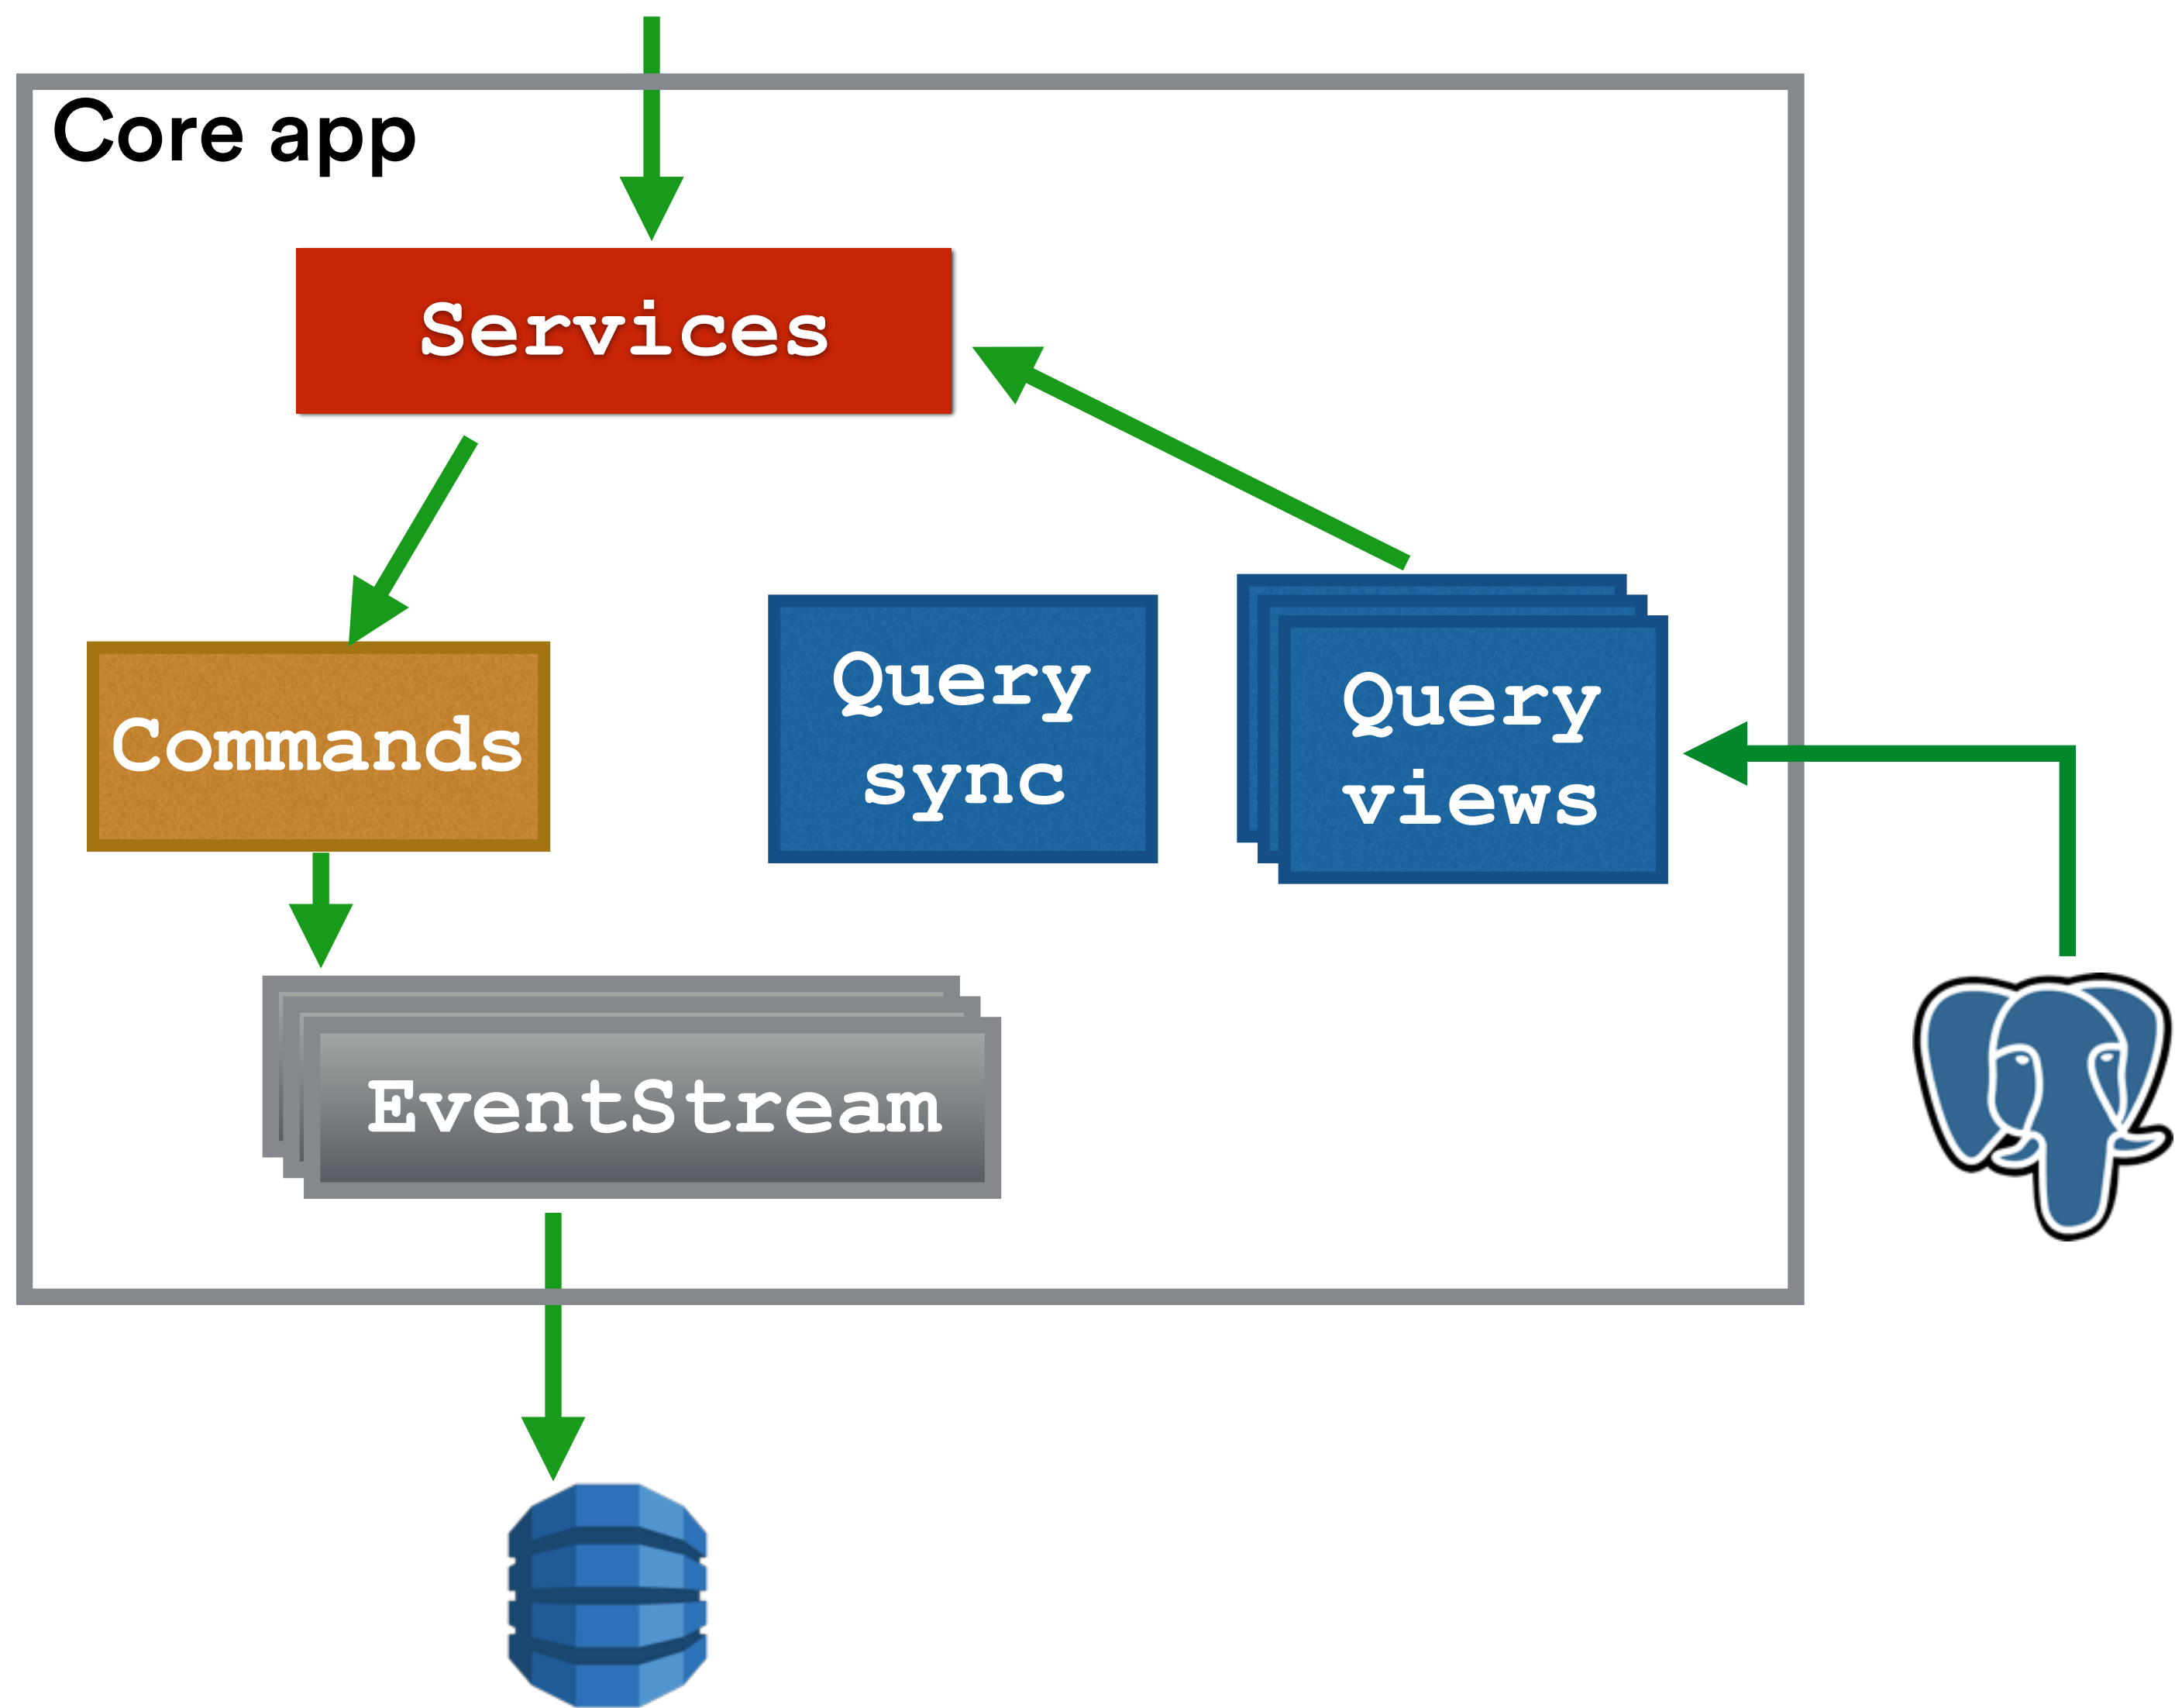


REST calls
e.g. Add User



REST calls

e.g. Add User

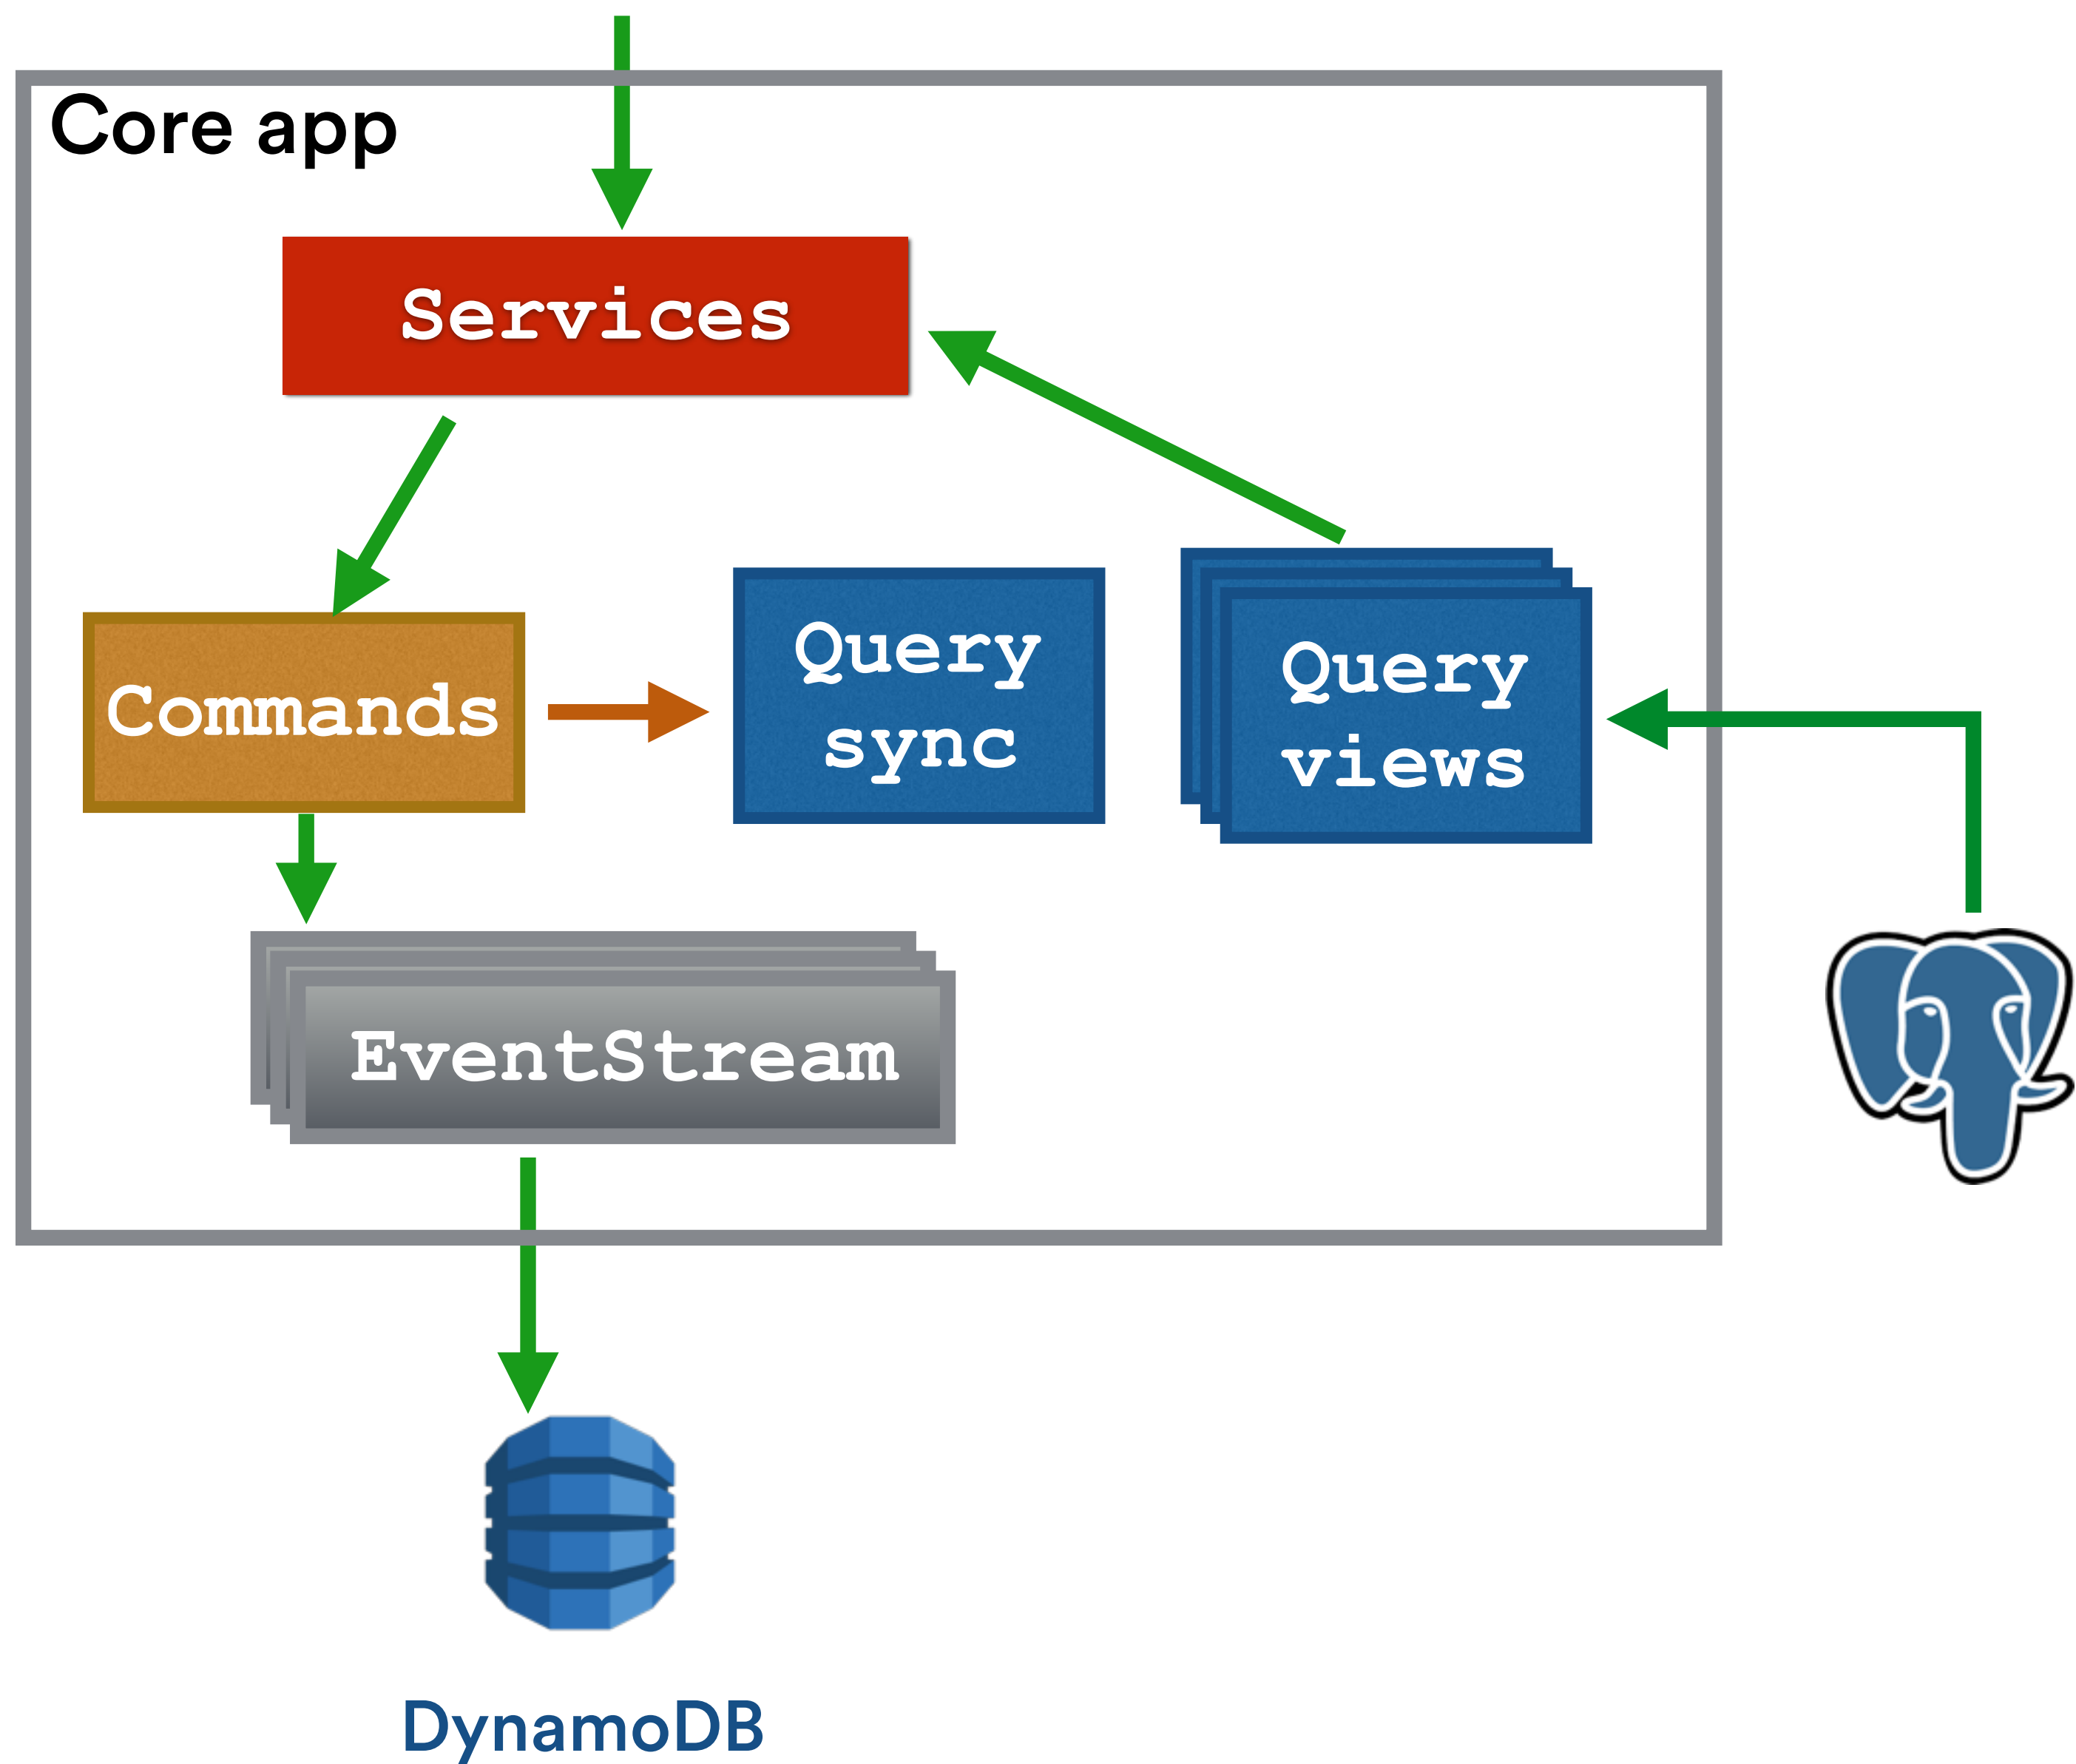


DynamoDB



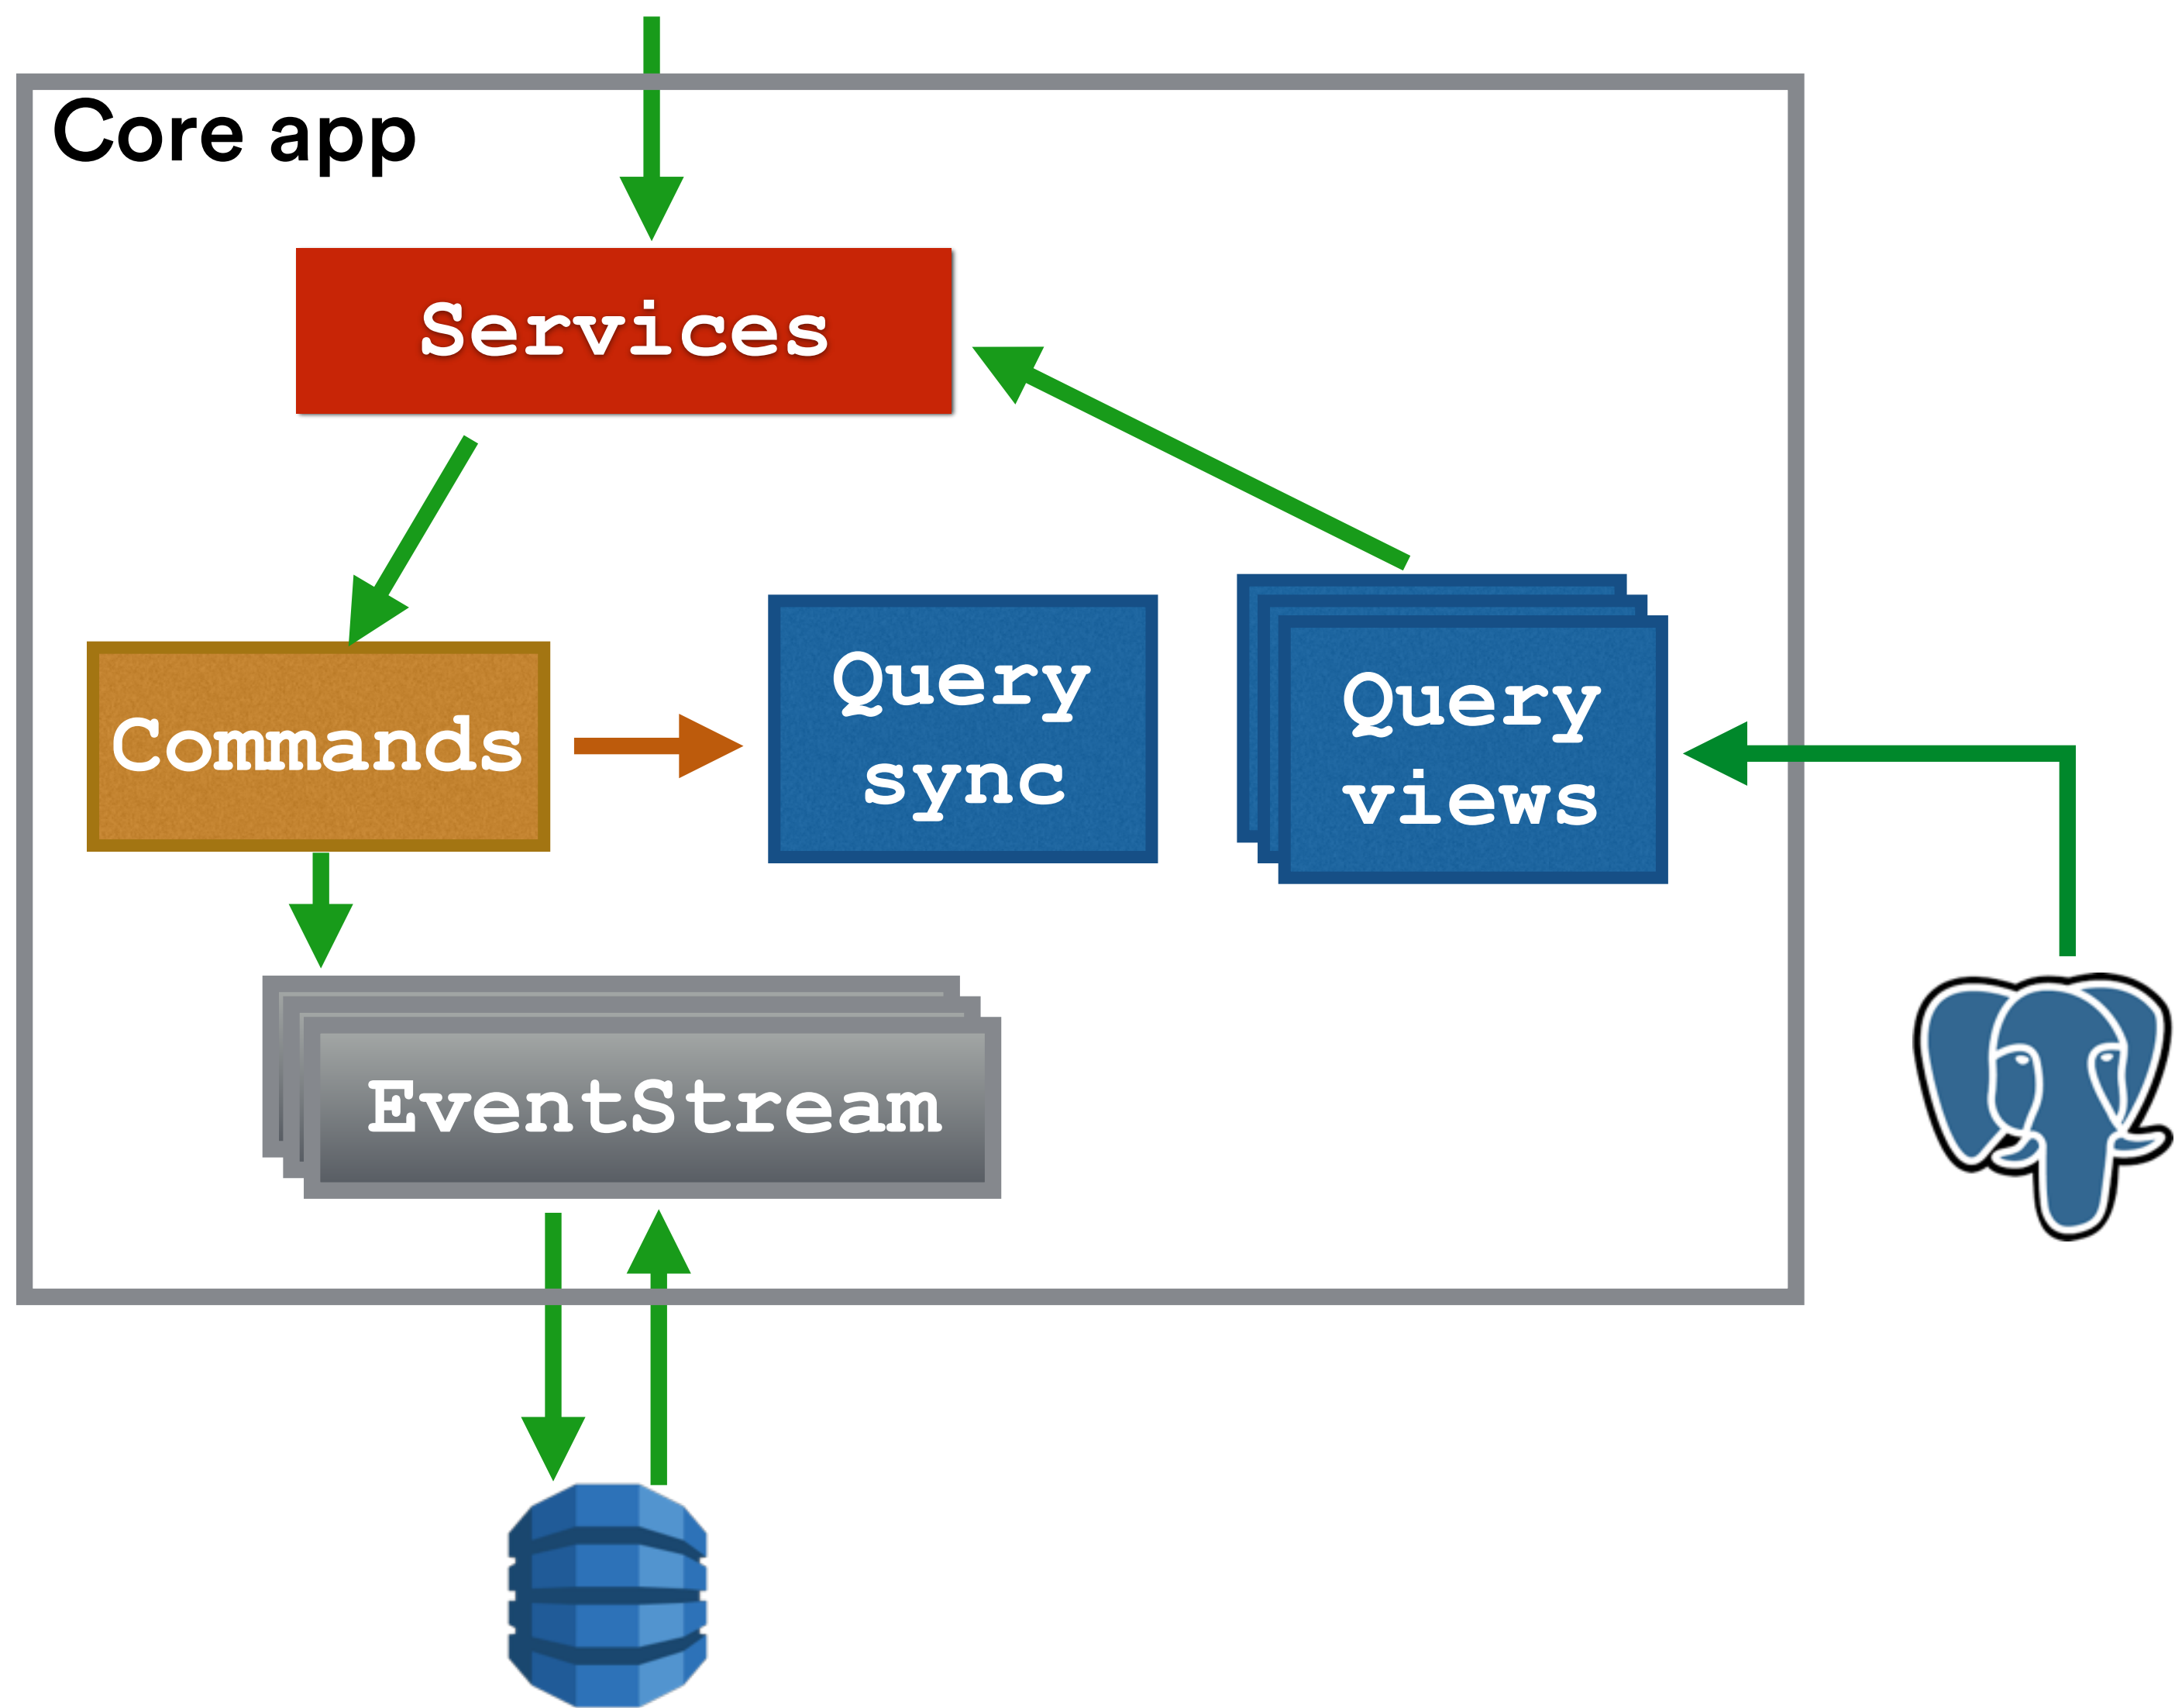
REST calls

e.g. Add User



REST calls

e.g. Add User

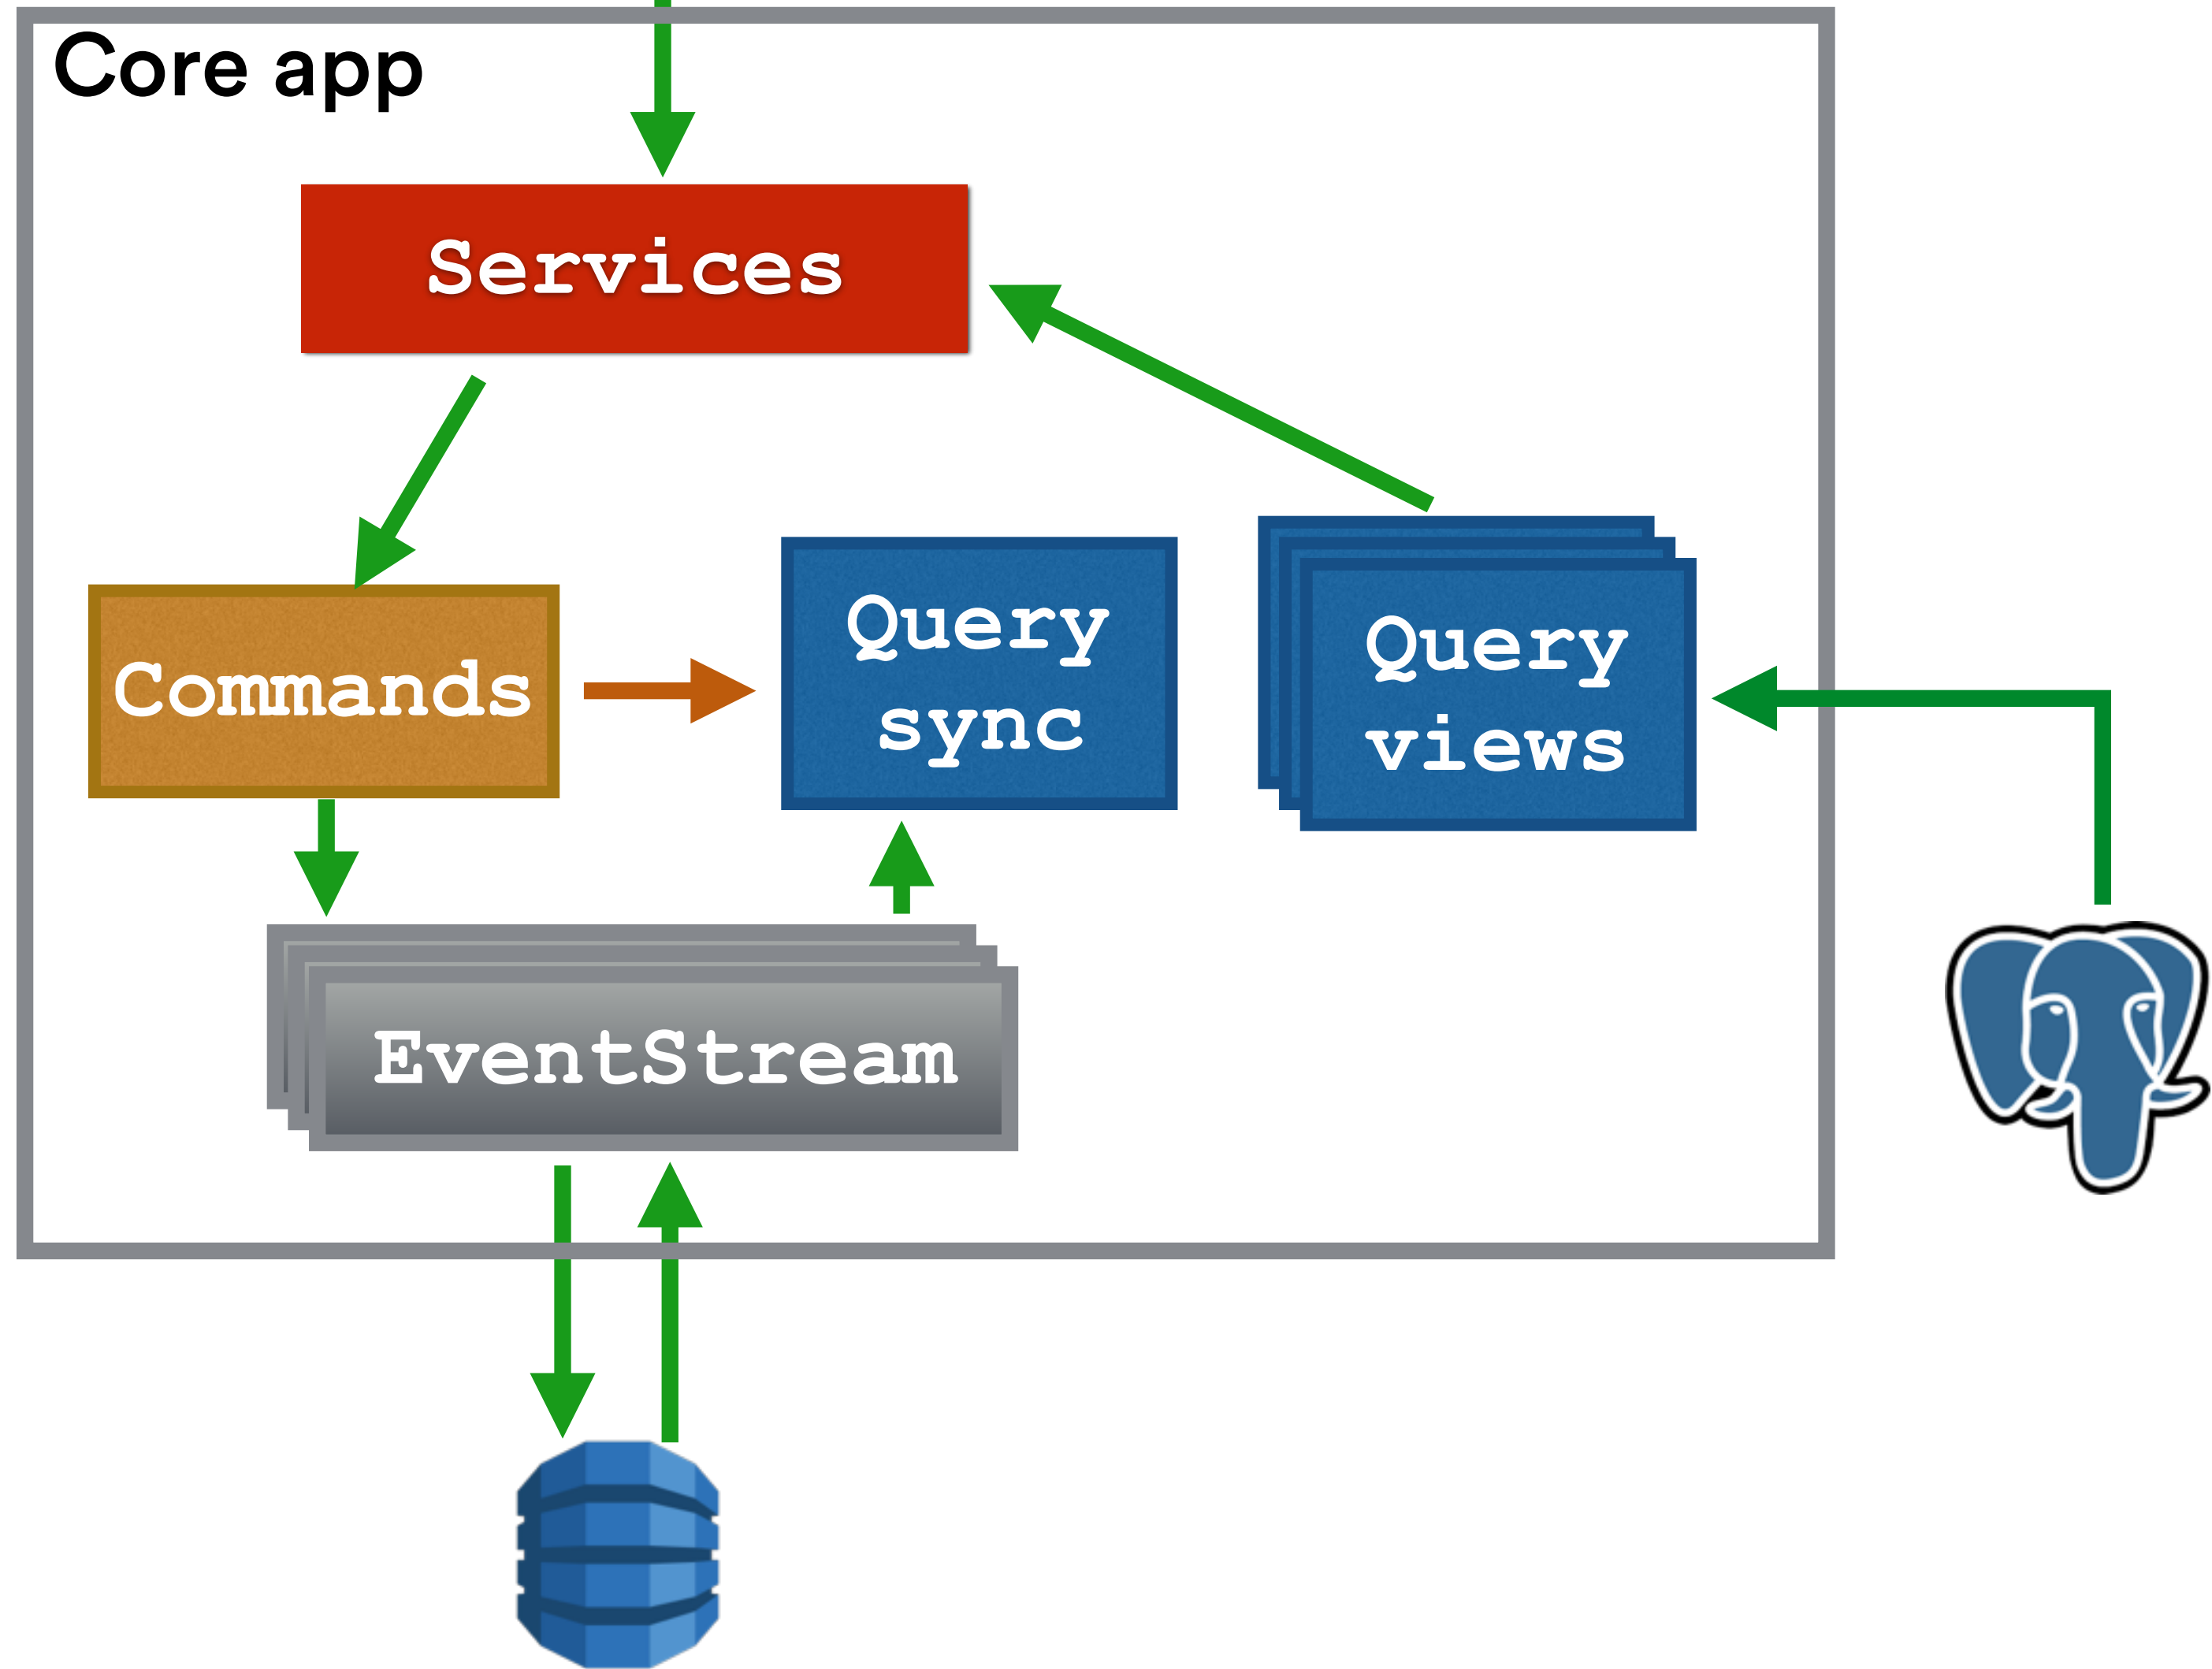


DynamoDB



REST calls

e.g. Add User

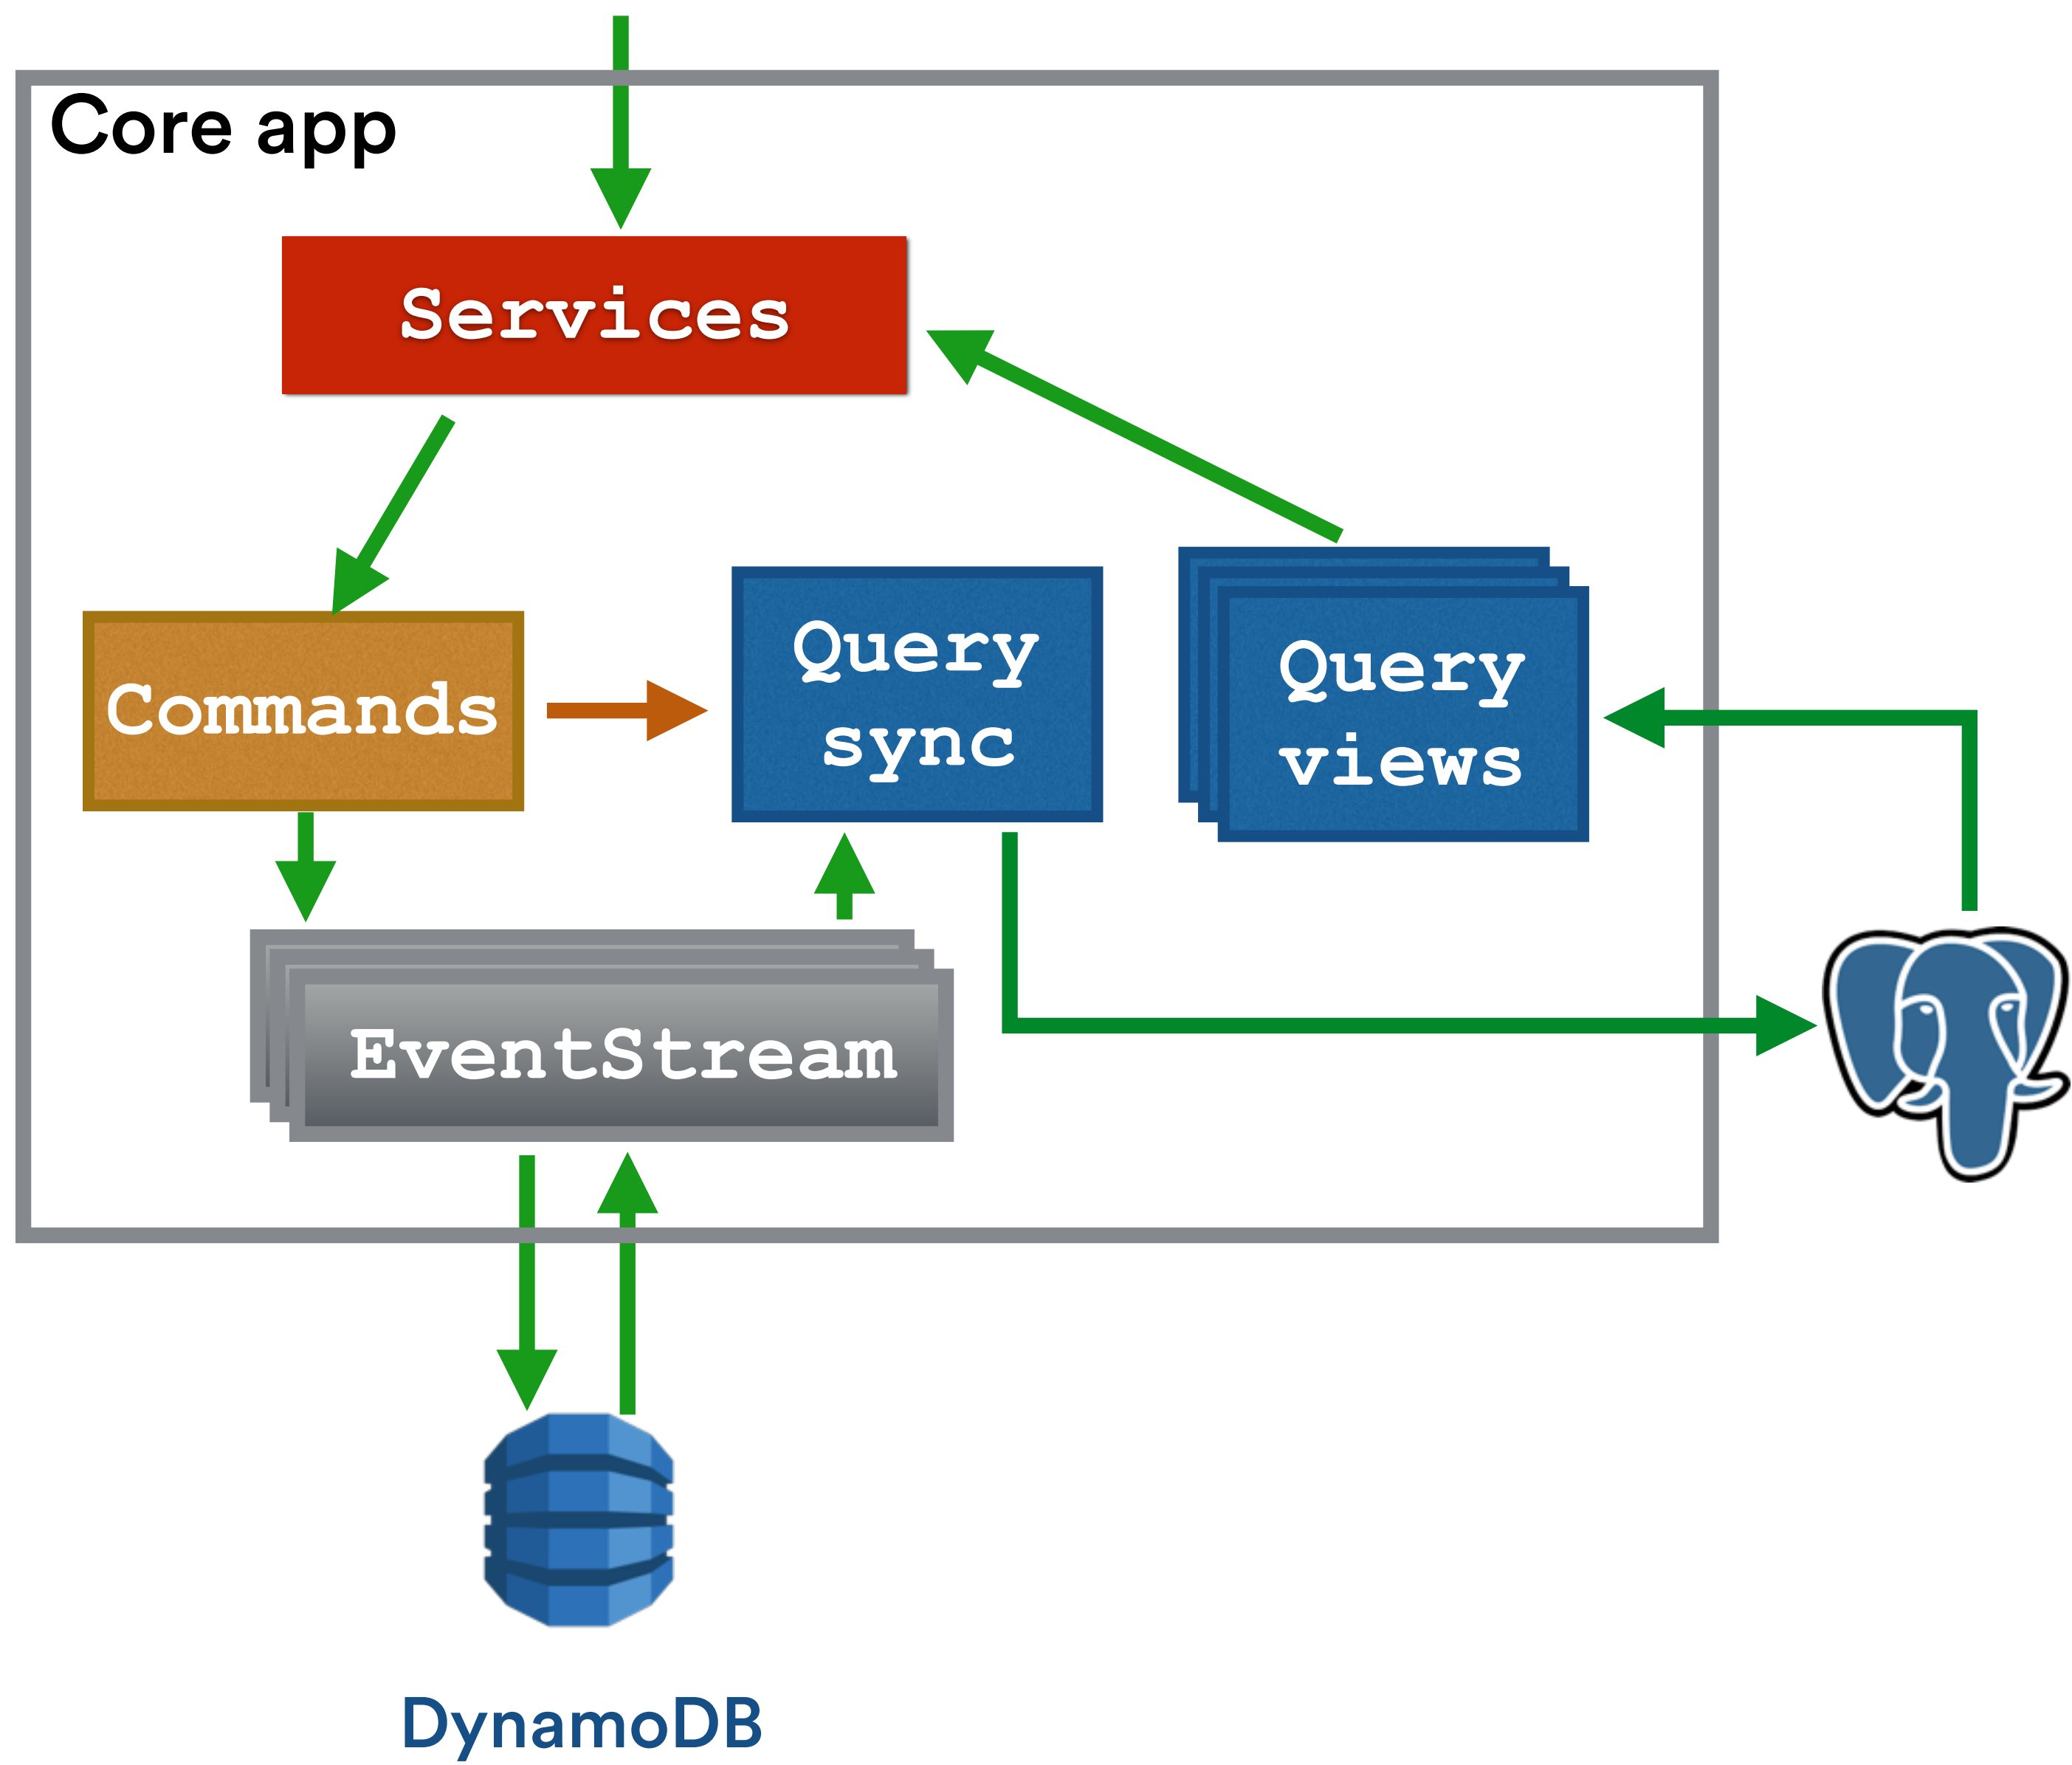


DynamoDB



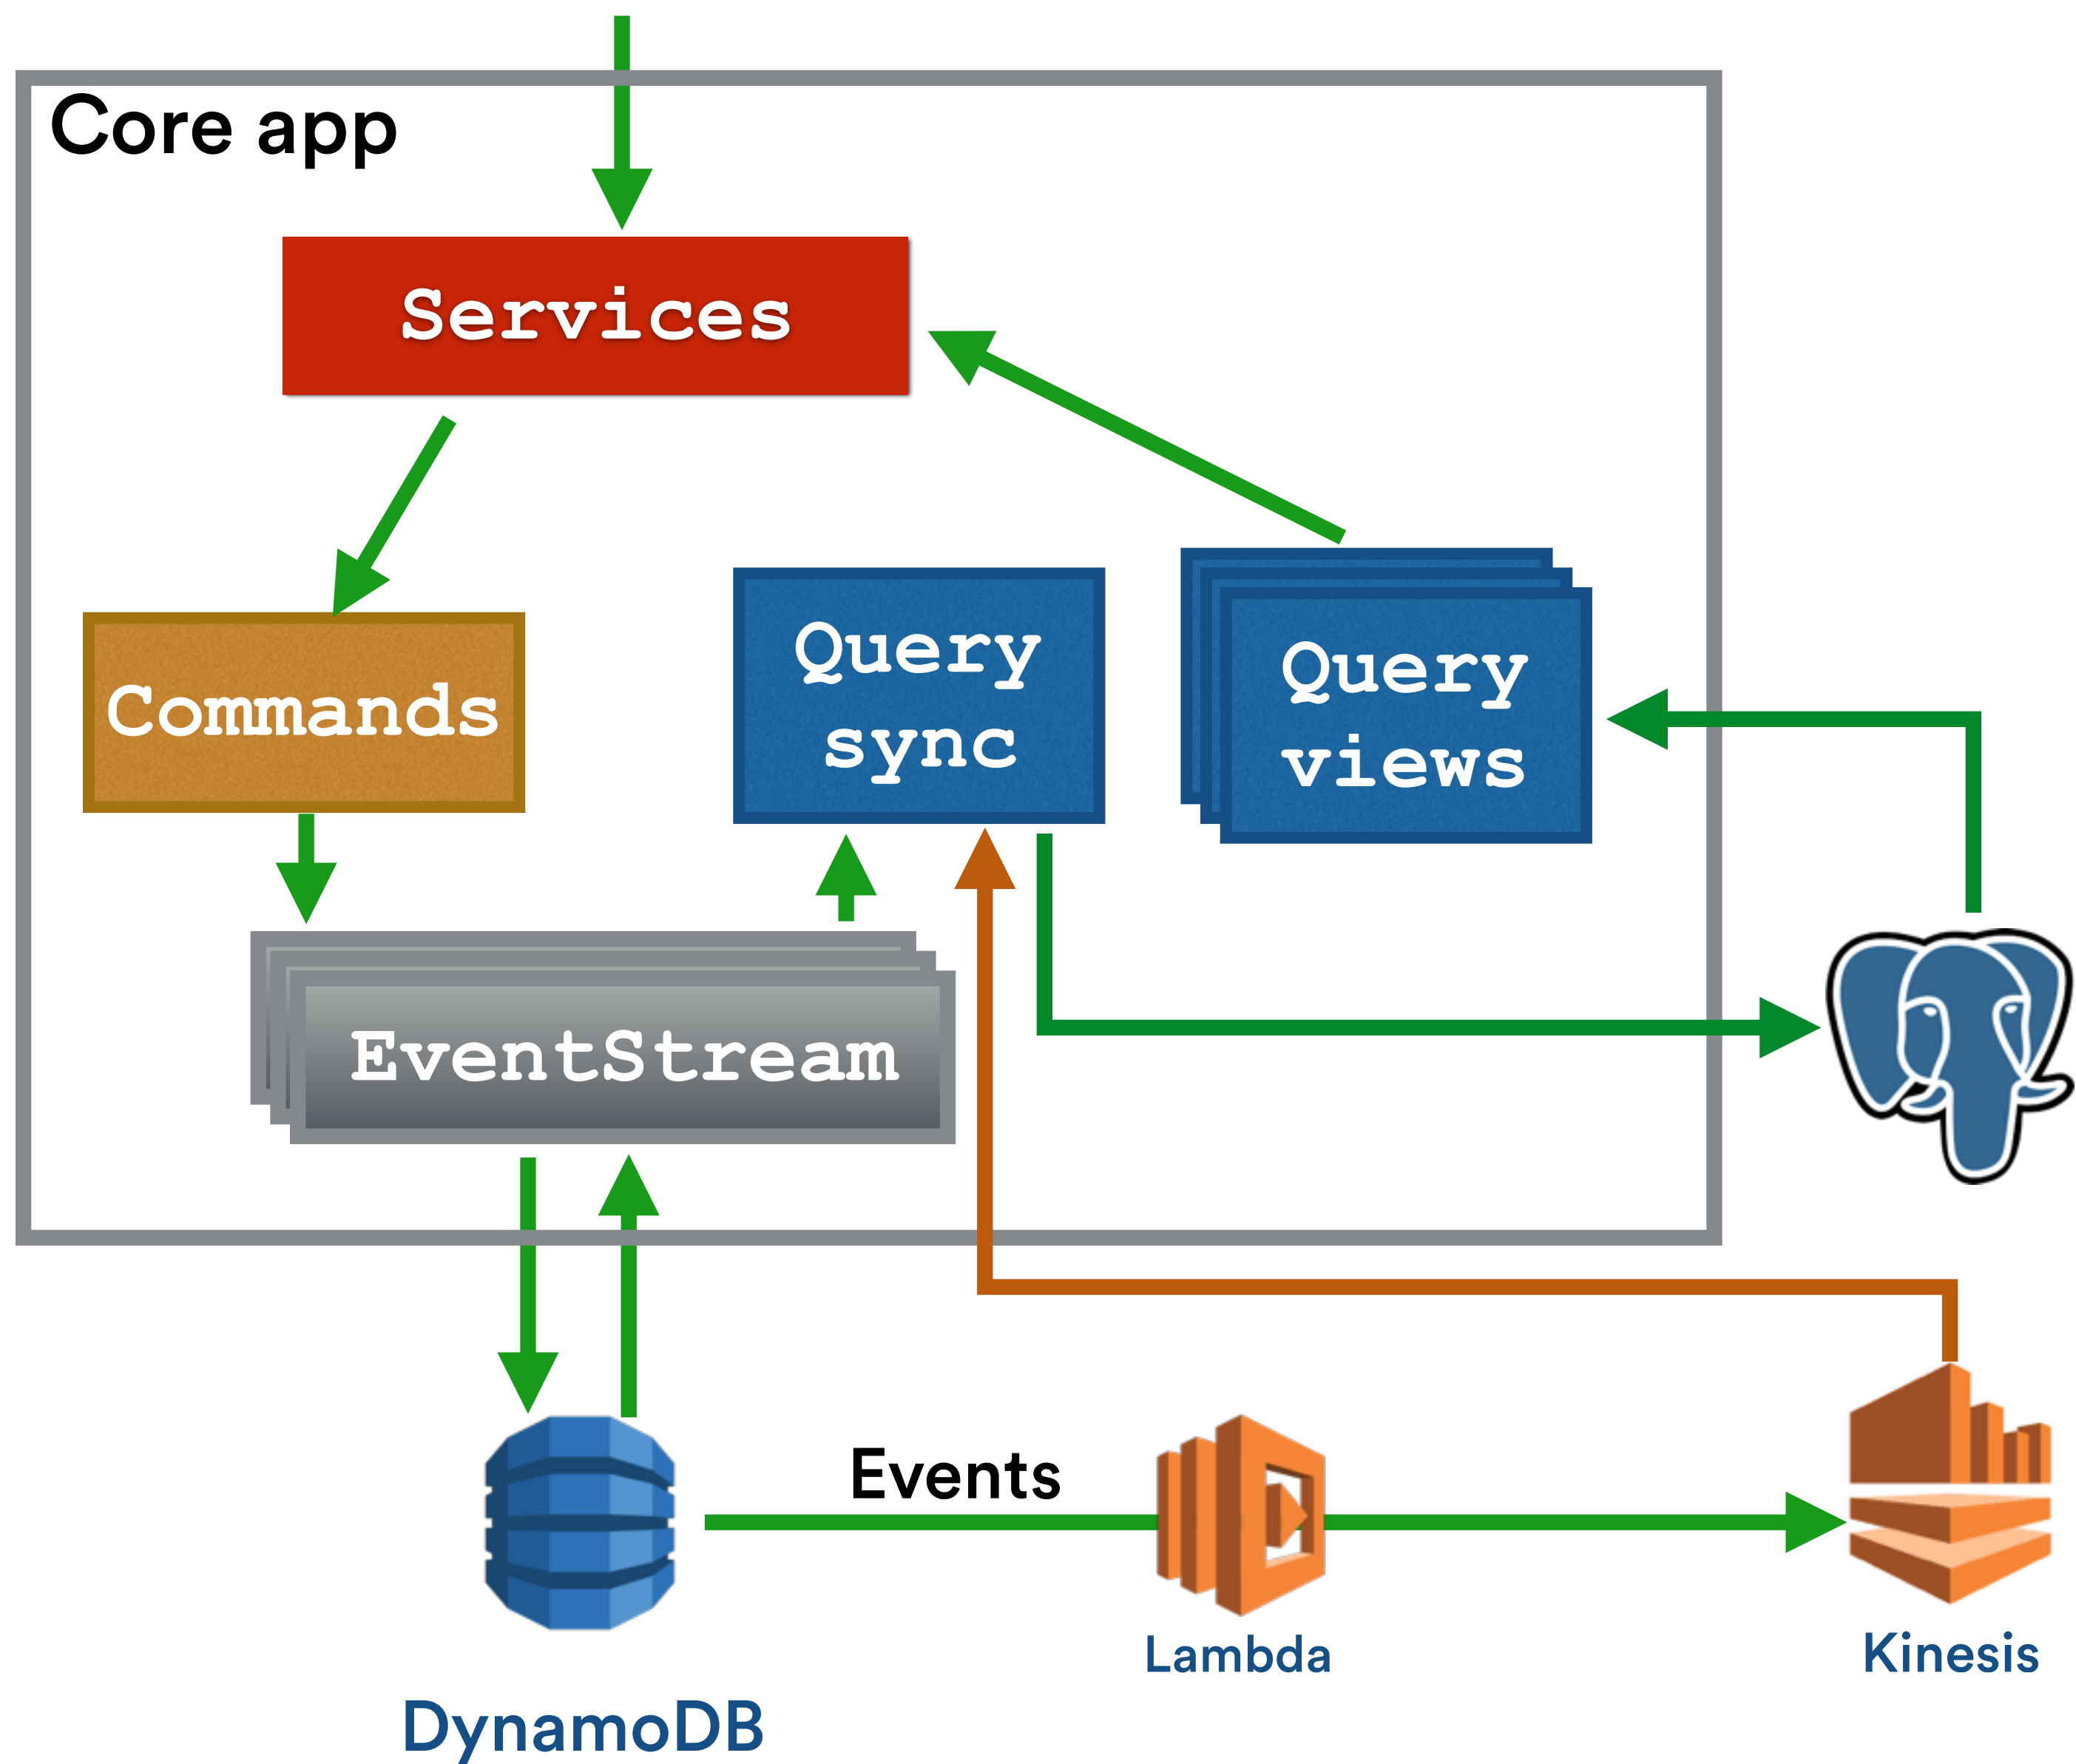
REST calls

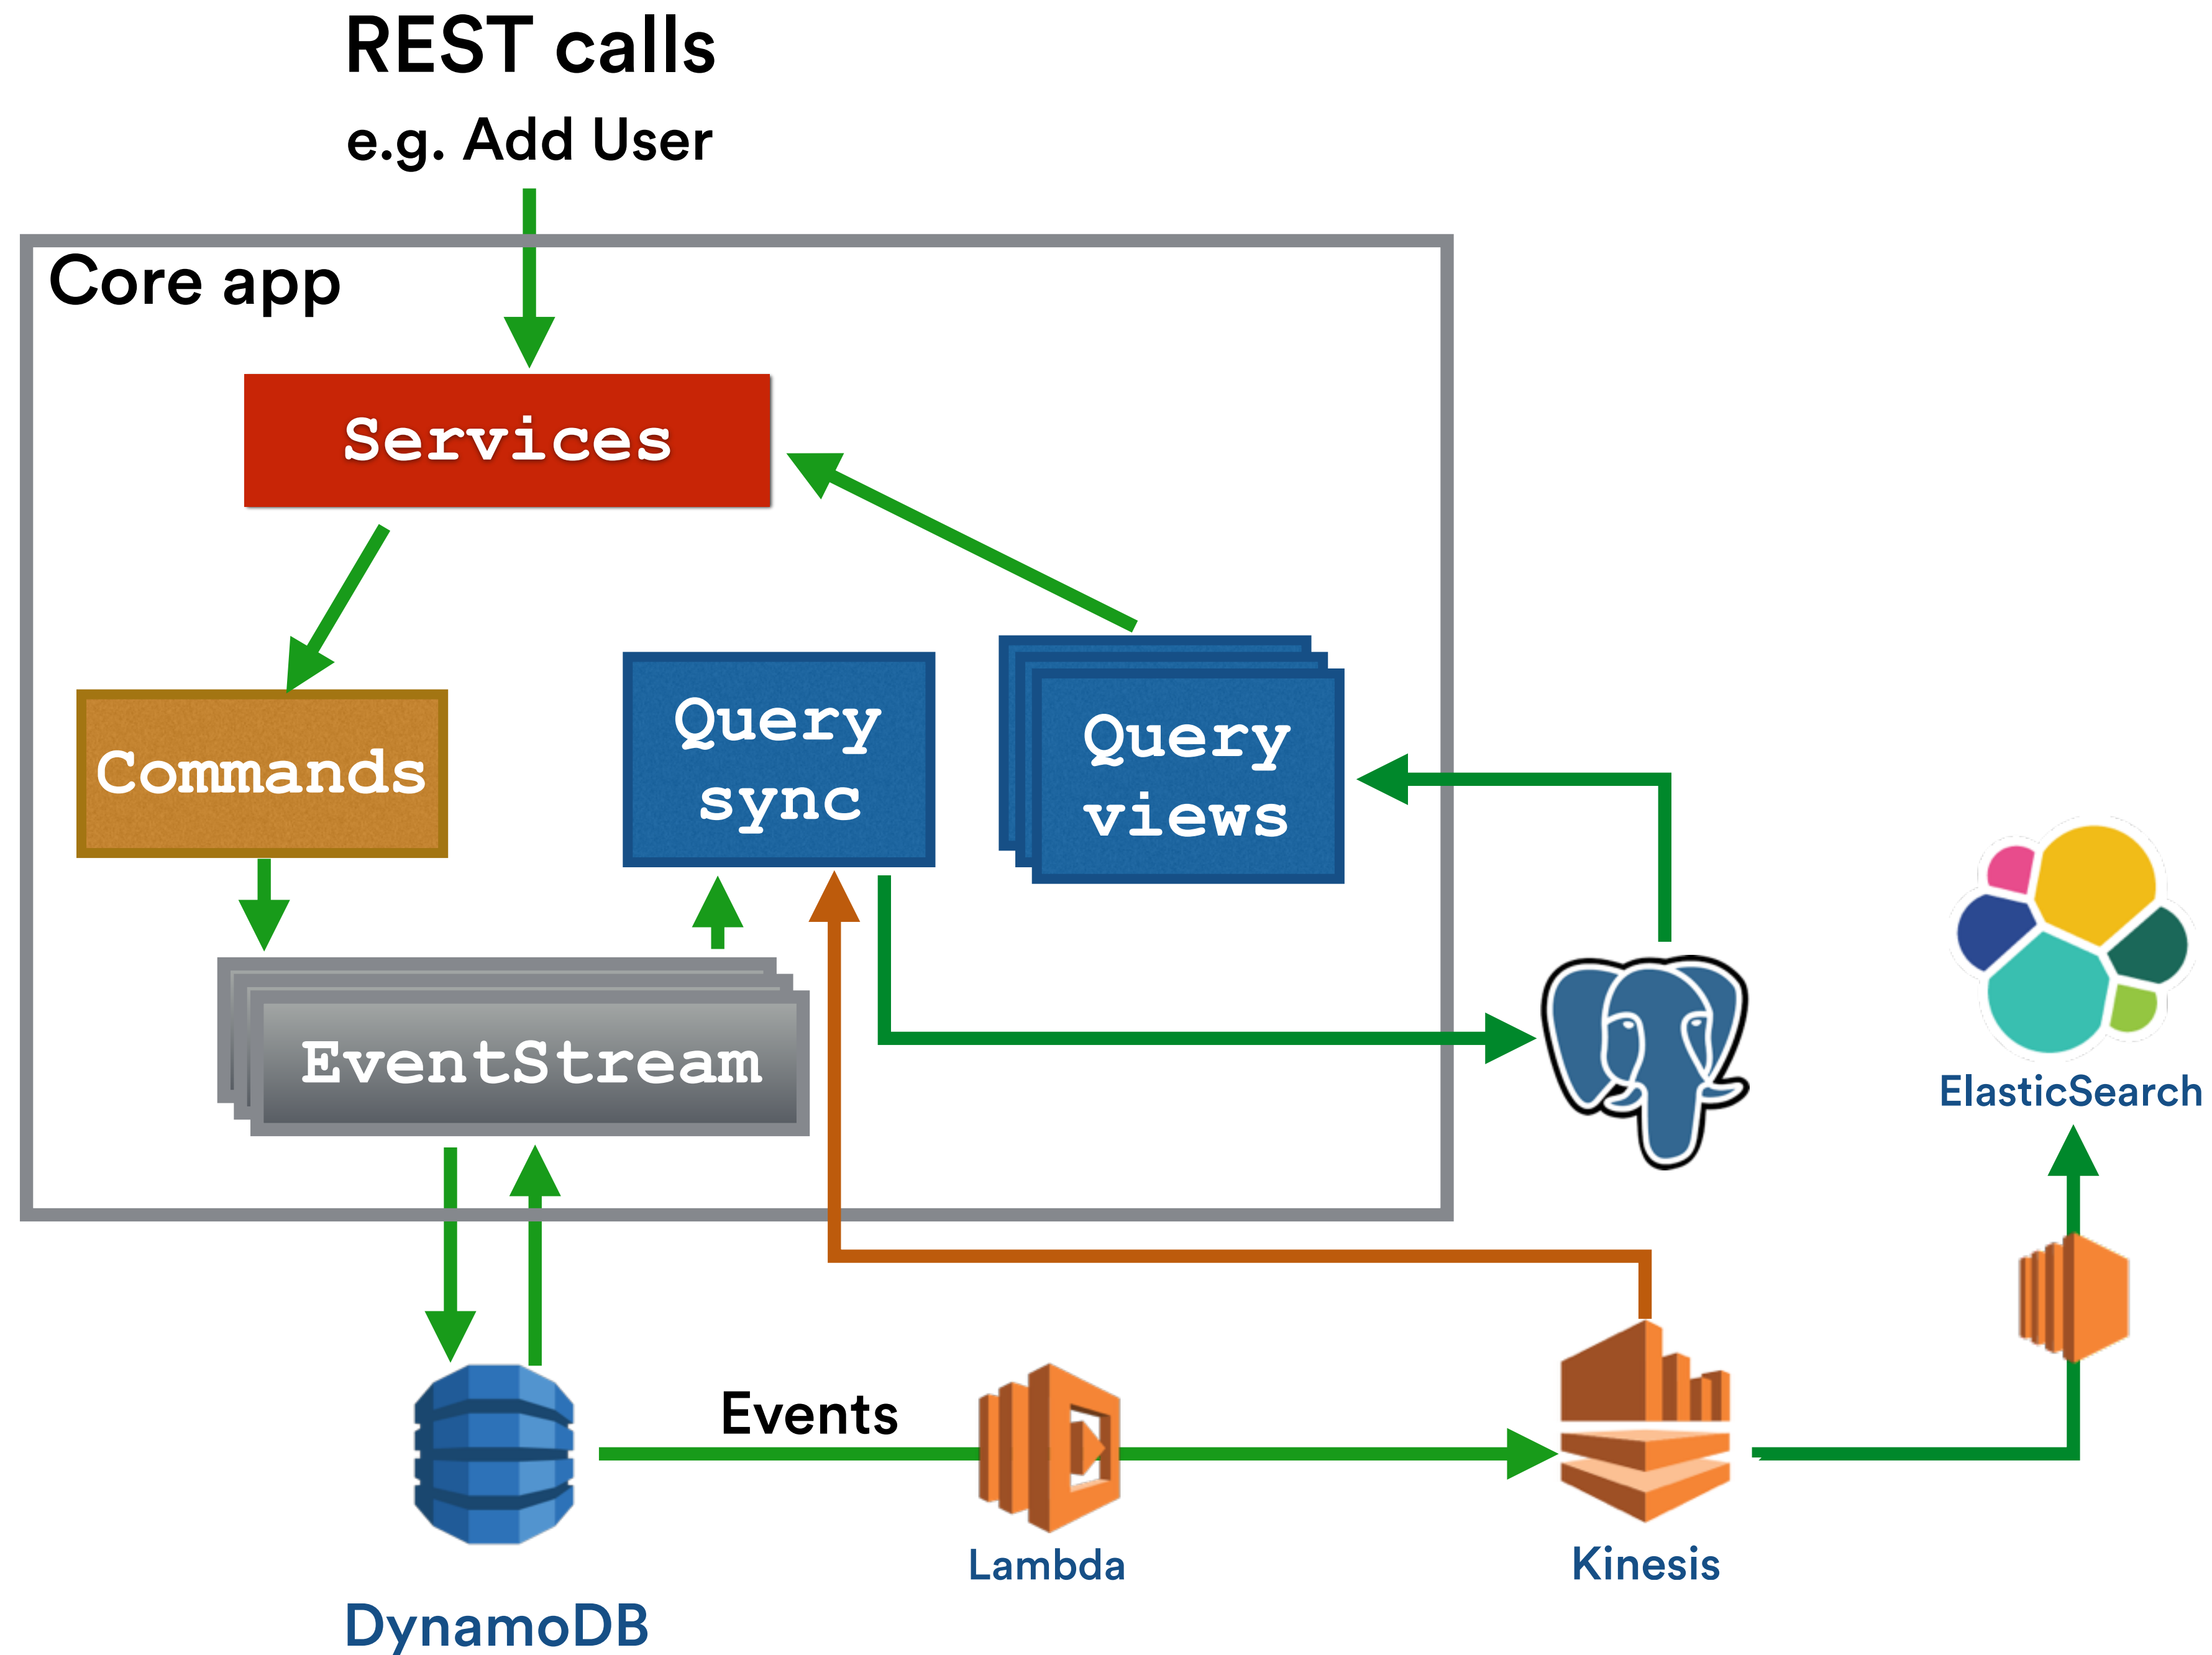
e.g. Add User



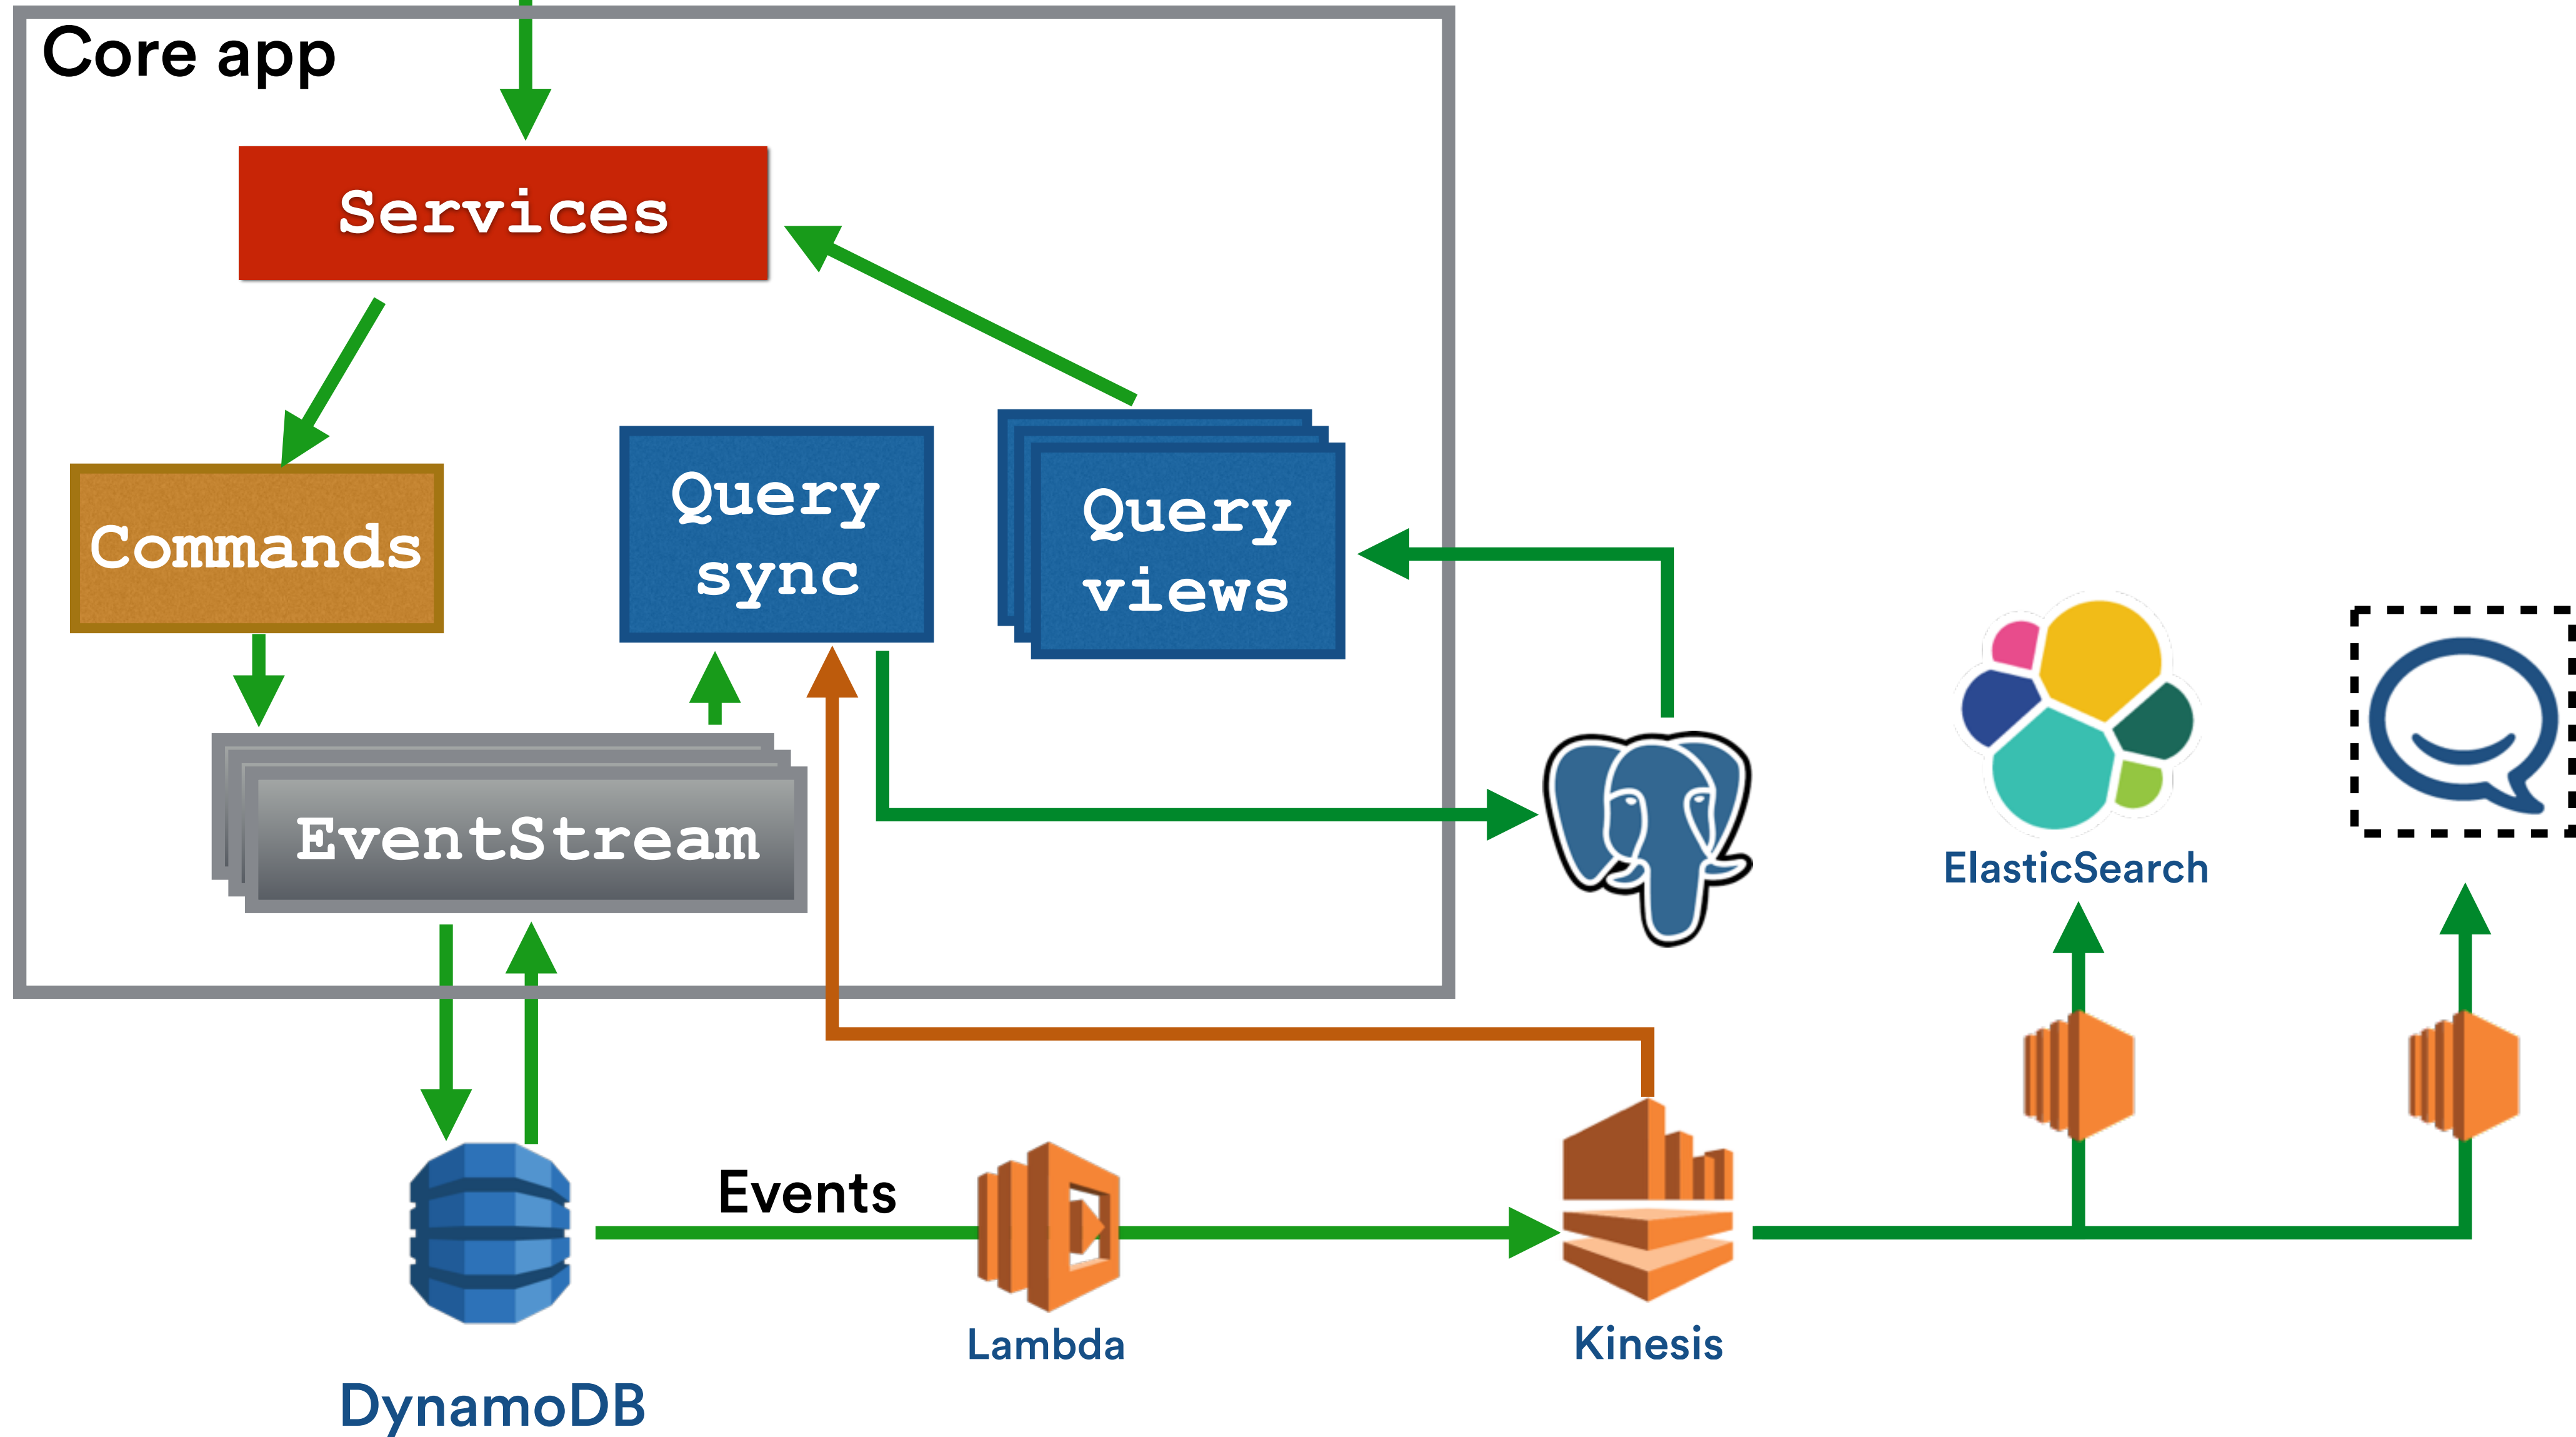
REST calls

e.g. Add User



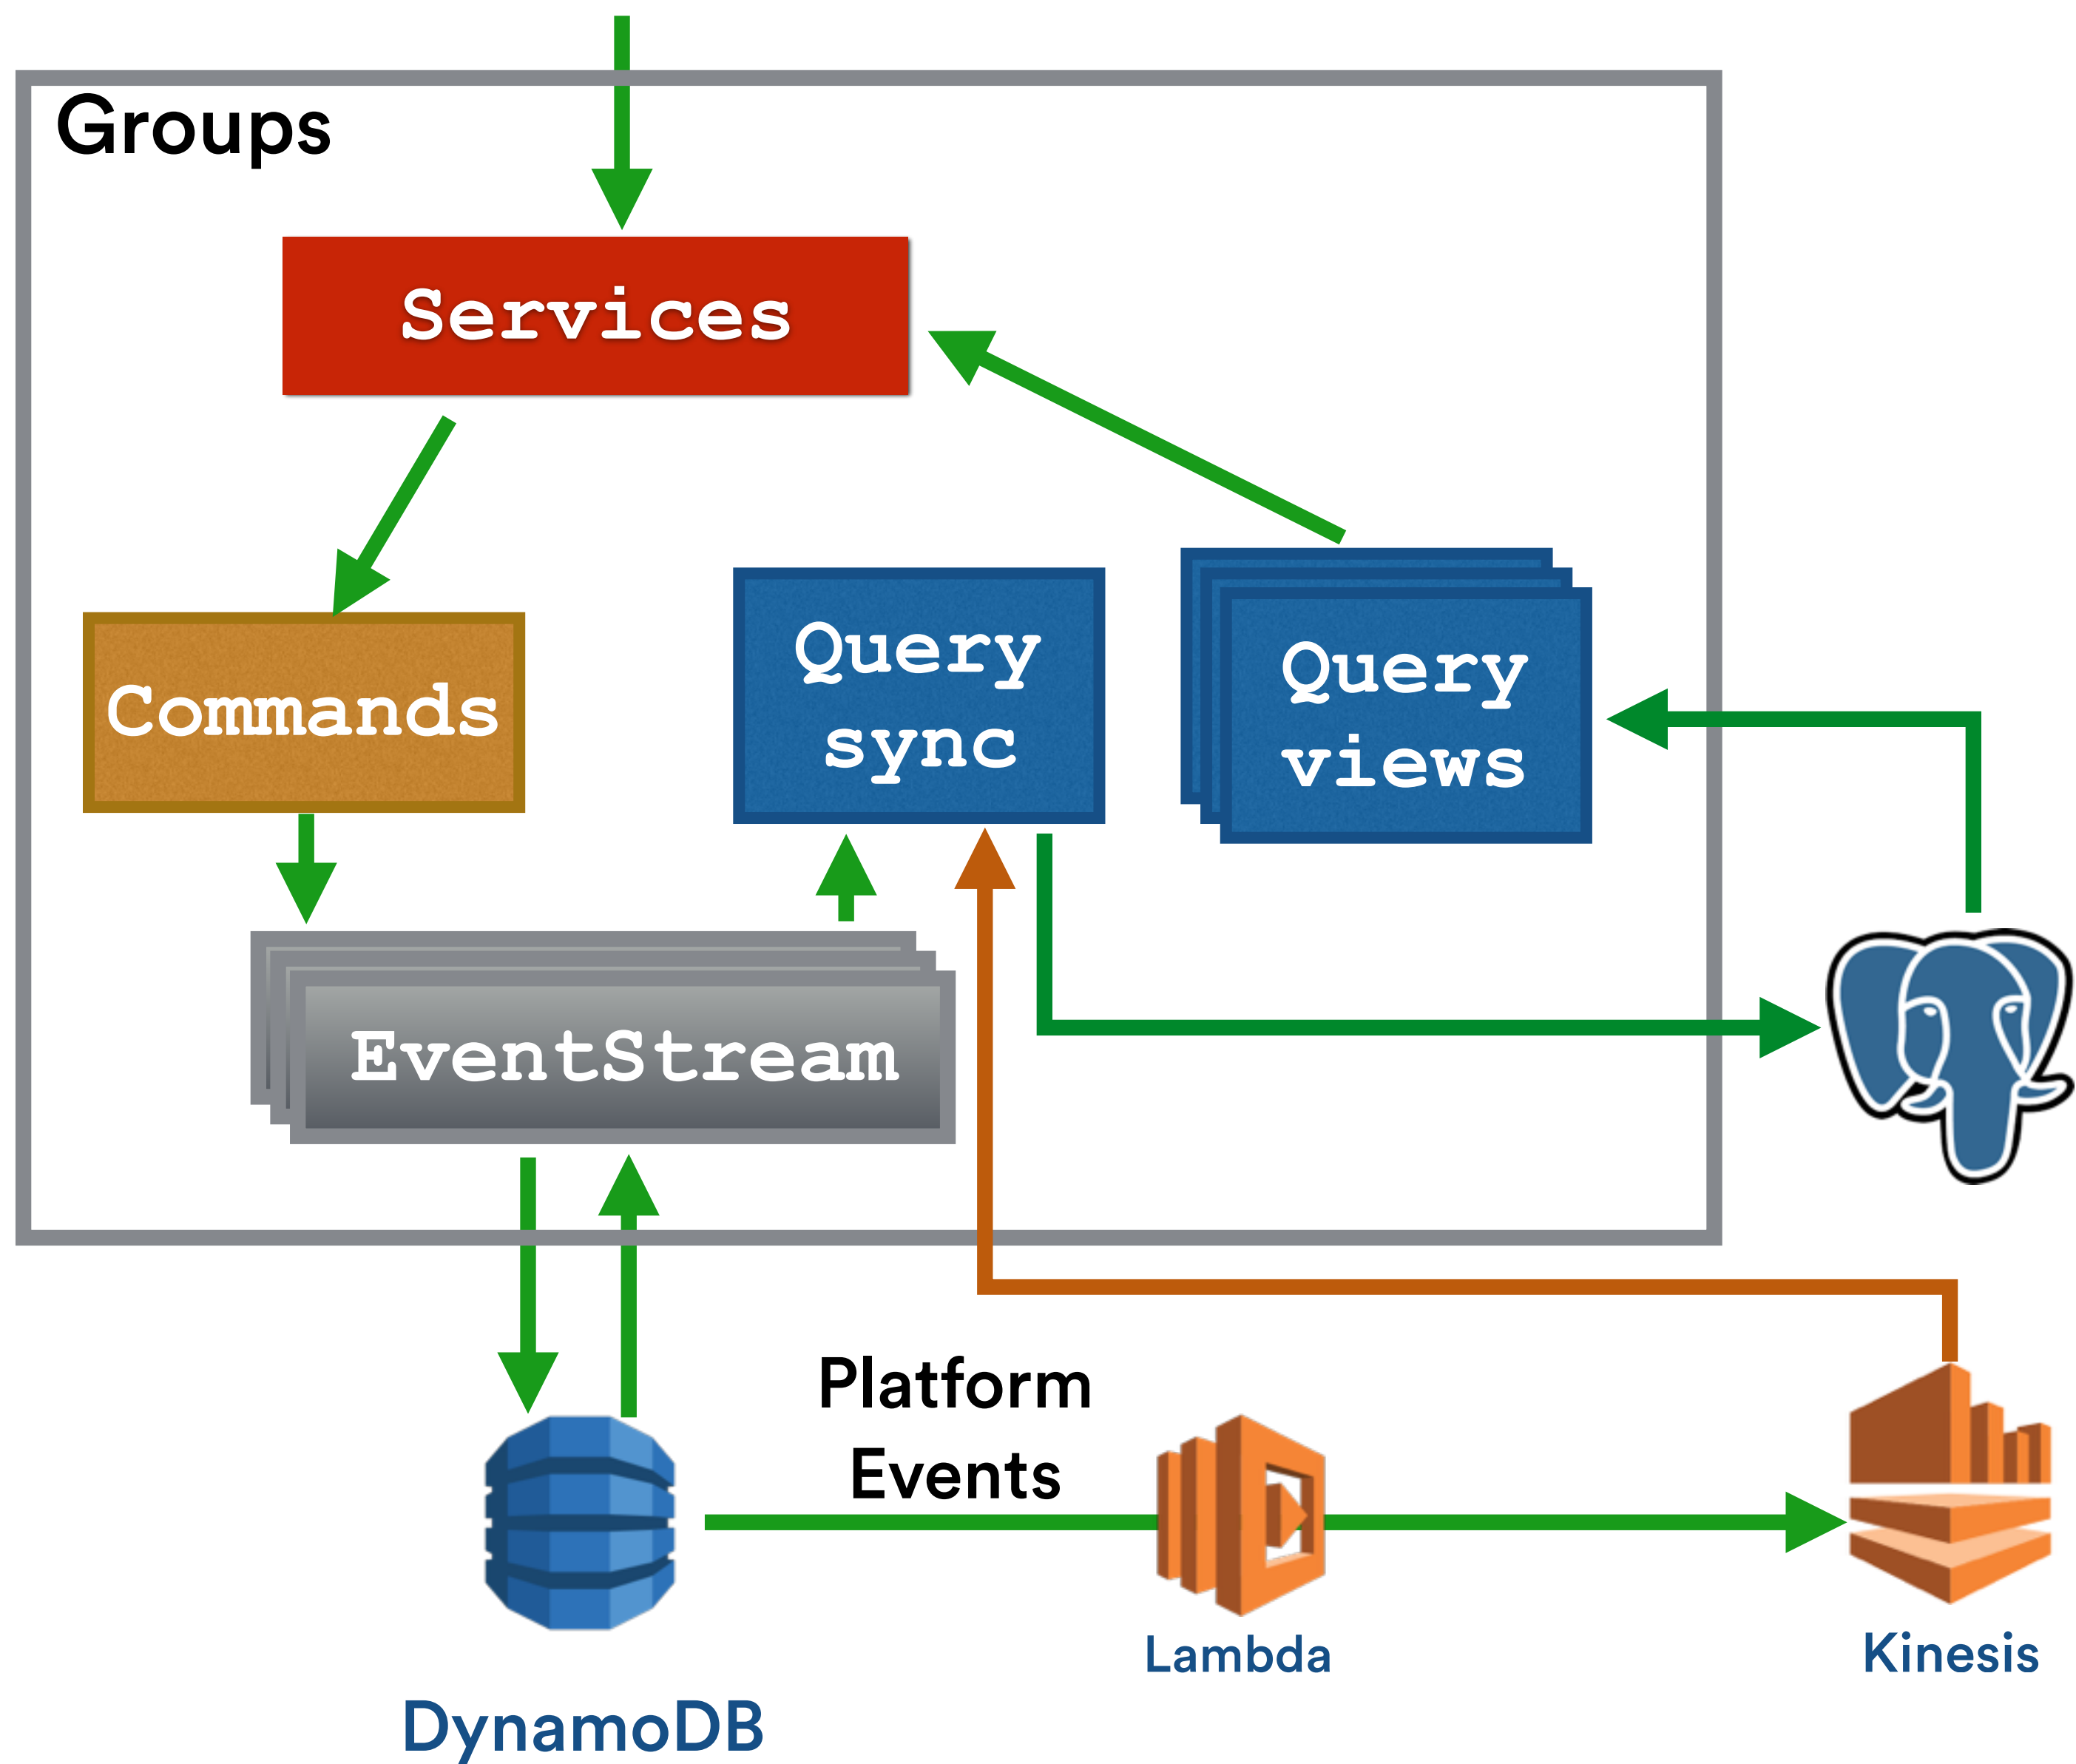


REST calls
e.g. Add User



REST calls

e.g. Add User



REST calls

e.g. Add User

Groups

Services

Commands

Query
sync

Query
views

EventStream

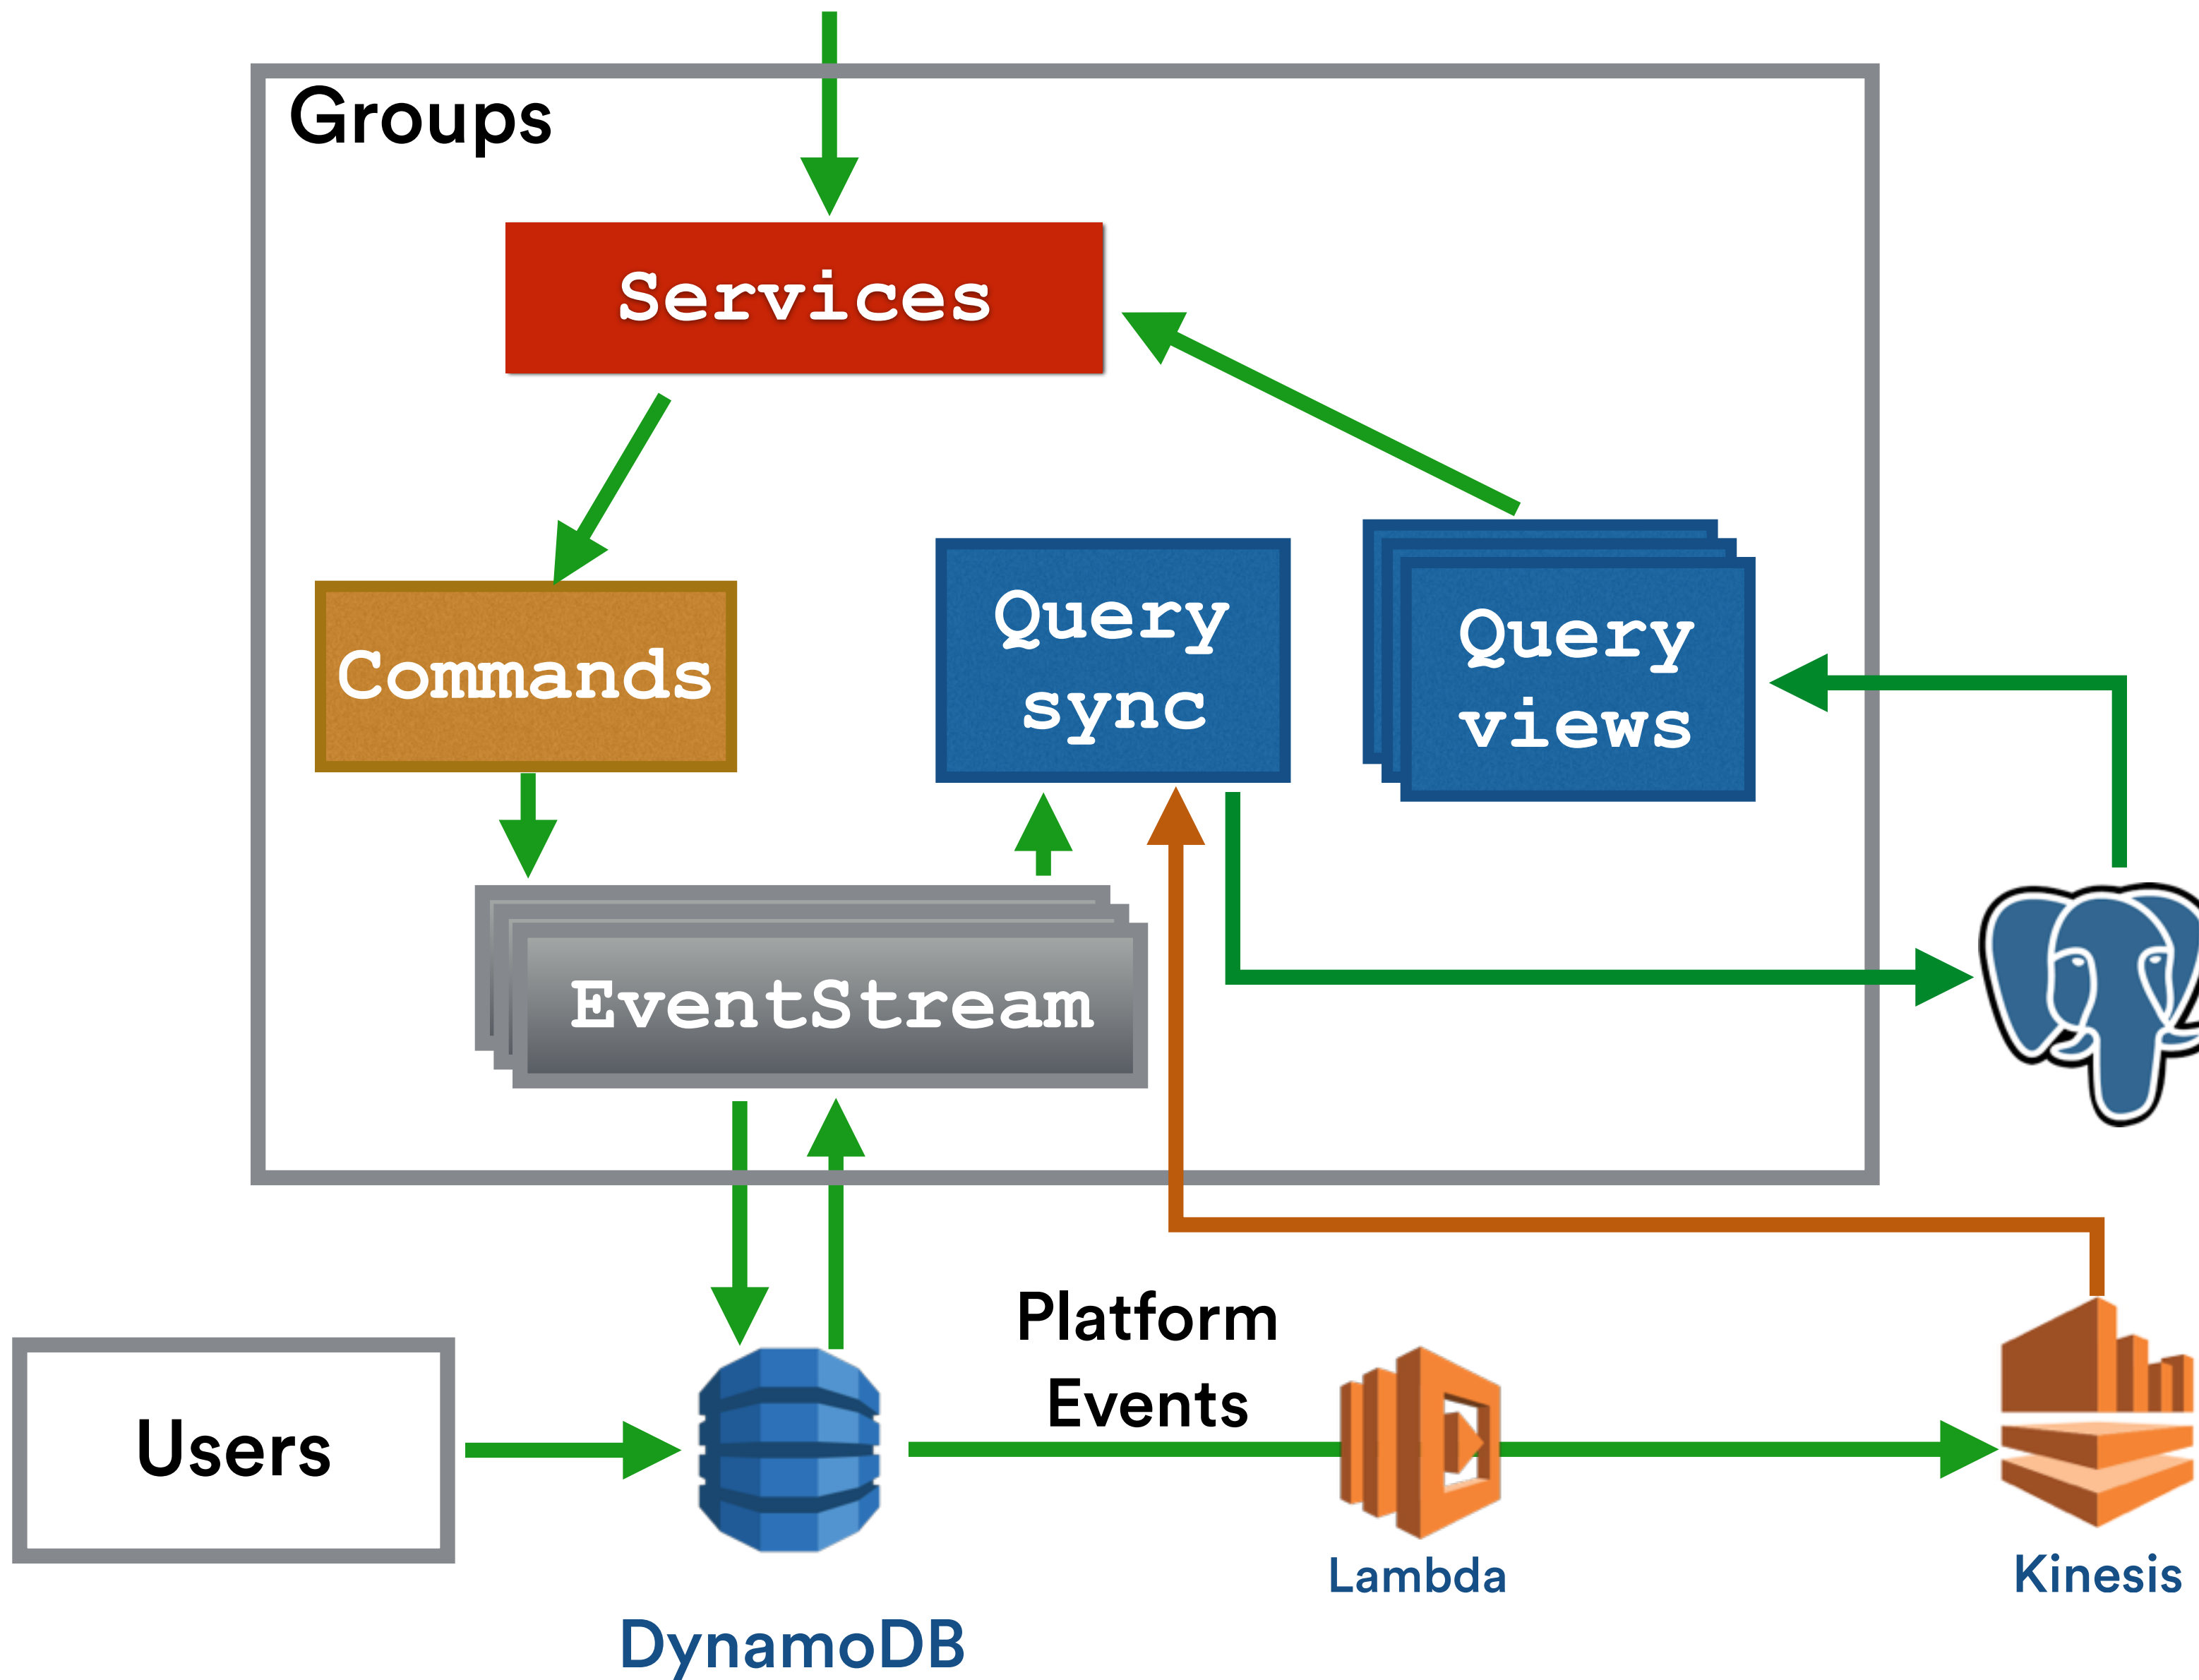
Users

Platform
Events

DynamoDB

Lambda

Kinesis



REST calls

e.g. Add User

Groups

Services

Commands

Query
sync

Query
views

EventStream

Users

DynamoDB

Platform
Events

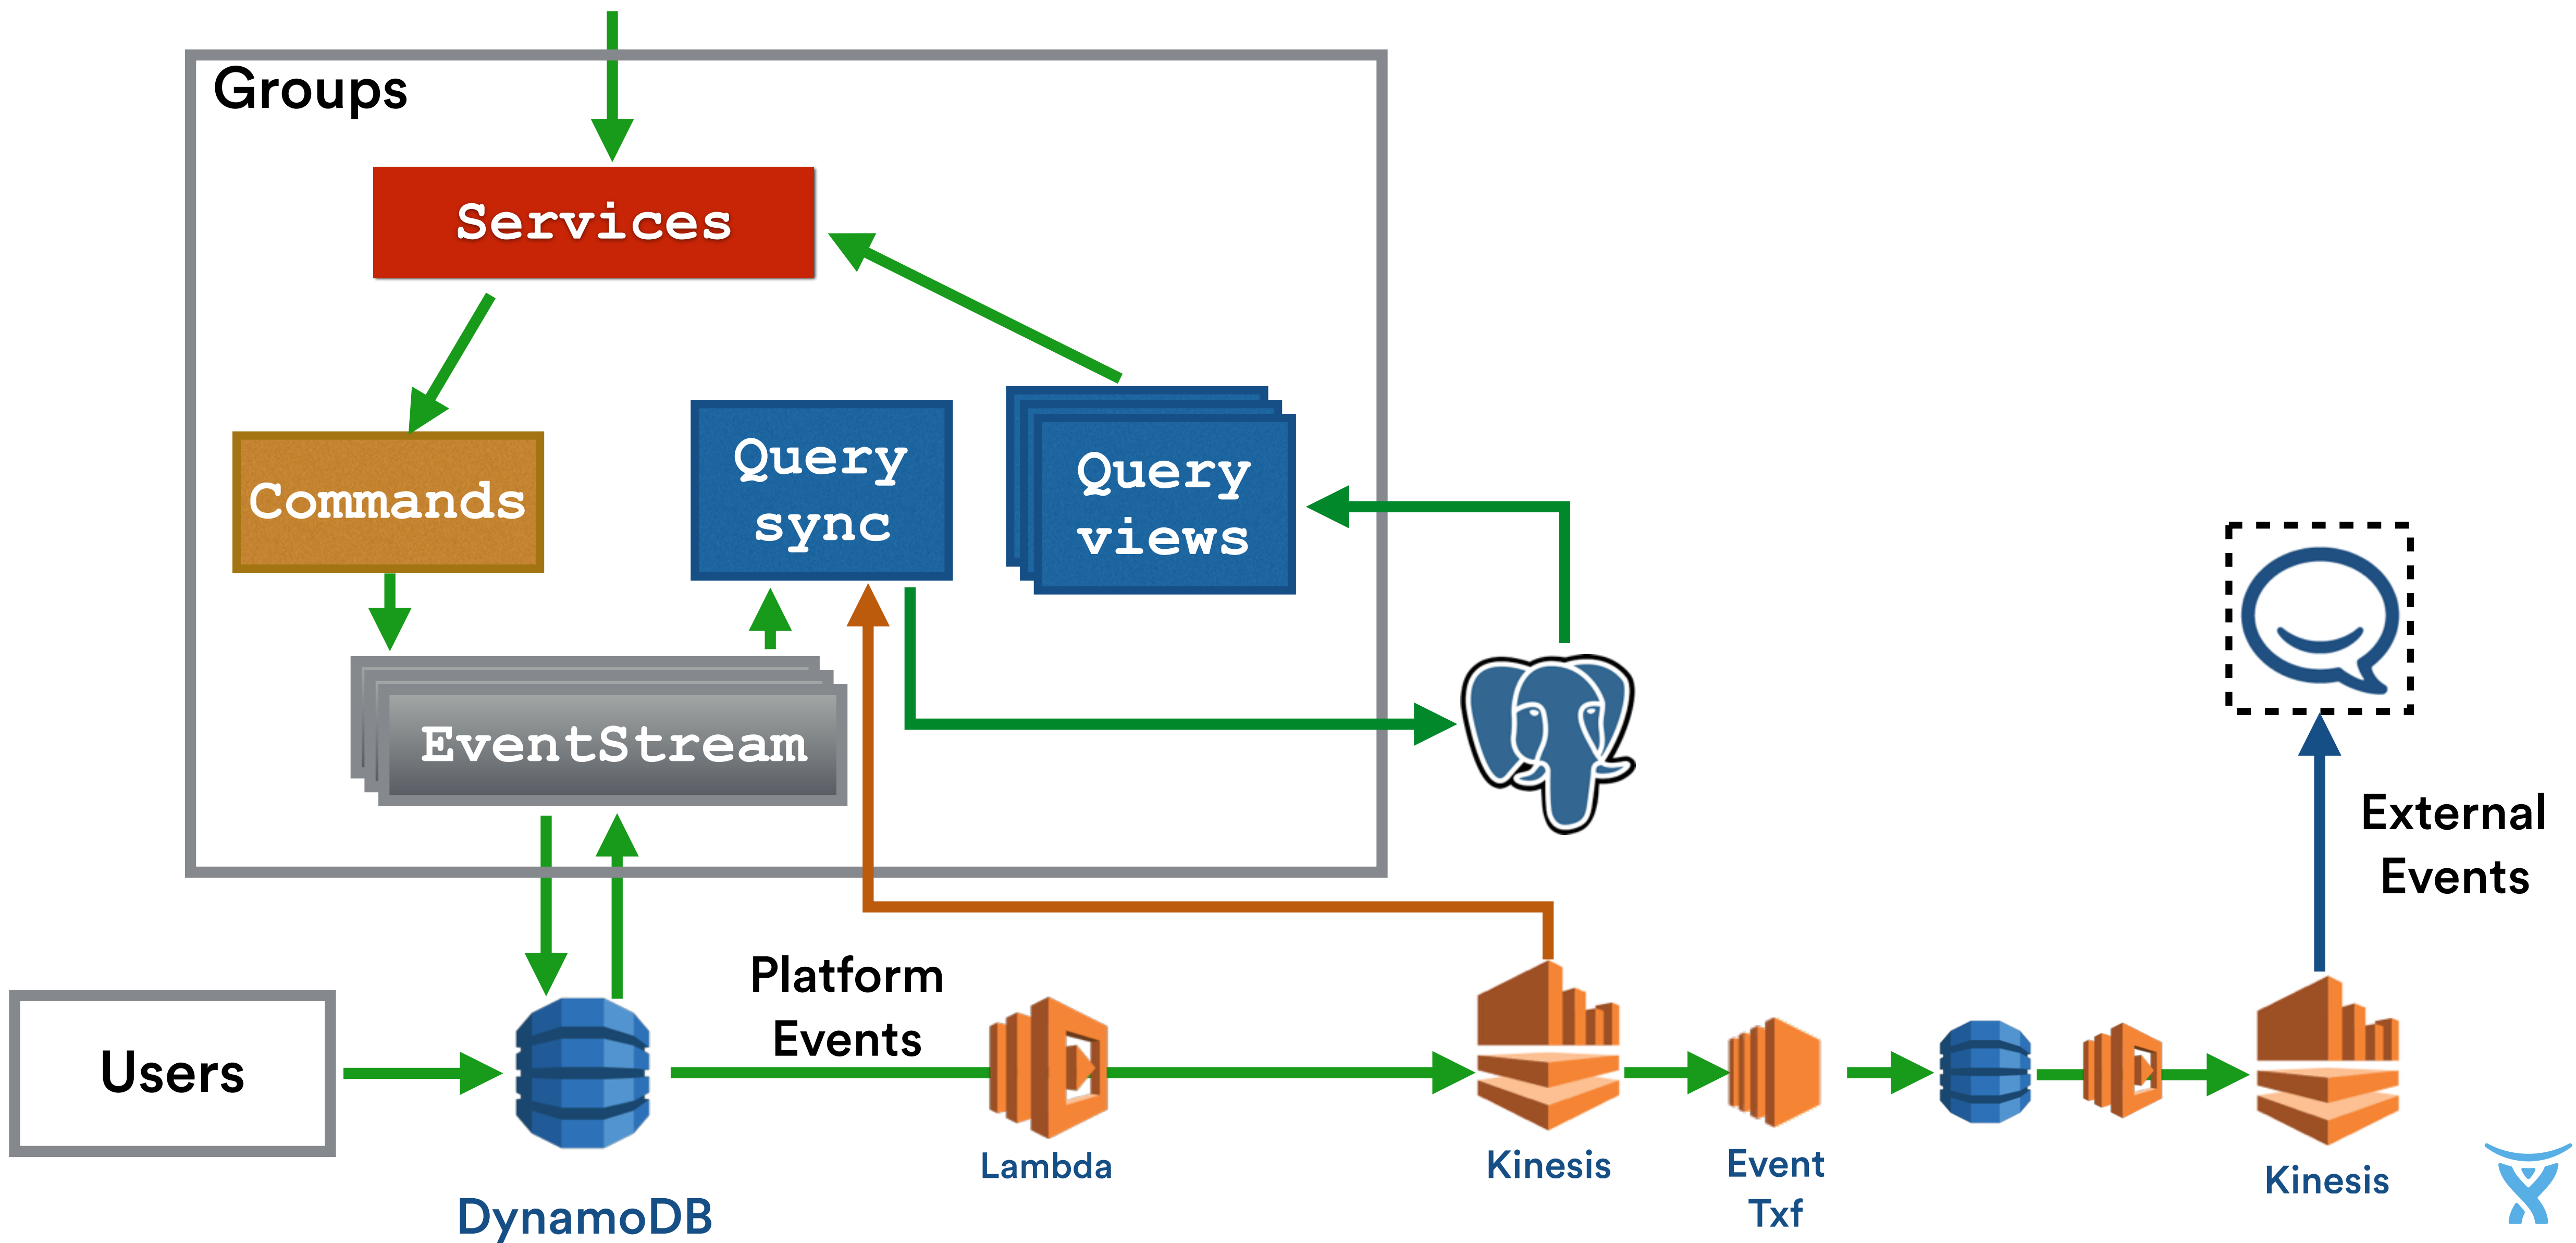
Lambda

Kinesis

Event
Txf

Kinesis

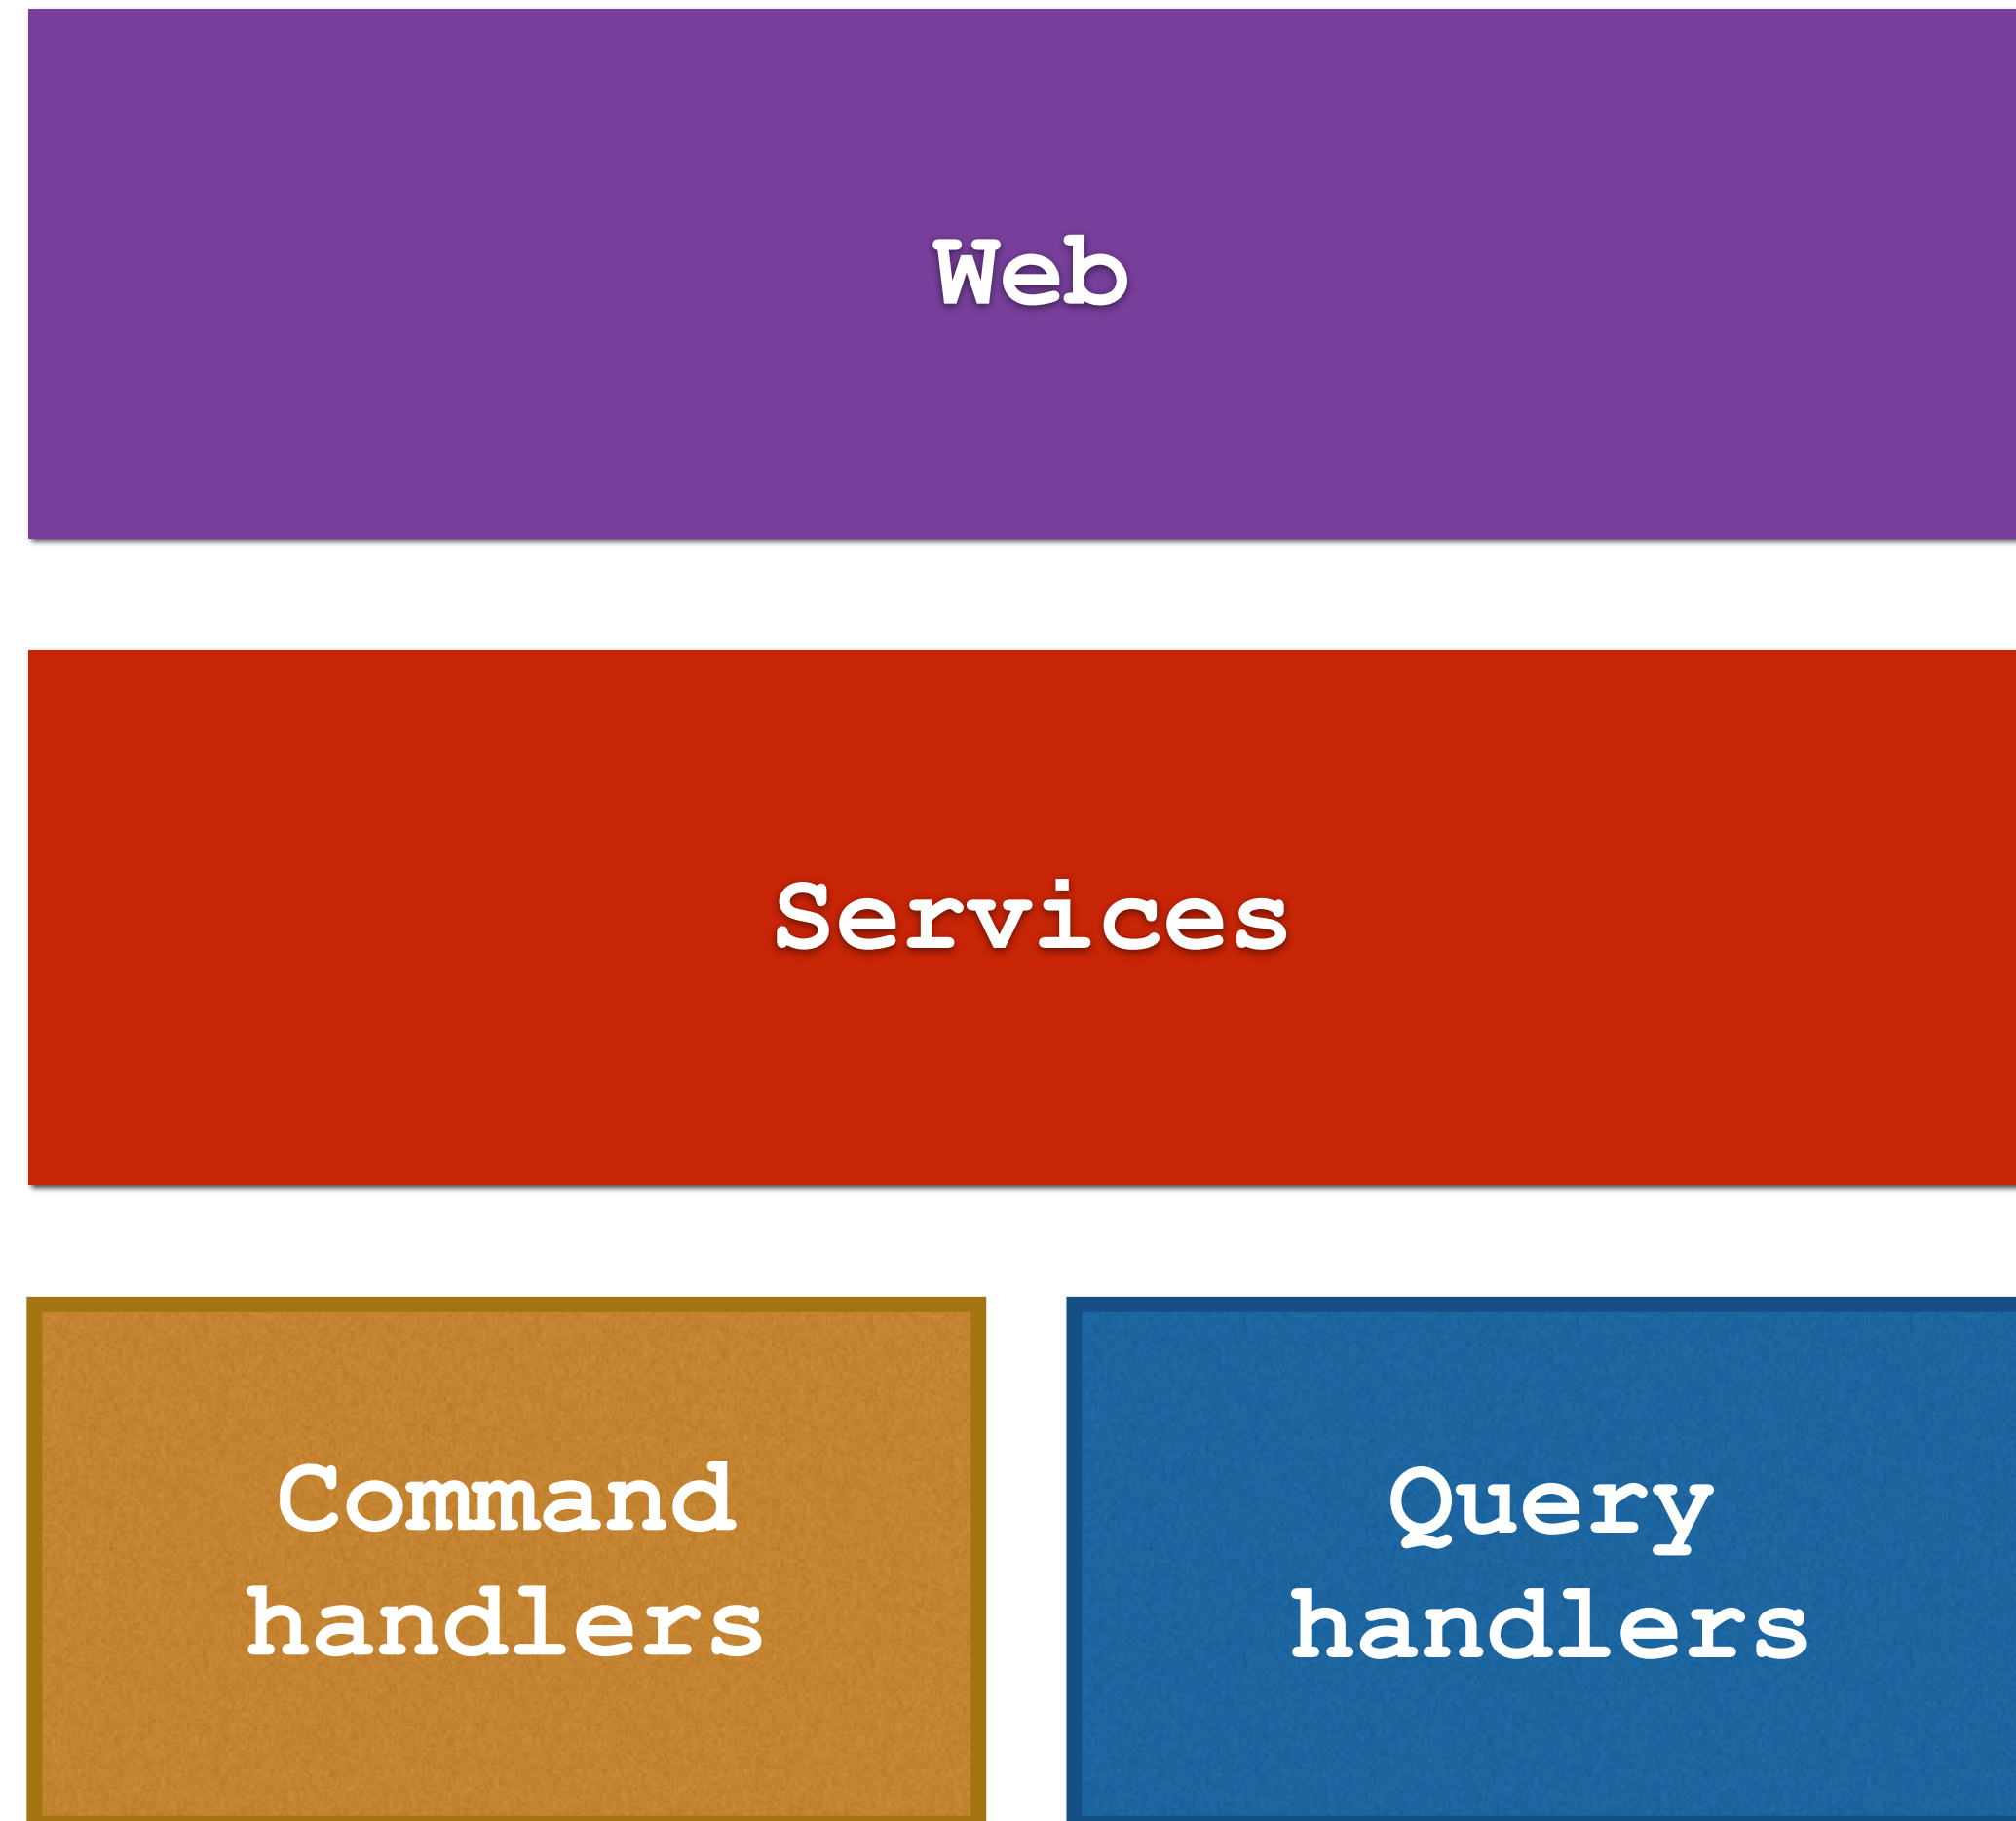
External
Events



Delving into code



Core app stack



Core app stack



`command.Op ~>`
`Result`

`query.Op ~>`
`Result`

```
case class Save[A](...)
  extends command.Op[A]
case class UserByName[A](...)
  extends query.Op[A]
```



Core app stack



command.Op ~>
Result

query.Op ~>
Result

```
def addUser: AppOp[A]
```

```
type AppOp[A] = Free[App, A]
```

```
type App[A] =
```

```
  Coproduct[command.Op, query.Op, A]
```

```
case class Save[A](...)
```

```
  extends command.Op[A]
```

```
case class UserByName[A](...)
```

```
  extends query.Op[A]
```



Core app stack

Web

```
post { entity(as[NewUser]) { u =>
  runApi(UserService.addUser(u))
}
```

```
val runApi: AppOp[A] => Route
```

Services

```
def addUser: AppOp[A]
```

```
type AppOp[A] = Free[App, A]
```

```
type App[A] =
  Coproduct[command.Op, query.Op, A]
```

Command
handlers

Query
handlers

```
case class Save[A](...)
  extends command.Op[A]
case class UserByName[A](...)
  extends query.Op[A]
```

command.Op ~>
Result

query.Op ~>
Result



Commands save events to streams

```
object UserCommandHandler {  
  def apply(streams: Streams): command.Op ~> Result =  
    new (command.Op ~> Result) {  
      def apply[X](a: command.Op[X]): Result[X] =  
        a match {  
          case Save(u) =>  
  
          case ...  
        }  
    }  
}
```



Commands save events to streams

```
object UserCommandHandler {  
  def apply(streams: Streams): command.Op ~> Result =  
    new (command.Op ~> Result) {  
      def apply[X](a: command.Op[X]): Result[X] =  
        a match {  
          case Save(u) =>  
            streams.allUsers.save(SingleStreamId, SetName(u.id, u.name))  
              .then(streams.userProfile.save(u.id, SetProfile(u))  
                .thenReturn(u.id)  
          case ...  
        }  
    }  
}
```



Query synchroniser is a scalaz-stream Channel...

[illegible]

Query synchroniser is a scalaz-stream Channel...

```
class QuerySynchroniser(db: DB) {  
  
  def synchronise(s: Process[Task, StreamEvent]) =  
    (s through writeEvent).run  
  
  private val writeEvent: Channel[Task, StreamEvent, Unit] =  
    Process.constant { e: StreamEvent =>  
  
    }  
}
```



Query synchroniser is a scalaz-stream Channel...

```
class QuerySynchroniser(db: DB) {  
  
  def synchronise(s: Process[Task, StreamEvent]) =  
    (s through writeEvent).run  
  
  private val writeEvent: Channel[Task, StreamEvent, Unit] =  
    Process.constant { e: StreamEvent =>  
      for {  
        r1 <- StreamState.transitionSeq(db, e.id.key, e.id.seq)  
  
      } yield ()  
    }  
}
```



Query synchroniser is a scalaz-stream Channel...

```
class QuerySynchroniser(db: DB) {  
  
  def synchronise(s: Process[Task, StreamEvent]) =  
    (s through writeEvent).run  
  
  private val writeEvent: Channel[Task, StreamEvent, Unit] =  
    Process.constant { e: StreamEvent =>  
      for {  
        r1 <- StreamState.transitionSeq(db, e.id.key, e.id.seq)  
        _ <- if (r1.success) db.process(e) else db.rollback  
      } yield ()  
    }  
}
```



...that you pass it an event stream

```
class KinesisProcessor(qs: QuerySynchroniser) extends IRecordProcessor {  
    override def processRecords(rs: ProcessRecordInput) = {  
  
    }  
}
```



...that you pass it an event stream

```
class KinesisProcessor(qs: QuerySynchroniser) extends IRecordProcessor {  
  
  override def processRecords(rs: ProcessRecordInput) = {  
    // Either emit events from Kinesis, or read events from Dynamo  
    val stream = Process.emitAll { rs.getRecords. ... }  
  
  }  
  
}
```



...that you pass it an event stream

```
class KinesisProcessor(qs: QuerySynchroniser) extends IRecordProcessor {  
  
  override def processRecords(rs: ProcessRecordInput) = {  
    // Either emit events from Kinesis, or read events from Dynamo  
    val stream = Process.emitAll { rs.getRecords. ... }  
  
    qs.synchronise(stream).attemptRunFor(...)   
  }  
  
}
```





Events as an API

Insert / Update Delta vs 'Set' events



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))

SetUserName(id, name)

SetUserEmail(id, email1)



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))

SetUserName(id, name)

SetUserEmail(id, email1)

SetUserEmail(id, email2)



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))

SetUserName(id, name)

SetUserEmail(id, email1)

SetUserEmail(id, email2)

Fits nicely with CRUD + PATCH



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))

SetUserName(id, name)

SetUserEmail(id, email1)

SetUserEmail(id, email2)

Fits nicely with CRUD + PATCH

Assume insert before update



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))

SetUserName(id, name)

SetUserEmail(id, email1)

SetUserEmail(id, email2)

Fits nicely with CRUD + PATCH

Assume insert before update

Encourages idempotent processing



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))

SetUserName(id, name)

SetUserEmail(id, email1)

SetUserEmail(id, email2)

Fits nicely with CRUD + PATCH

Assume insert before update

Encourages idempotent processing

Single code path for query sync



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))

SetUserName(id, name)

SetUserEmail(id, email1)

SetUserEmail(id, email2)

Fits nicely with CRUD + PATCH

Assume insert before update

Encourages idempotent processing

Single code path for query sync

Minimally sized events to avoid conflict



Insert / Update Delta

vs

‘Set’ events

AddUser(id, name, email1)

UpdateUser(id, email = *Some*(email))

Fits nicely with CRUD + PATCH

Assume insert before update

SetUserName(id, name)

SetUserEmail(id, email1)

SetUserEmail(id, email2)

Encourages idempotent processing

Single code path for query sync

Minimally sized events to avoid conflict



Single stream

Multiple streams



Single stream

**Transactions and
consistent data
resolution**

Multiple streams



Single stream

**Transactions and
consistent data
resolution**

Multiple streams

**Sharding for
throughput**



Rules for splitting streams

1. Place independent events on different streams



Rules for splitting streams

1. Place independent events on different streams
2. Split streams by event type and unique Id



Rules for splitting streams

1. Place independent events on different streams
2. Split streams by event type and unique Id
3. Identify the 'transactions' you really need



Rules for splitting streams

1. Place independent events on different streams
2. Split streams by event type and unique Id
3. Identify the 'transactions' you really need
4. Use hierarchical streams to maximise number of streams



Rules for splitting streams

1. Place independent events on different streams
2. Split streams by event type and unique Id
3. Identify the 'transactions' you really need
4. Use hierarchical streams to maximise number of streams
5. Splitting and joining streams later is possible





**Let go of
transactions
and consistency**

Query views get populated eventually



Query views get populated eventually

- No guaranteed order between events on different streams



Query views get populated eventually

- No guaranteed order between events on different streams
- A field should only be updated by a single event stream



Query views get populated eventually

- No guaranteed order between events on different streams
- A field should only be updated by a single event stream
- No foreign key constraints



Query views get populated eventually

- No guaranteed order between events on different streams
- A field should only be updated by a single event stream
- No foreign key constraints
- Unique or data constraints 'enforced' on write



Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag



Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag
- Pass as request header for:



Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag
- Pass as request header for:
 - Condition write ('transaction')



Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag
- Pass as request header for:
 - Condition write ('transaction')
 - Force query view update ('consistency')



Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag
- Pass as request header for:
 - Condition write ('transaction')
 - Force query view update ('consistency')
 - Caching



Using tokens to enforce state

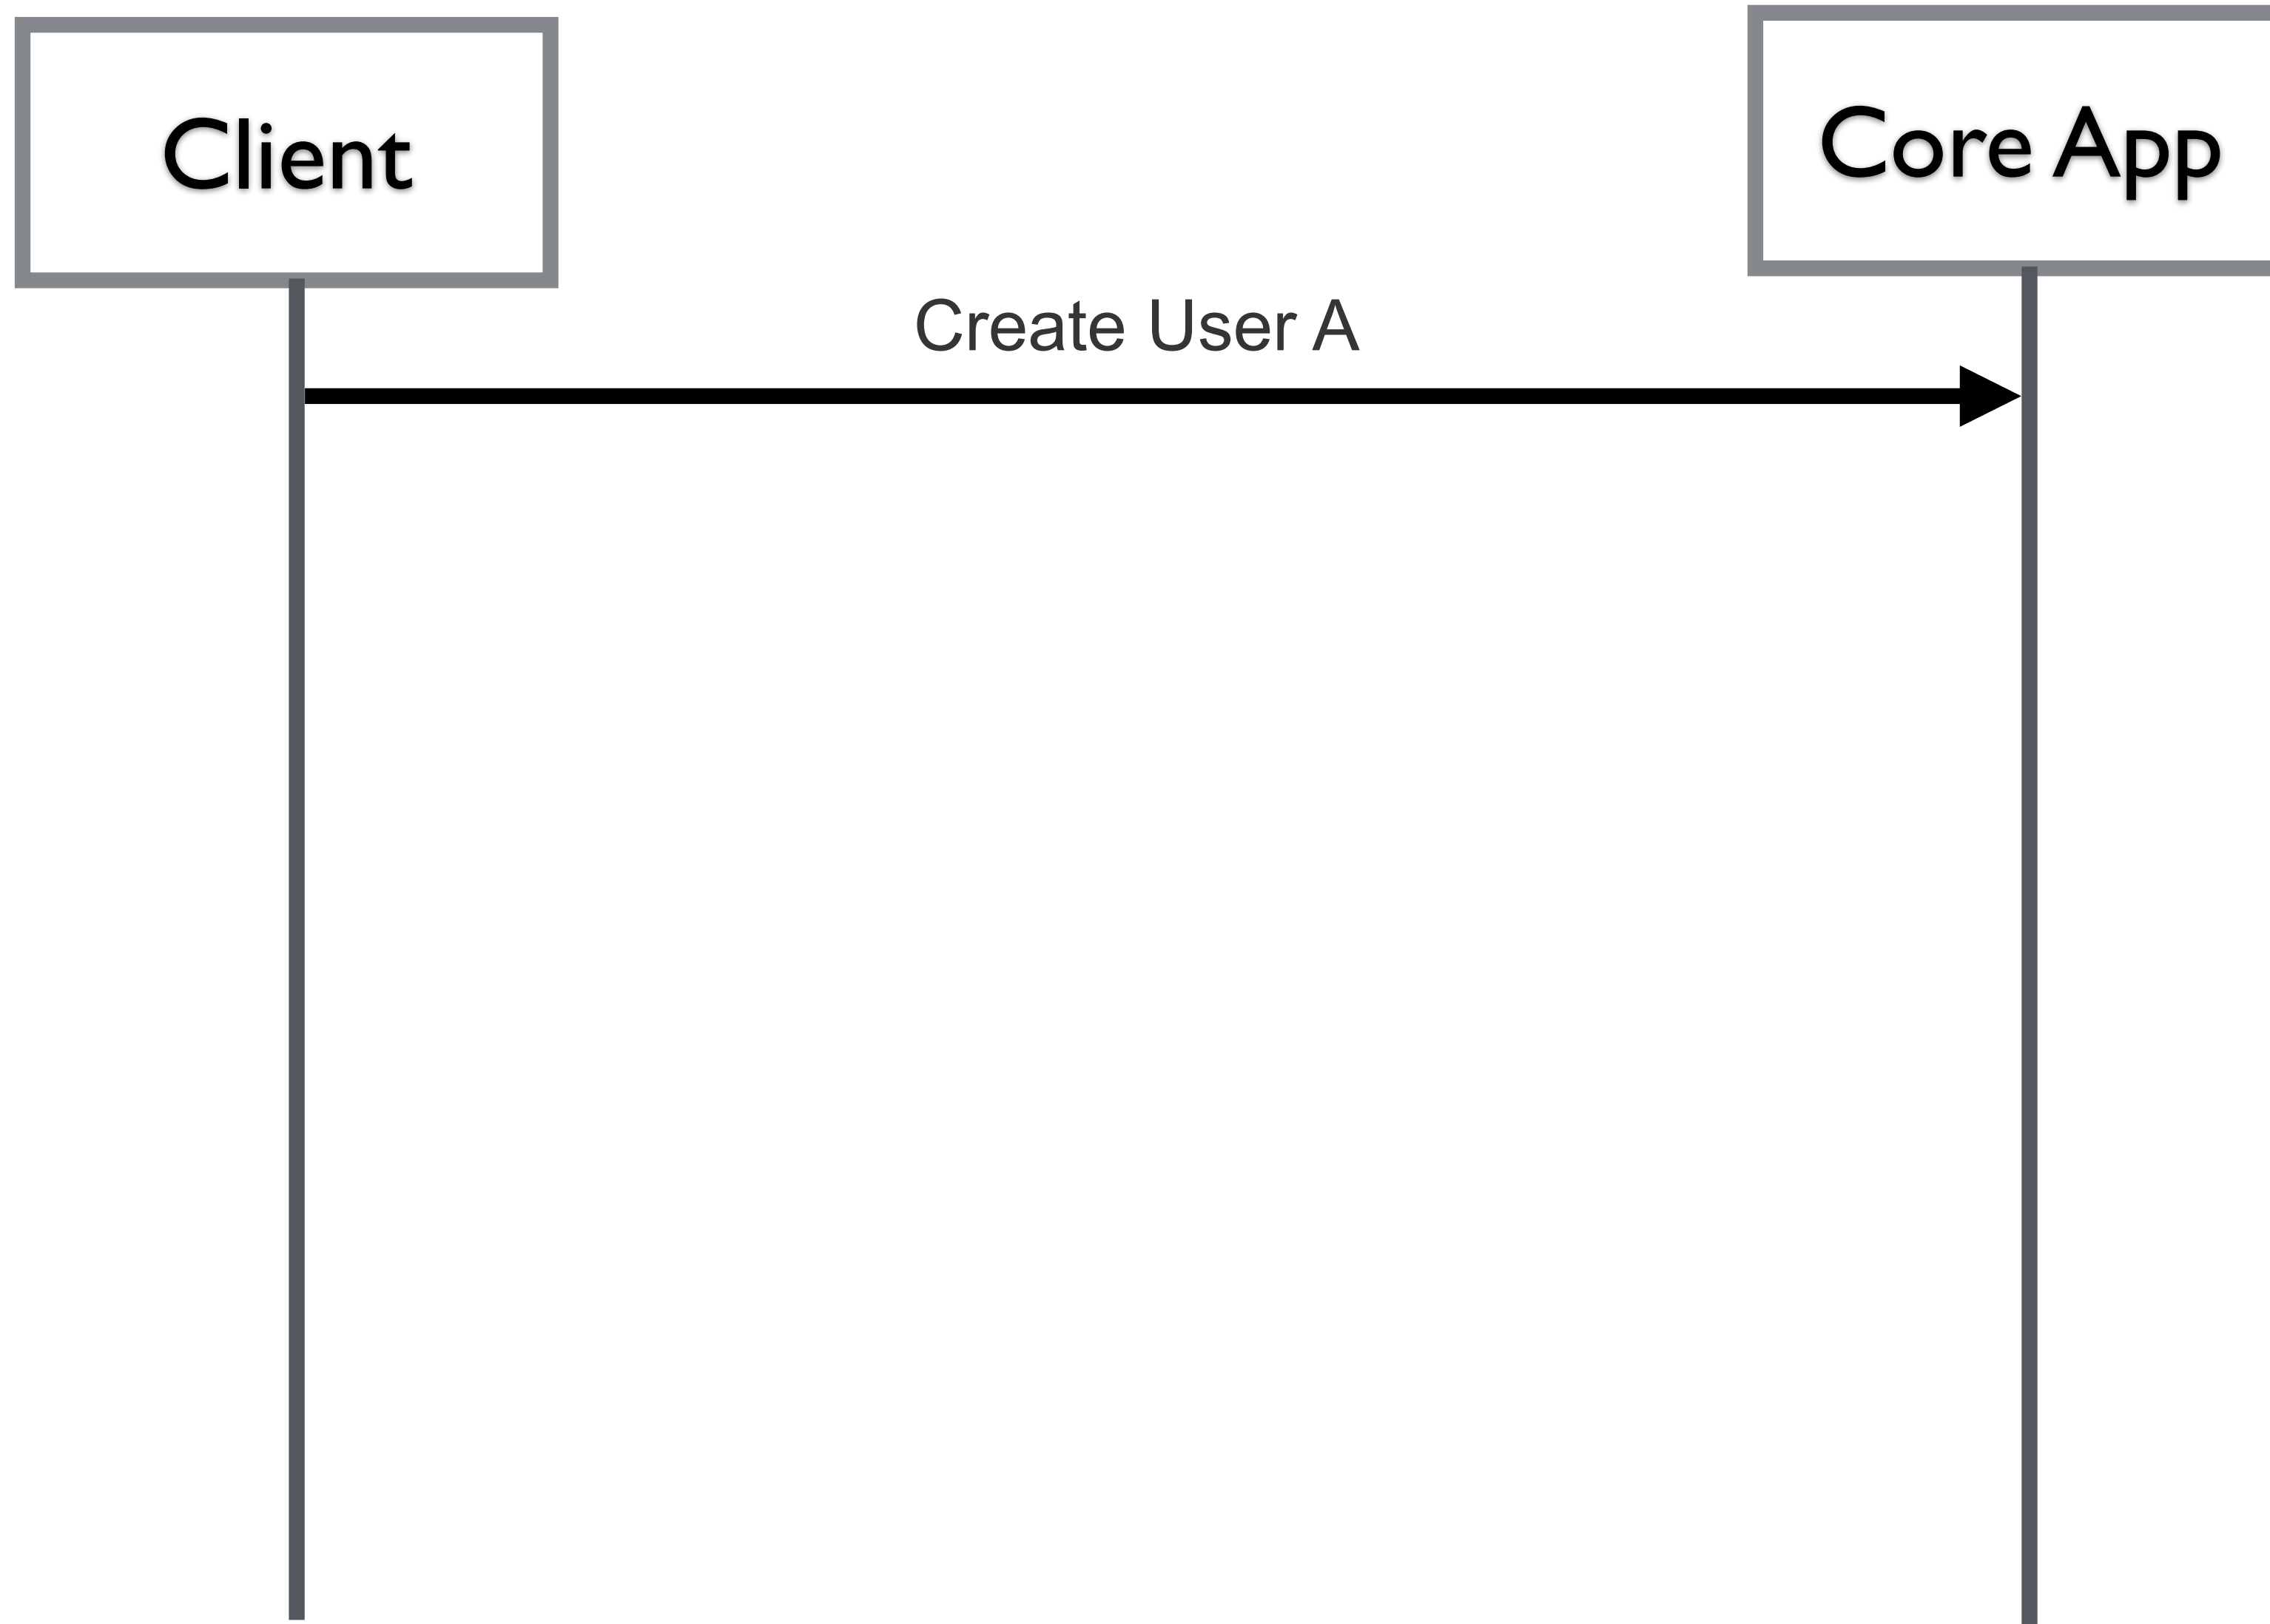
Client



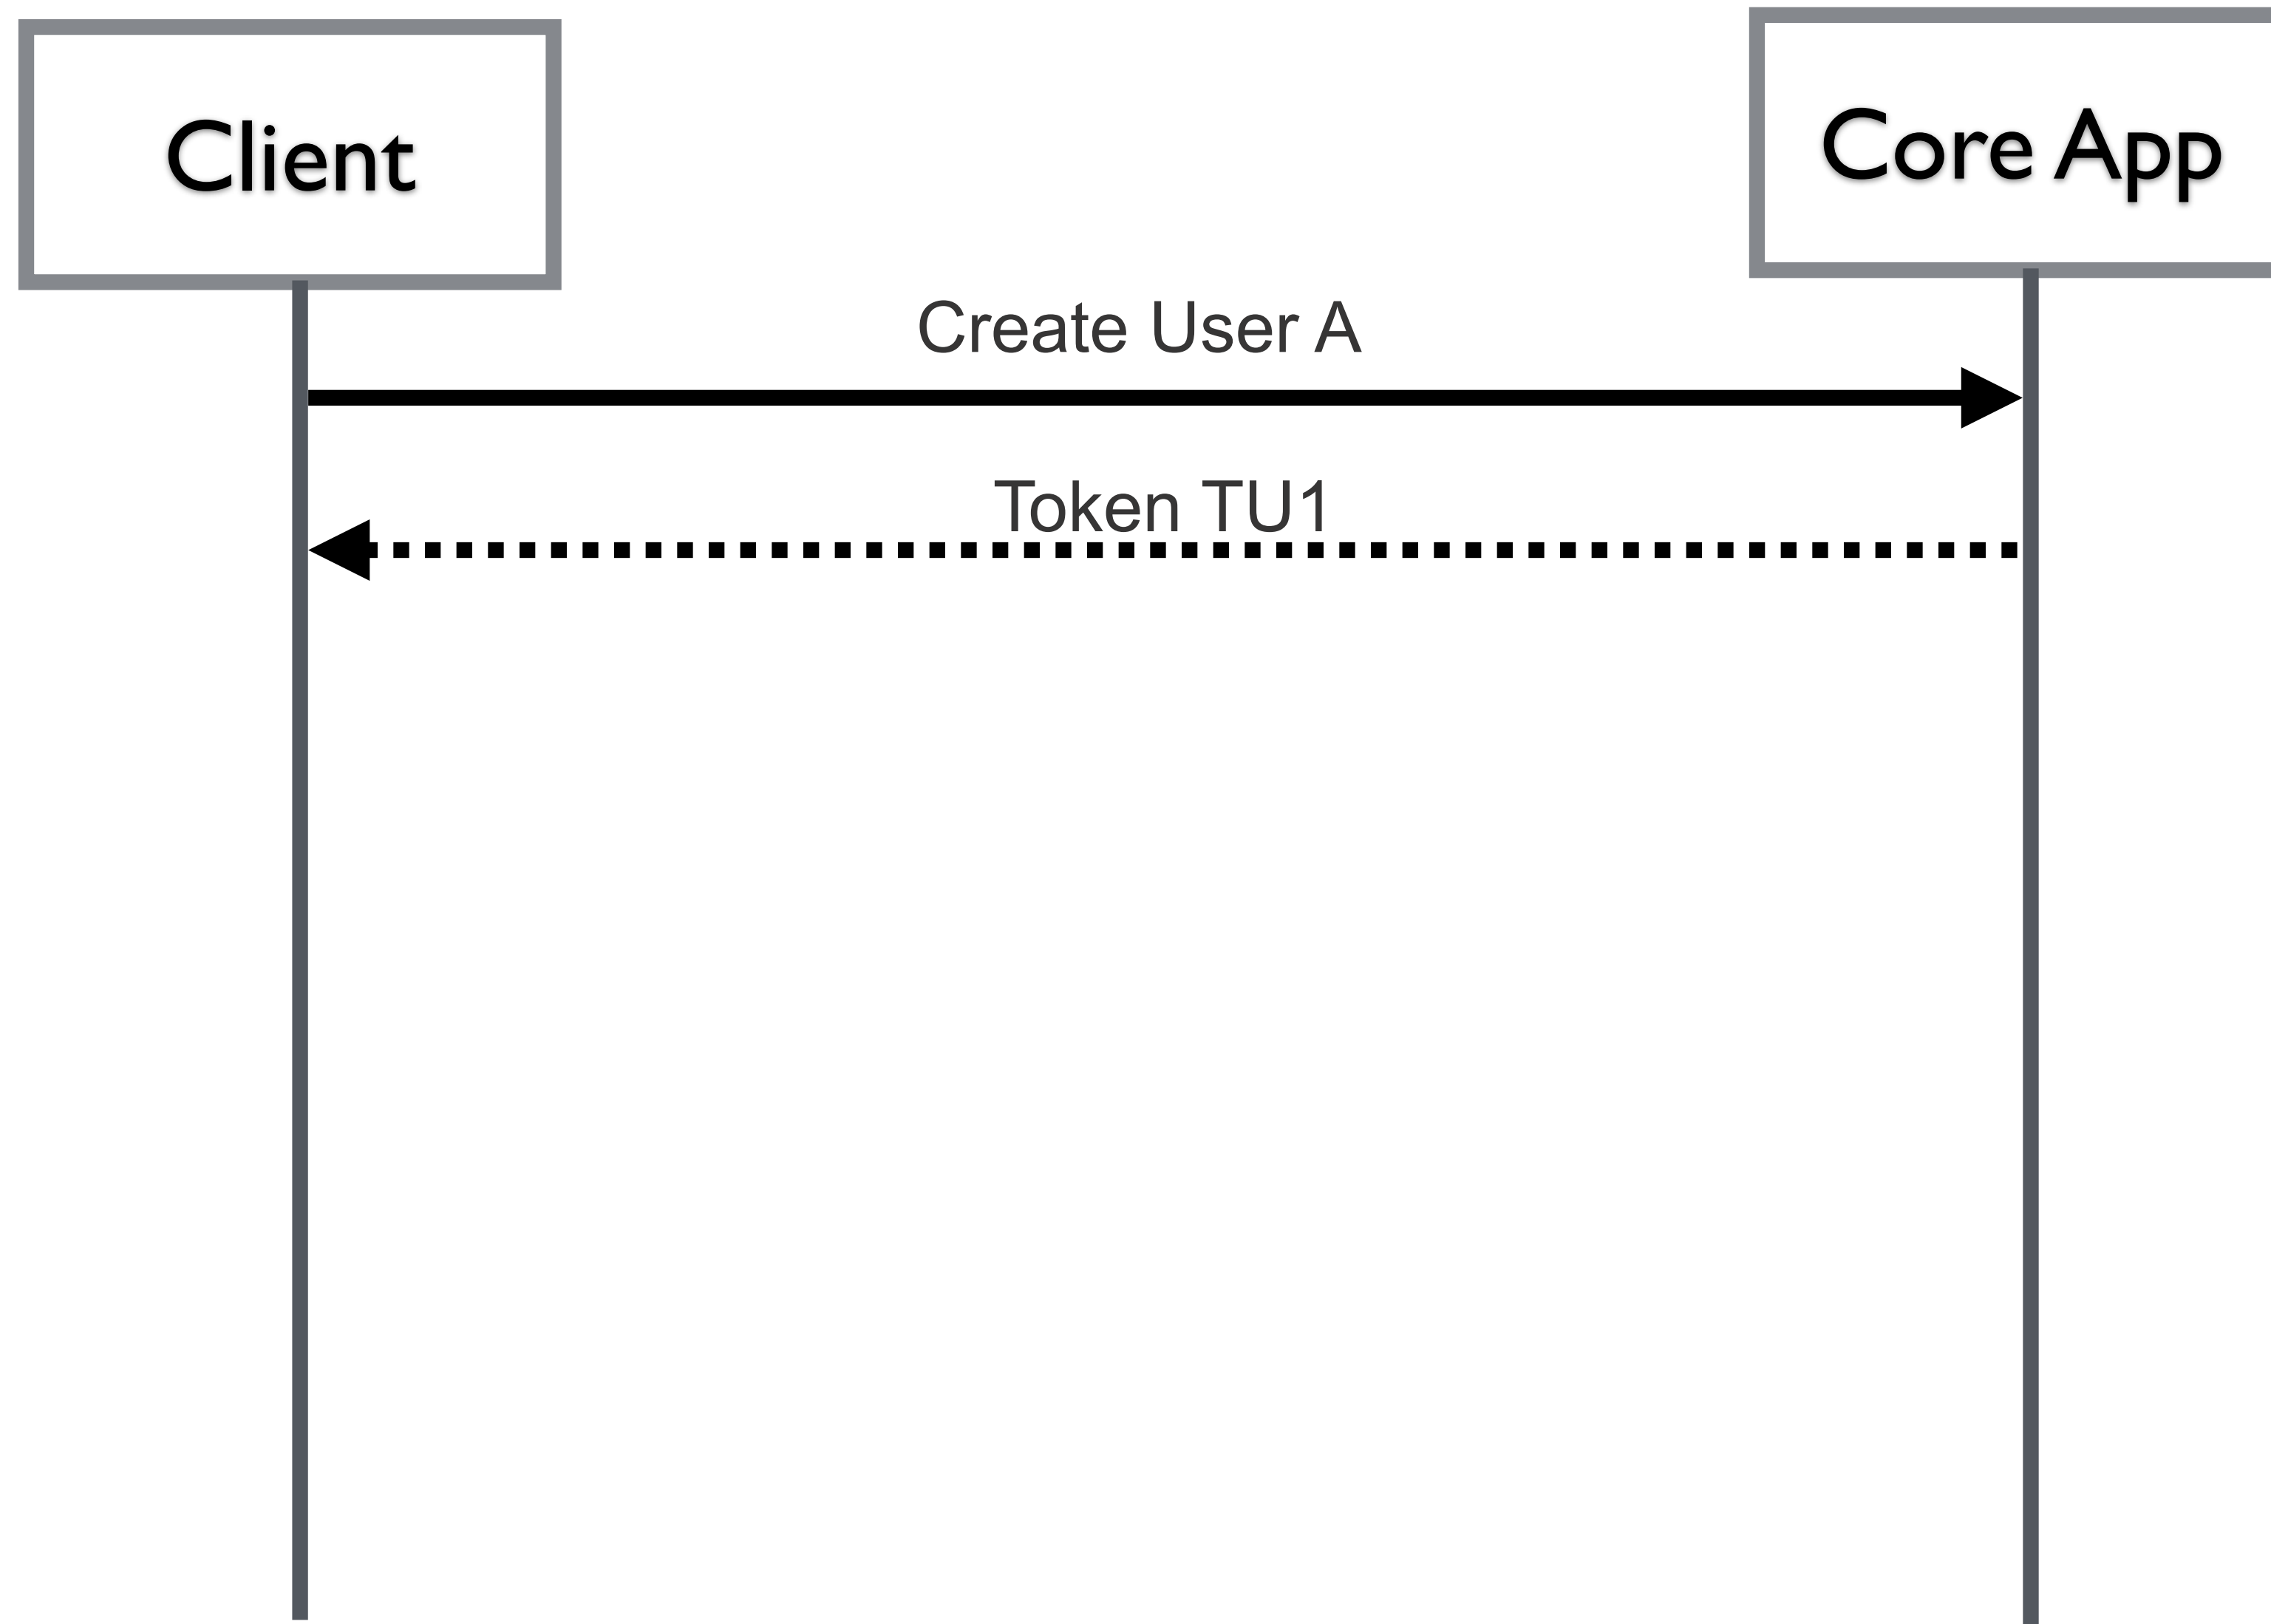
Core App



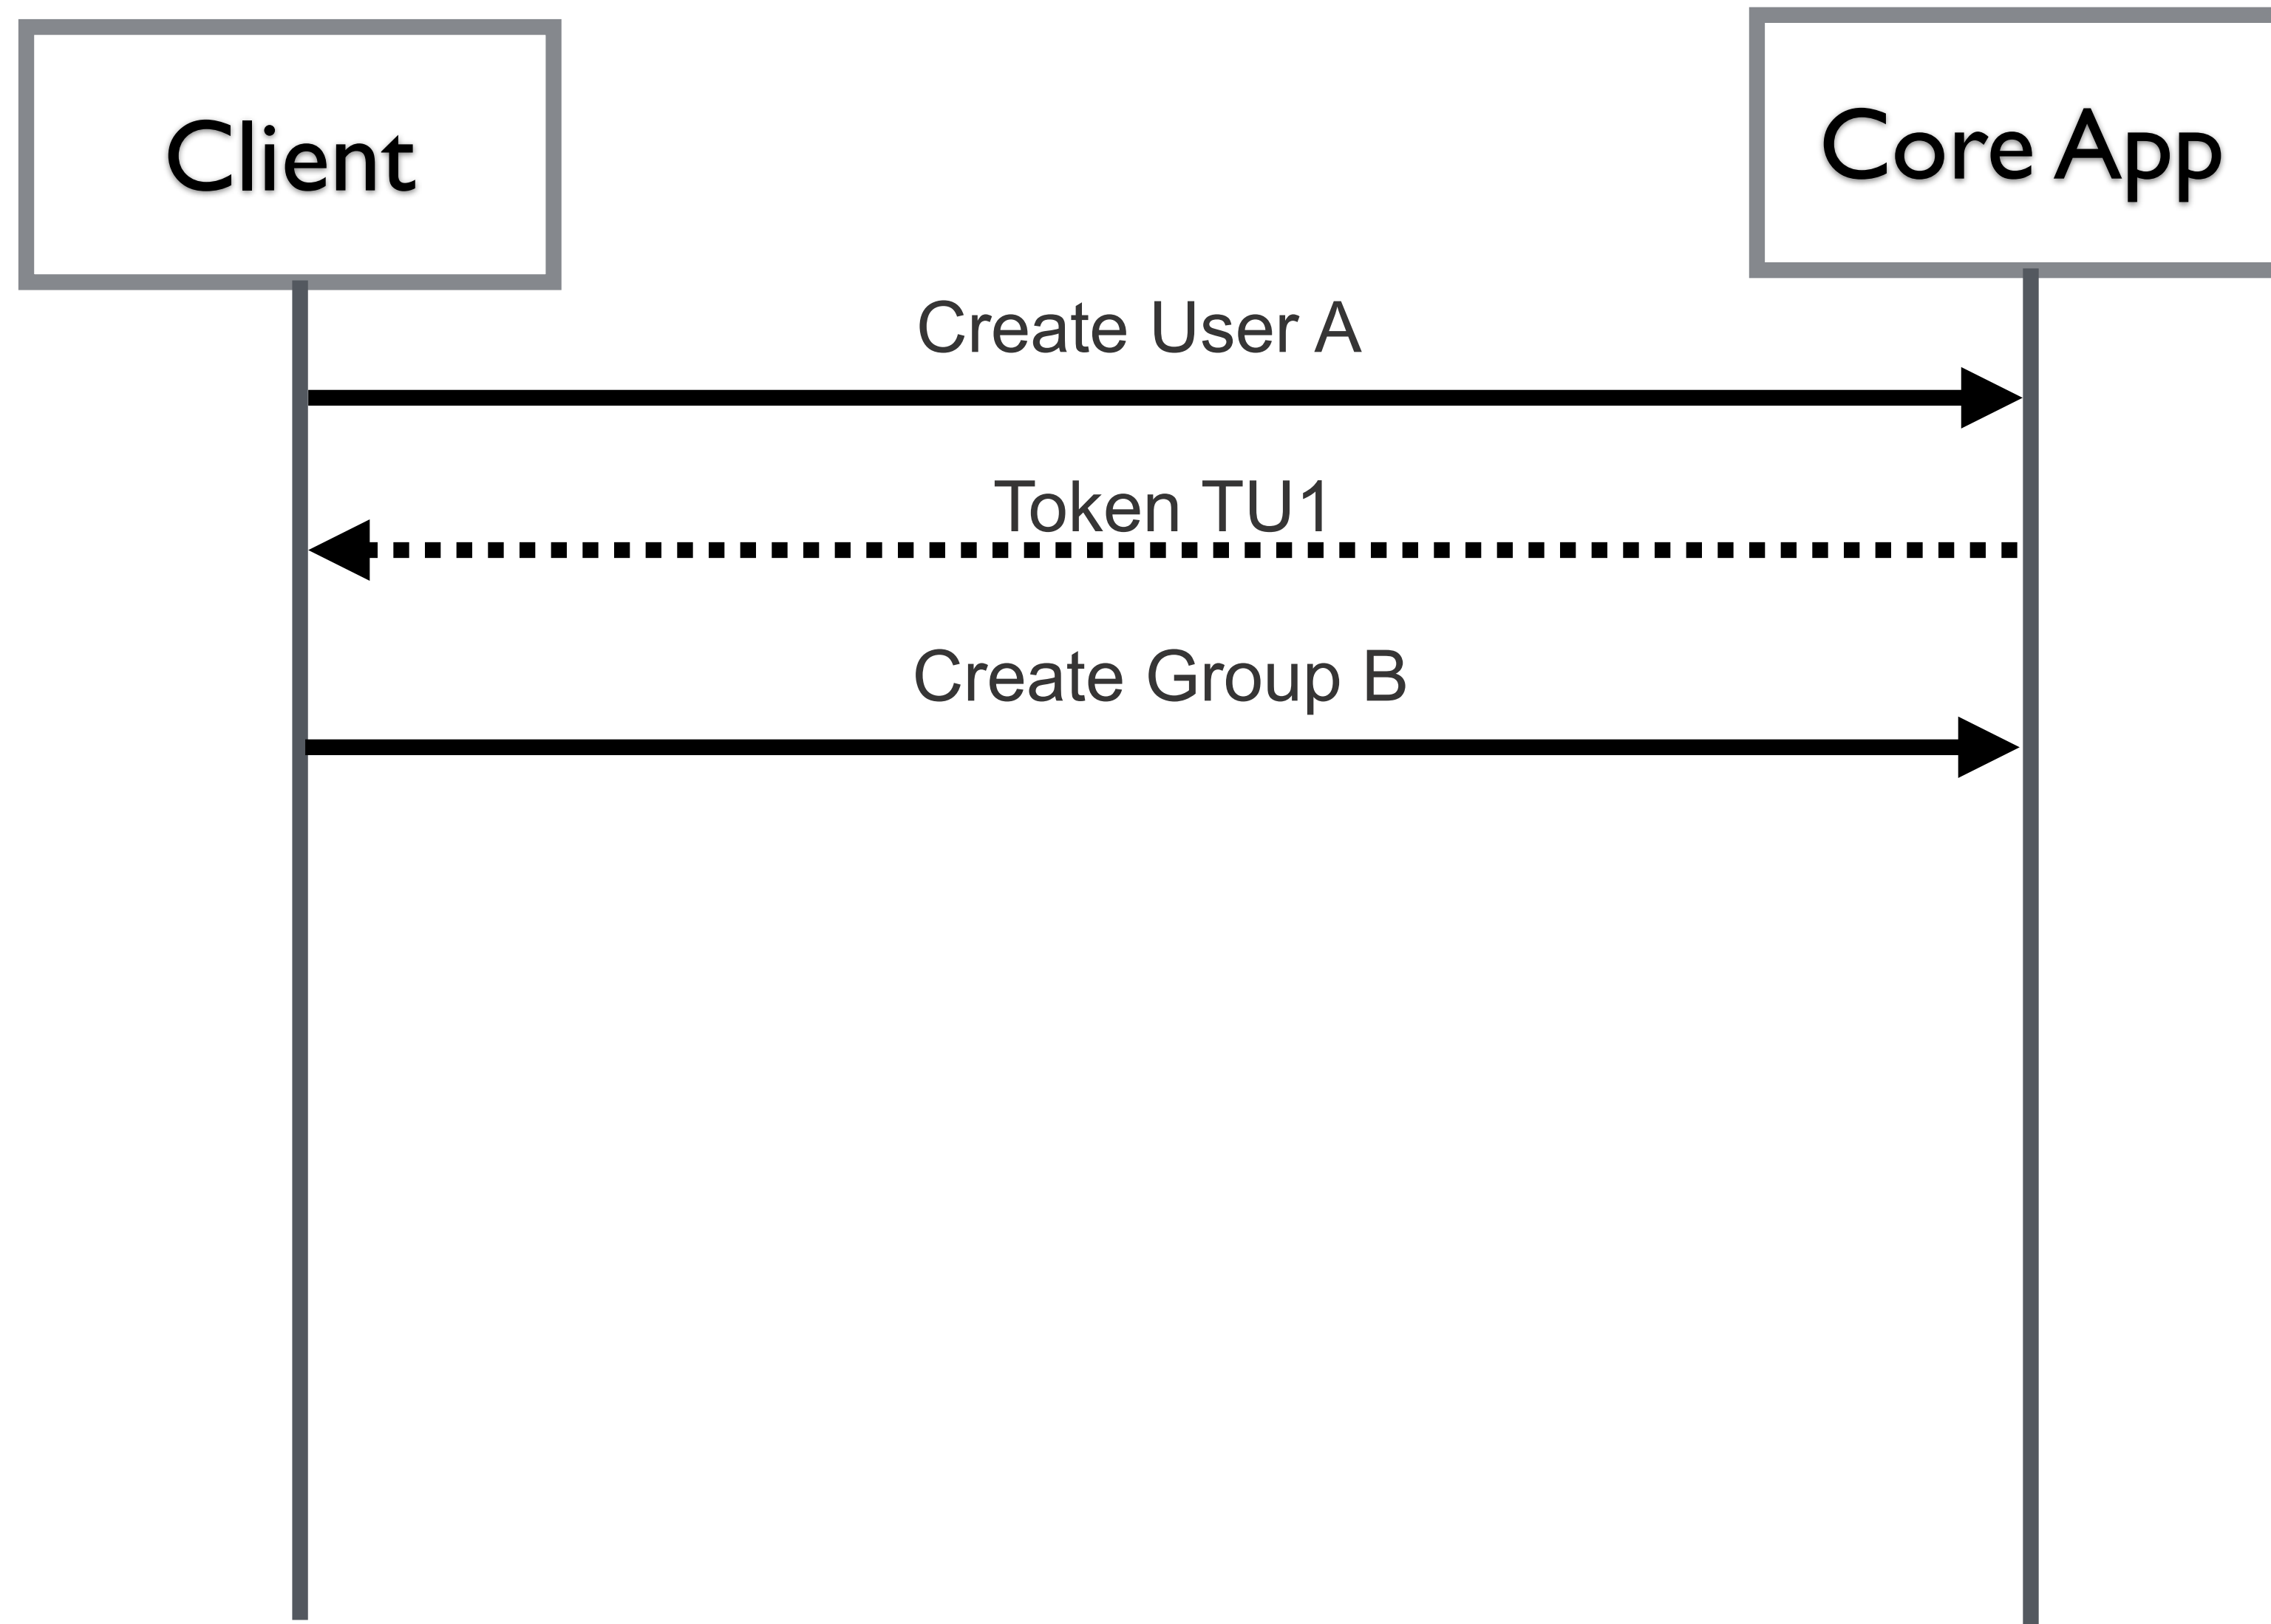
Using tokens to enforce state



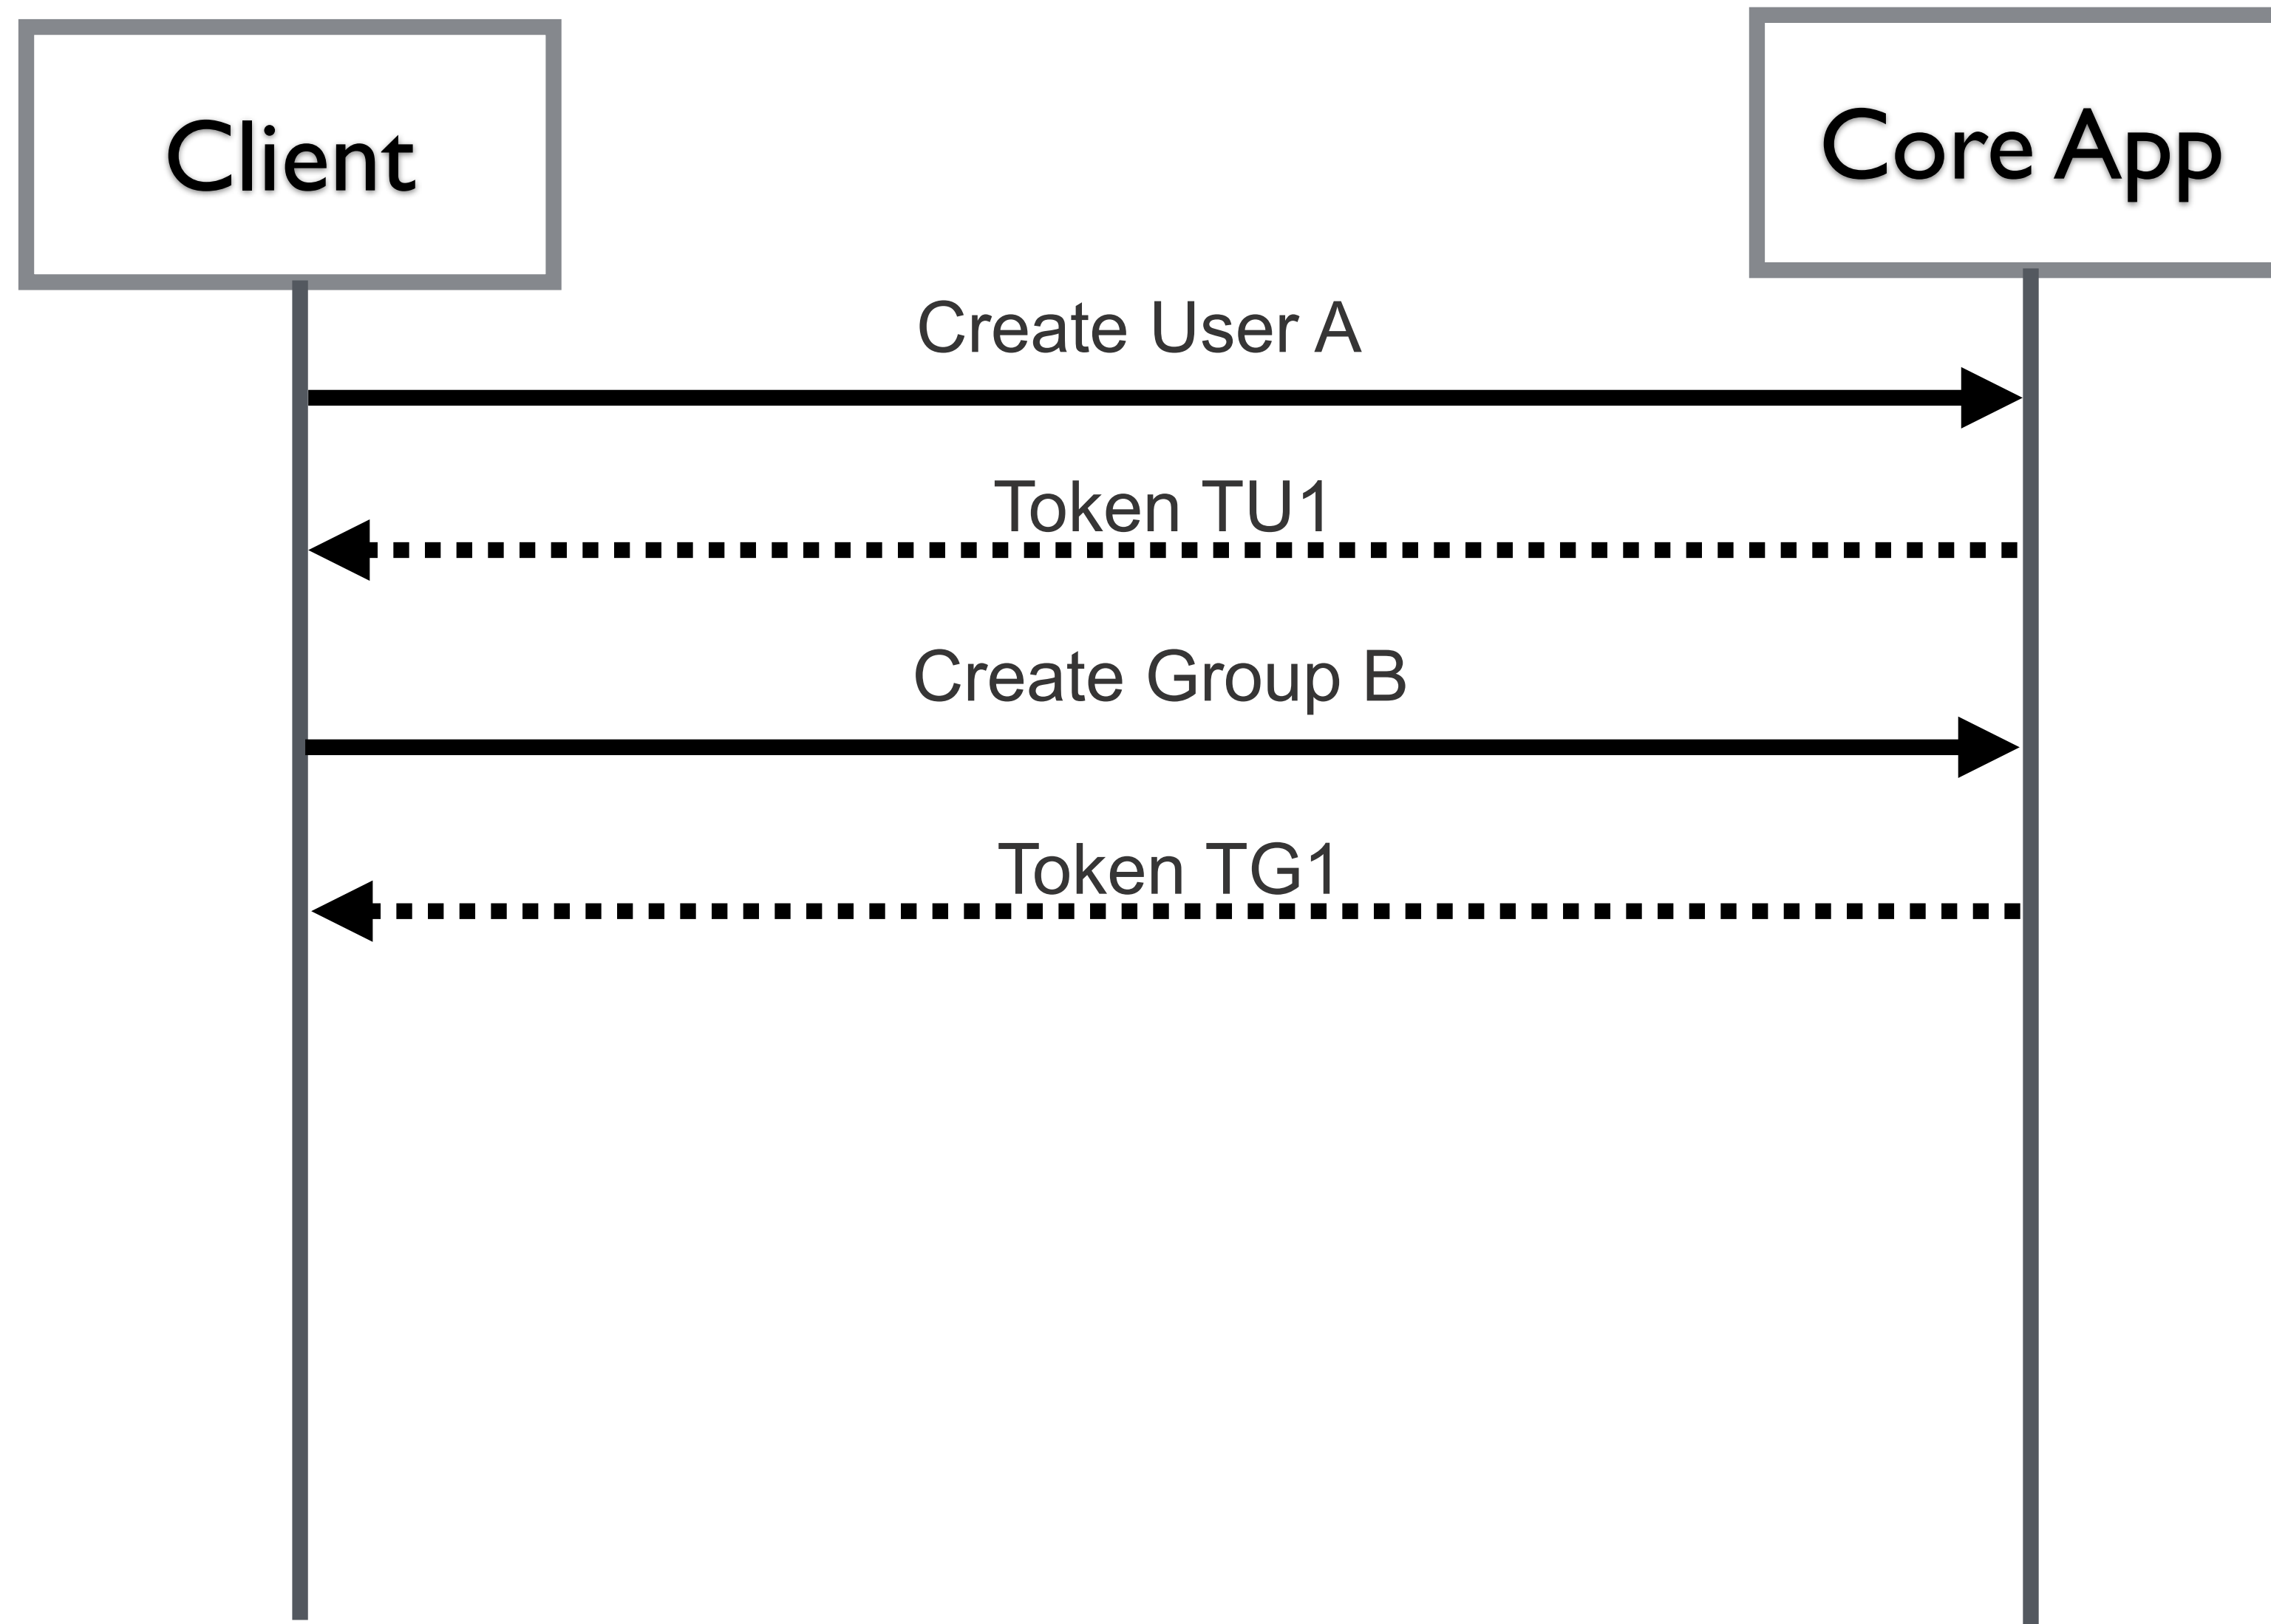
Using tokens to enforce state



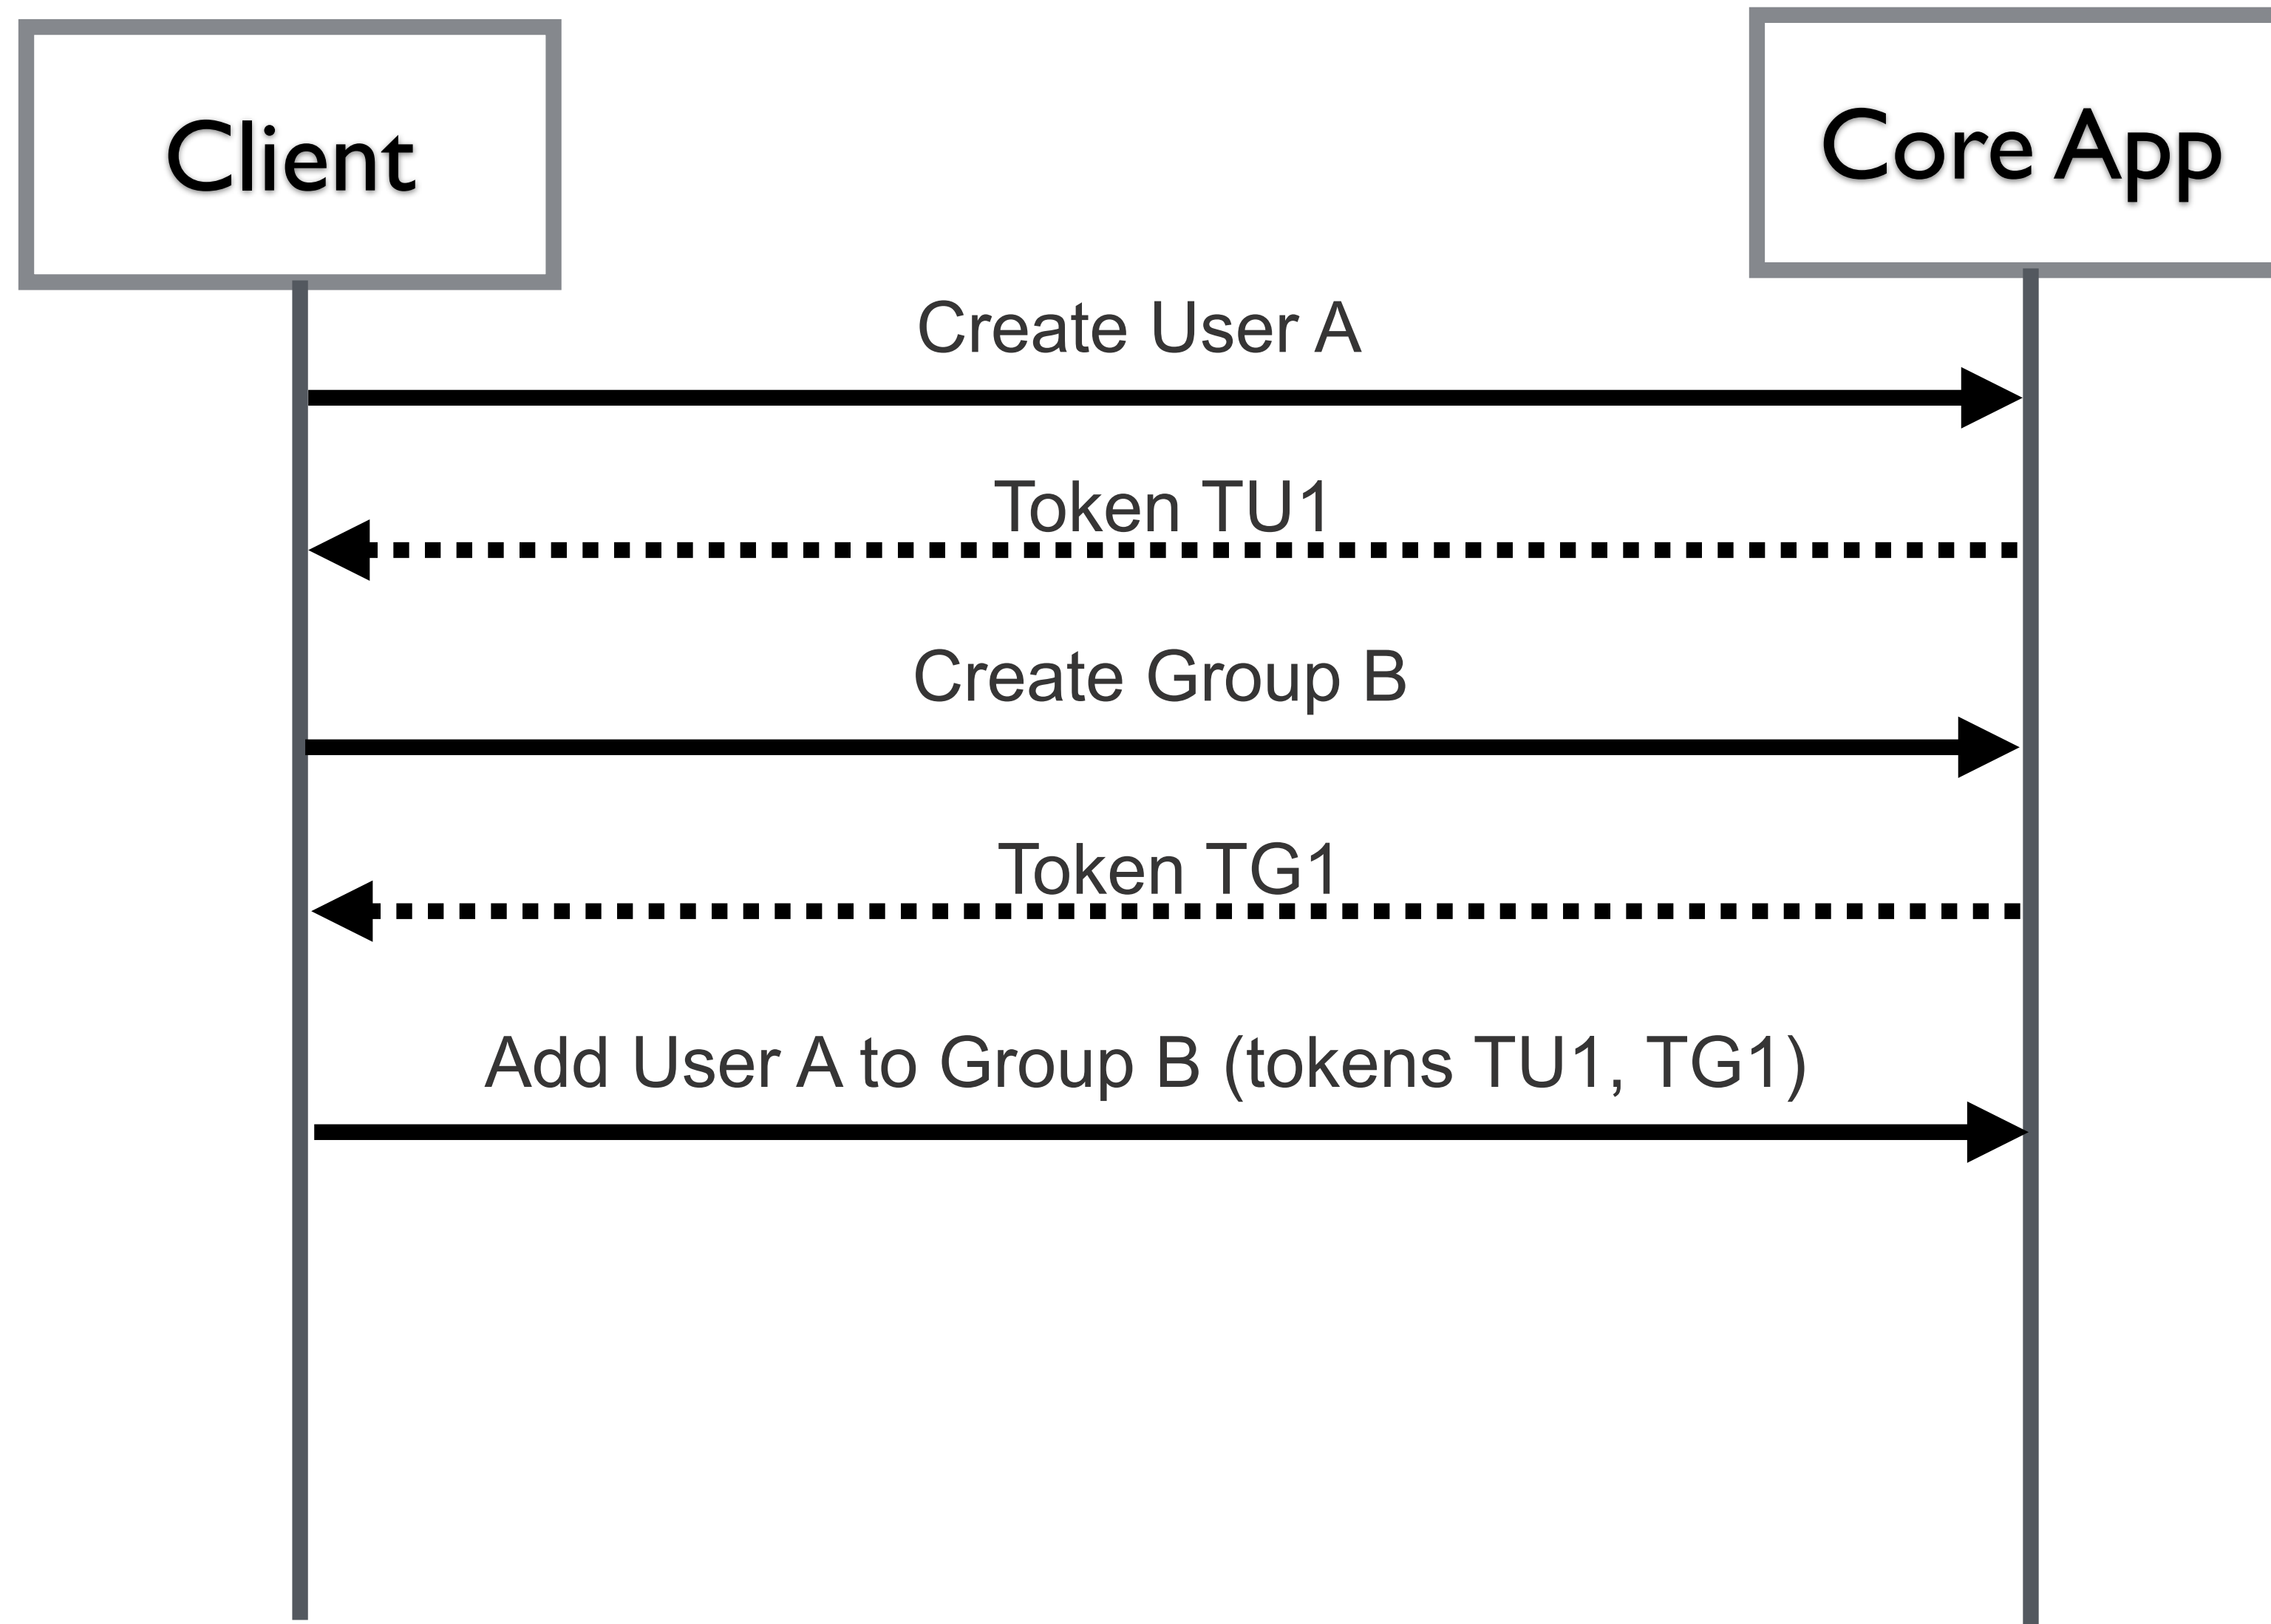
Using tokens to enforce state



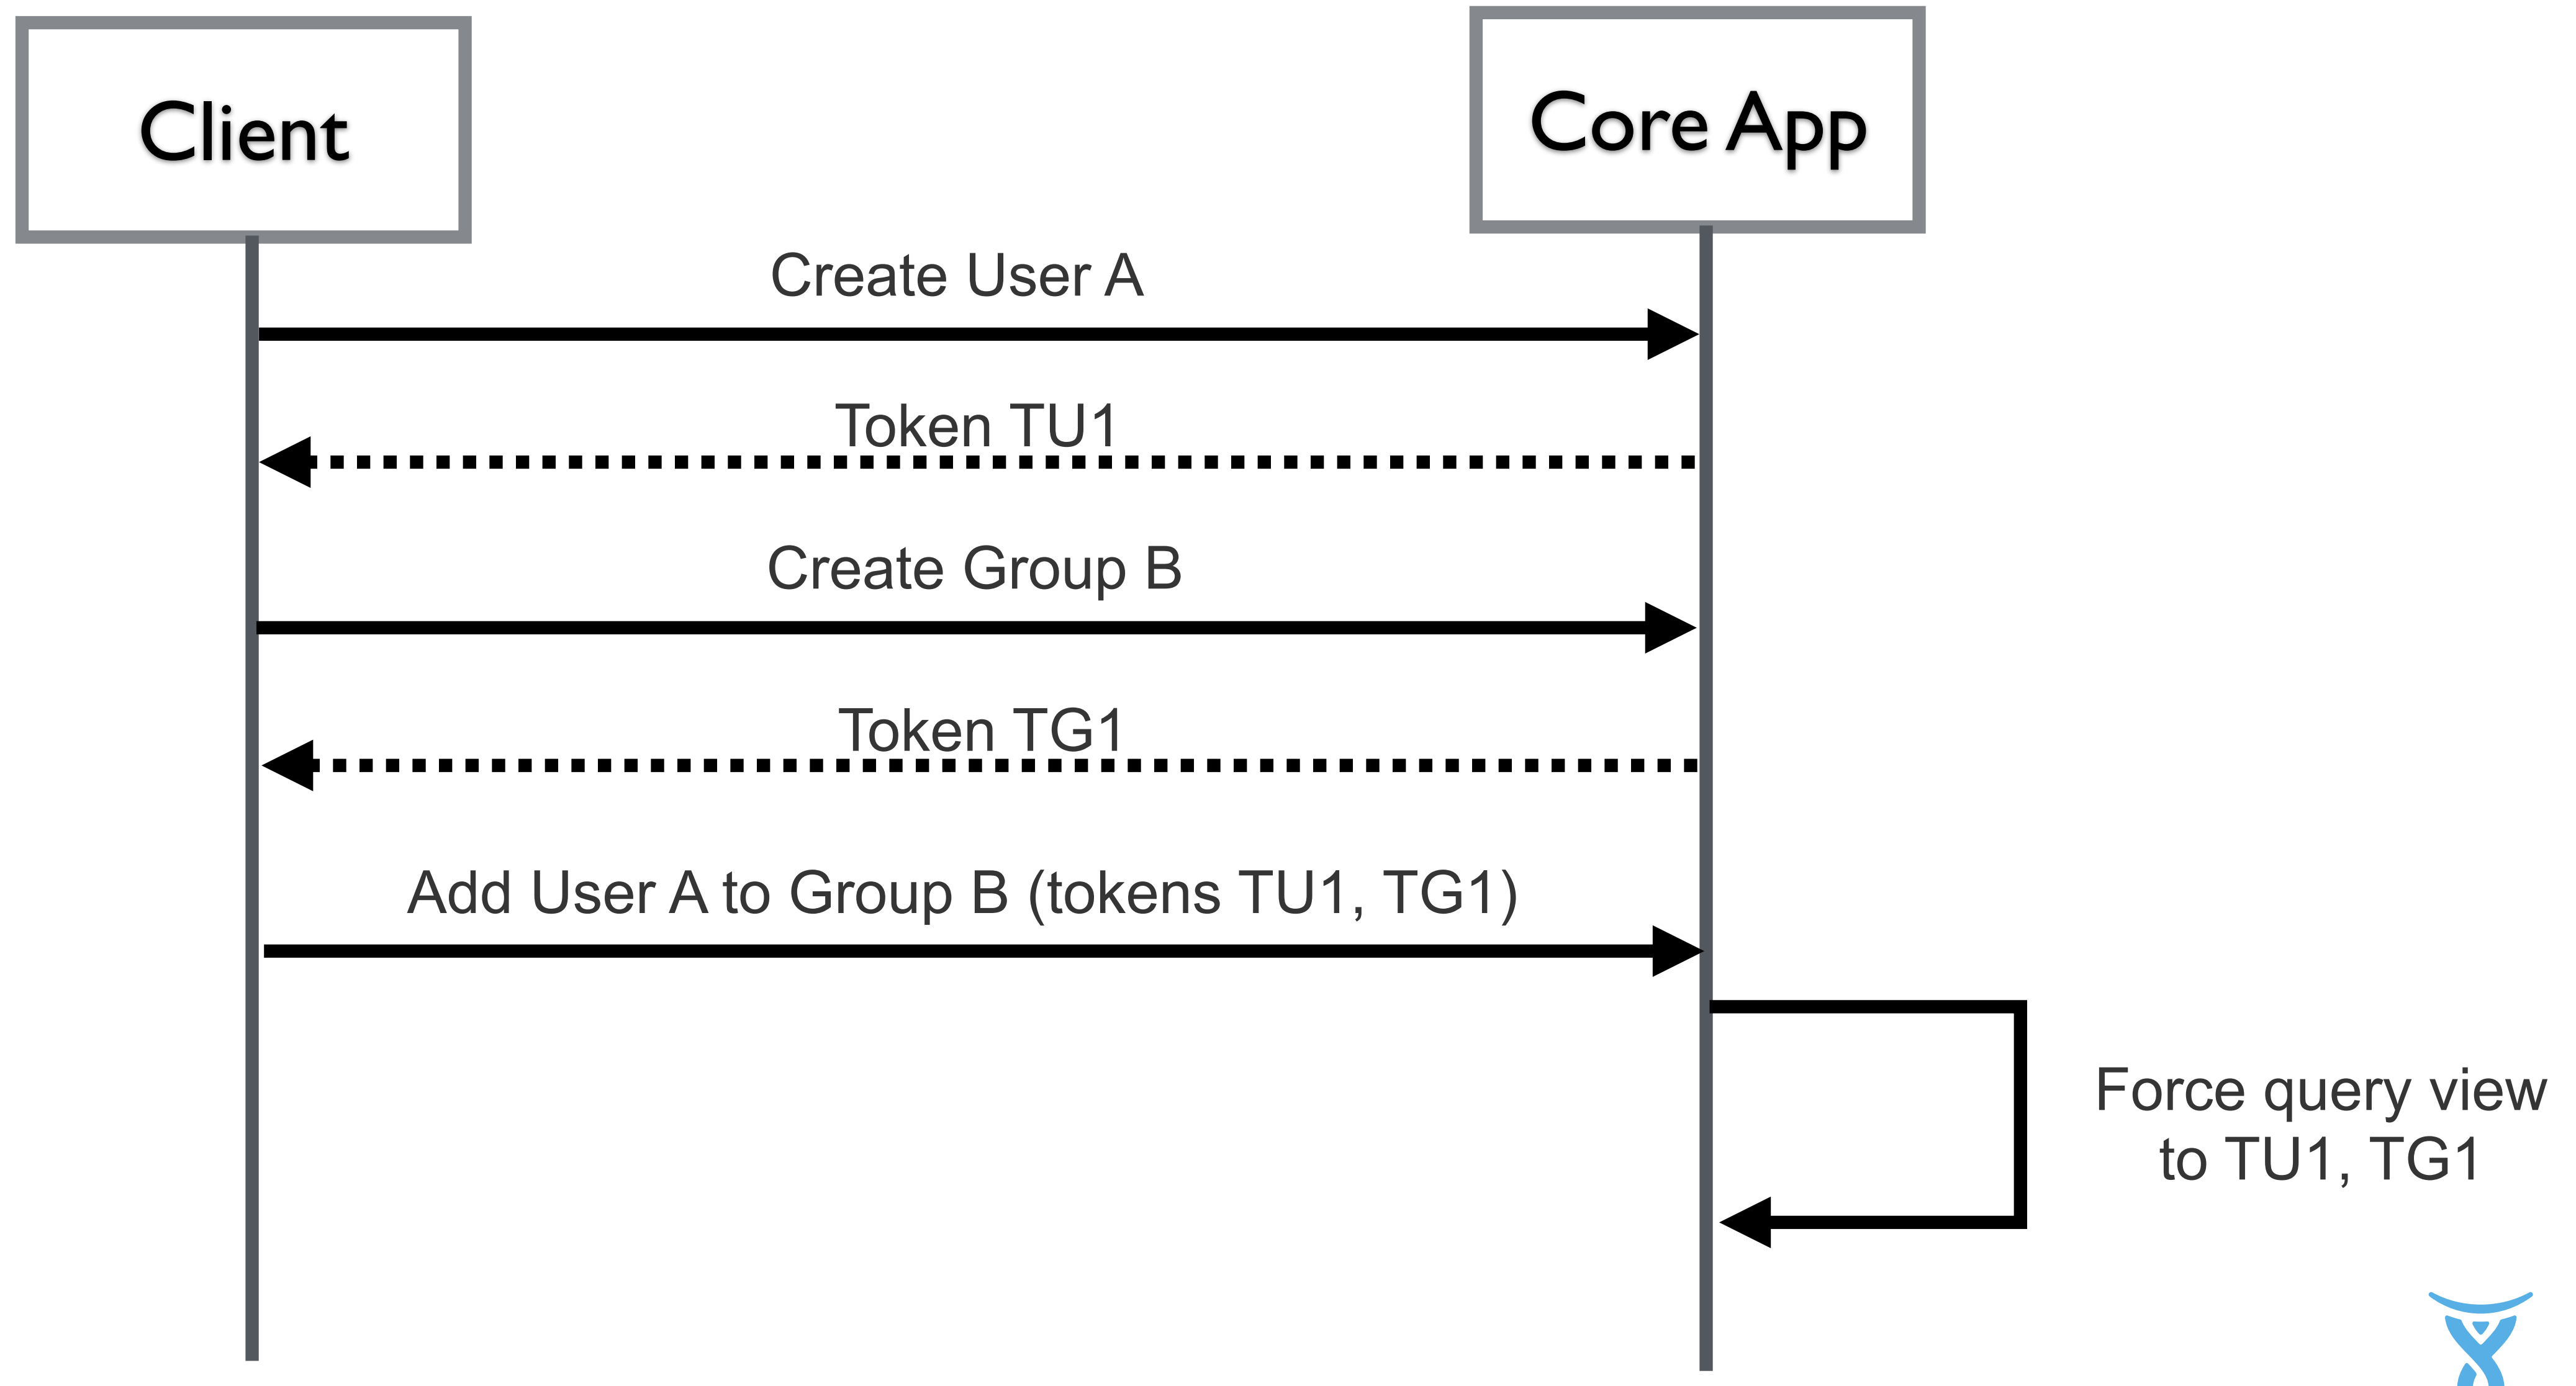
Using tokens to enforce state



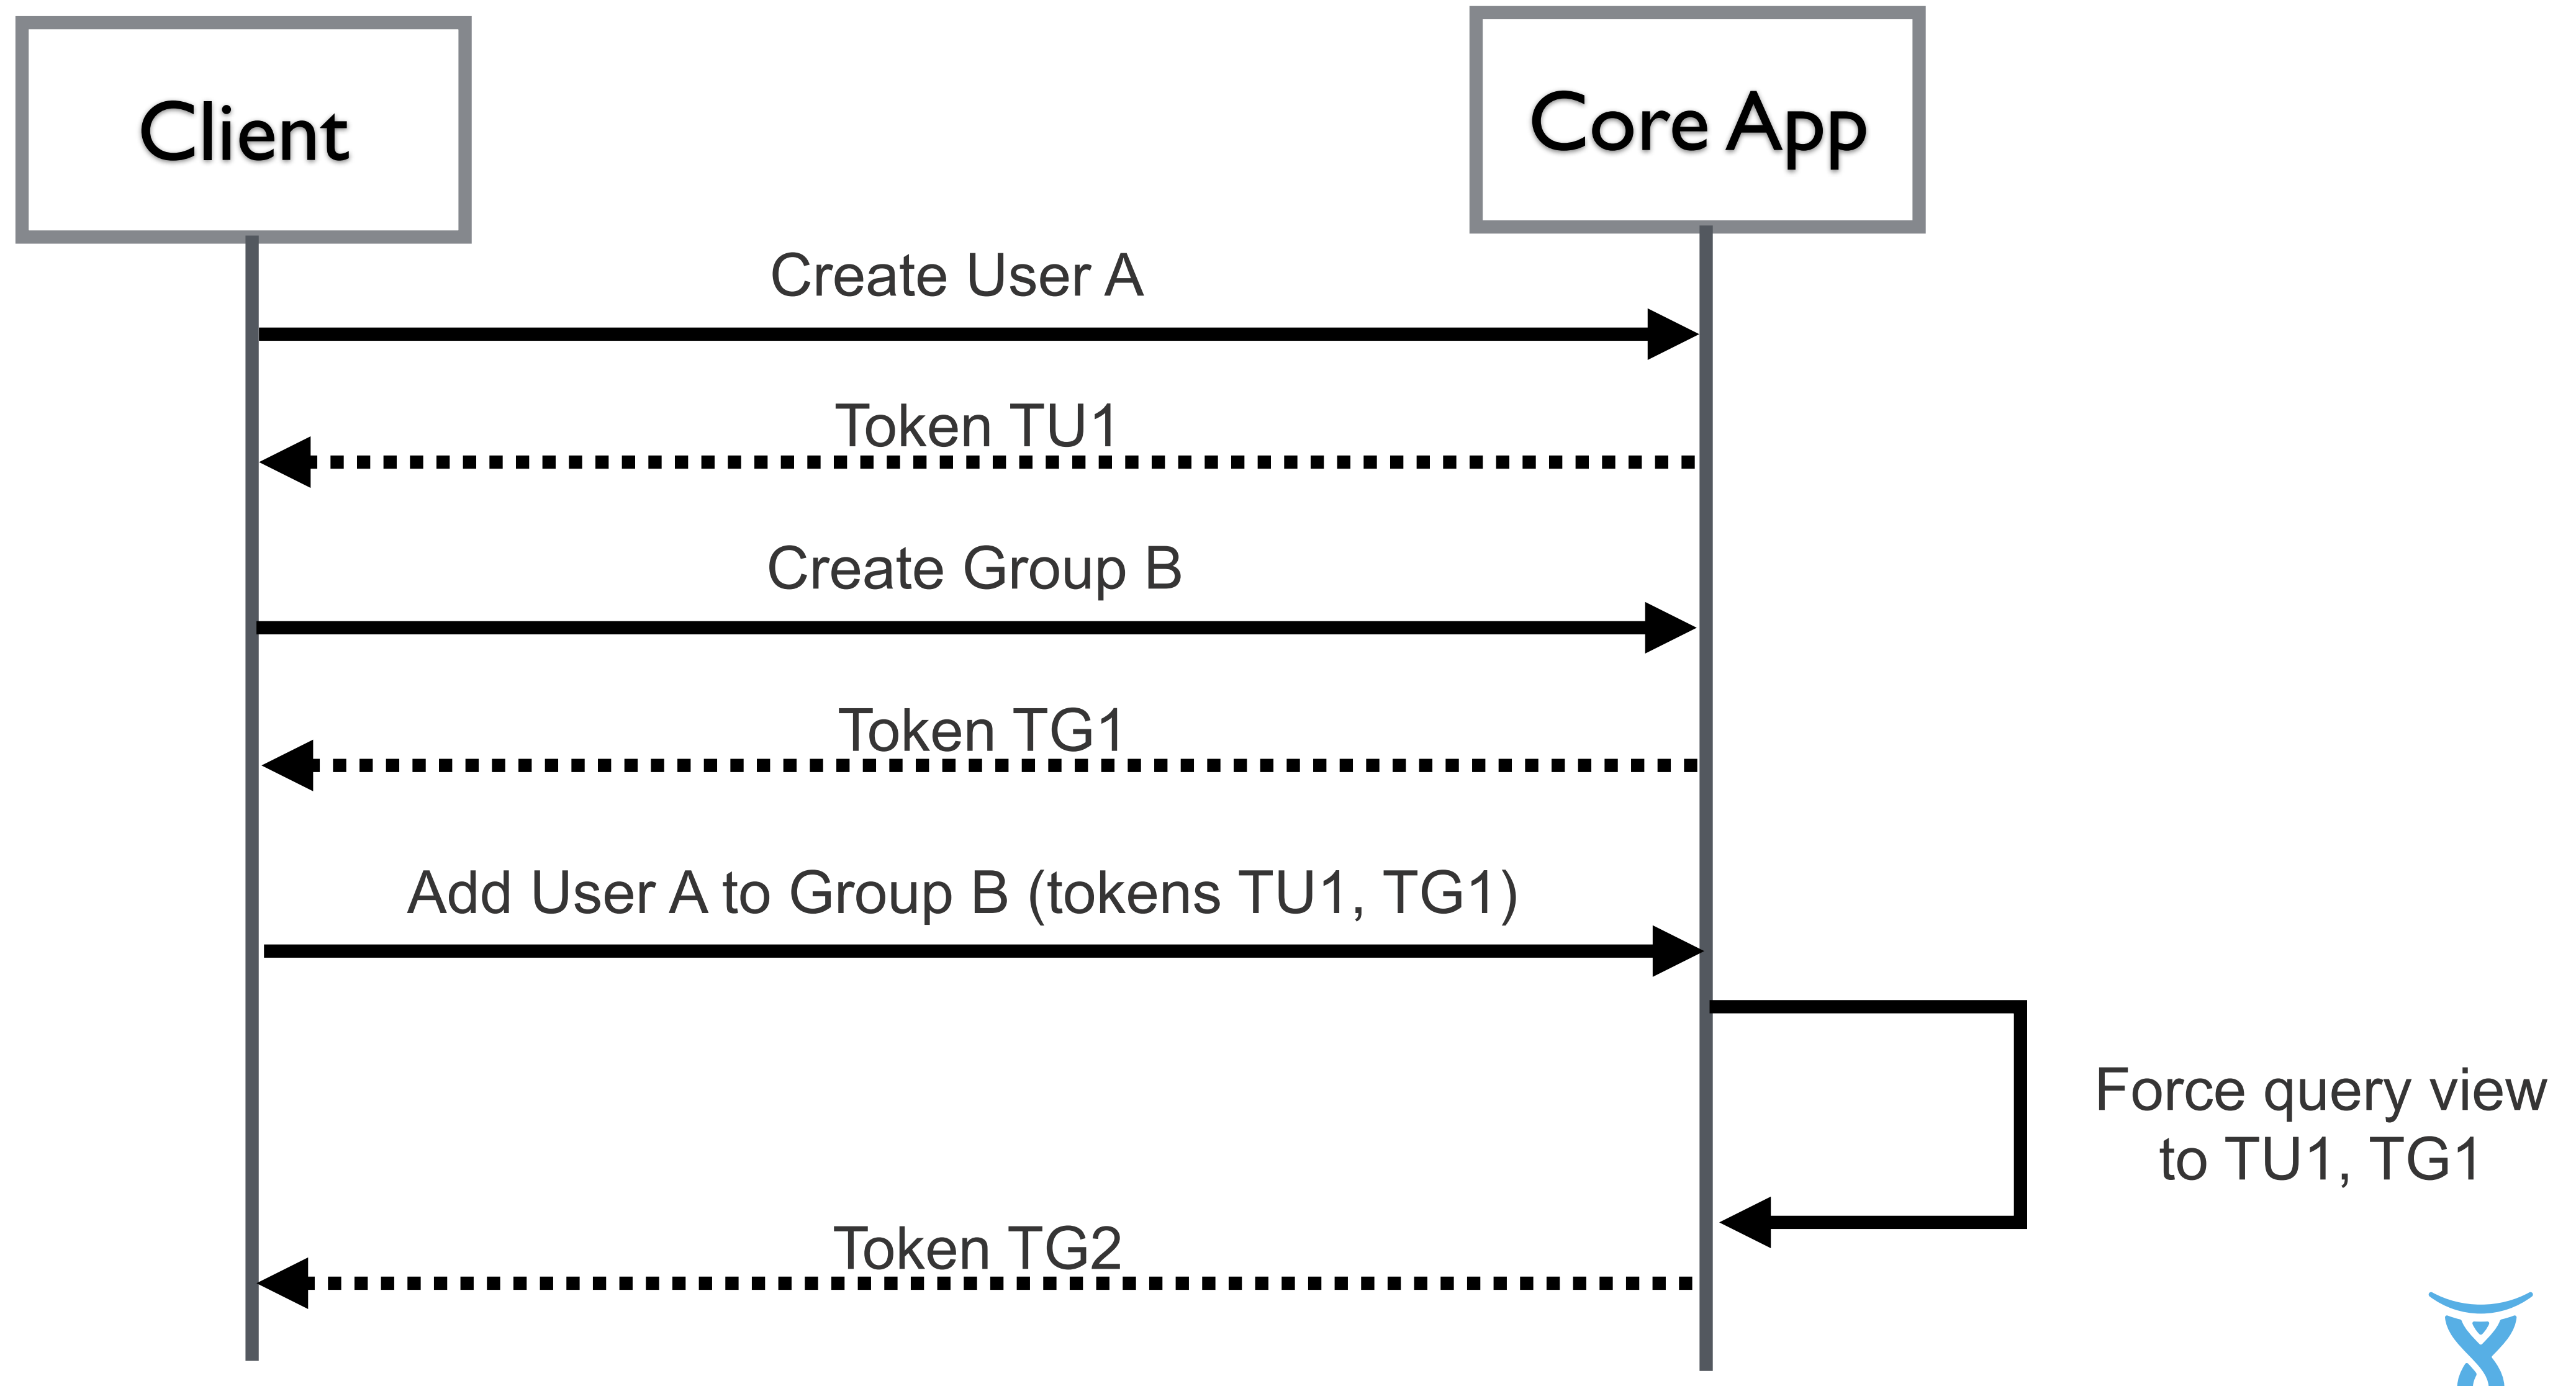
Using tokens to enforce state



Using tokens to enforce state



Using tokens to enforce state





Key takeaways

Key takeaways

- Start small and challenge everything!



Key takeaways

- Start small and challenge everything!
- Incremental architecture for incremental demos



Key takeaways

- Start small and challenge everything!
- Incremental architecture for incremental demos
- Think “Events as an API”



Key takeaways

- Start small and challenge everything!
- Incremental architecture for incremental demos
- Think “Events as an API”
- Accept weaker transactions and eventual consistency



*“We should using
event sourcing more
than we do”*

Martin Fowler (very loosely paraphrased)



event sourcing lib:
bitbucket.org/atlassianlabs/eventsrc

