



# Immutable data stores and CQRS for safety, flexibility and profit



SIDNEY SHEK • ARCHITECT • @SIDNEYSHEK • #SCALADAYS

# Our exploration

WHAT AND WHY

EVOLVING THE ARCHITECTURE

DELVING INTO CODE

3.5 DESIGN PRINCIPLES



# Universe of Users

users

<b>Id</b>	<b>Name</b>	<b>Username</b>	<b>APIKey</b>
1	Homer	homer	d0a
2	Bart	bart	f00
3	Maggie	maggie	baa



# Universe of Users

users

<b>Id</b>	<b>Name</b>	<b>Username</b>	<b>APIKey</b>
1	Homer	homer	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa



# Universe of Users

users

<b>Id</b>	<b>Name</b>	<b>Username</b>	<b>APIKey</b>
1	Homer	homers	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa



# Universe of Users

users			
<b>Id</b>	<b>Name</b>	<b>Username</b>	<b>APIKey</b>
1	Homer	homer	d0a
2	Bart	bart	f00
3	Maggie	maggie	baa

↑

<b>Seq</b>	<b>Event</b>	<b>Time</b>
123	<b>SetUsername(3, Maggie)</b>	0



# Universe of Users

users			
Id	Name	Username	APIKey
1	Homer	homer	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa

↑

Seq	Event	Time
123	SetUsername(3, Maggie)	0
124	SetName(3, Lisa Jr)	10



# Universe of Users

users			
Id	Name	Username	APIKey
1	Homer	homers	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa

↑

Seq	Event	Time
123	SetUsername(3, Maggie)	0
124	SetName(3, Lisa Jr)	10
125	SetUsername(1, homers)	15



# Universe of Users

users_new		
Id	Name	Derived
1	Homer	Homer1
2	Bart	Bart2
3	Lisa Jr	Lisa Jr3

users

Id	Name	Username	APIKey
1	Homer	homers	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa

events

Seq	Event	Time
123	SetUsername(3, Maggie)	0
124	SetName(3, Lisa Jr)	10
125	SetUsername(1, homers)	15



# Universe of Users

users_new		
Id	Name	Derived
1	Homer	Homer1
2	Bart	Bart2
3	Lisa Jr	Lisa Jr3

users			
Id	Name	Username	APIKey
1	Homer	homers	d0a
2	Bart	bart	f00
3	Lisa Jr	maggie	baa

Seq	Event	Time
123	SetUsername(3, Maggie)	0
124	SetName(3, Lisa Jr)	10
125	SetUsername(1, homers)	15



# Our Identity System requirements



# Our Identity System requirements

- Users, groups and memberships



# Our Identity System requirements

- Users, groups and memberships
  - Searching for users



# Our Identity System requirements

- Users, groups and memberships
  - Searching for users
  - Retrieve by email



# Our Identity System requirements

- Users, groups and memberships
  - Searching for users
  - Retrieve by email
- High volume low latency reads



# Our Identity System requirements

- Users, groups and memberships
  - Searching for users
  - Retrieve by email
  - Incremental synchronisation
- High volume low latency reads



# Our Identity System requirements

- Users, groups and memberships
  - Searching for users
  - Retrieve by email
  - Incremental synchronisation
- Audit trails for changes
- High volume low latency reads



# Our Identity System requirements

- Users, groups and memberships
  - Searching for users
  - Retrieve by email
  - Incremental synchronisation
- Audit trails for changes
- High volume low latency reads
- Highly available
  - Disaster recovery
  - Zero-downtime upgrades



# Our Identity System requirements

- Users, groups and memberships
  - Searching for users
  - Retrieve by email
  - Incremental synchronisation
- Audit trails for changes
  - High volume low latency reads
  - Highly available
    - Disaster recovery
    - Zero-downtime upgrades
  - Testing with production-like data

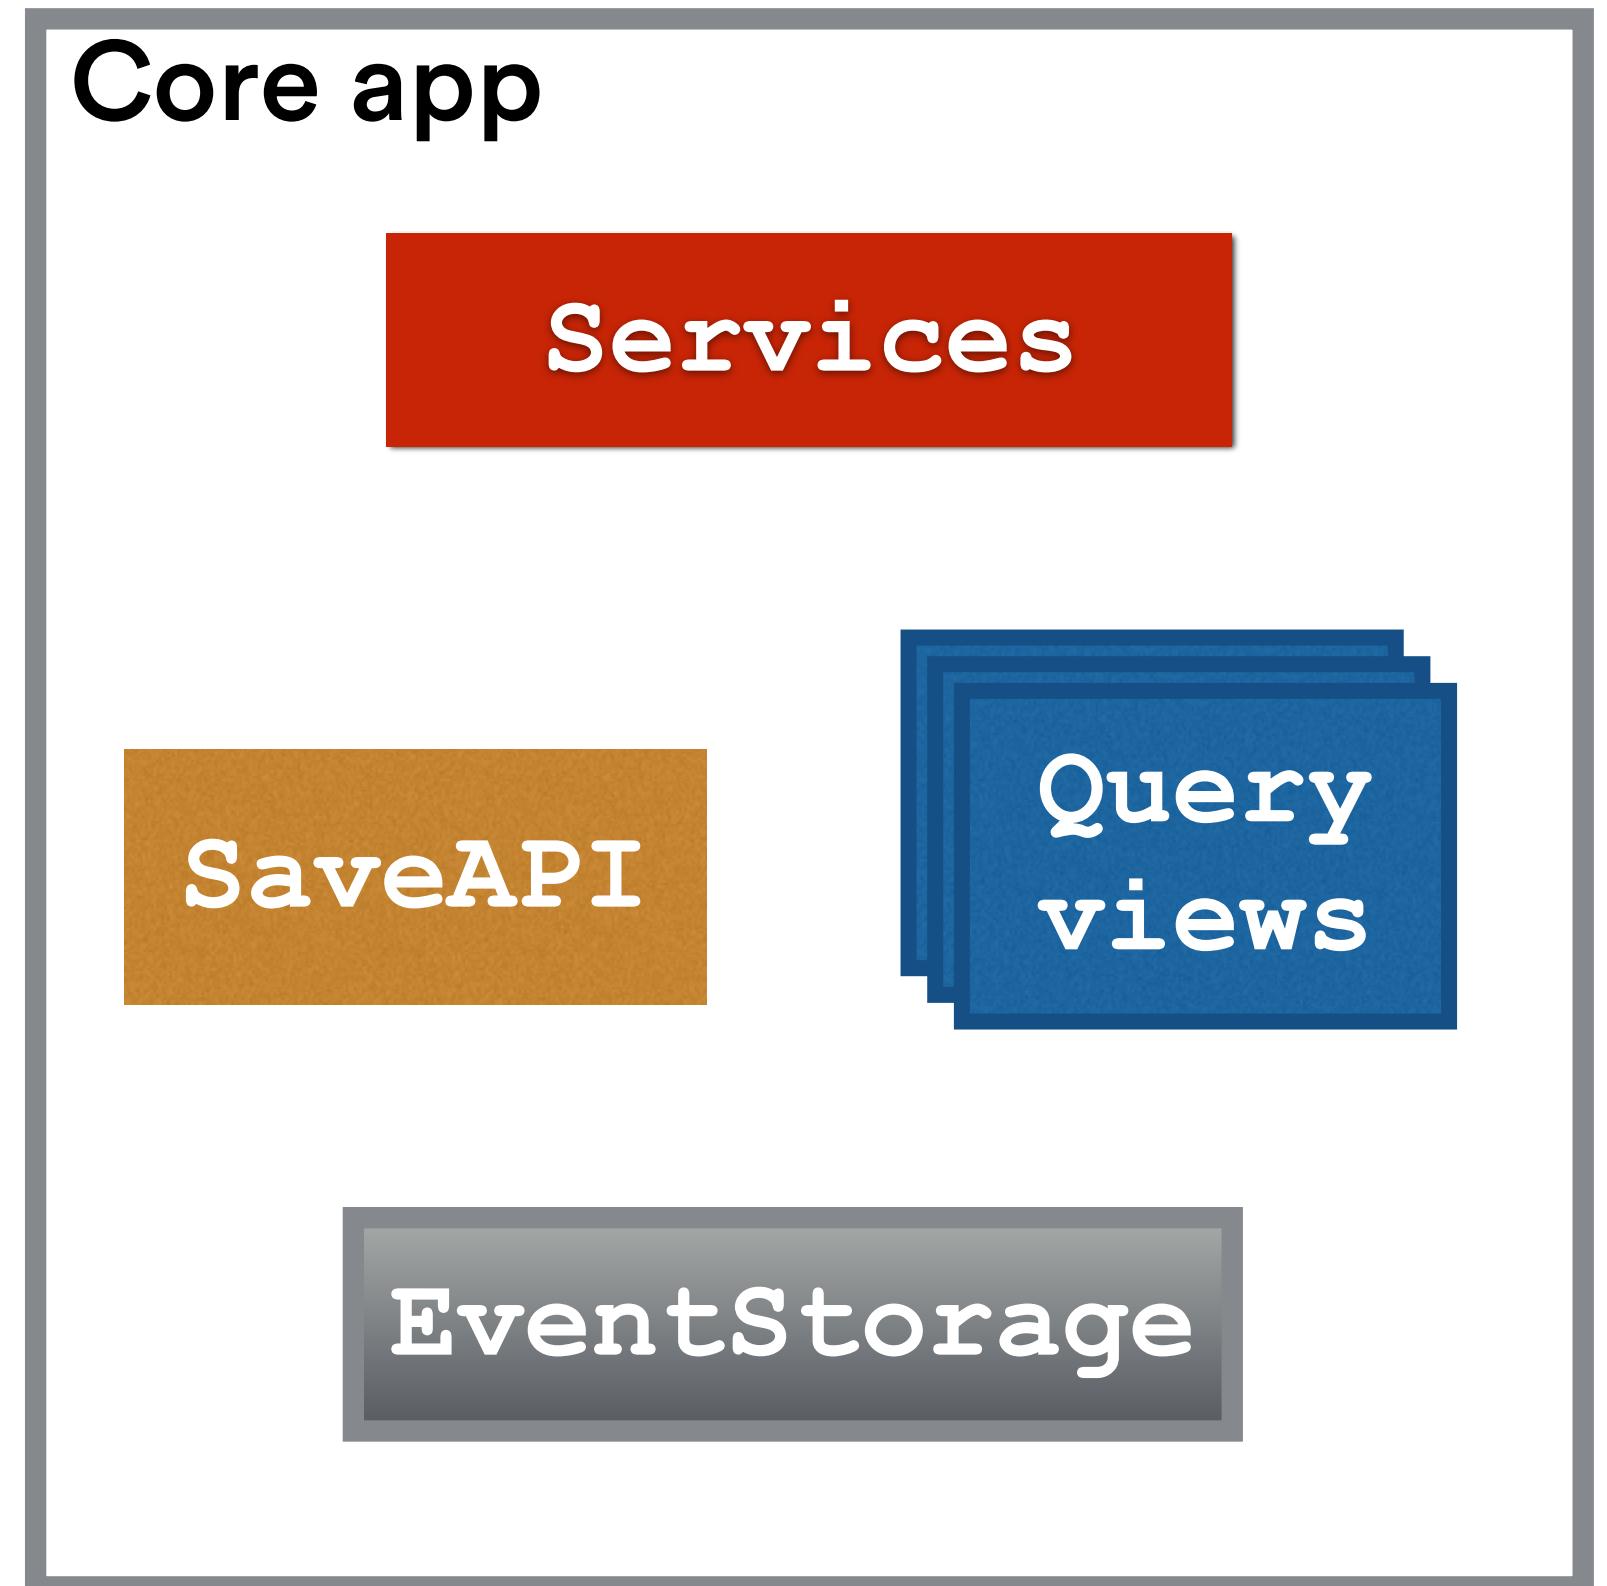




Evolving the  
architecture

# REST calls

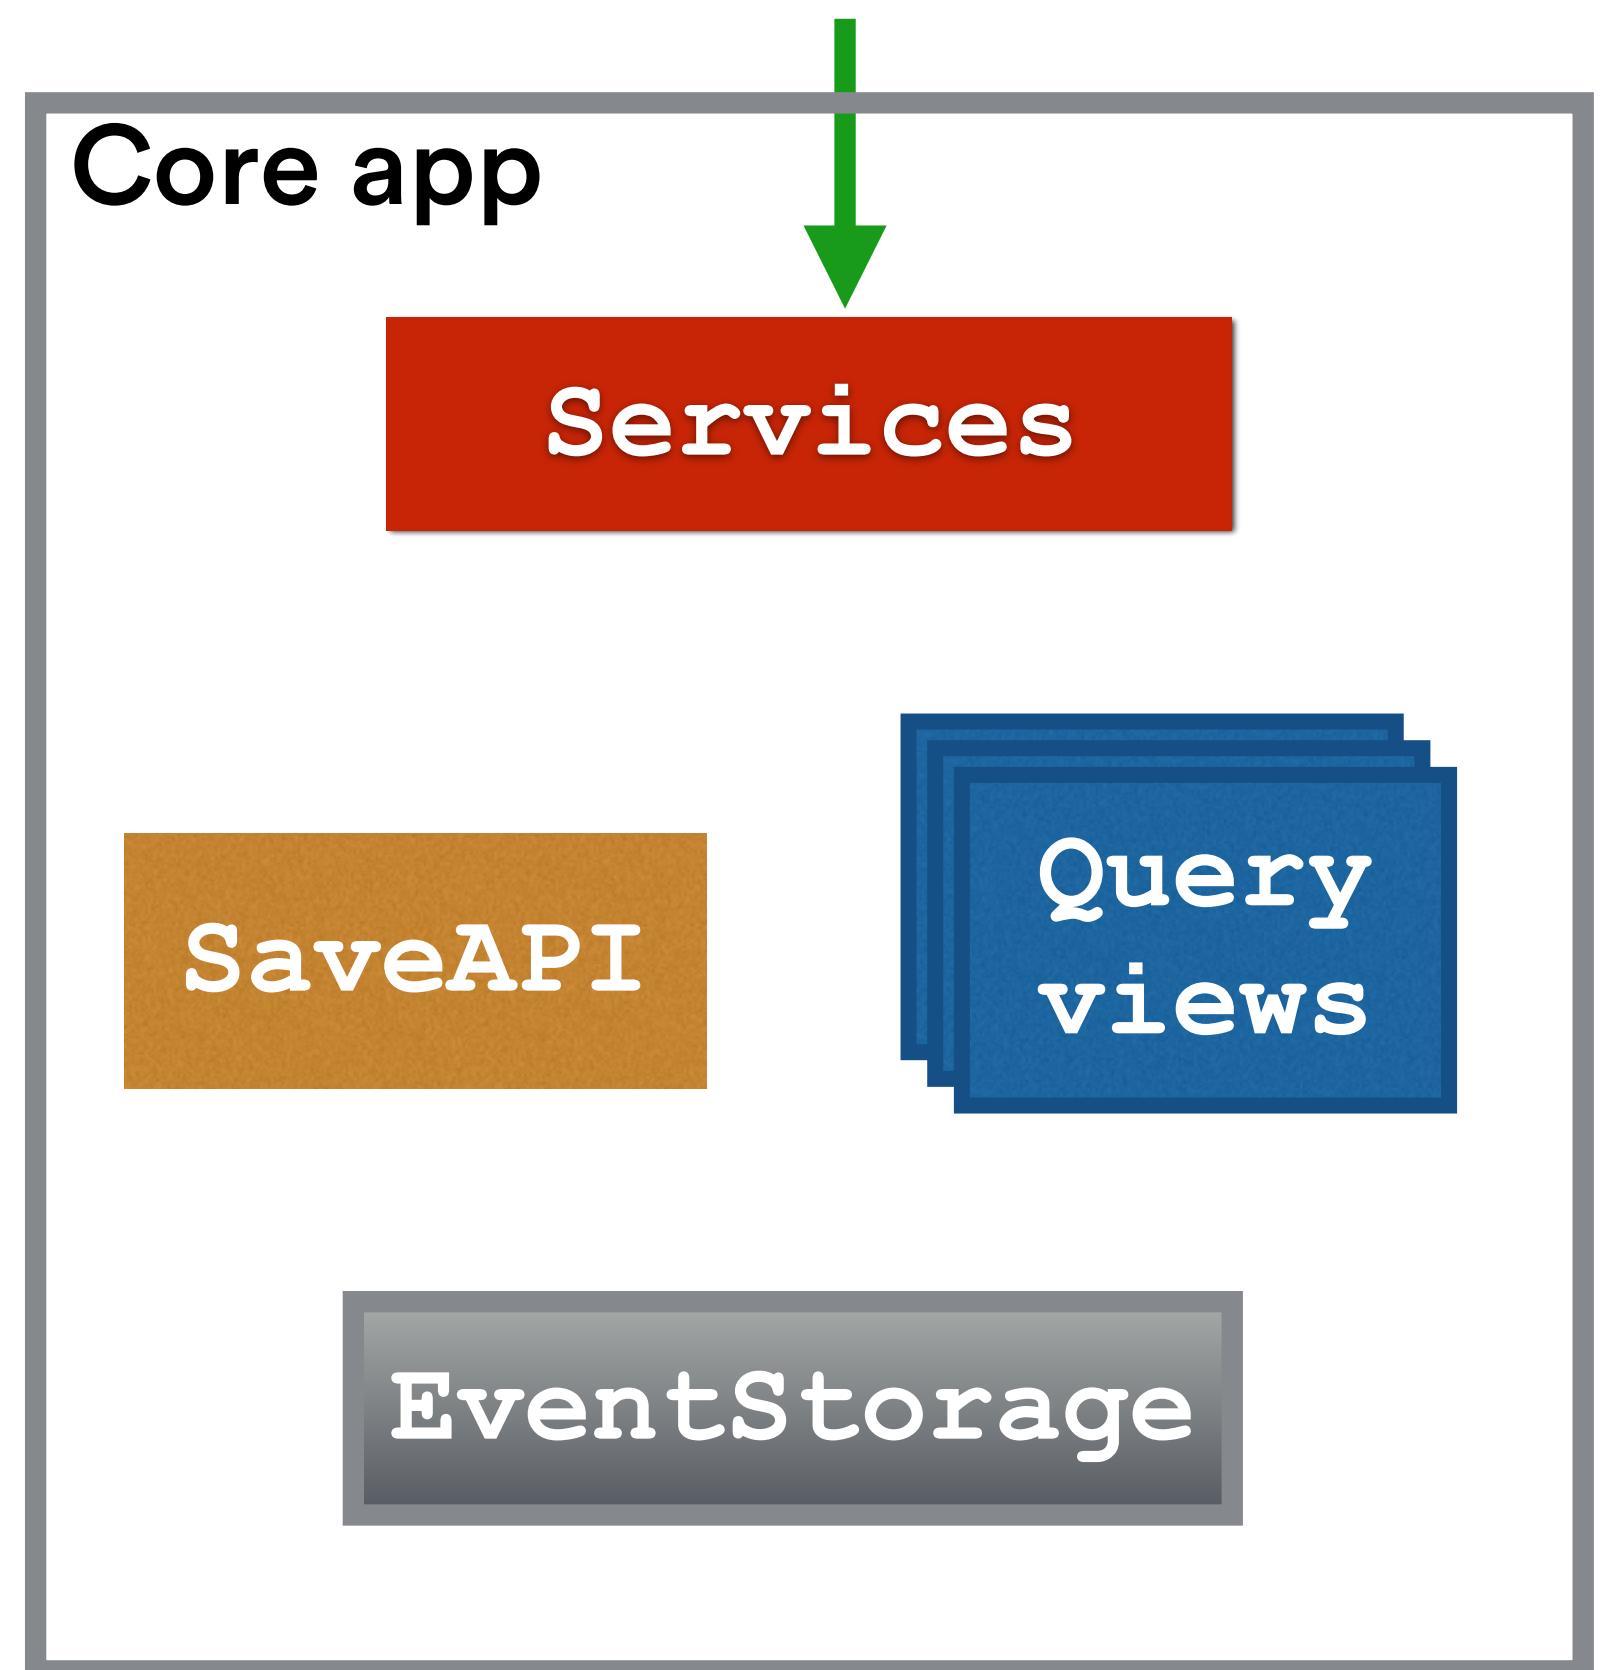
e.g. Add User



DynamoDB



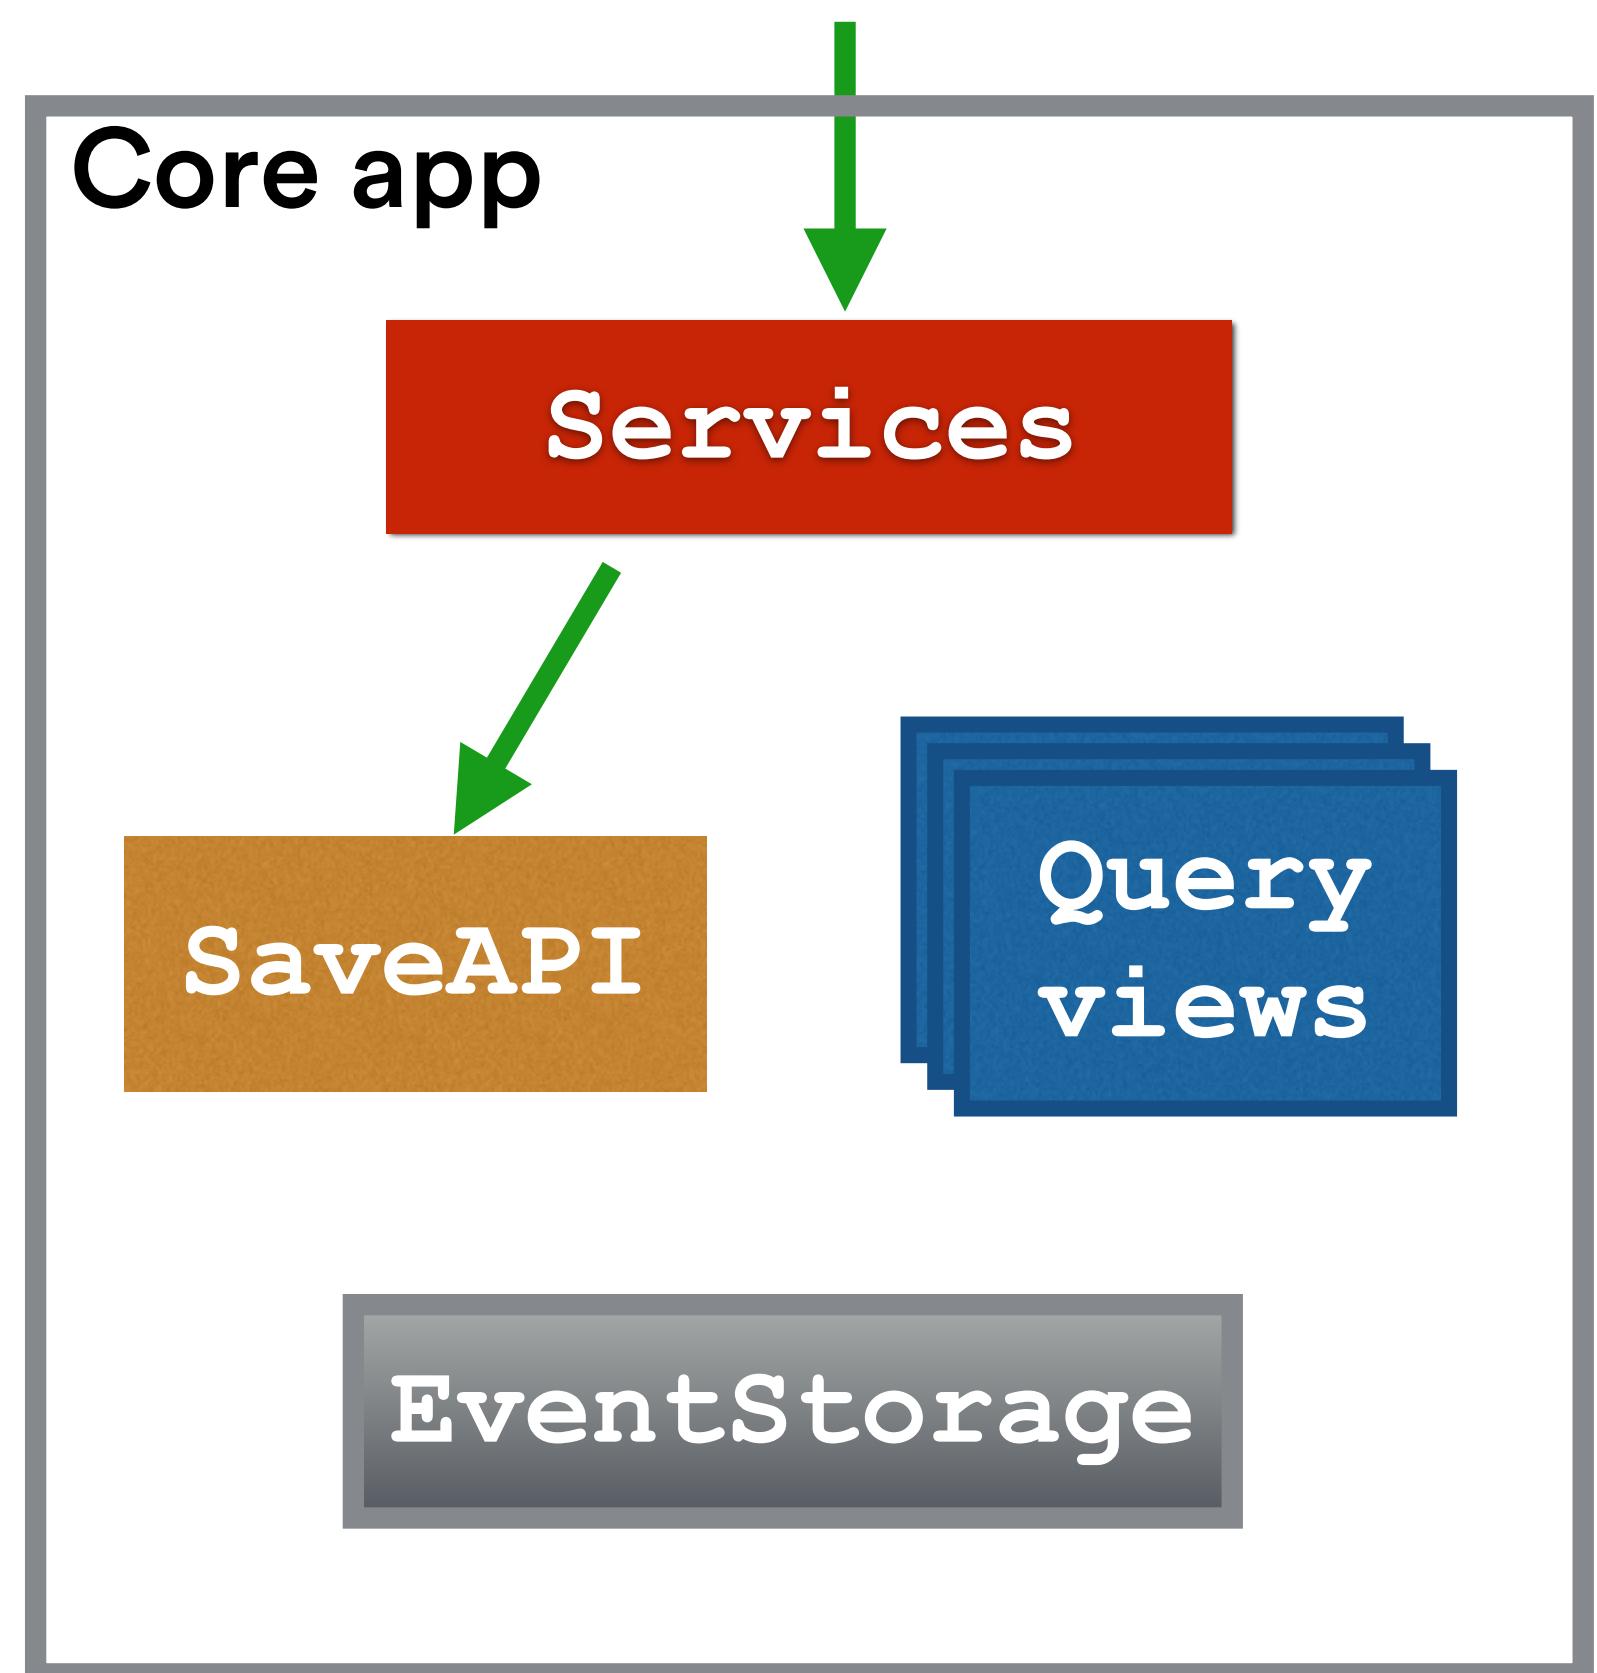
**REST calls**  
e.g. Add User



DynamoDB



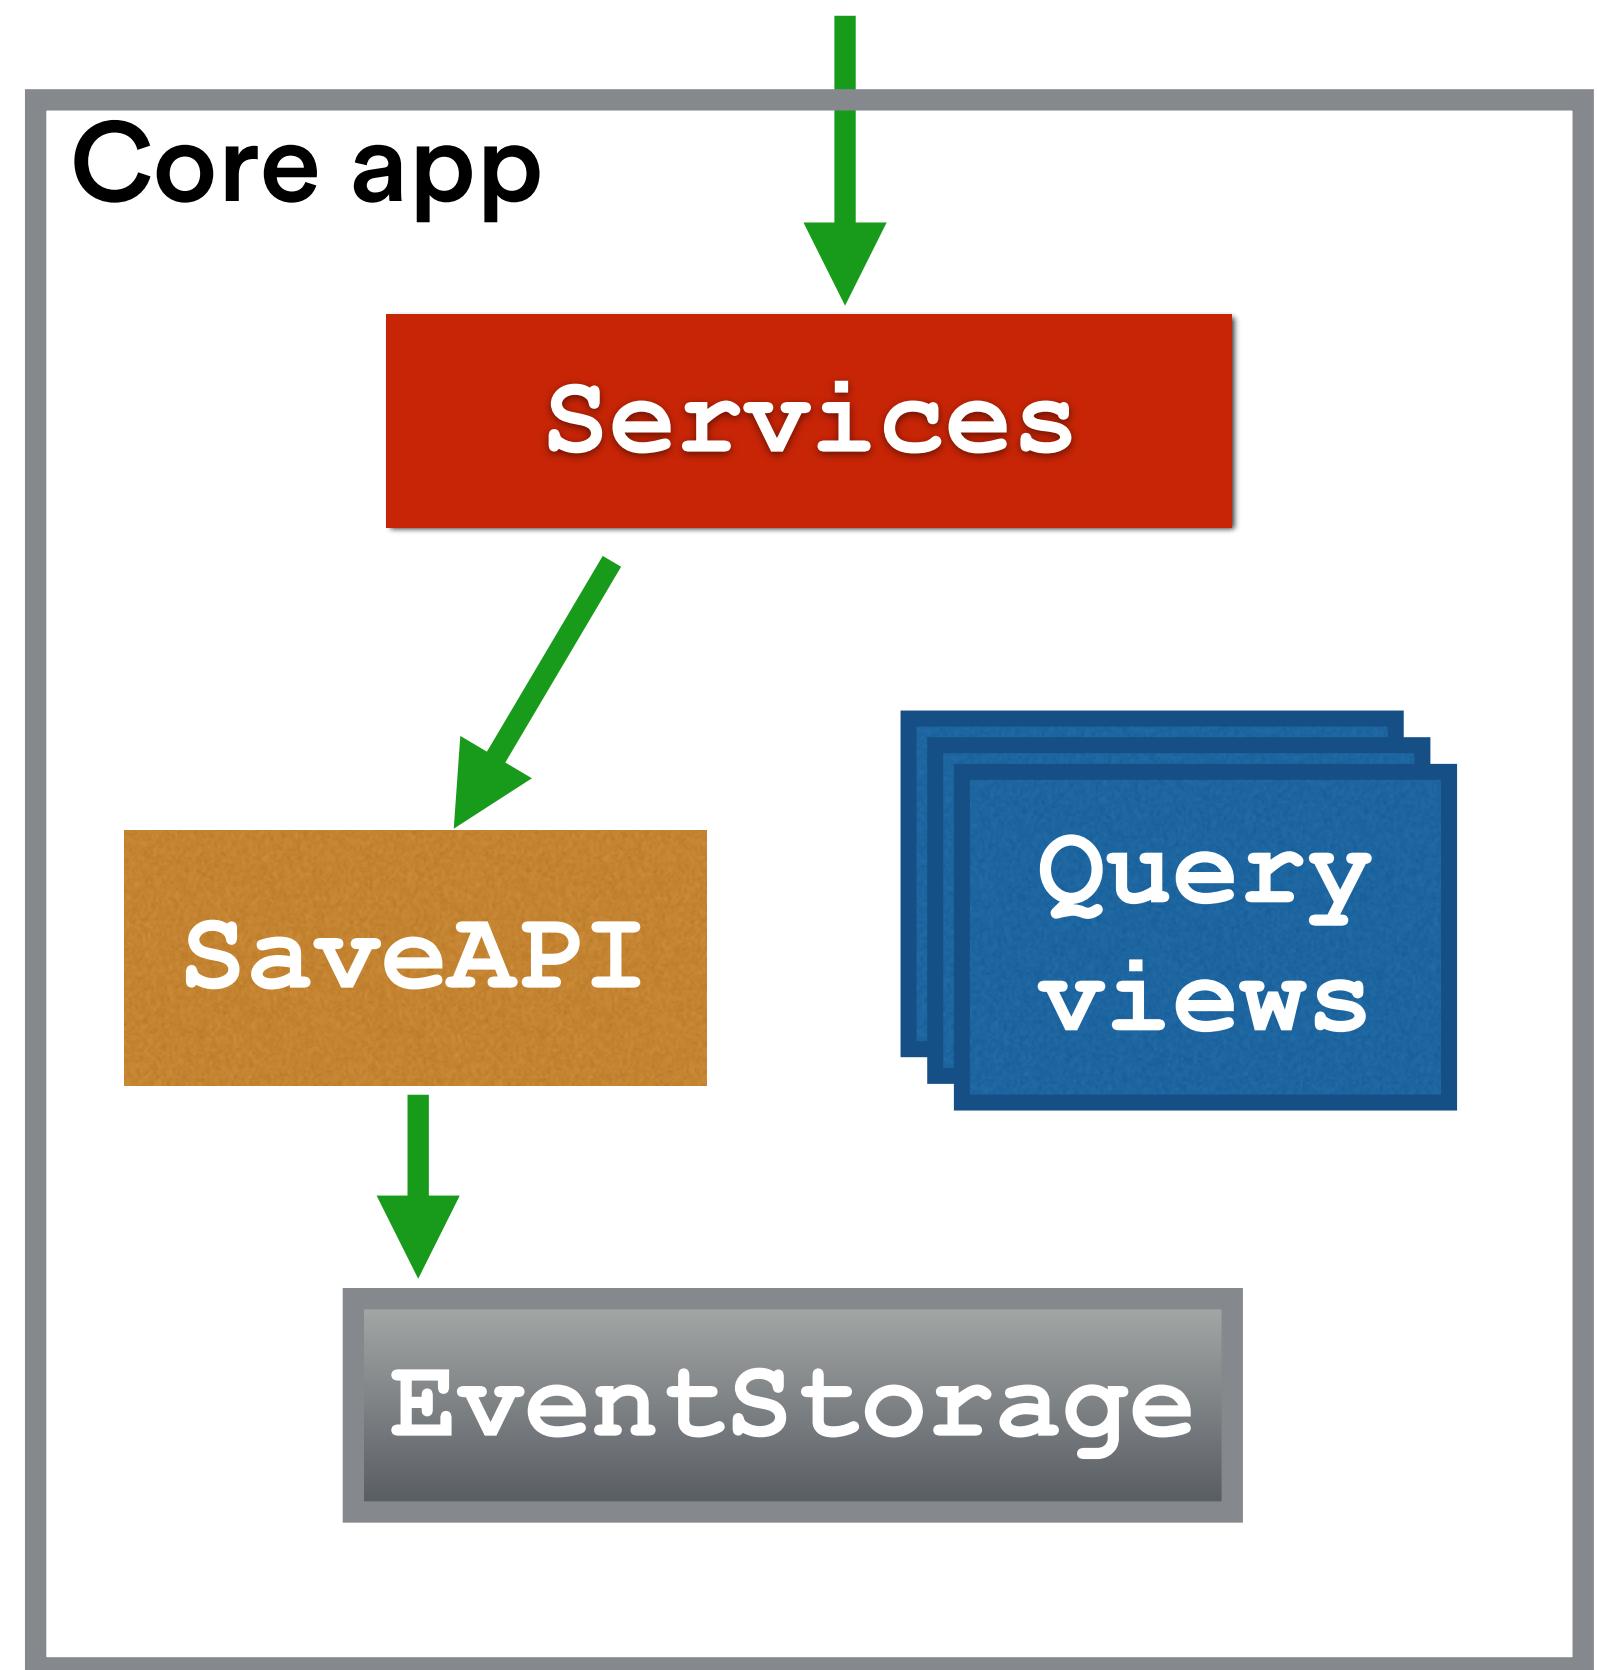
**REST calls**  
e.g. Add User



DynamoDB



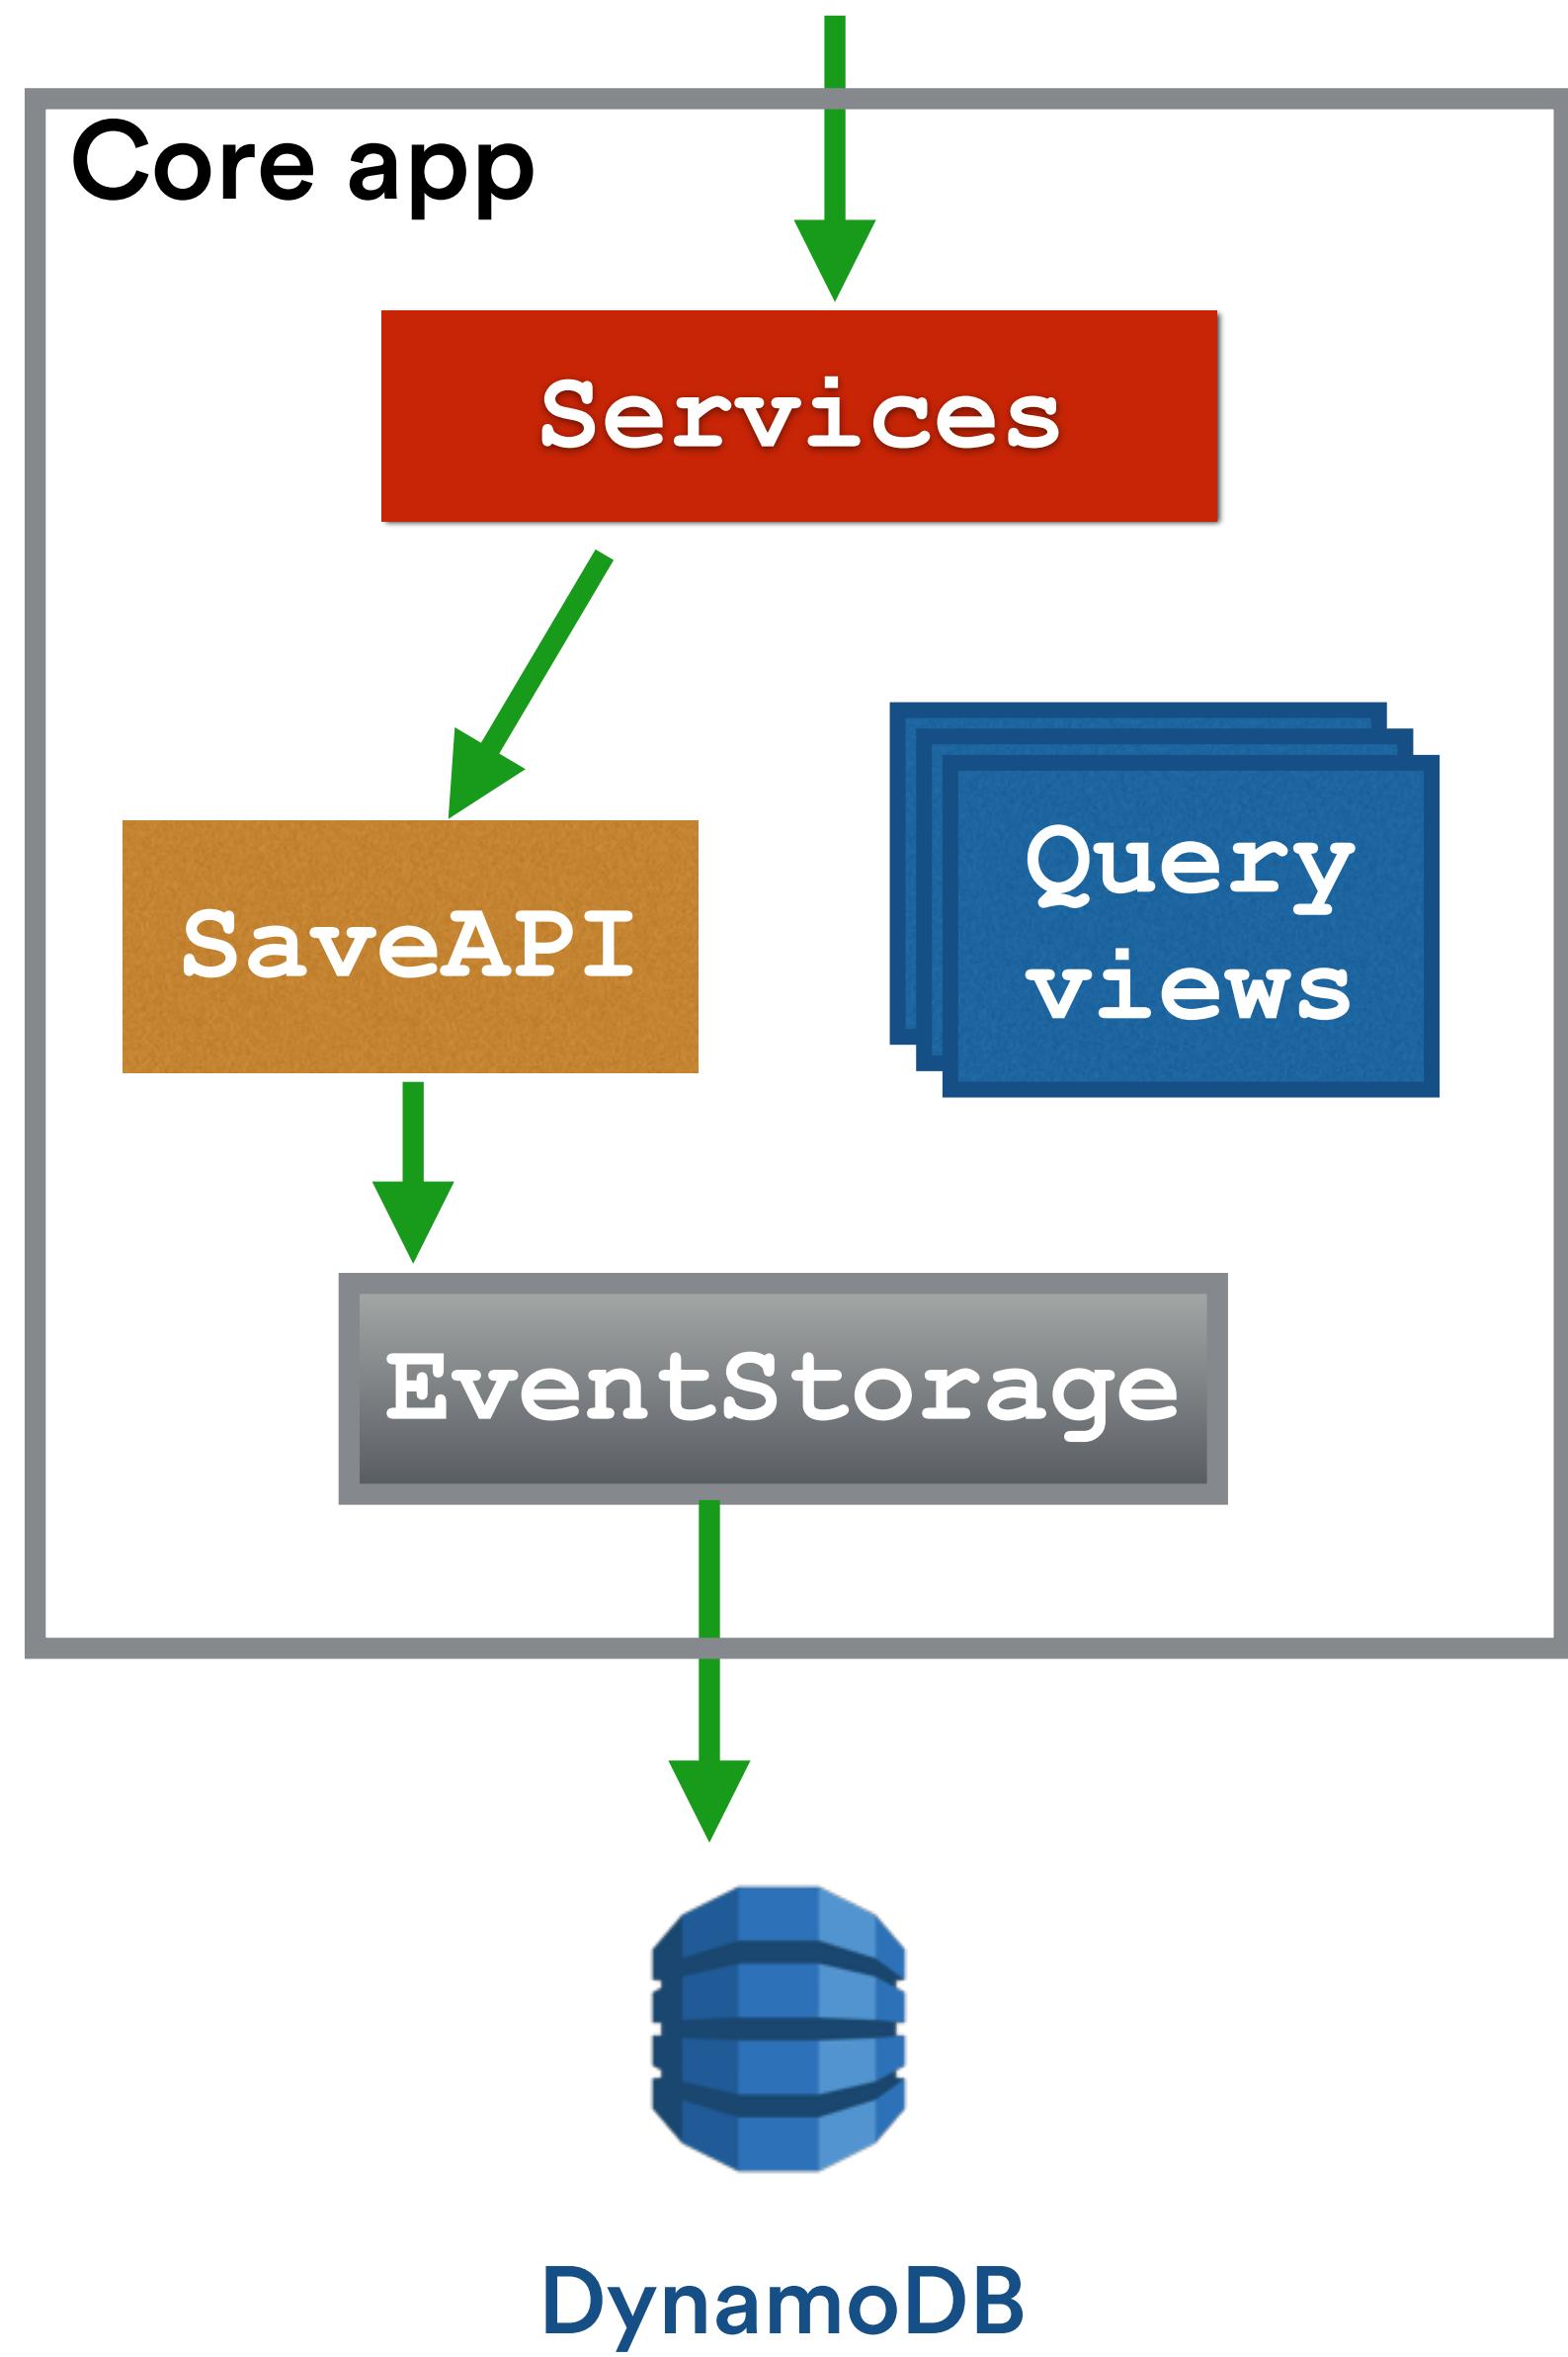
**REST calls**  
e.g. Add User



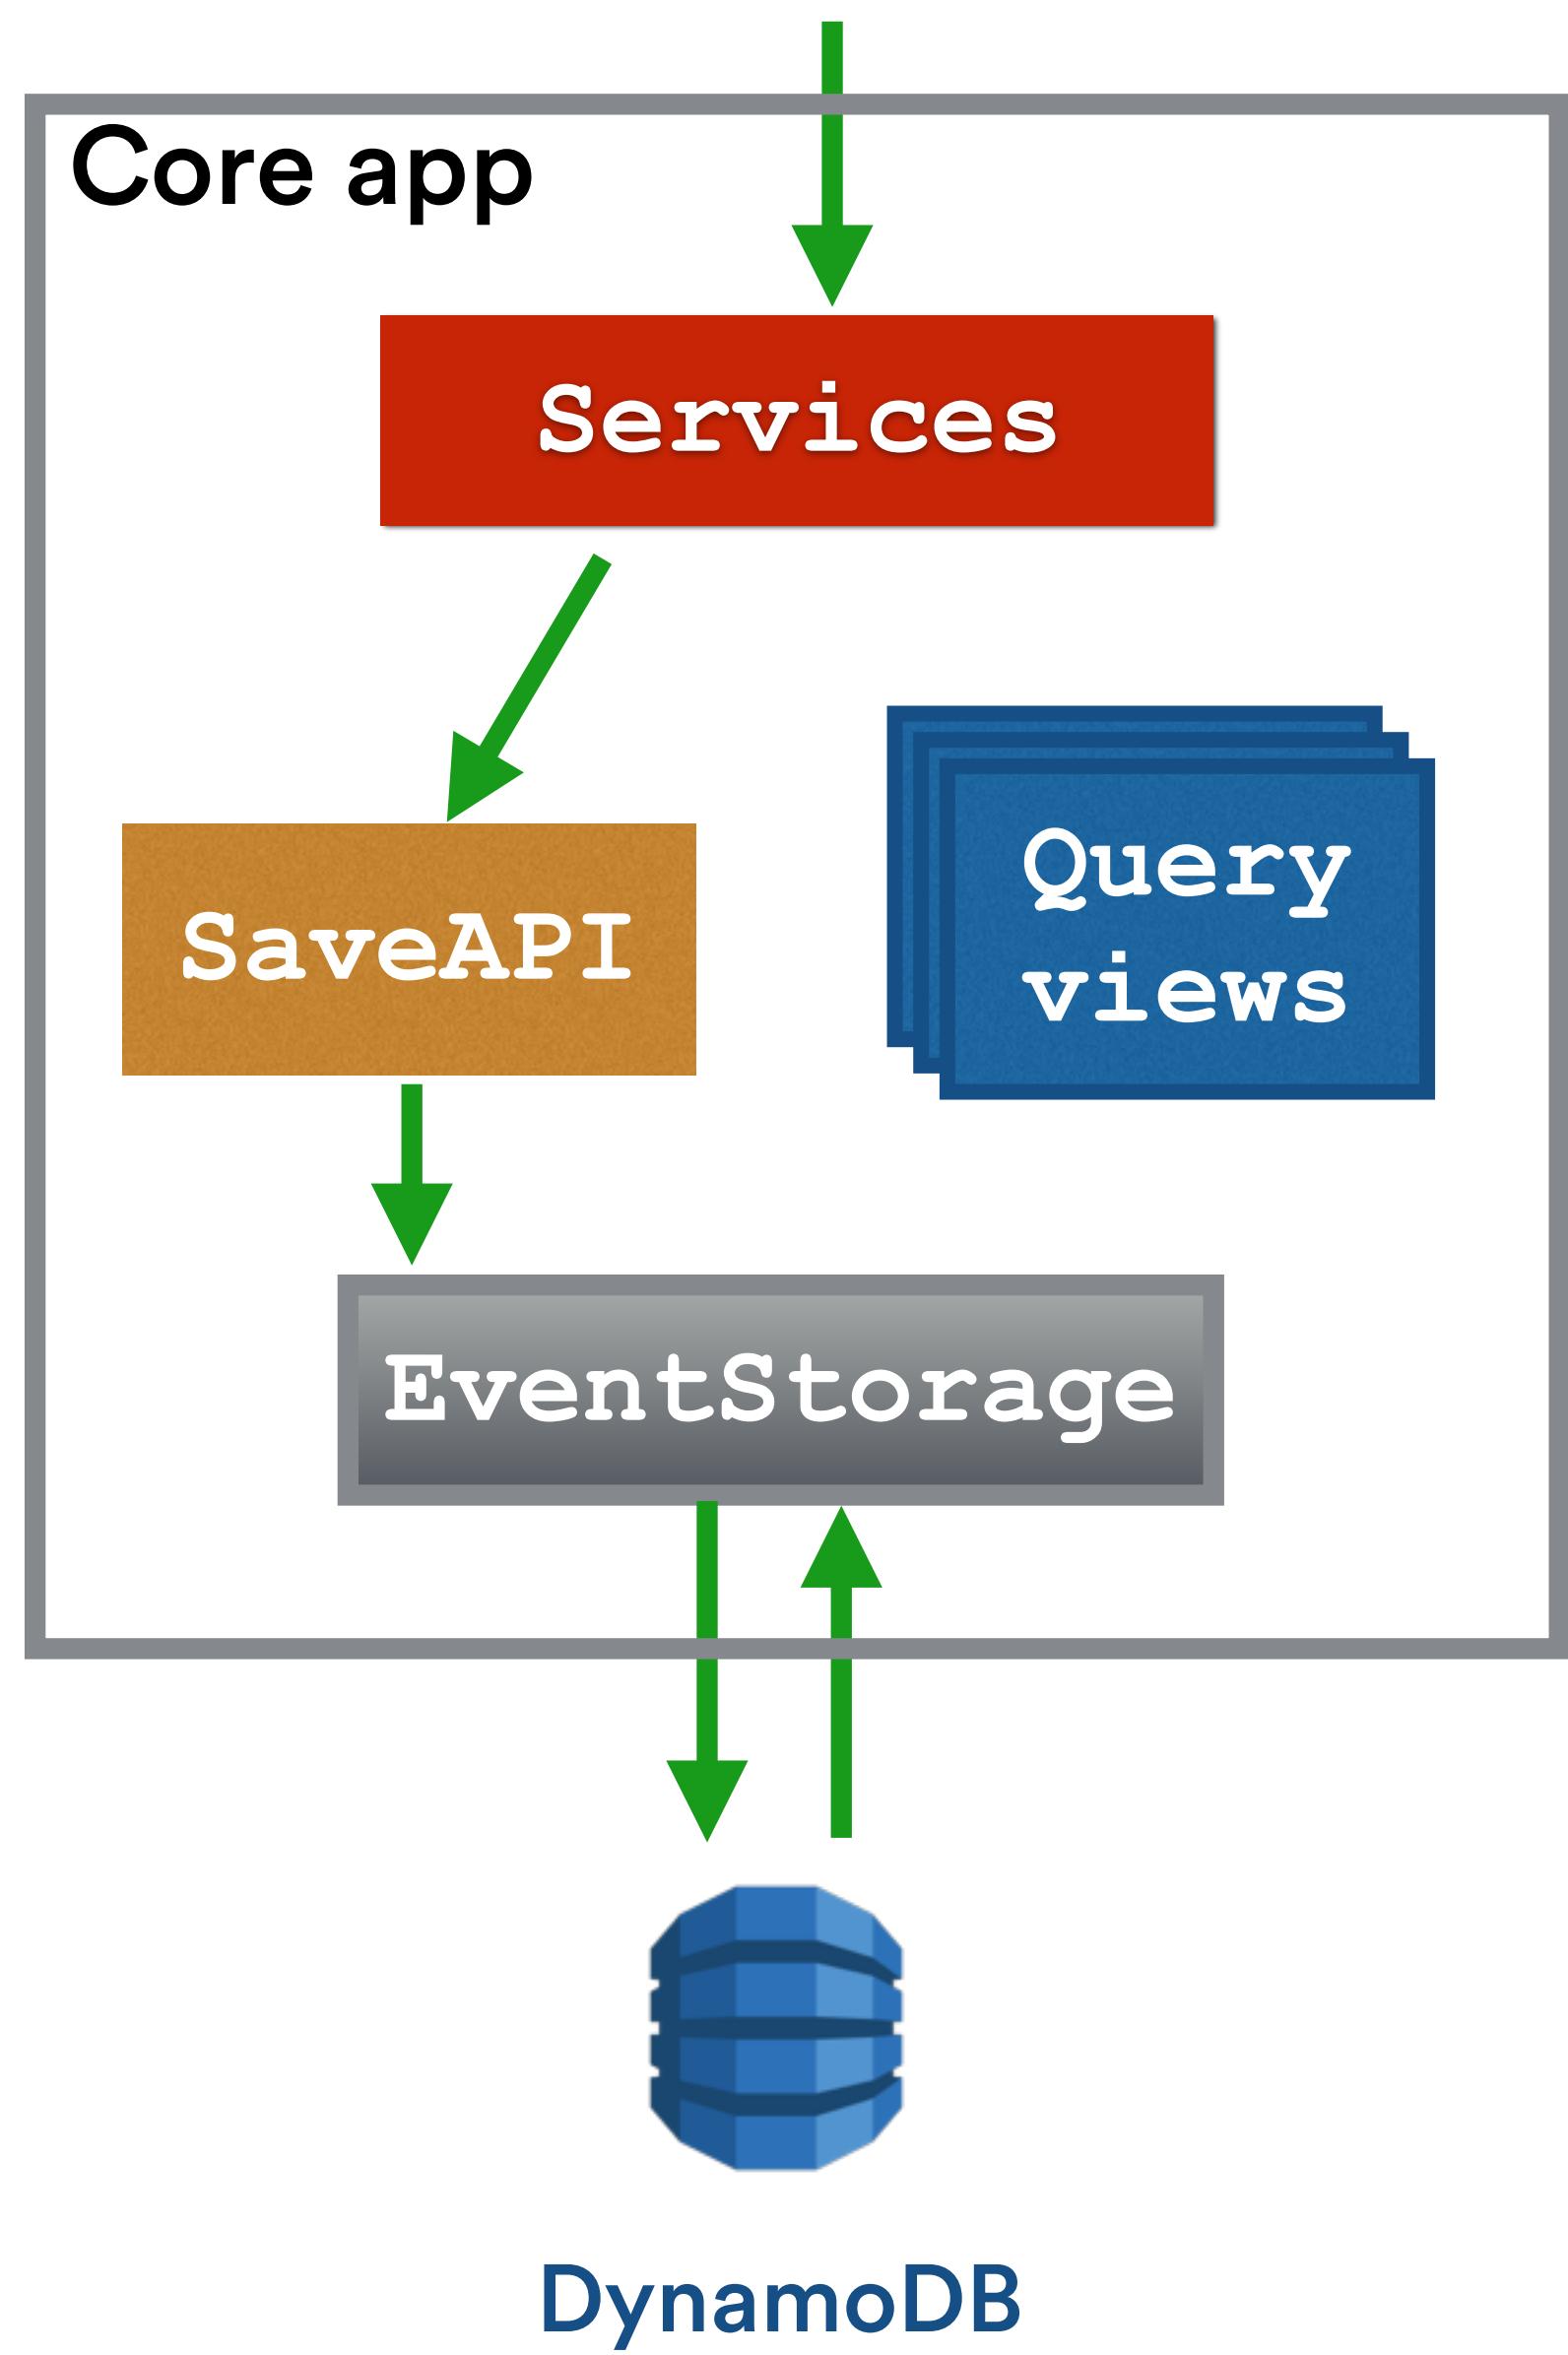
DynamoDB



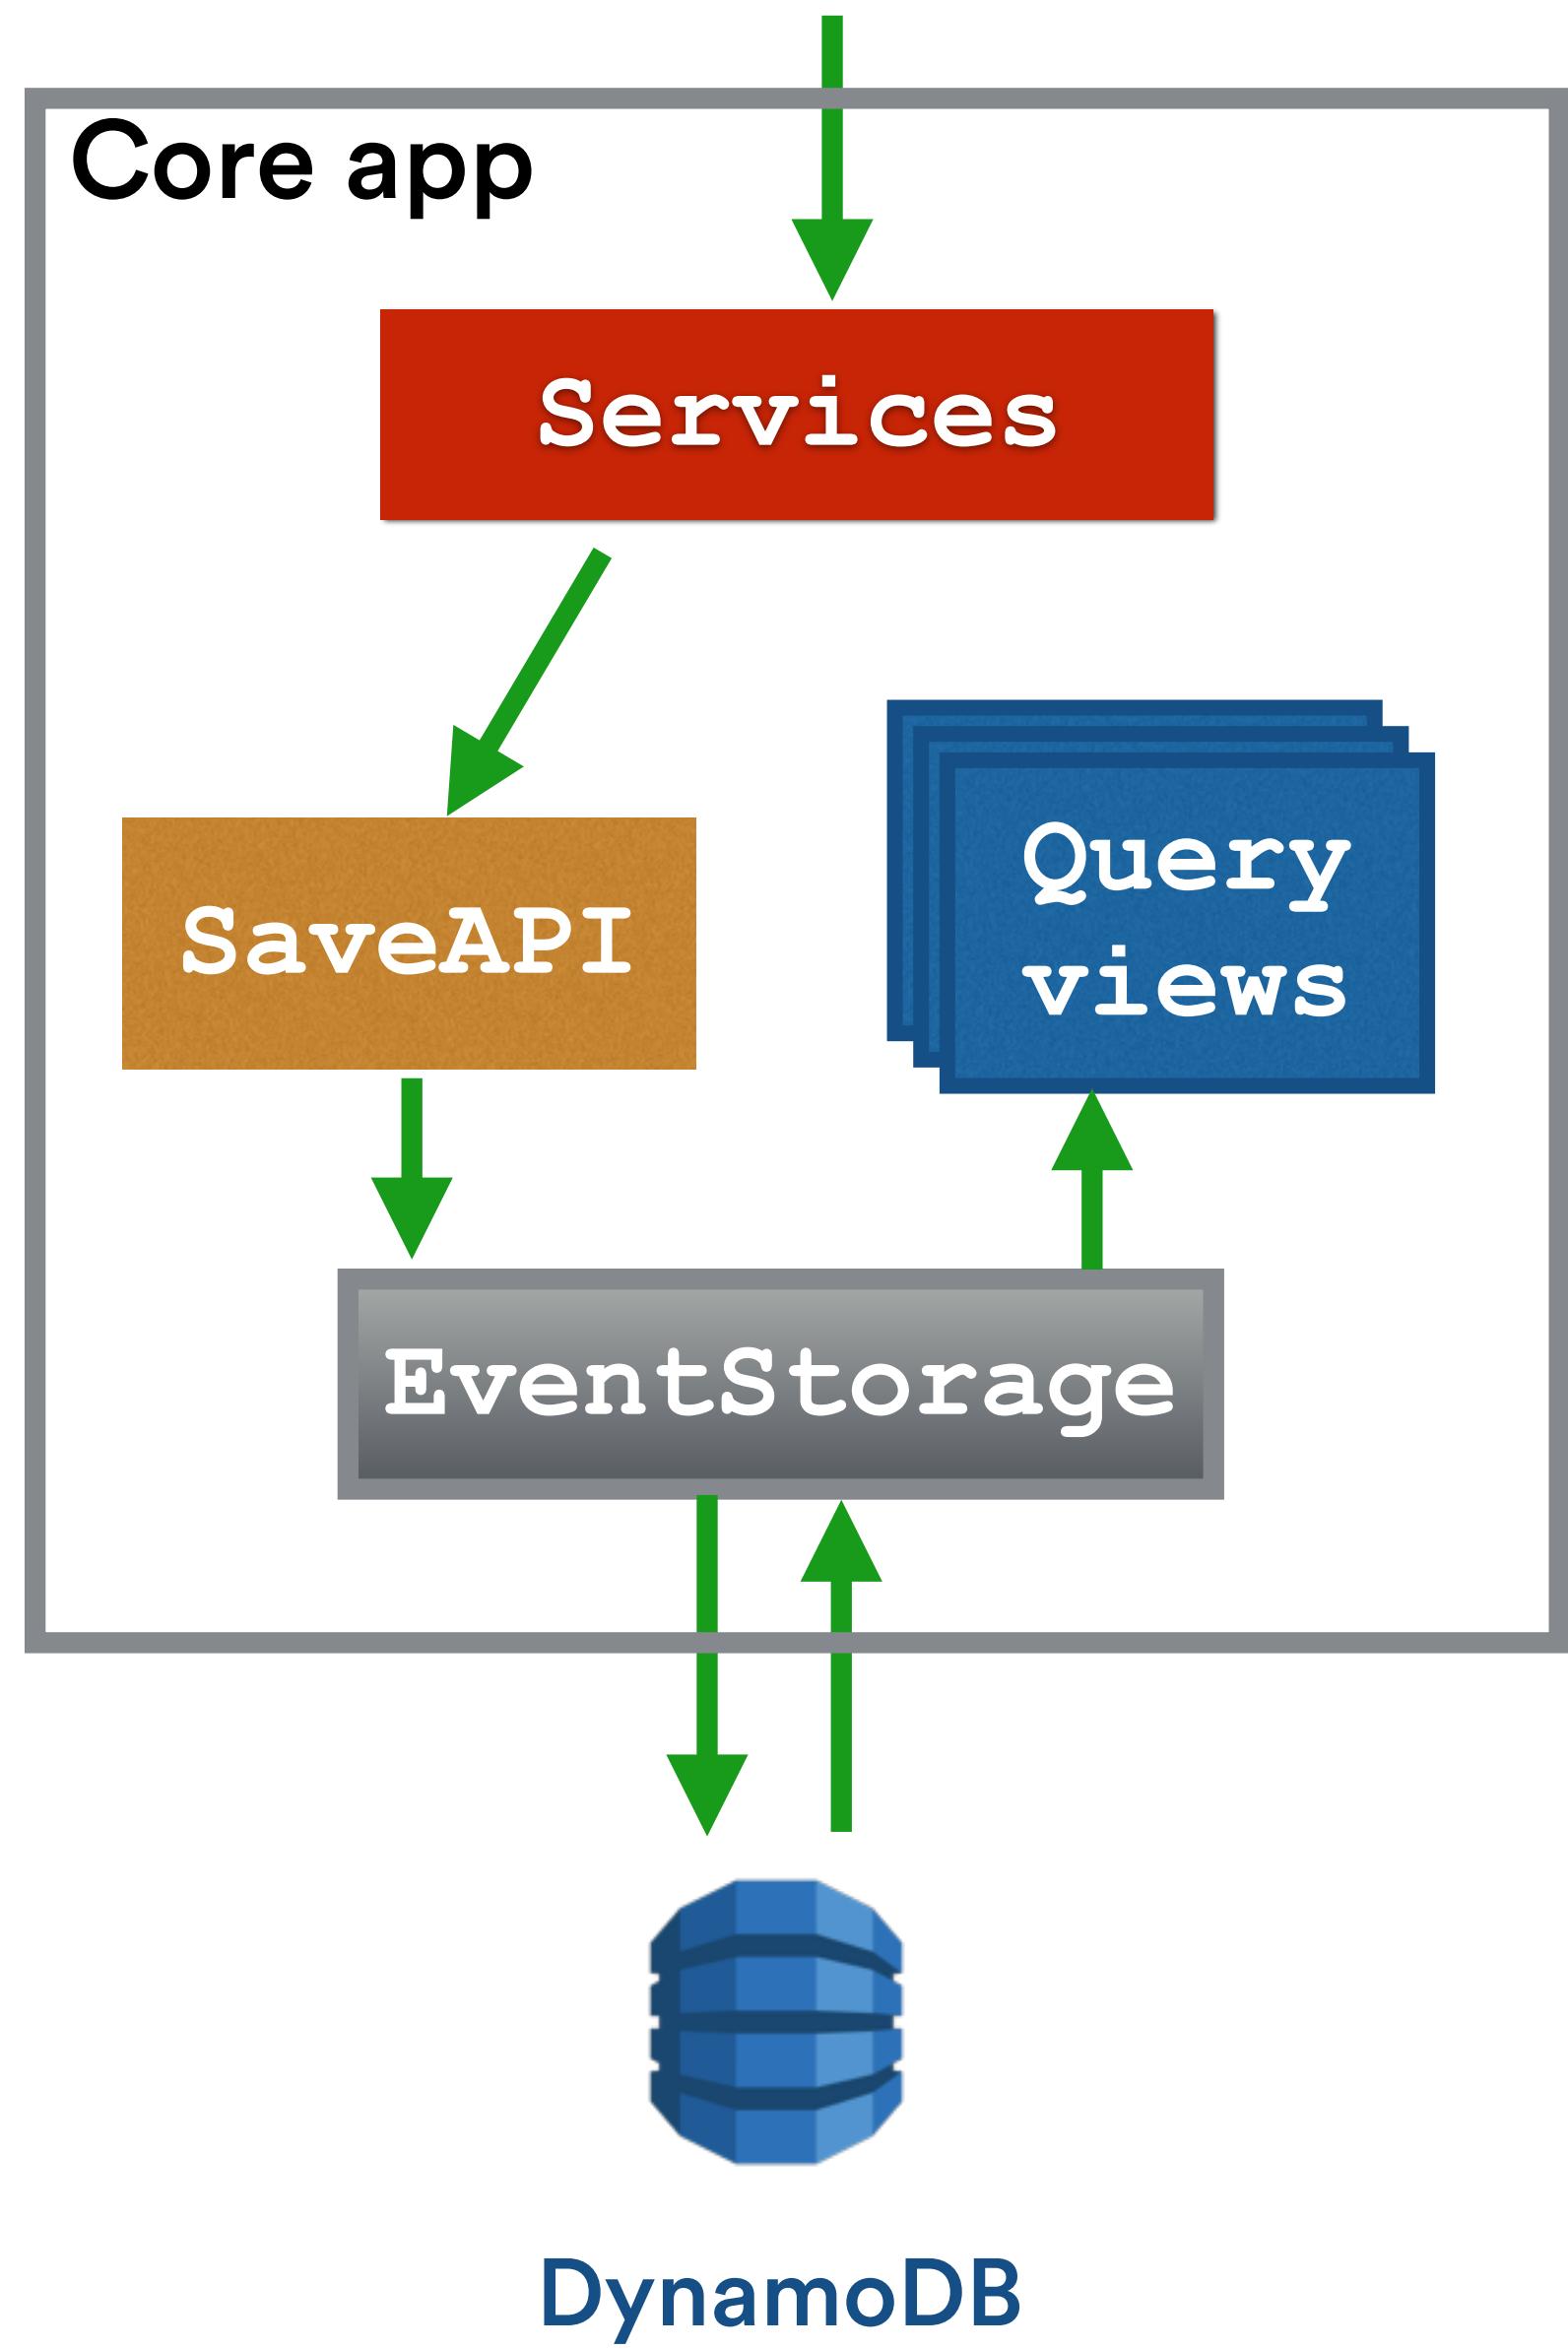
**REST calls**  
e.g. Add User



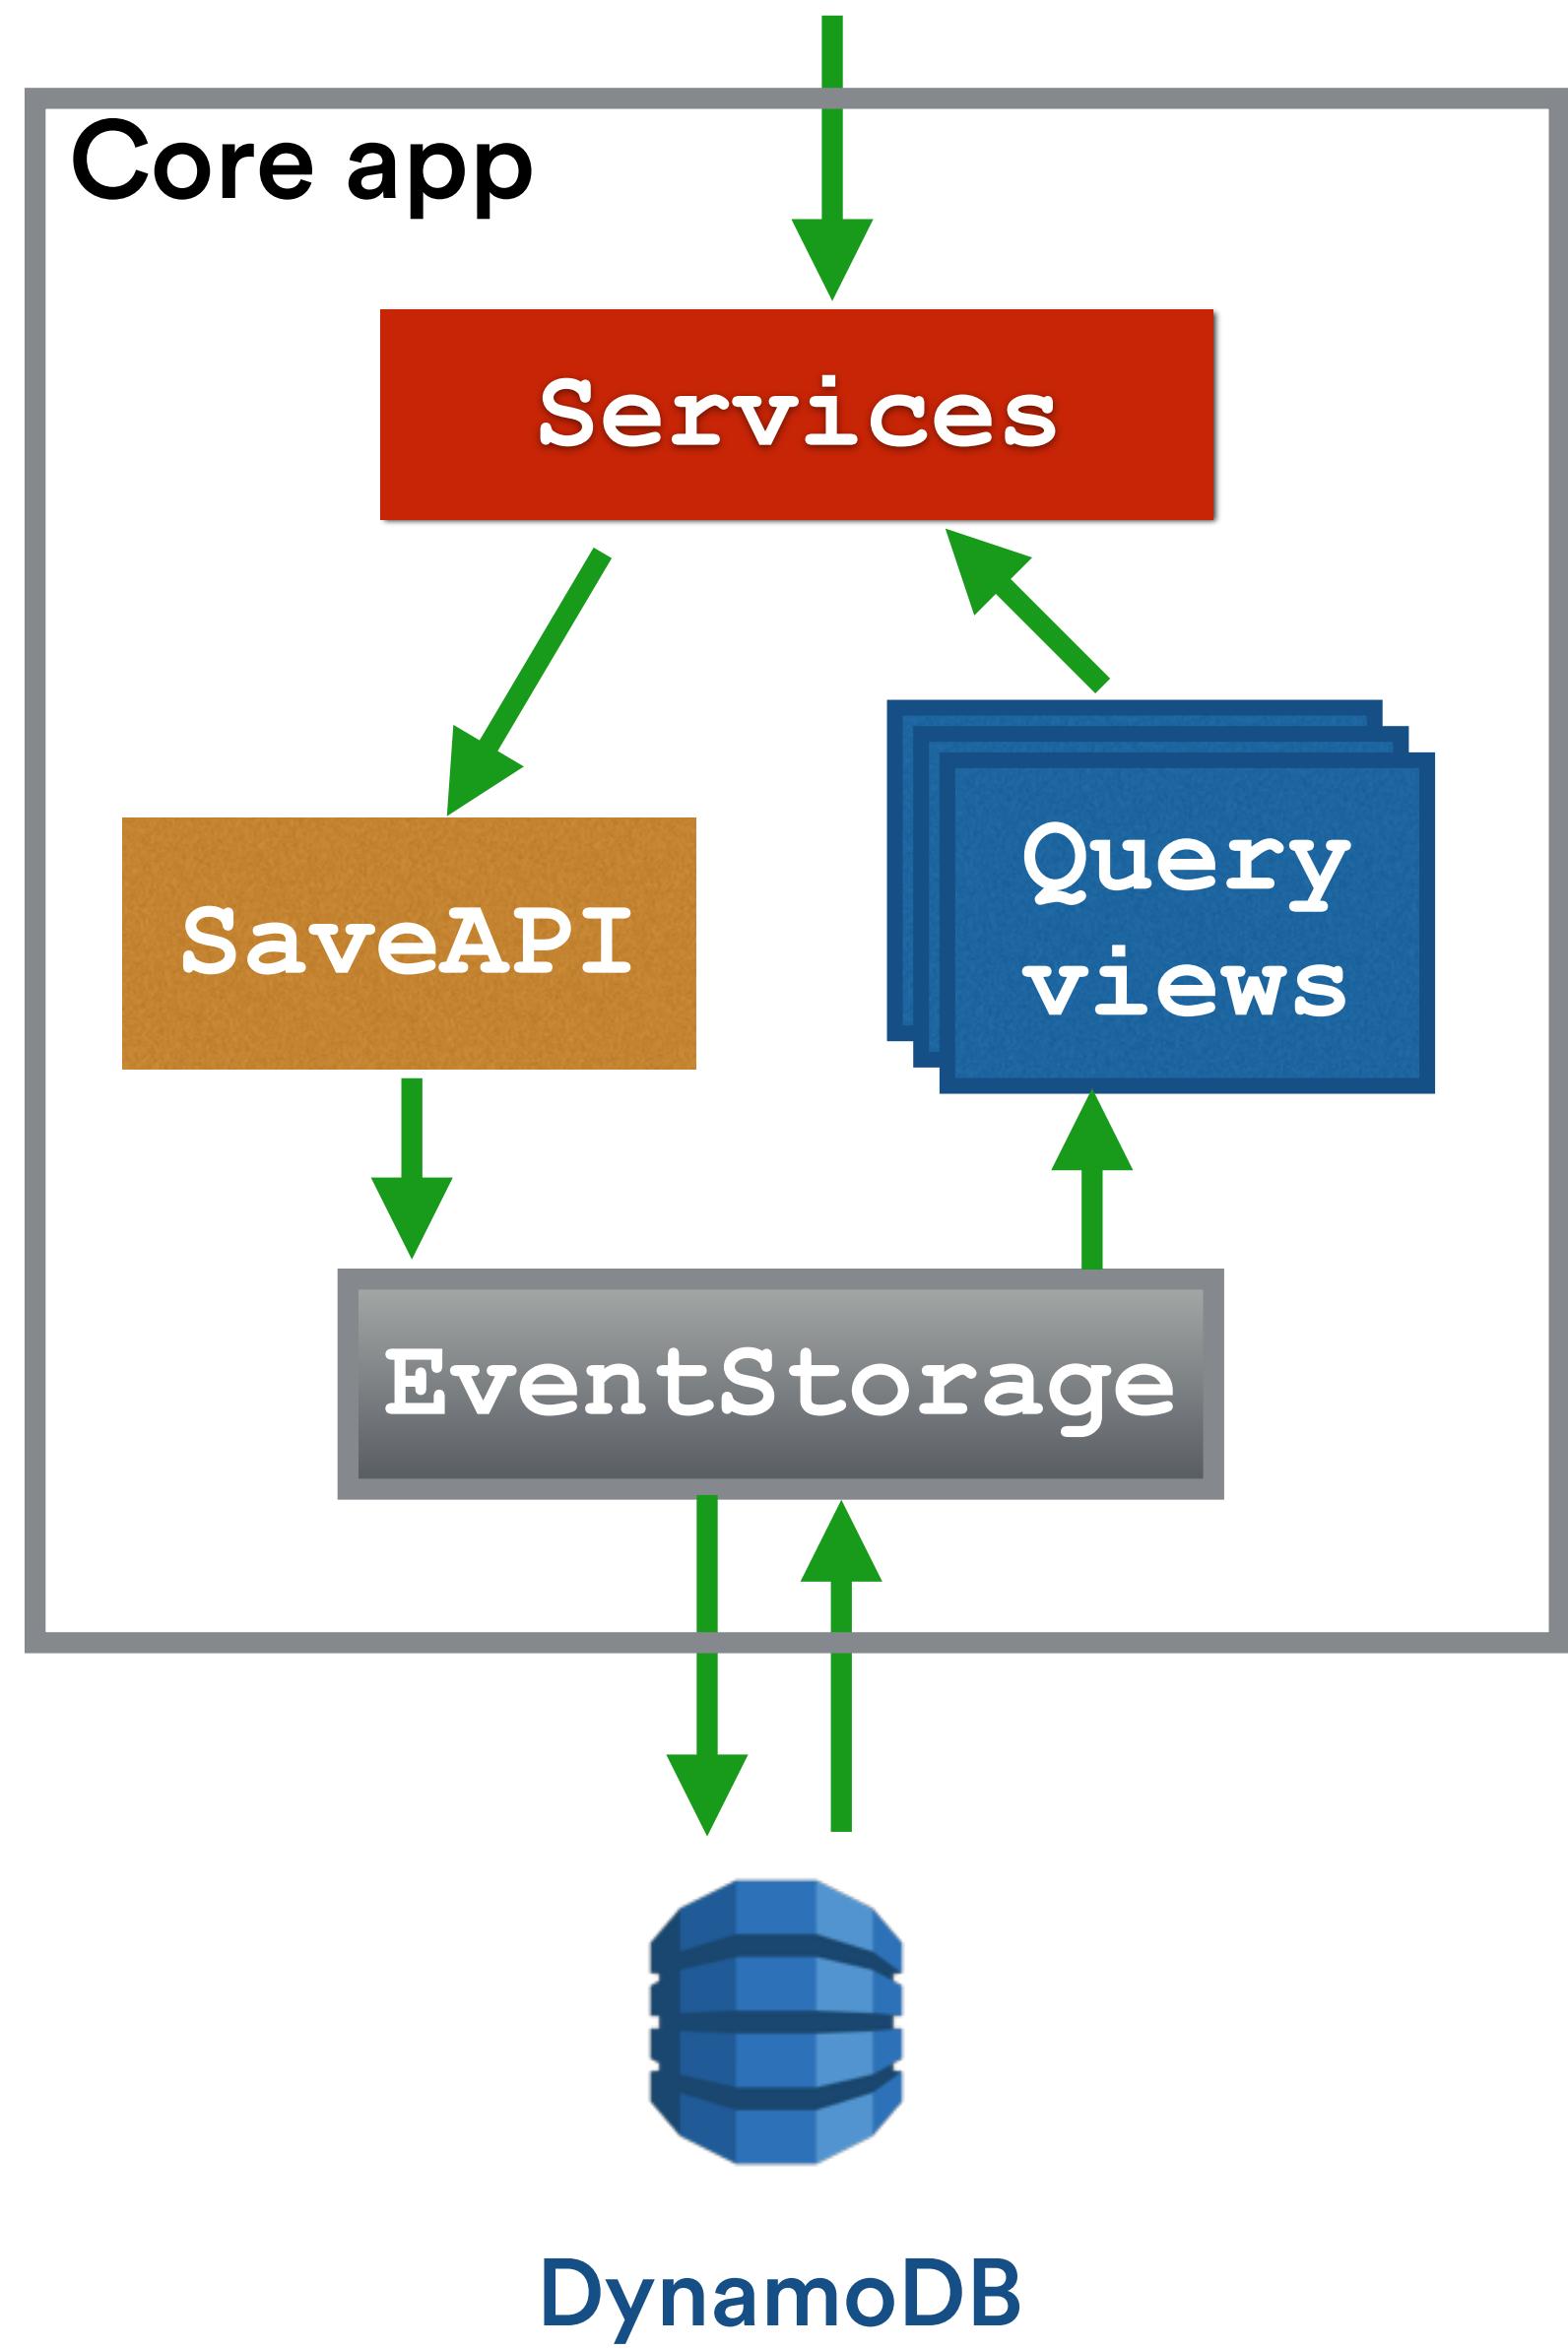
**REST calls**  
e.g. Add User



**REST calls**  
e.g. Add User

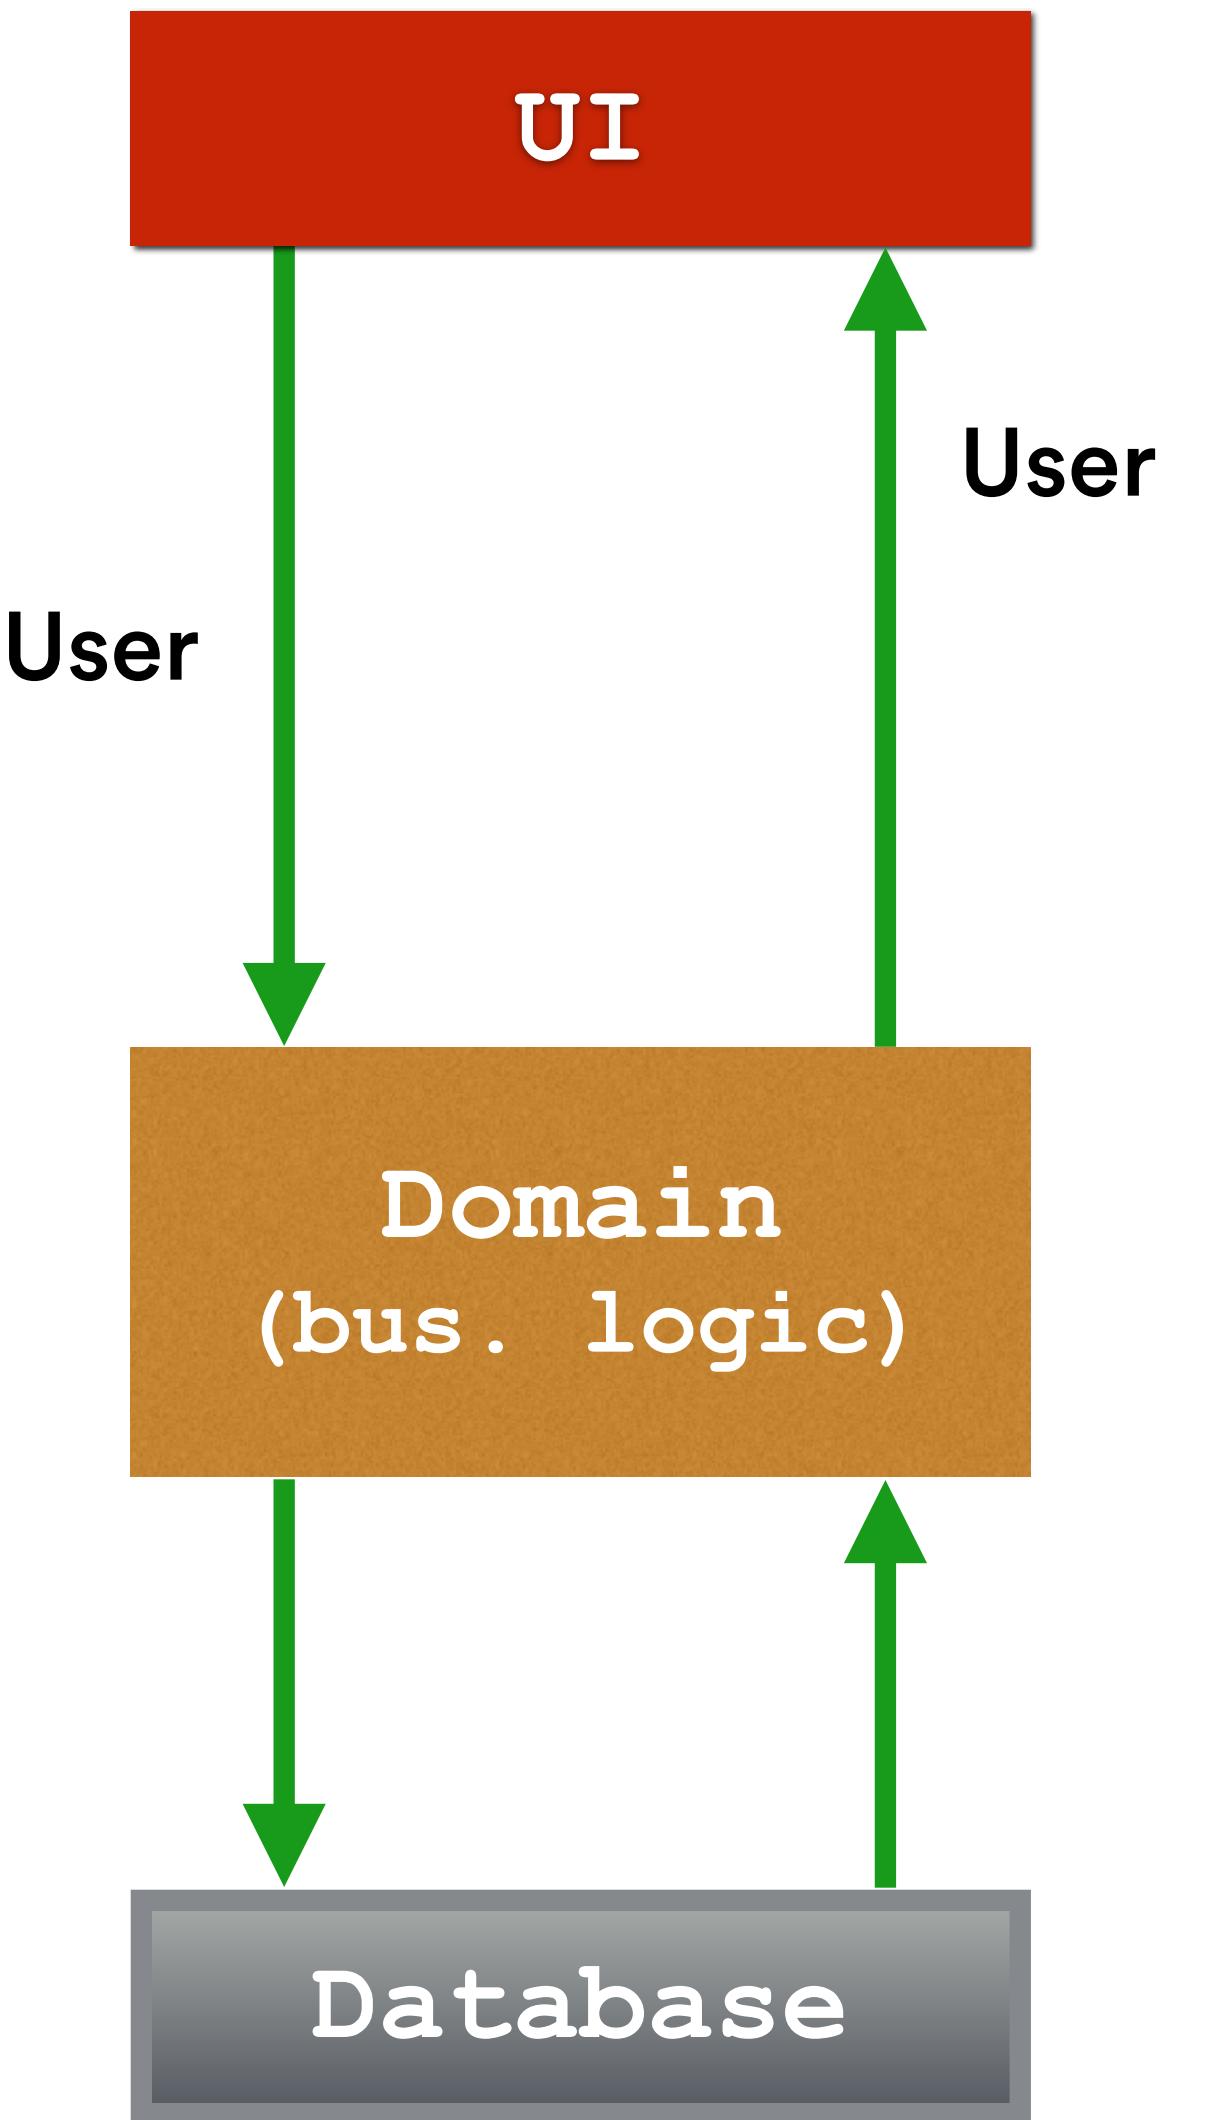


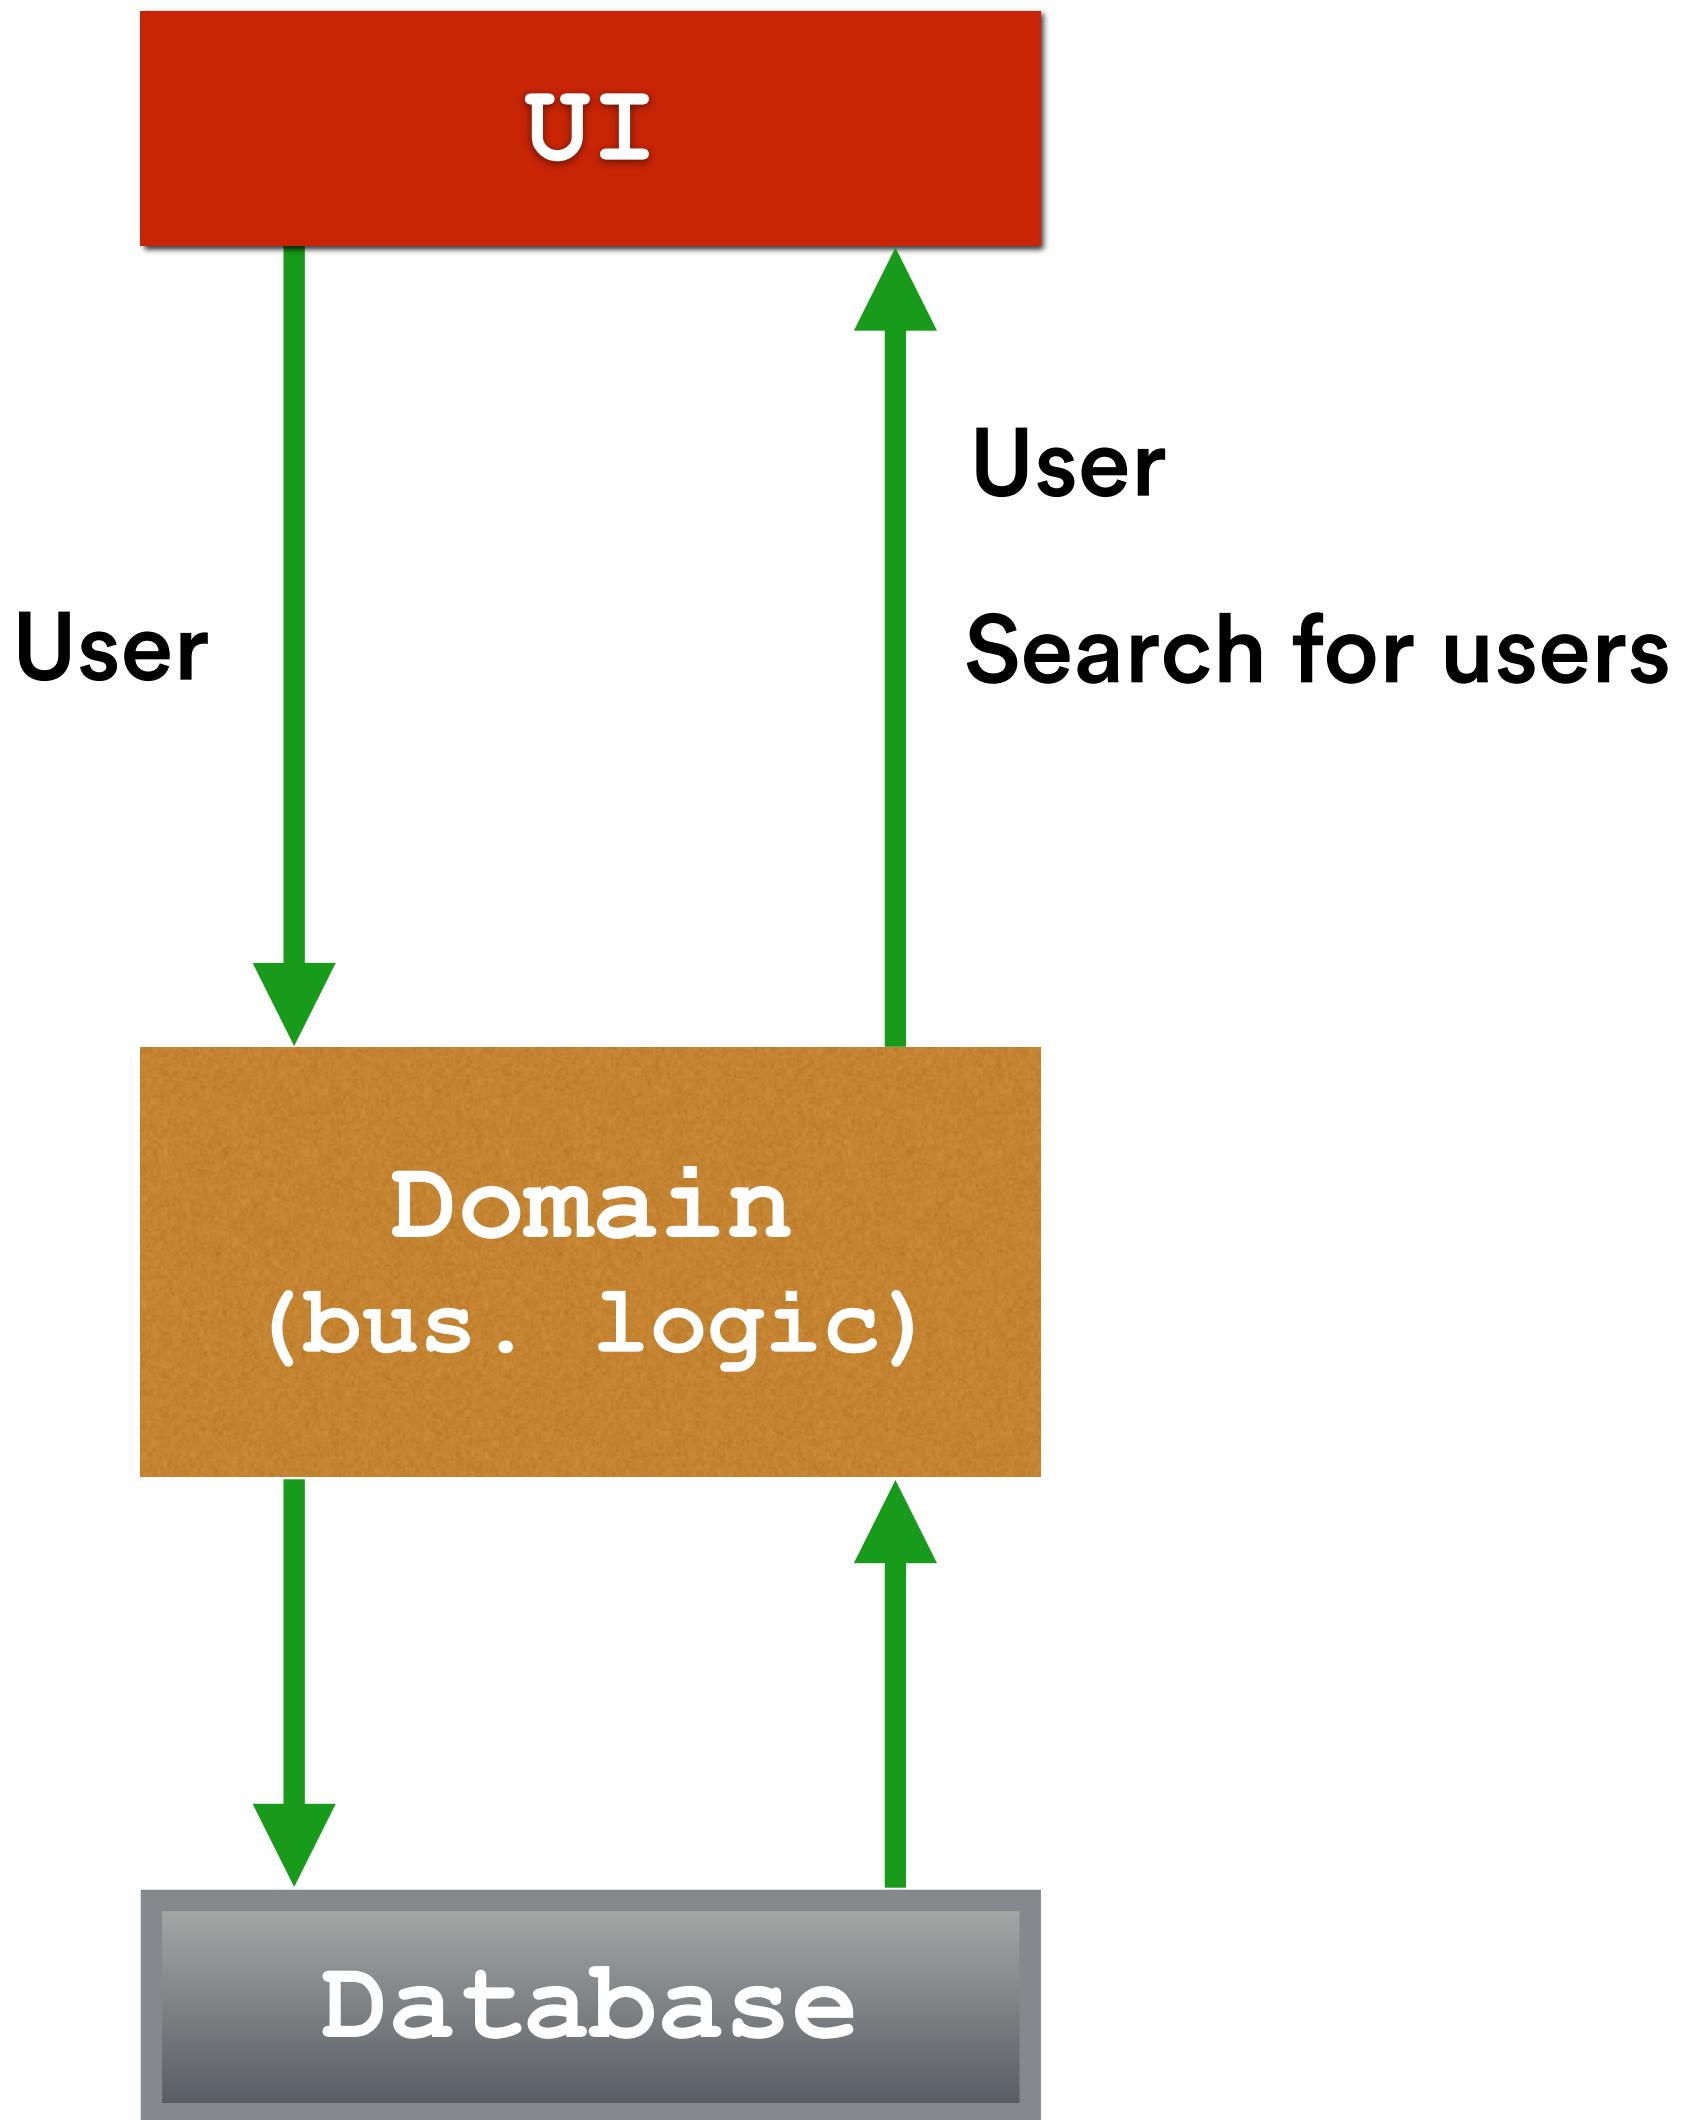
**REST calls**  
e.g. Add User

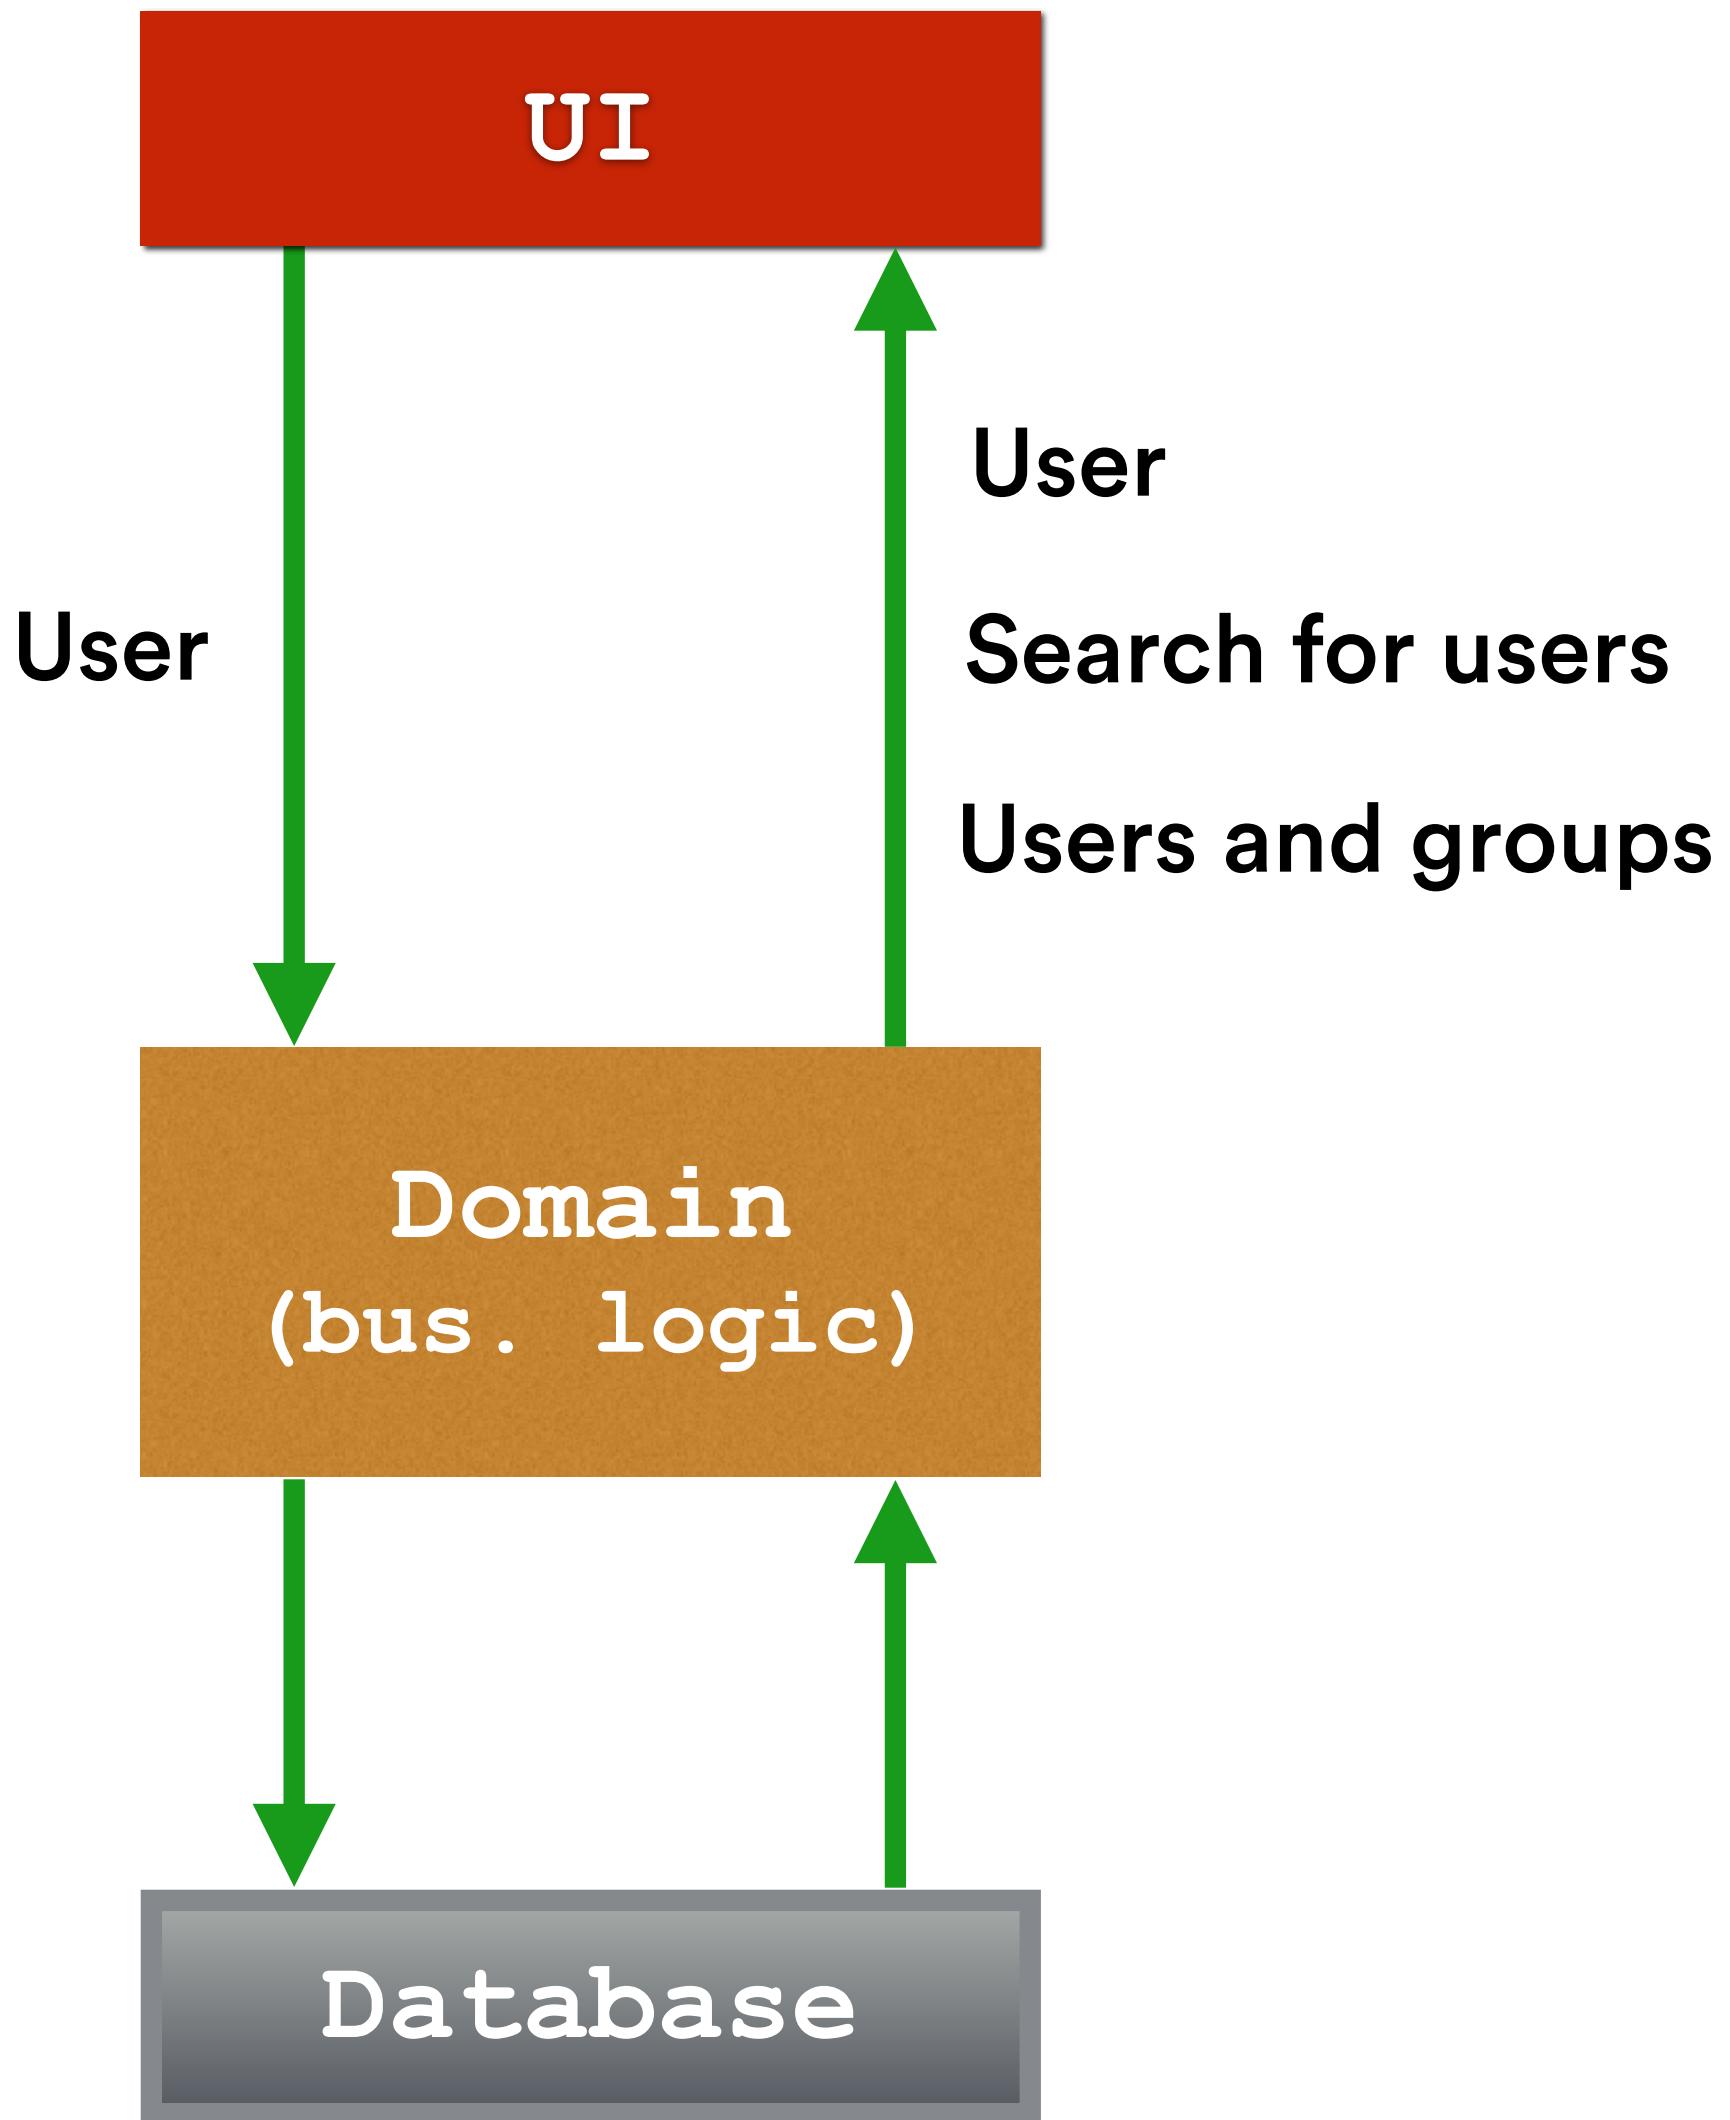


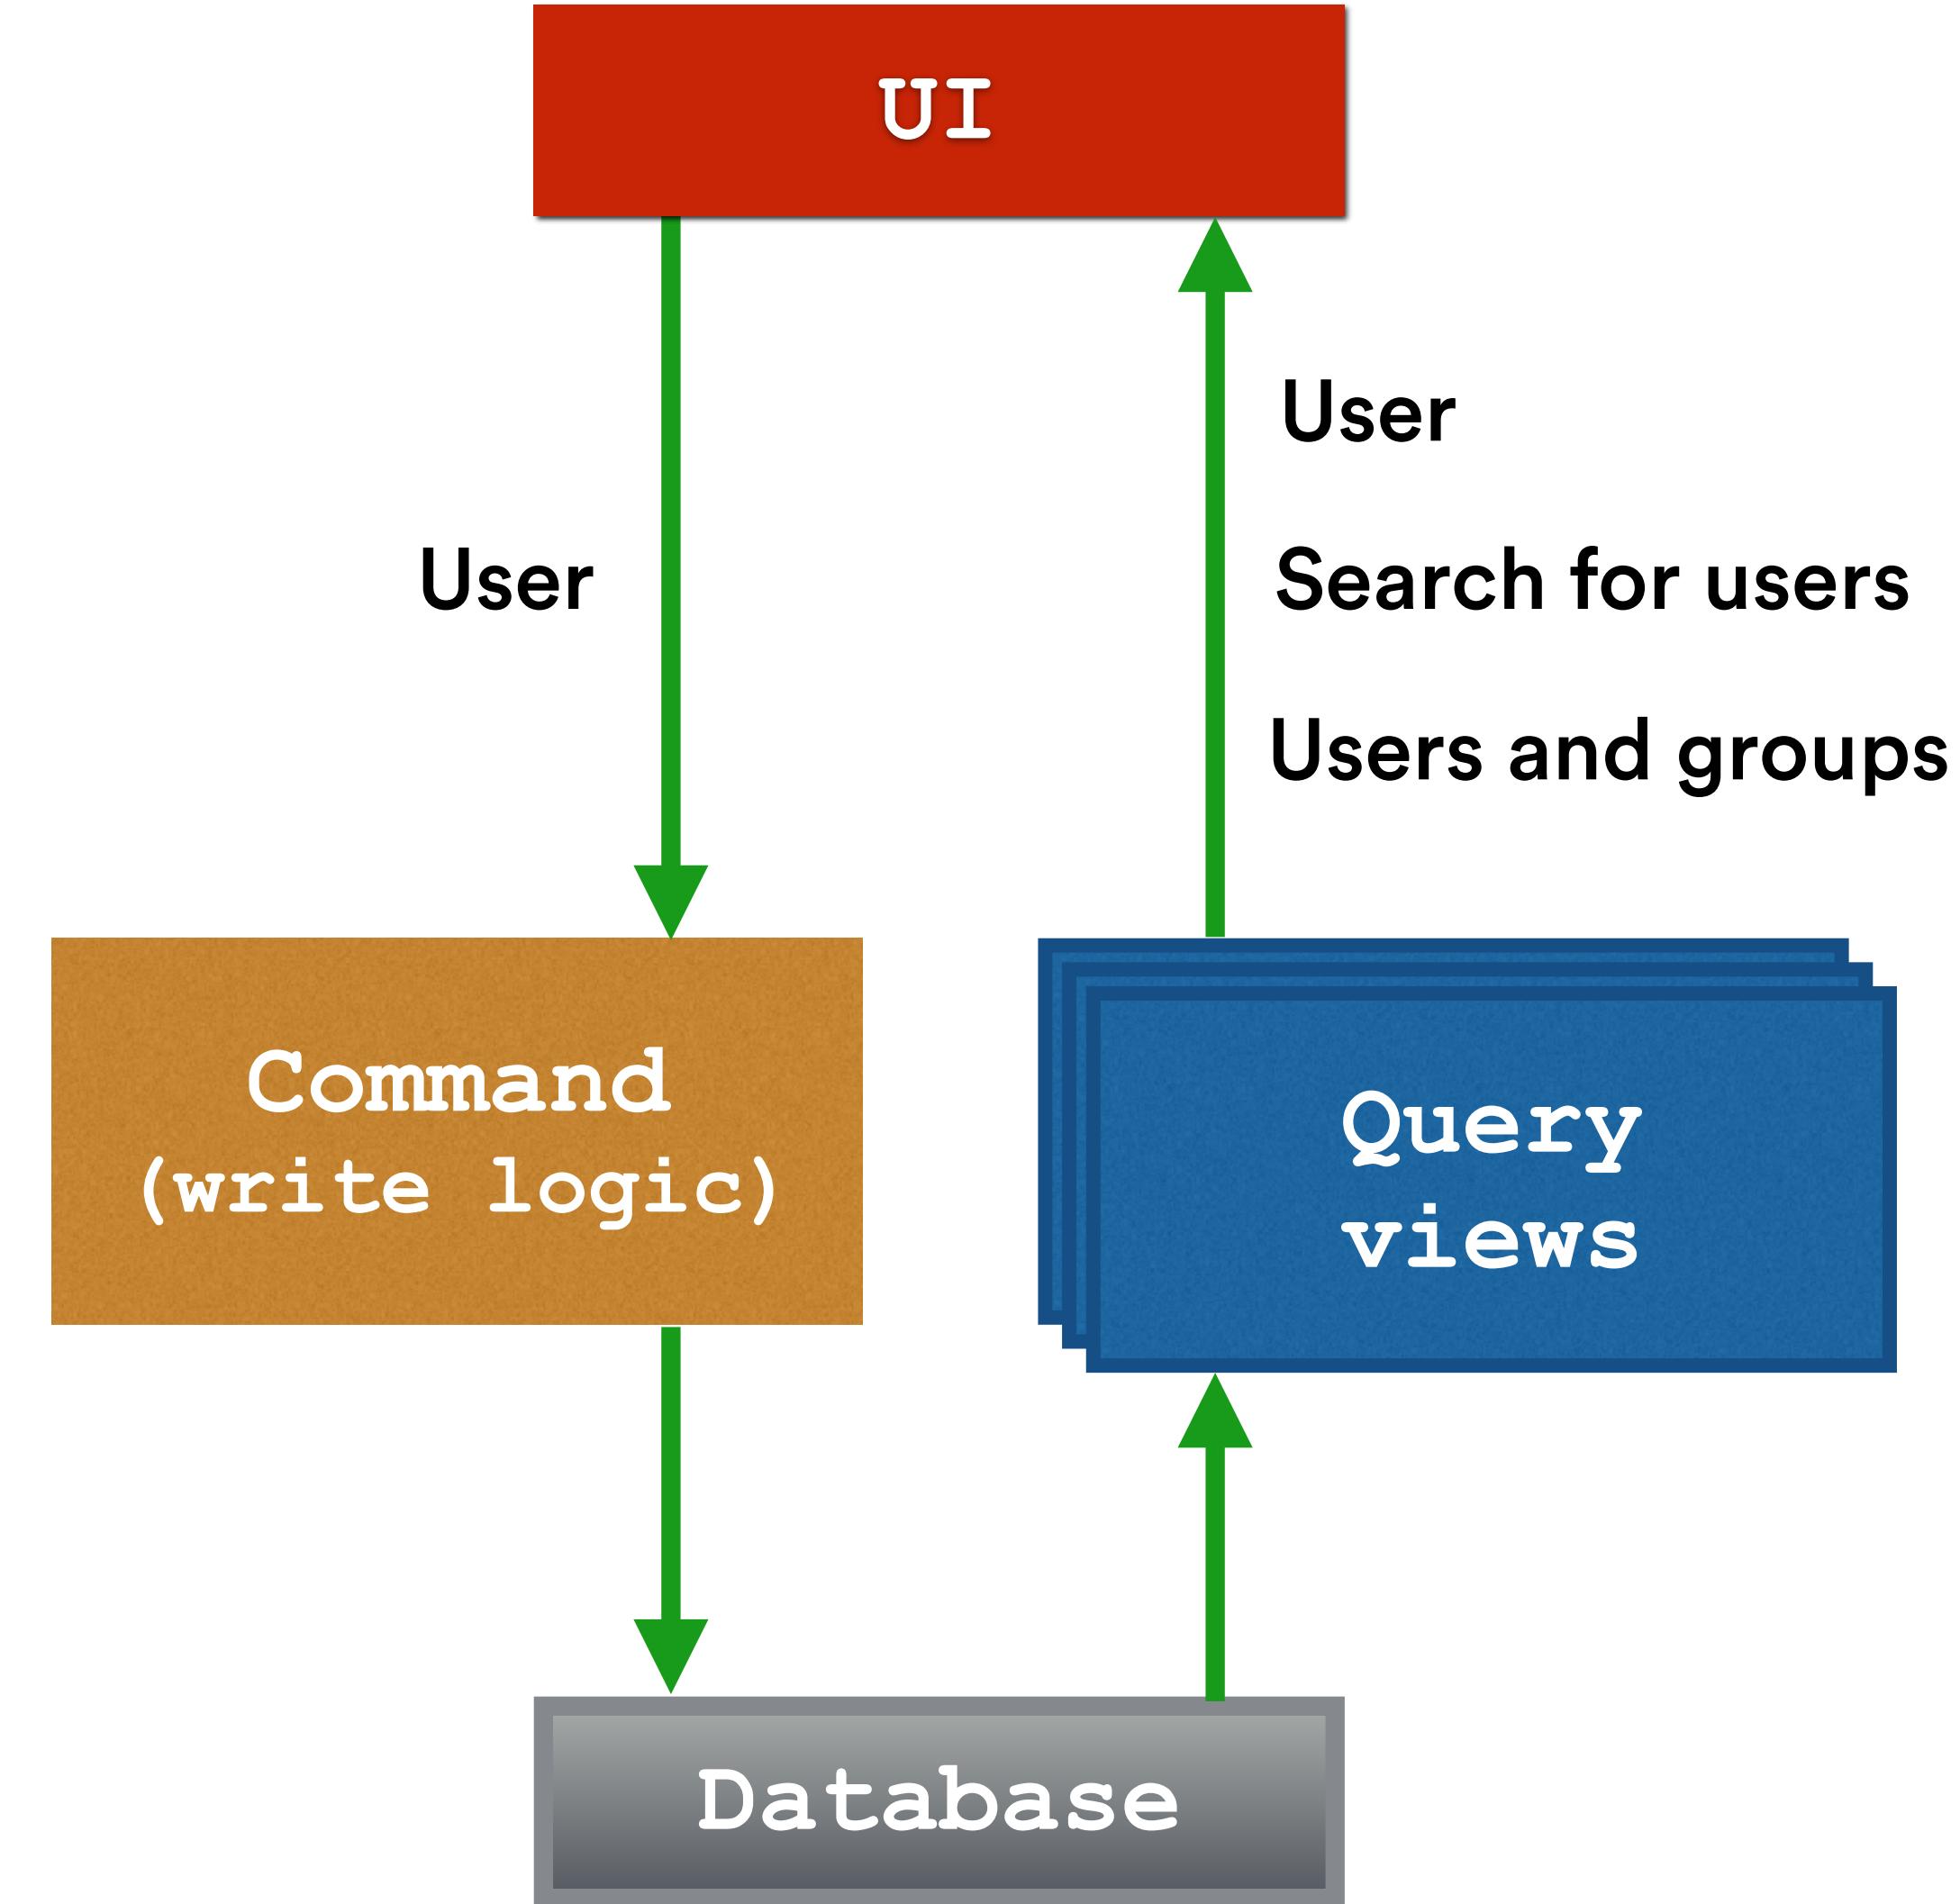


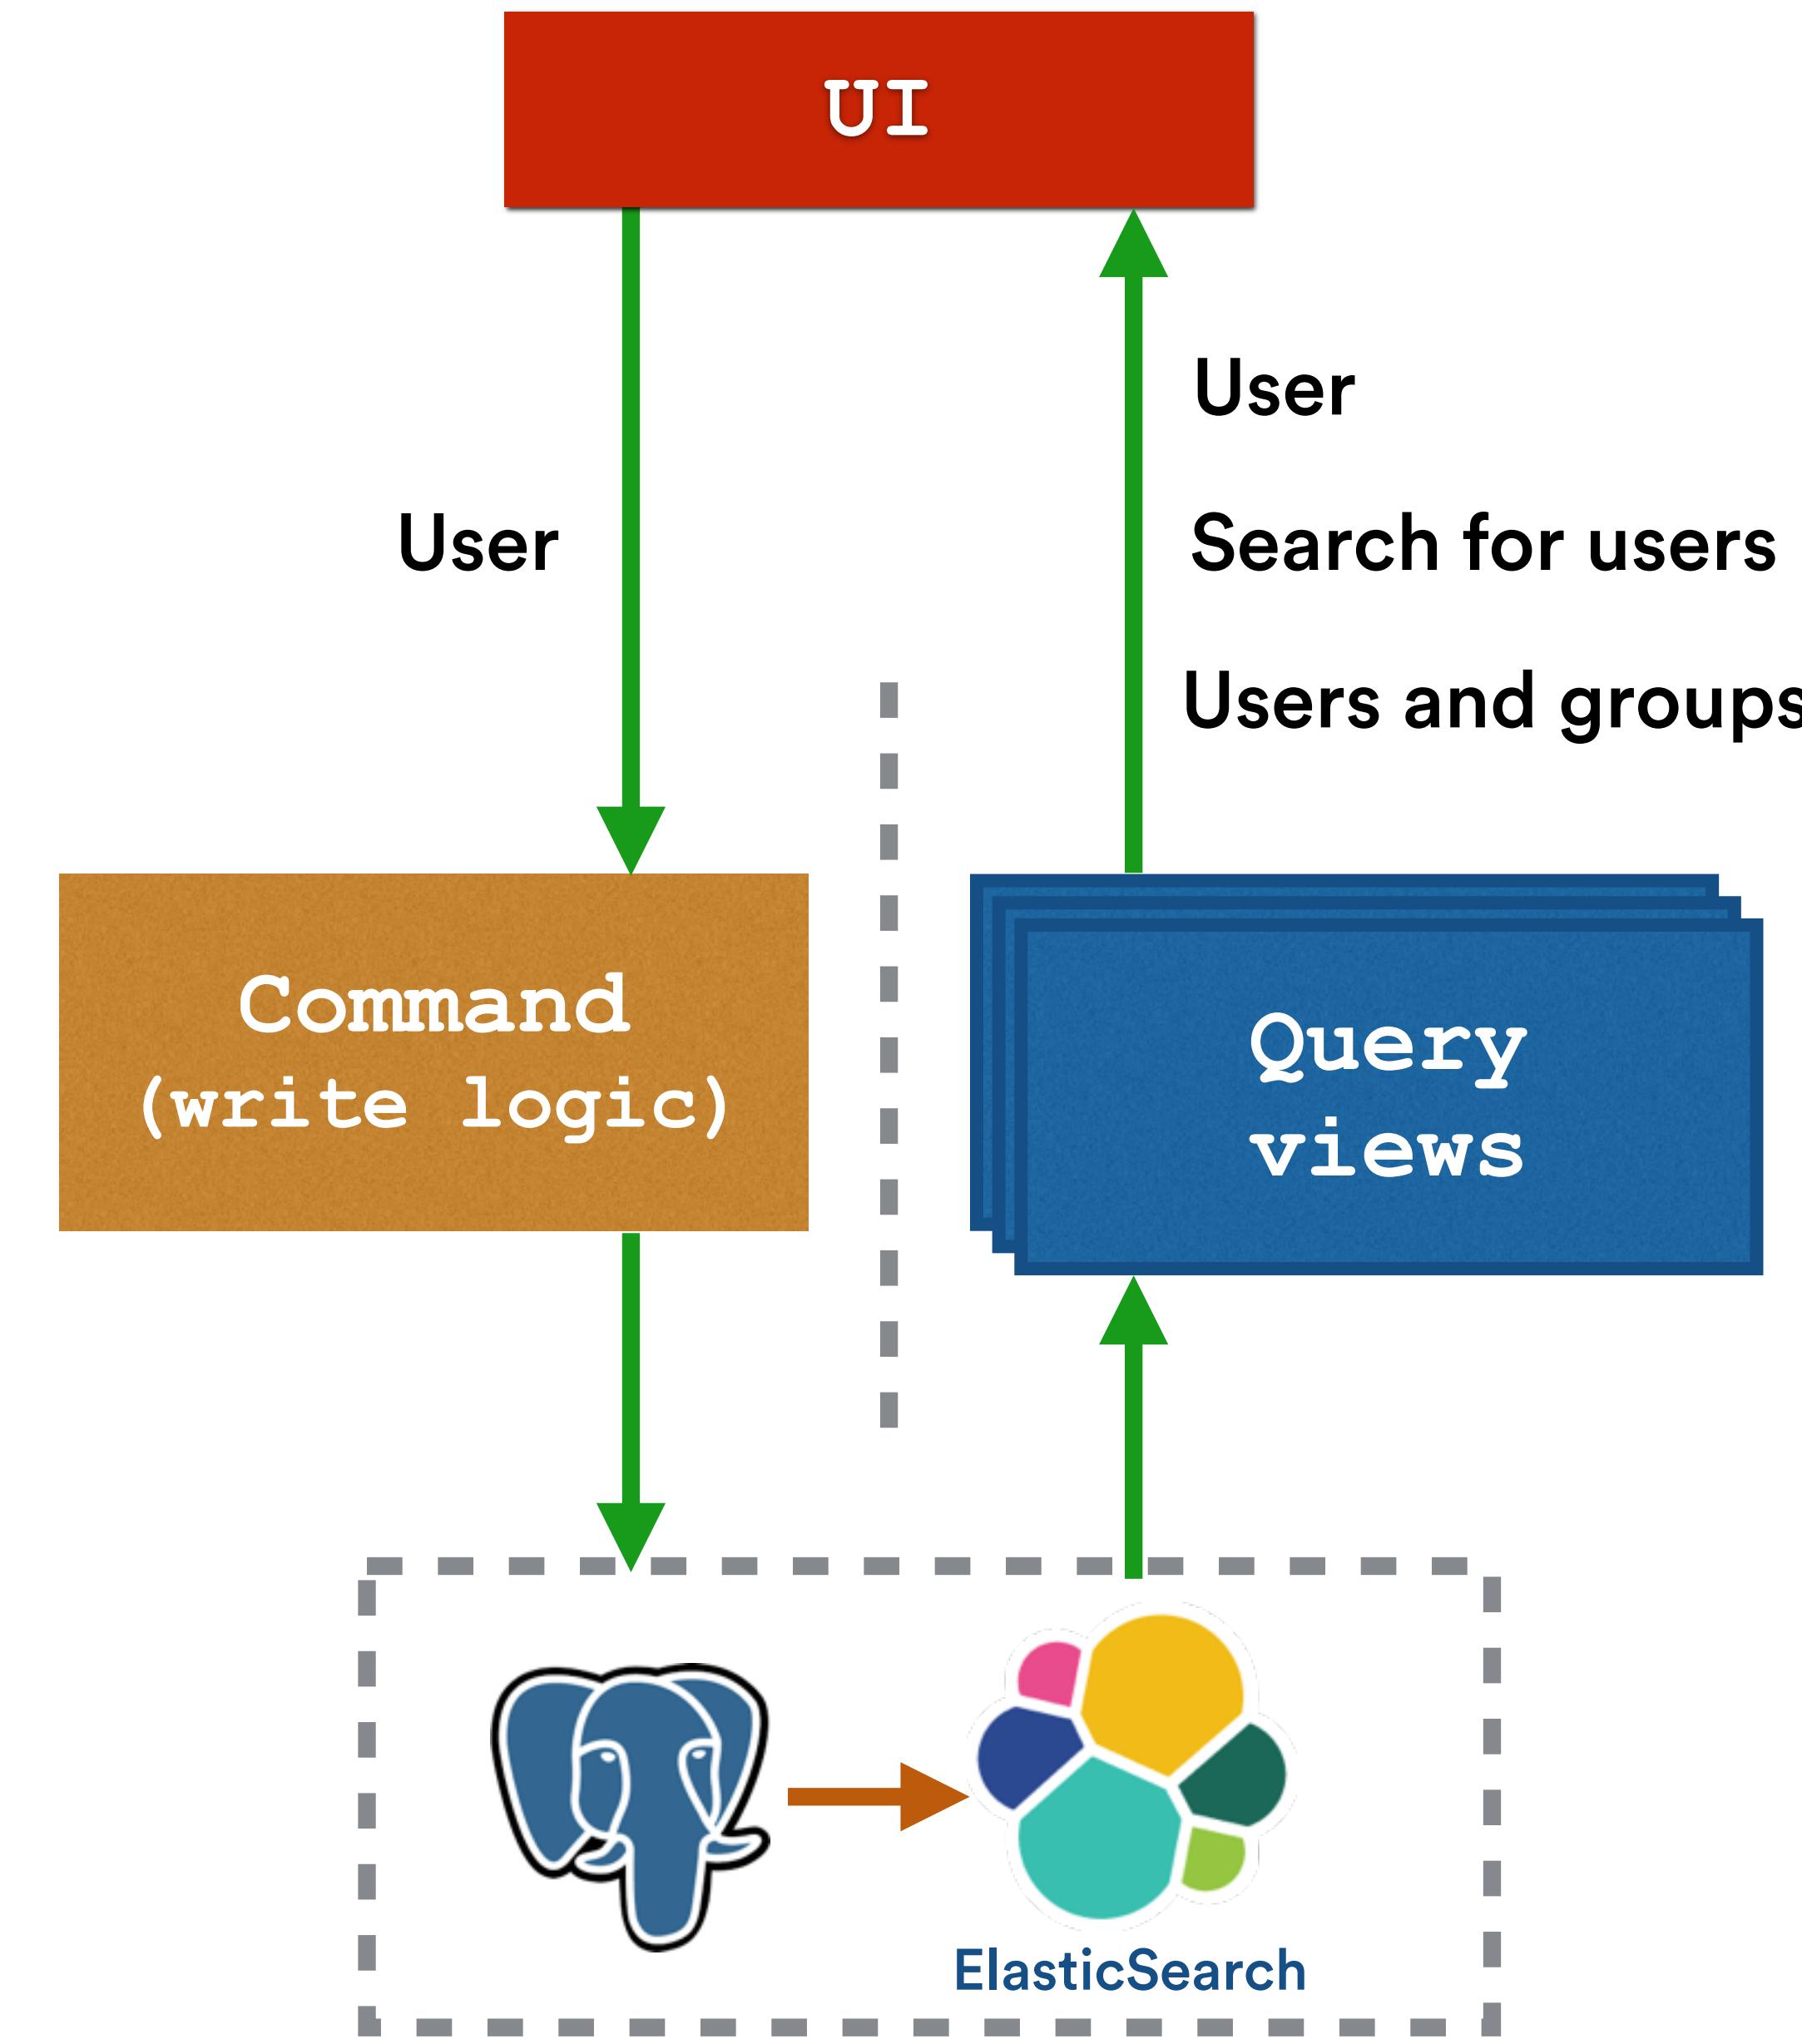
# Command Query Responsibility Segregation



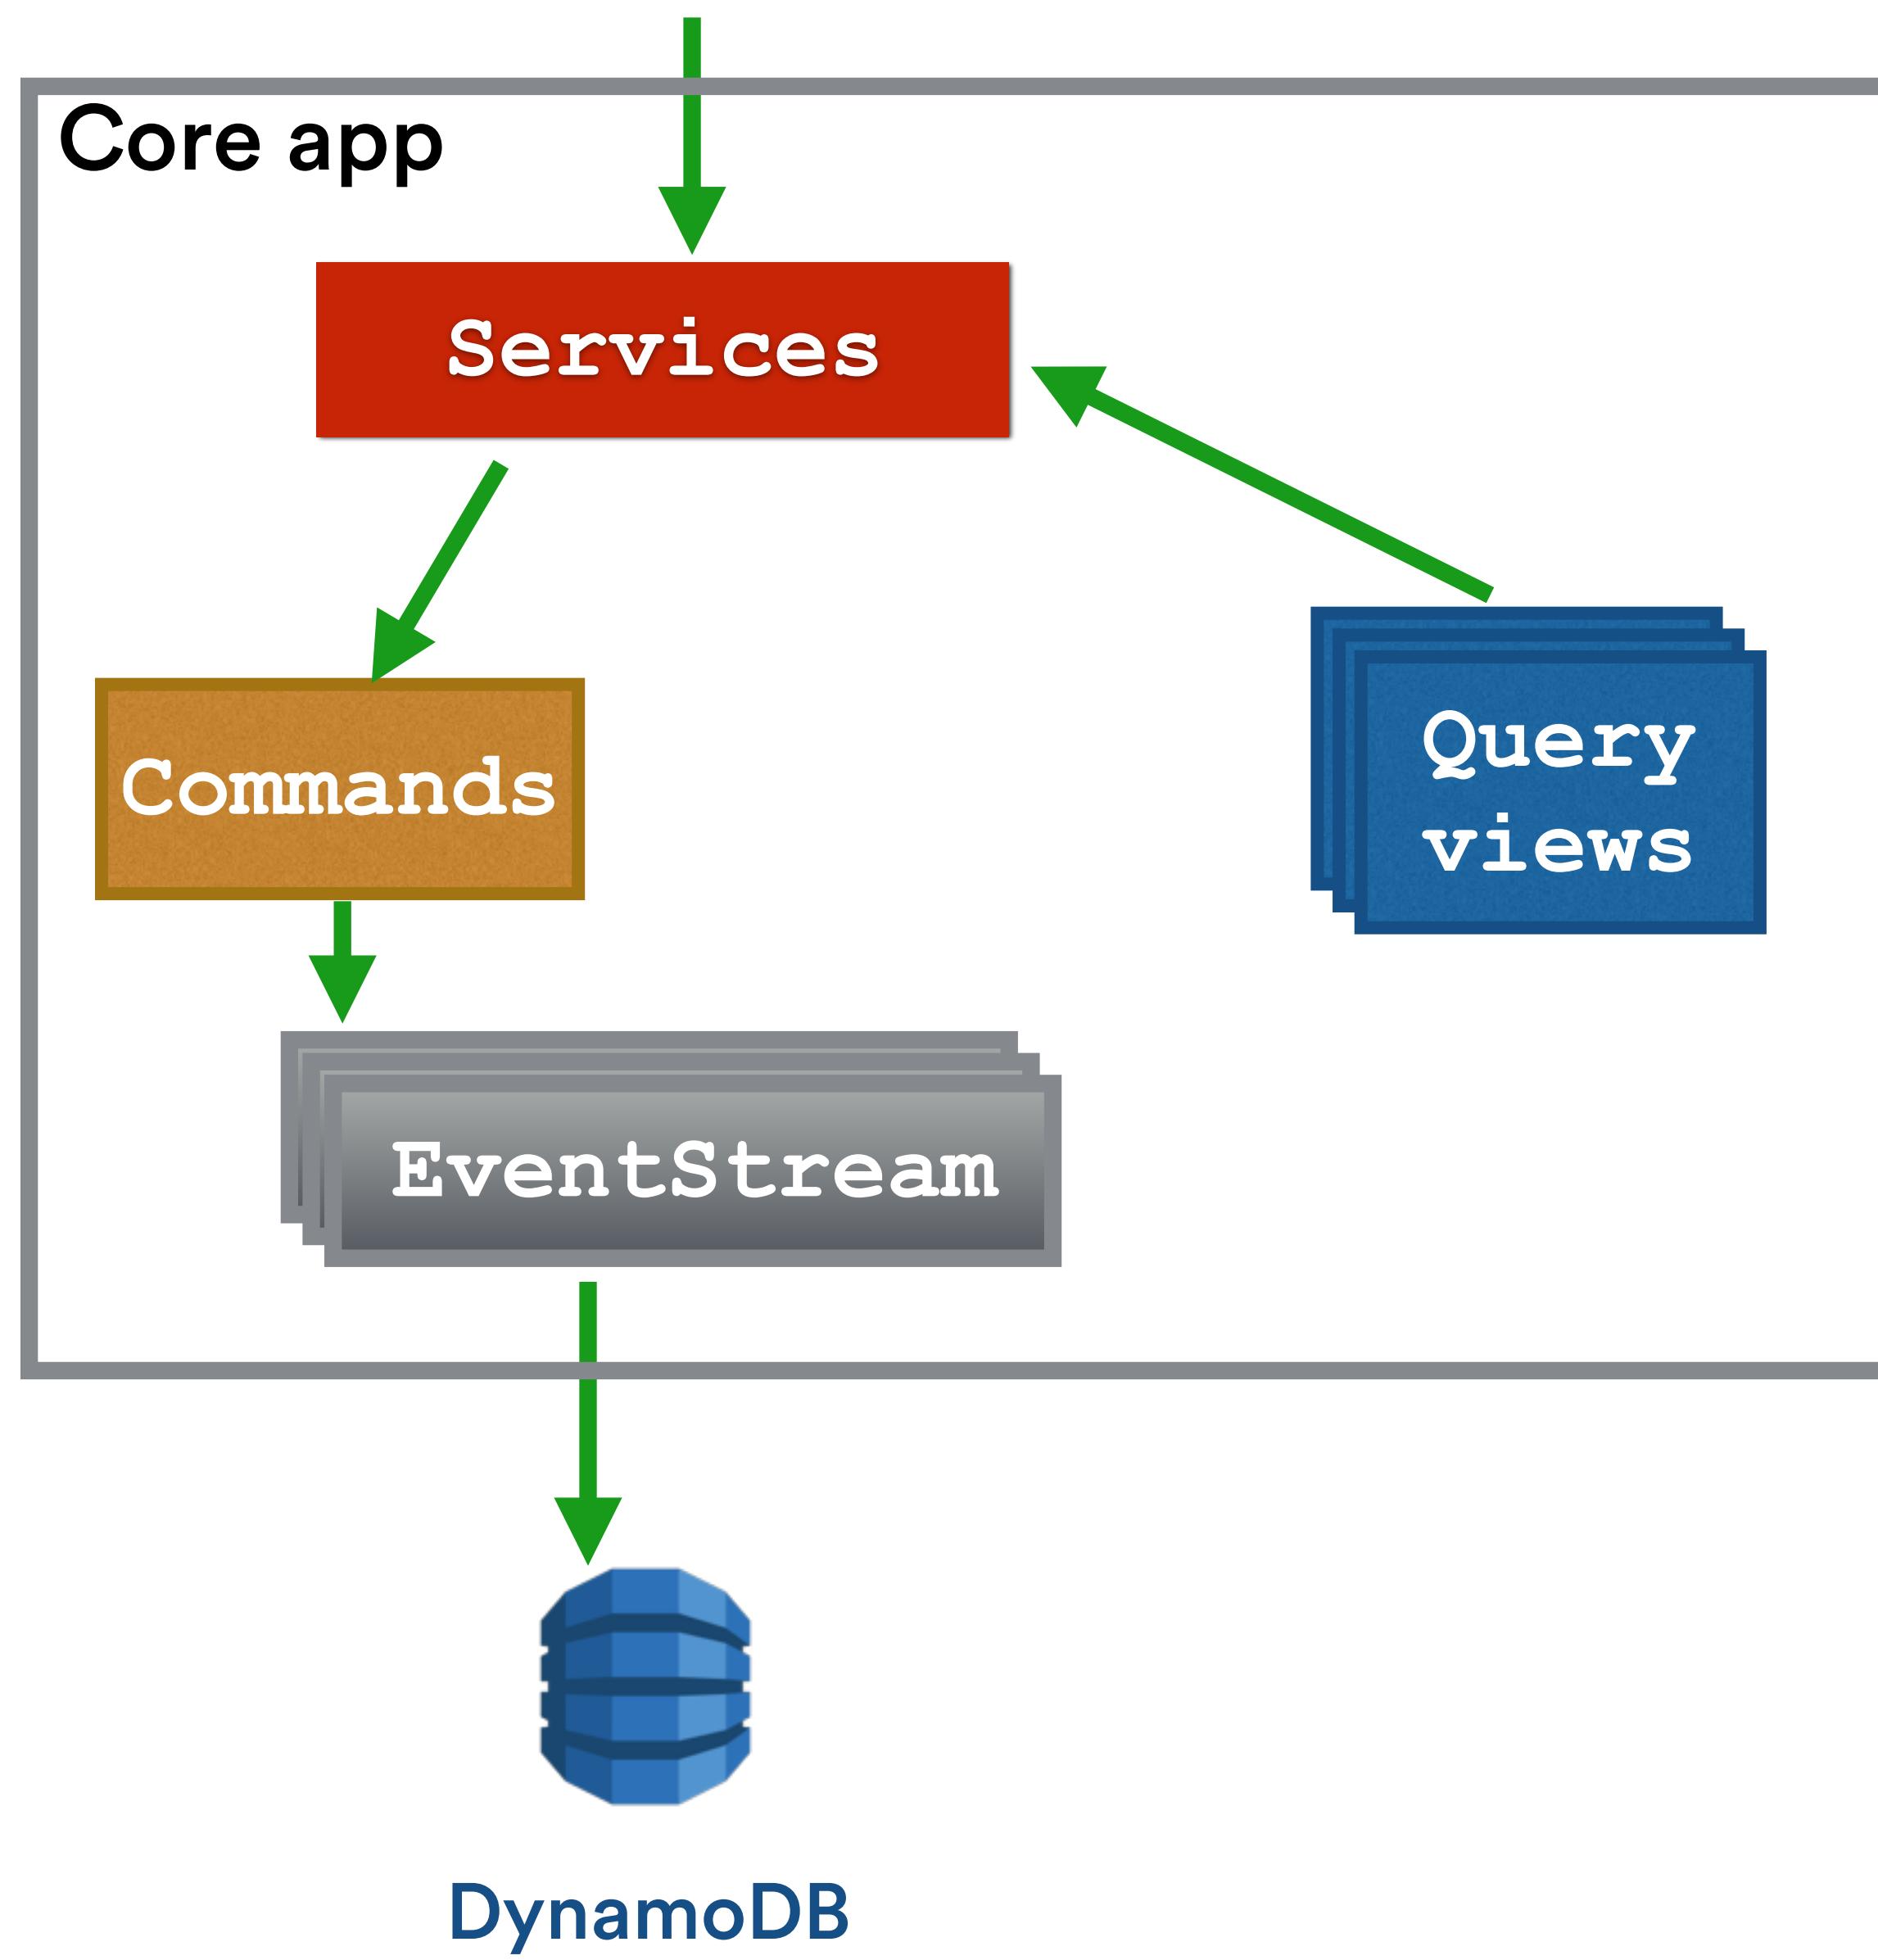






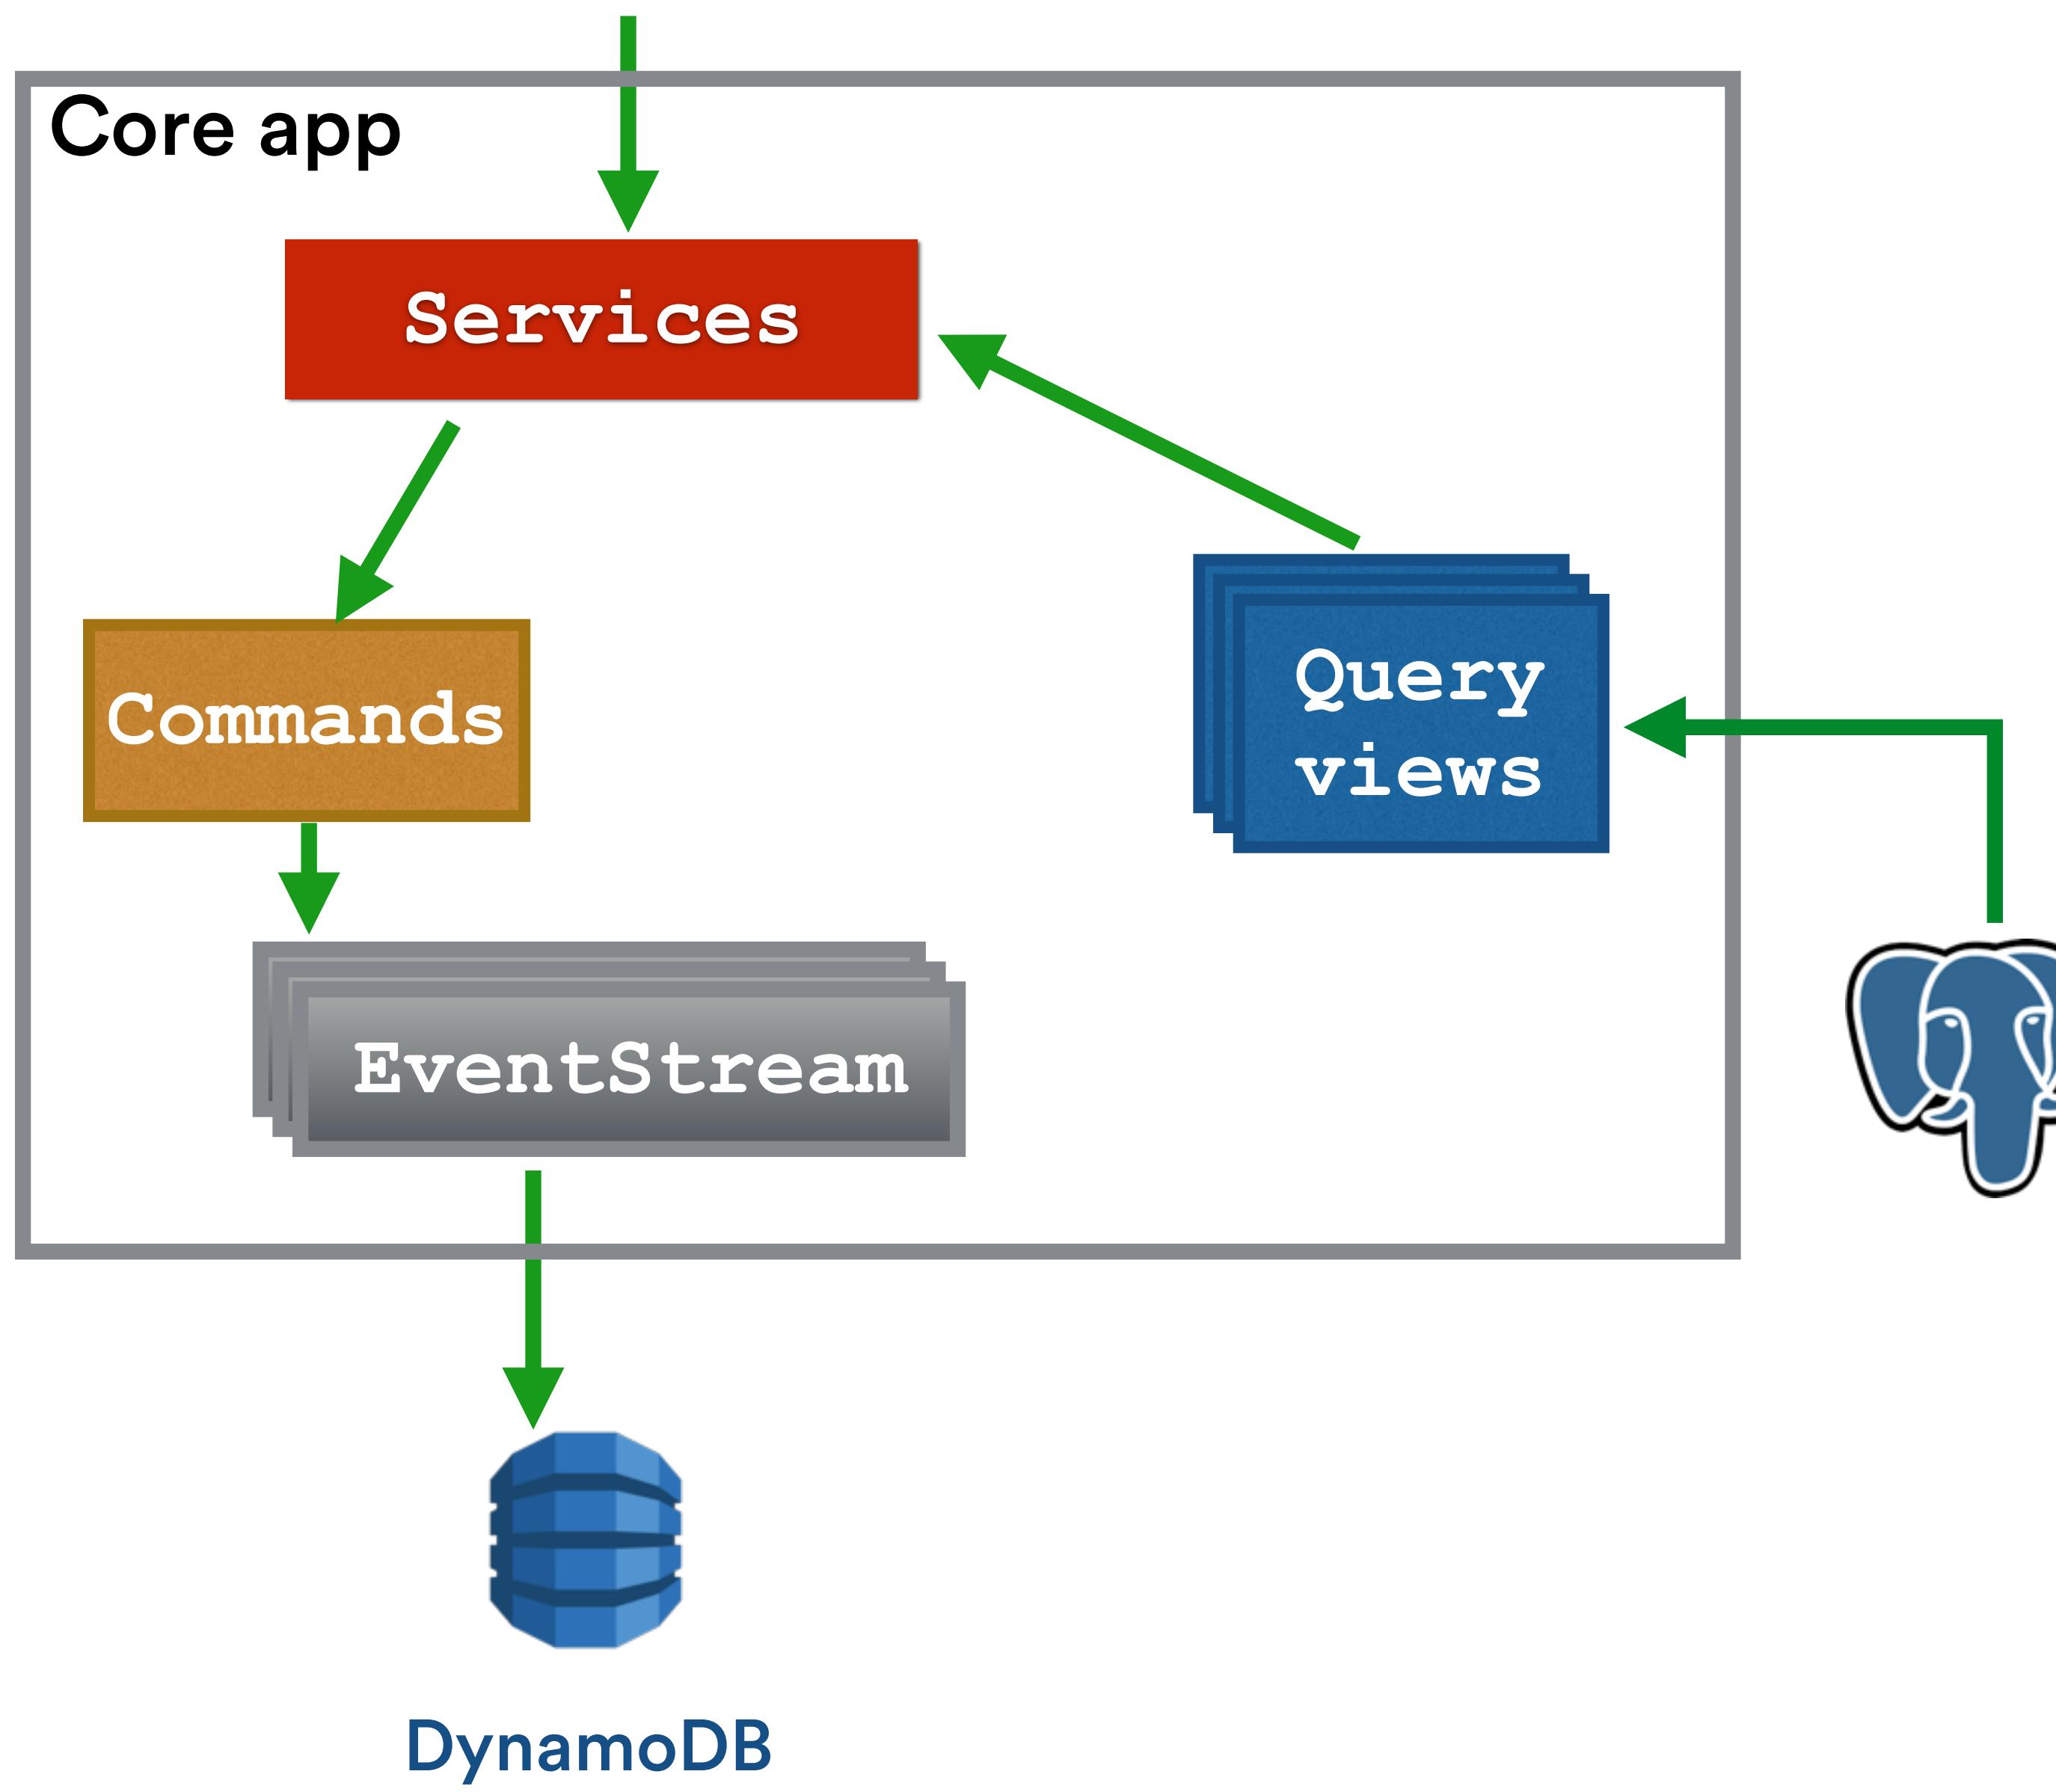


**REST calls**  
e.g. Add User

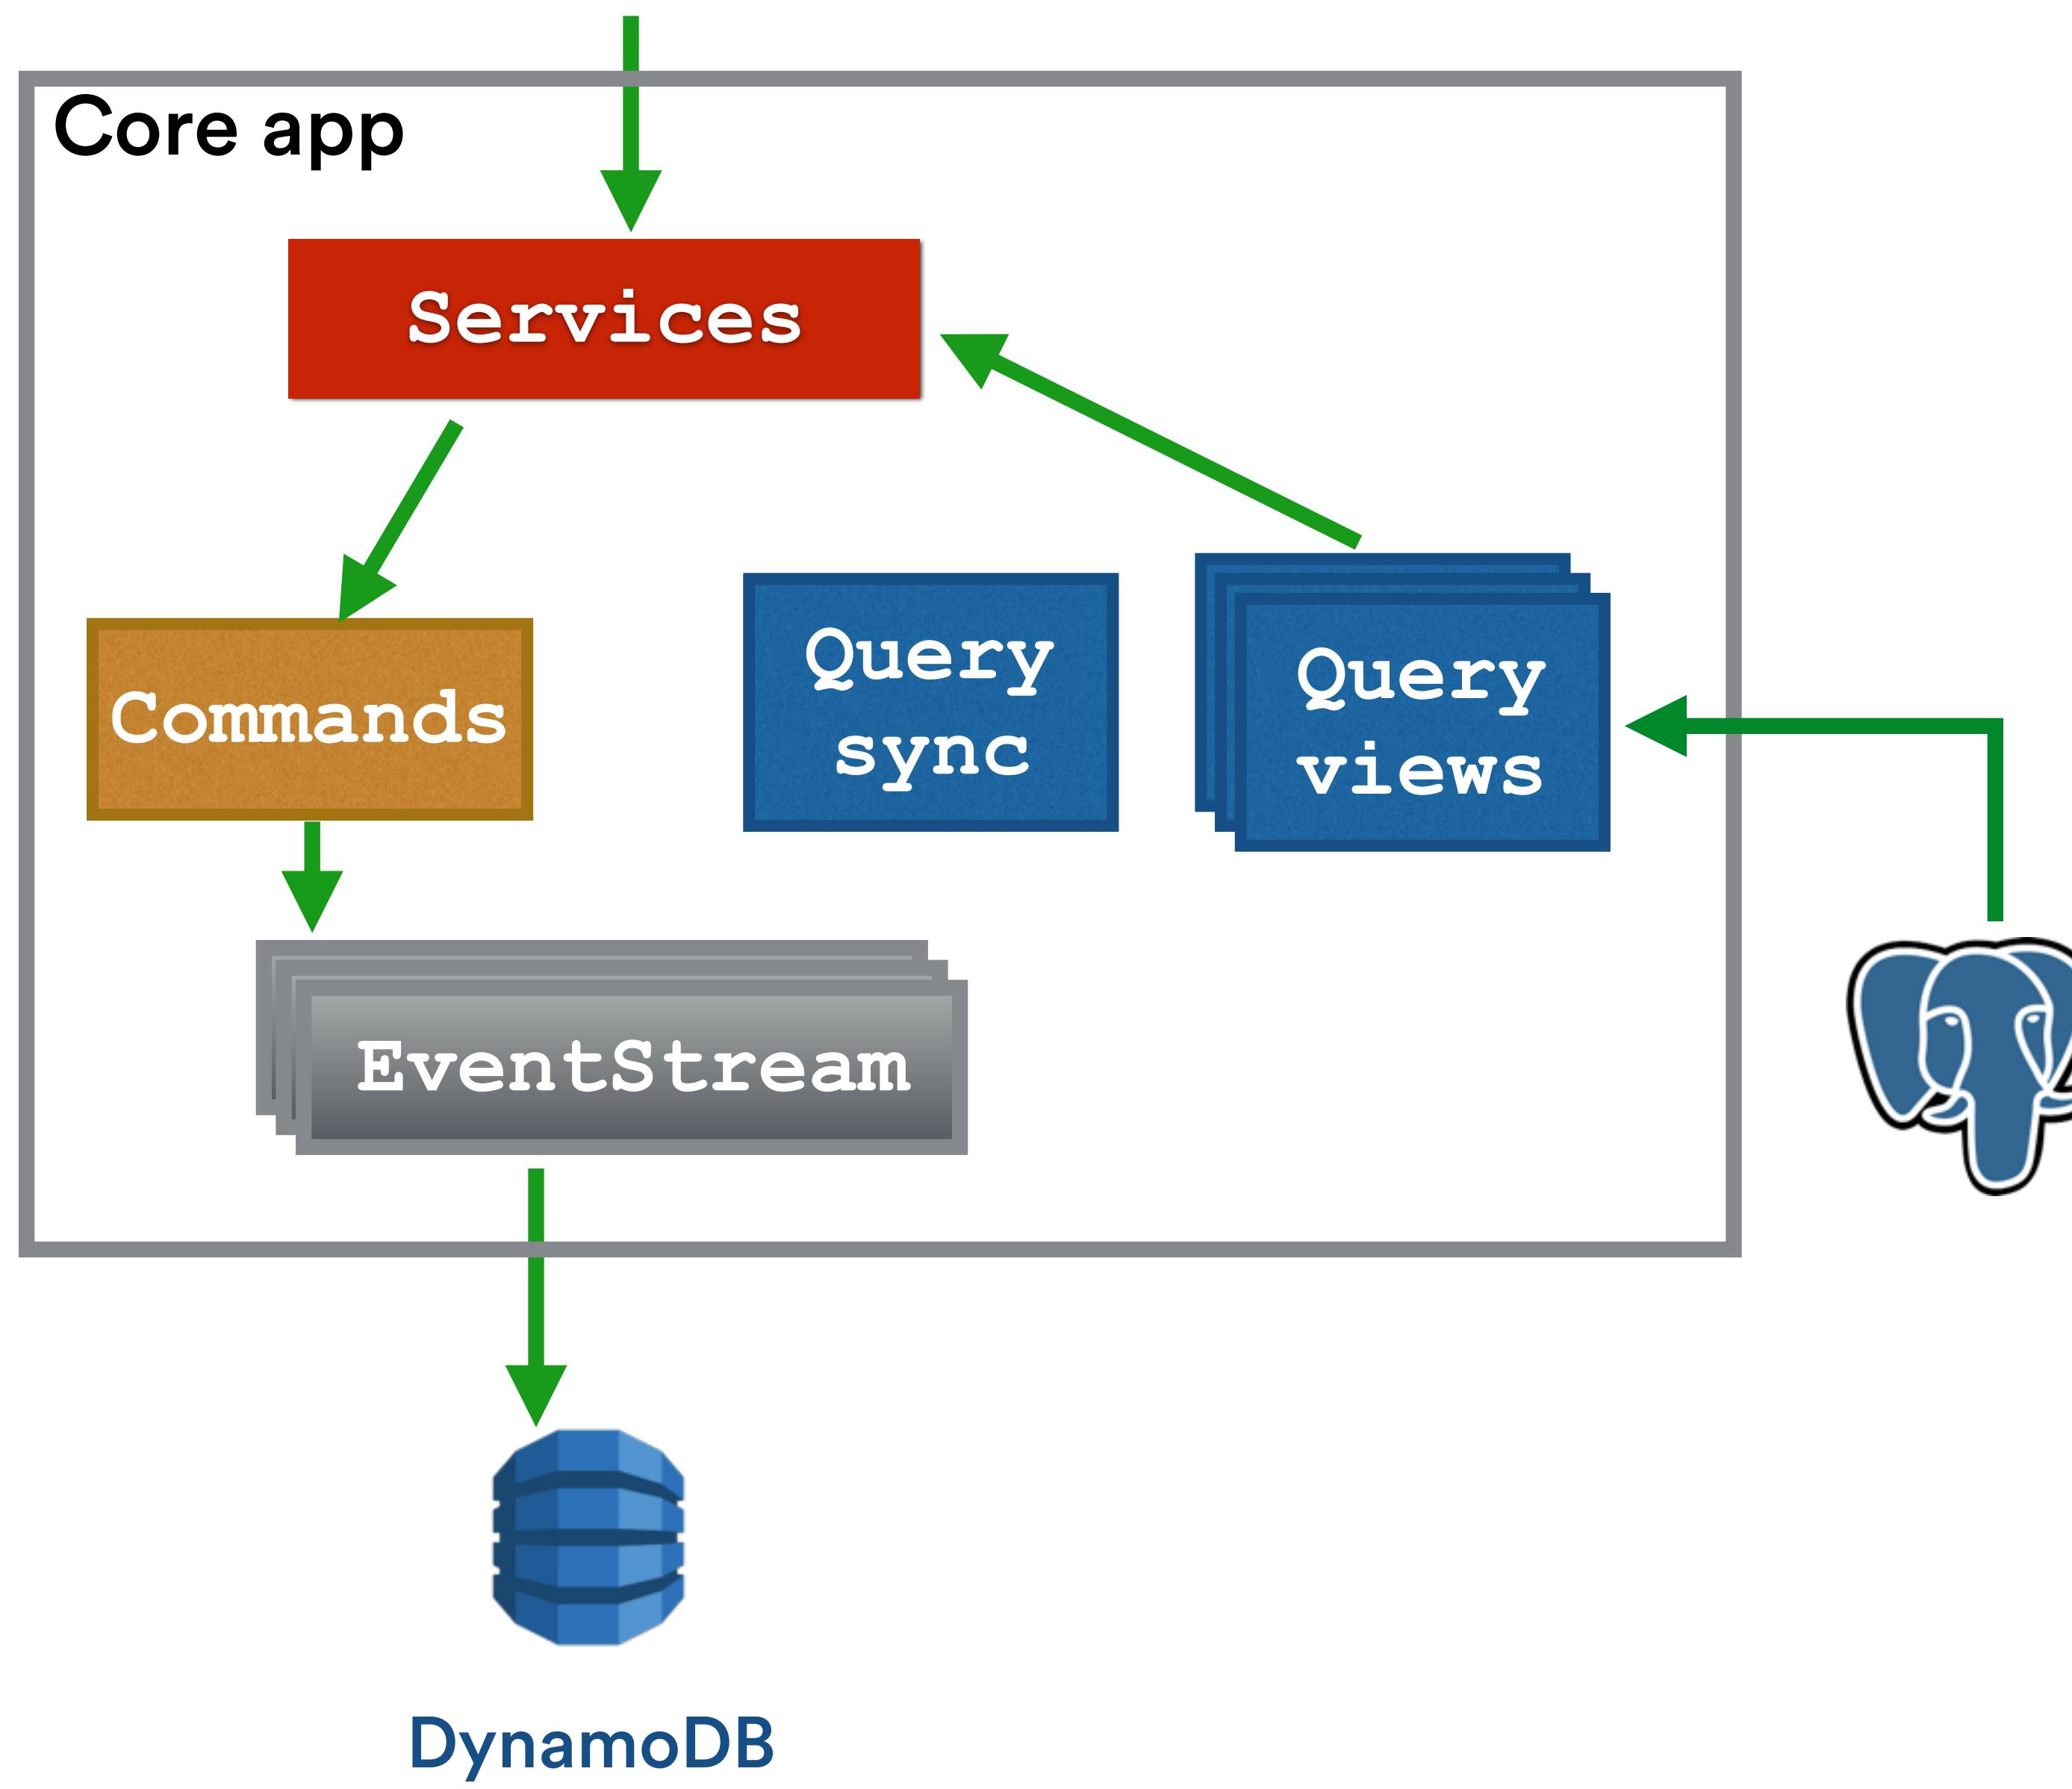


**REST calls**

e.g. Add User

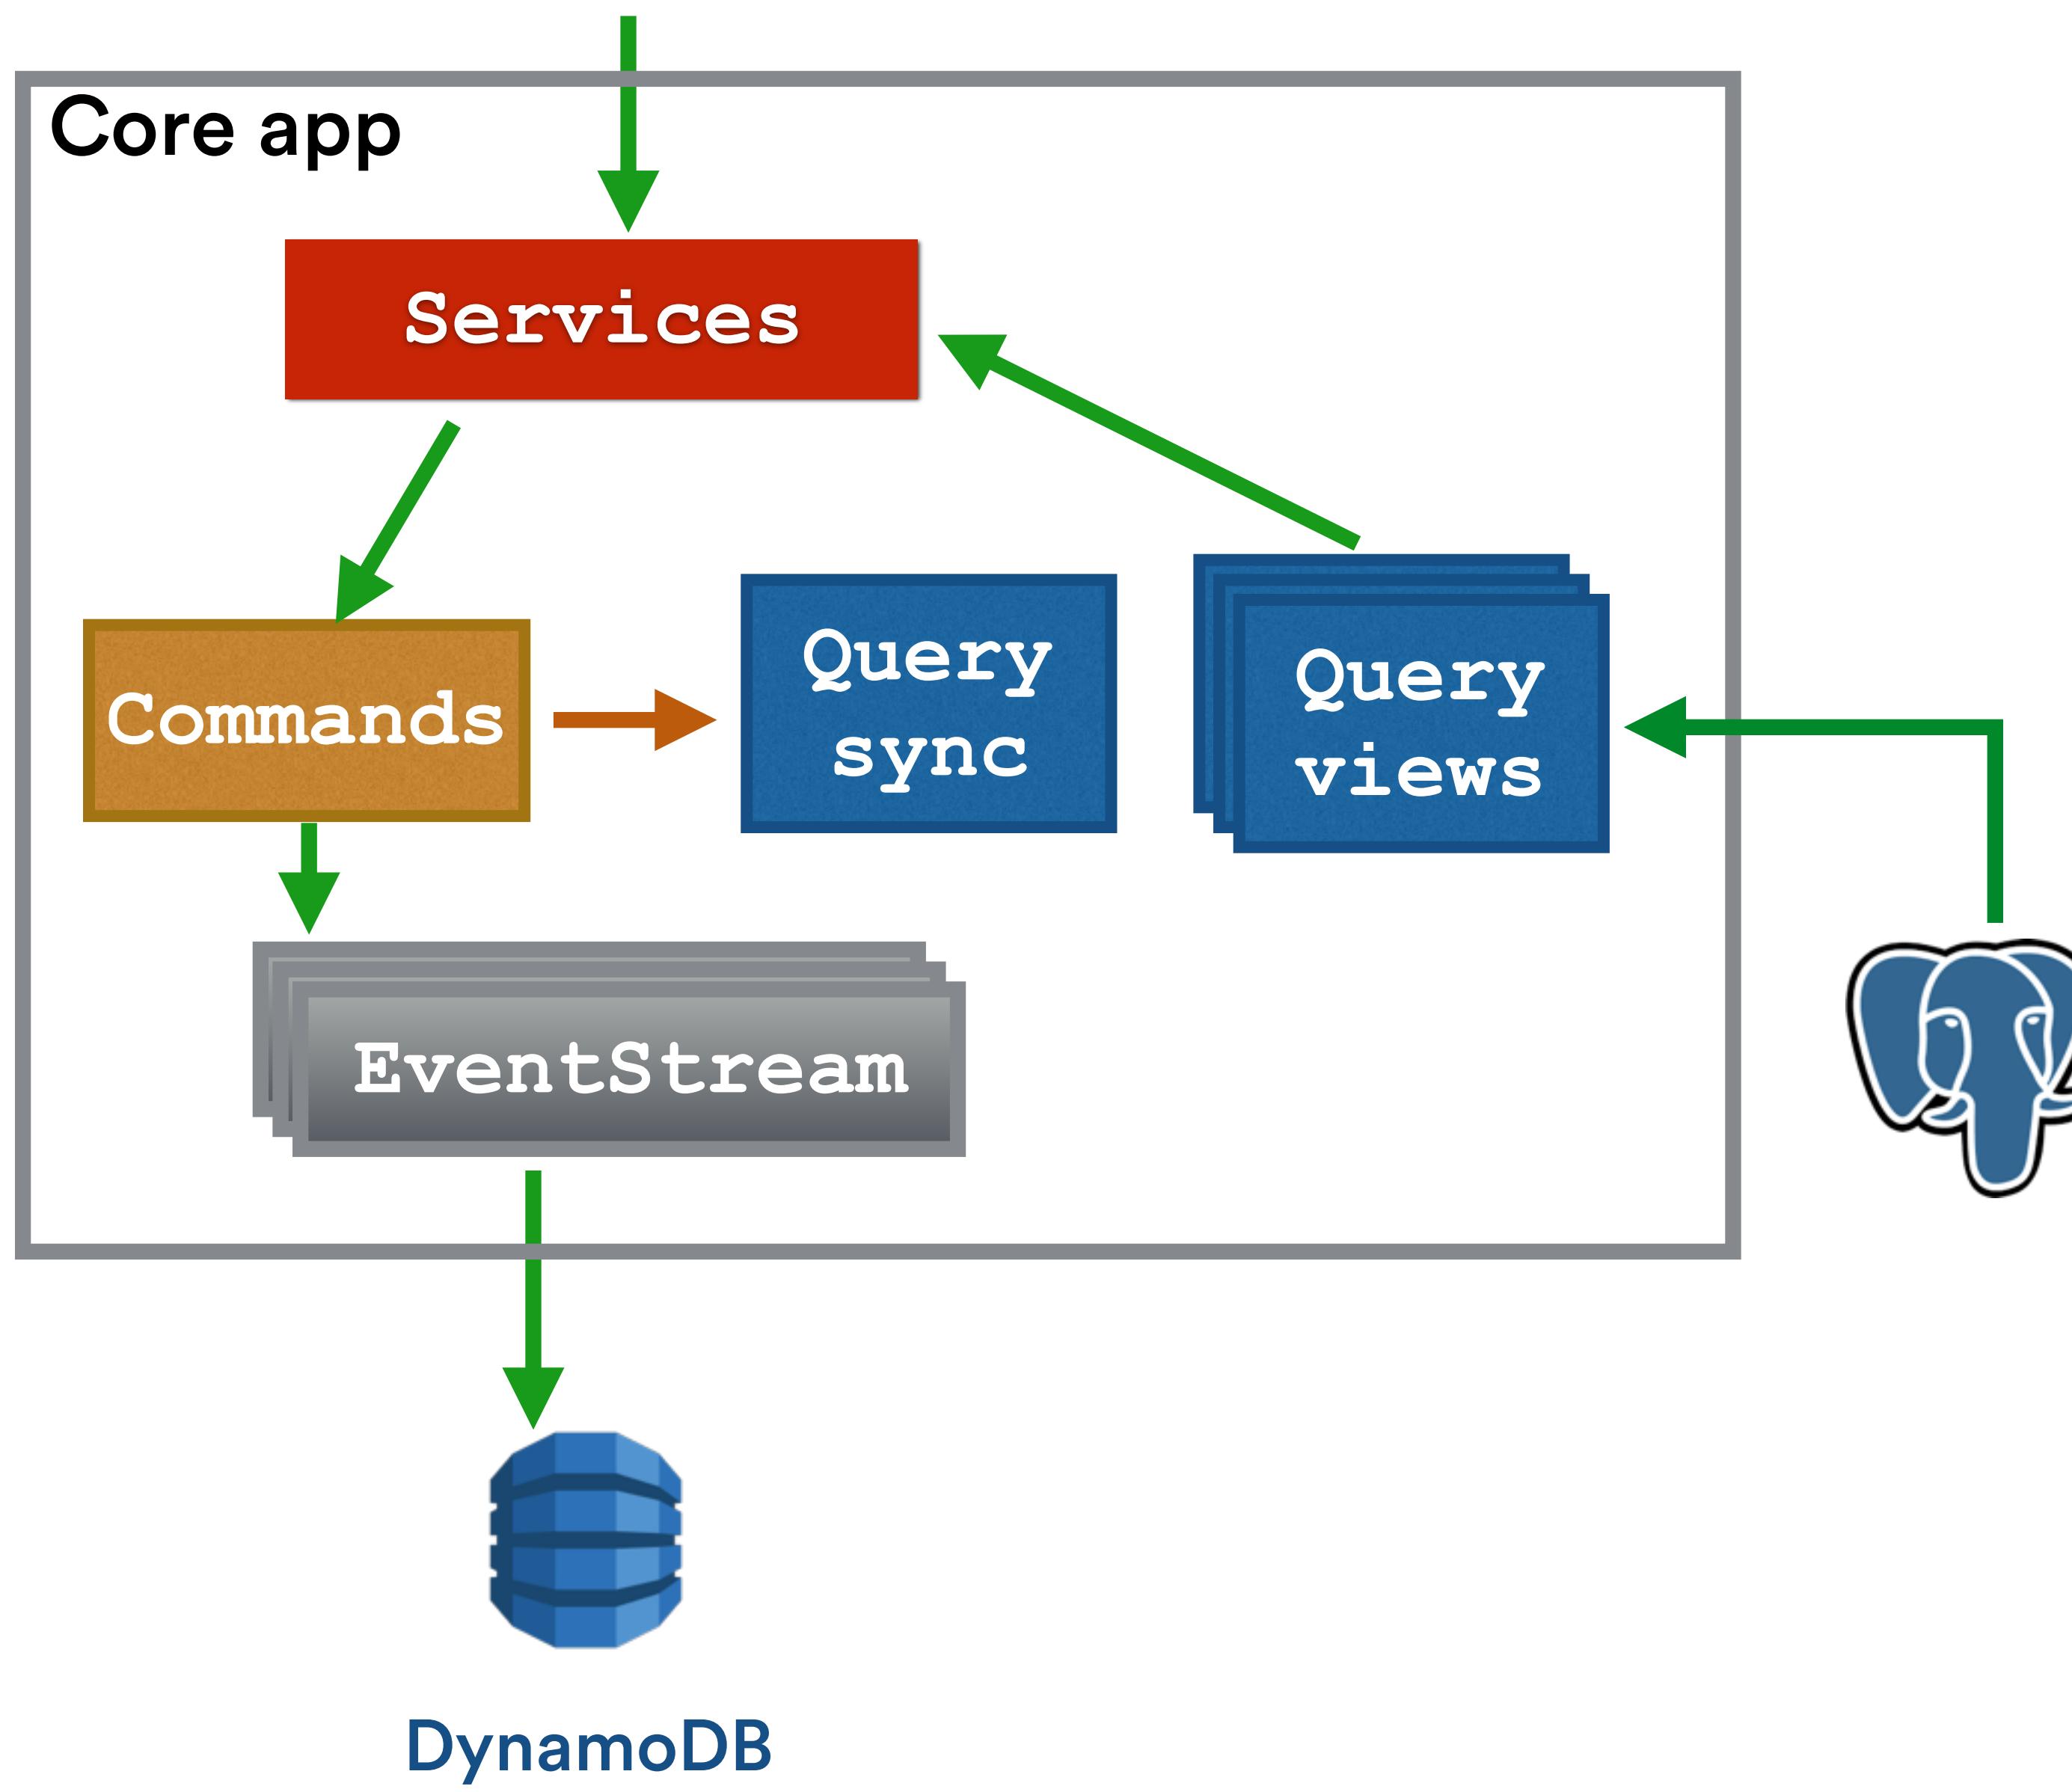


**REST calls**  
e.g. Add User

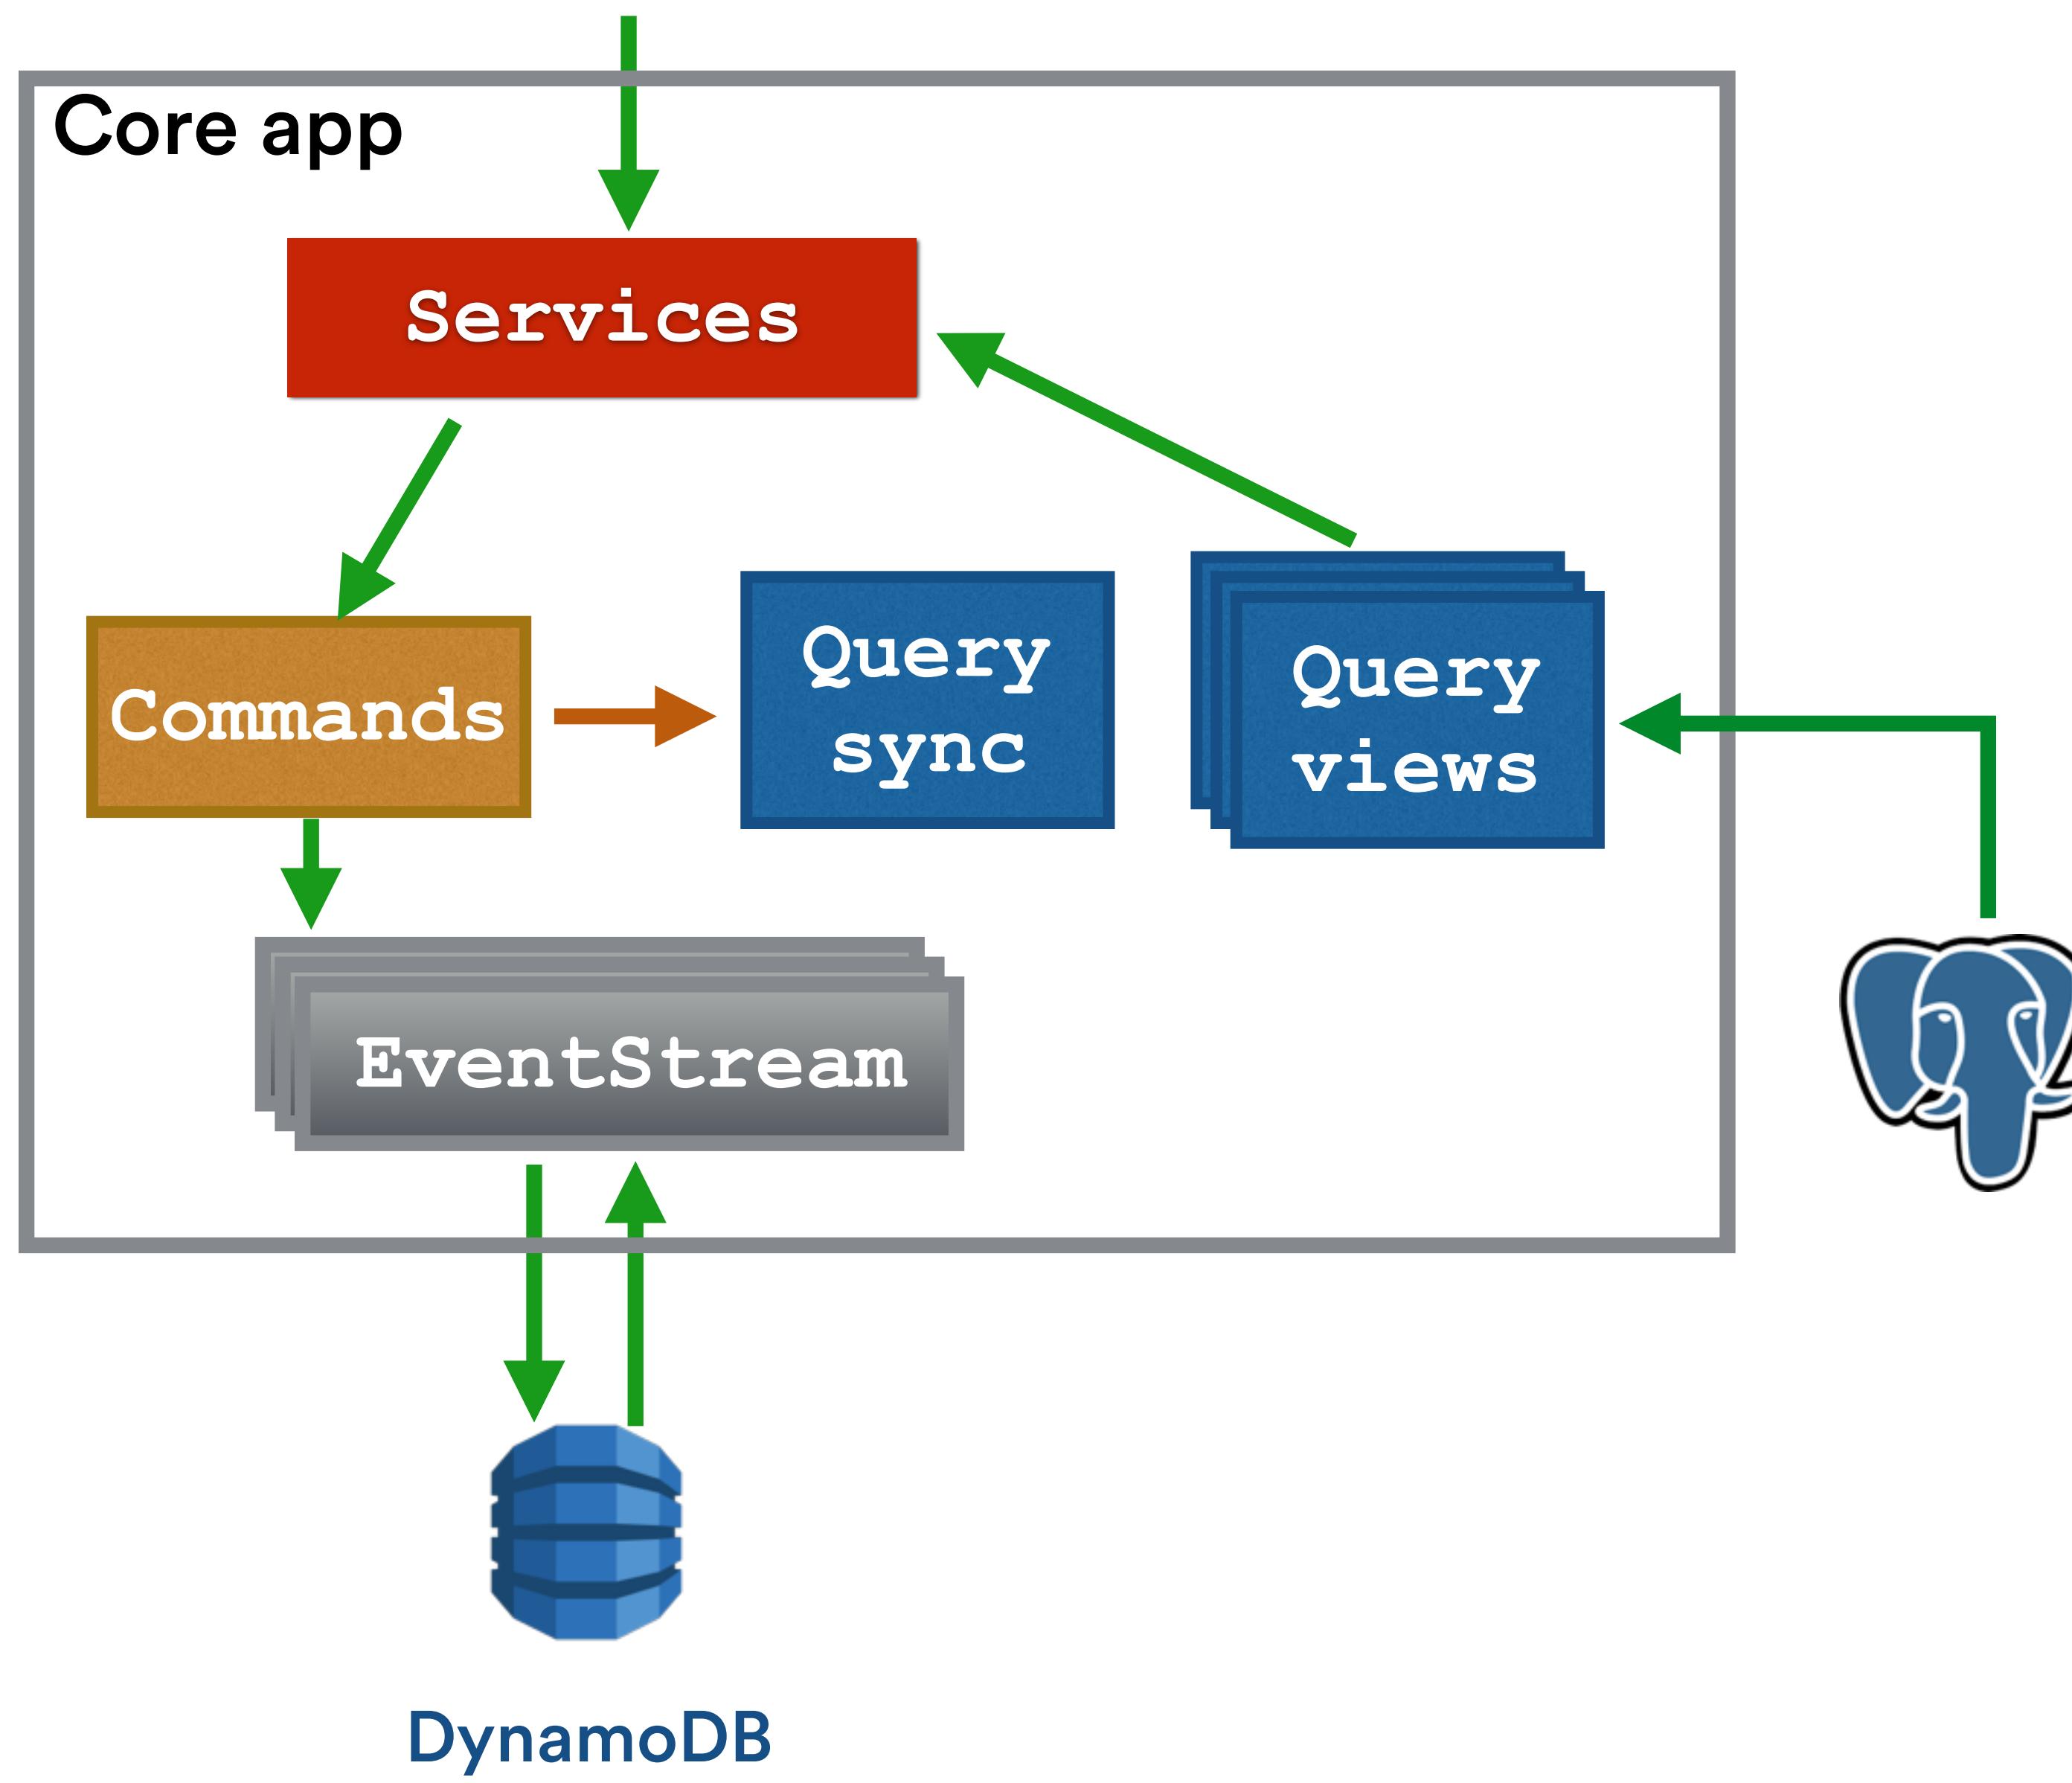


# REST calls

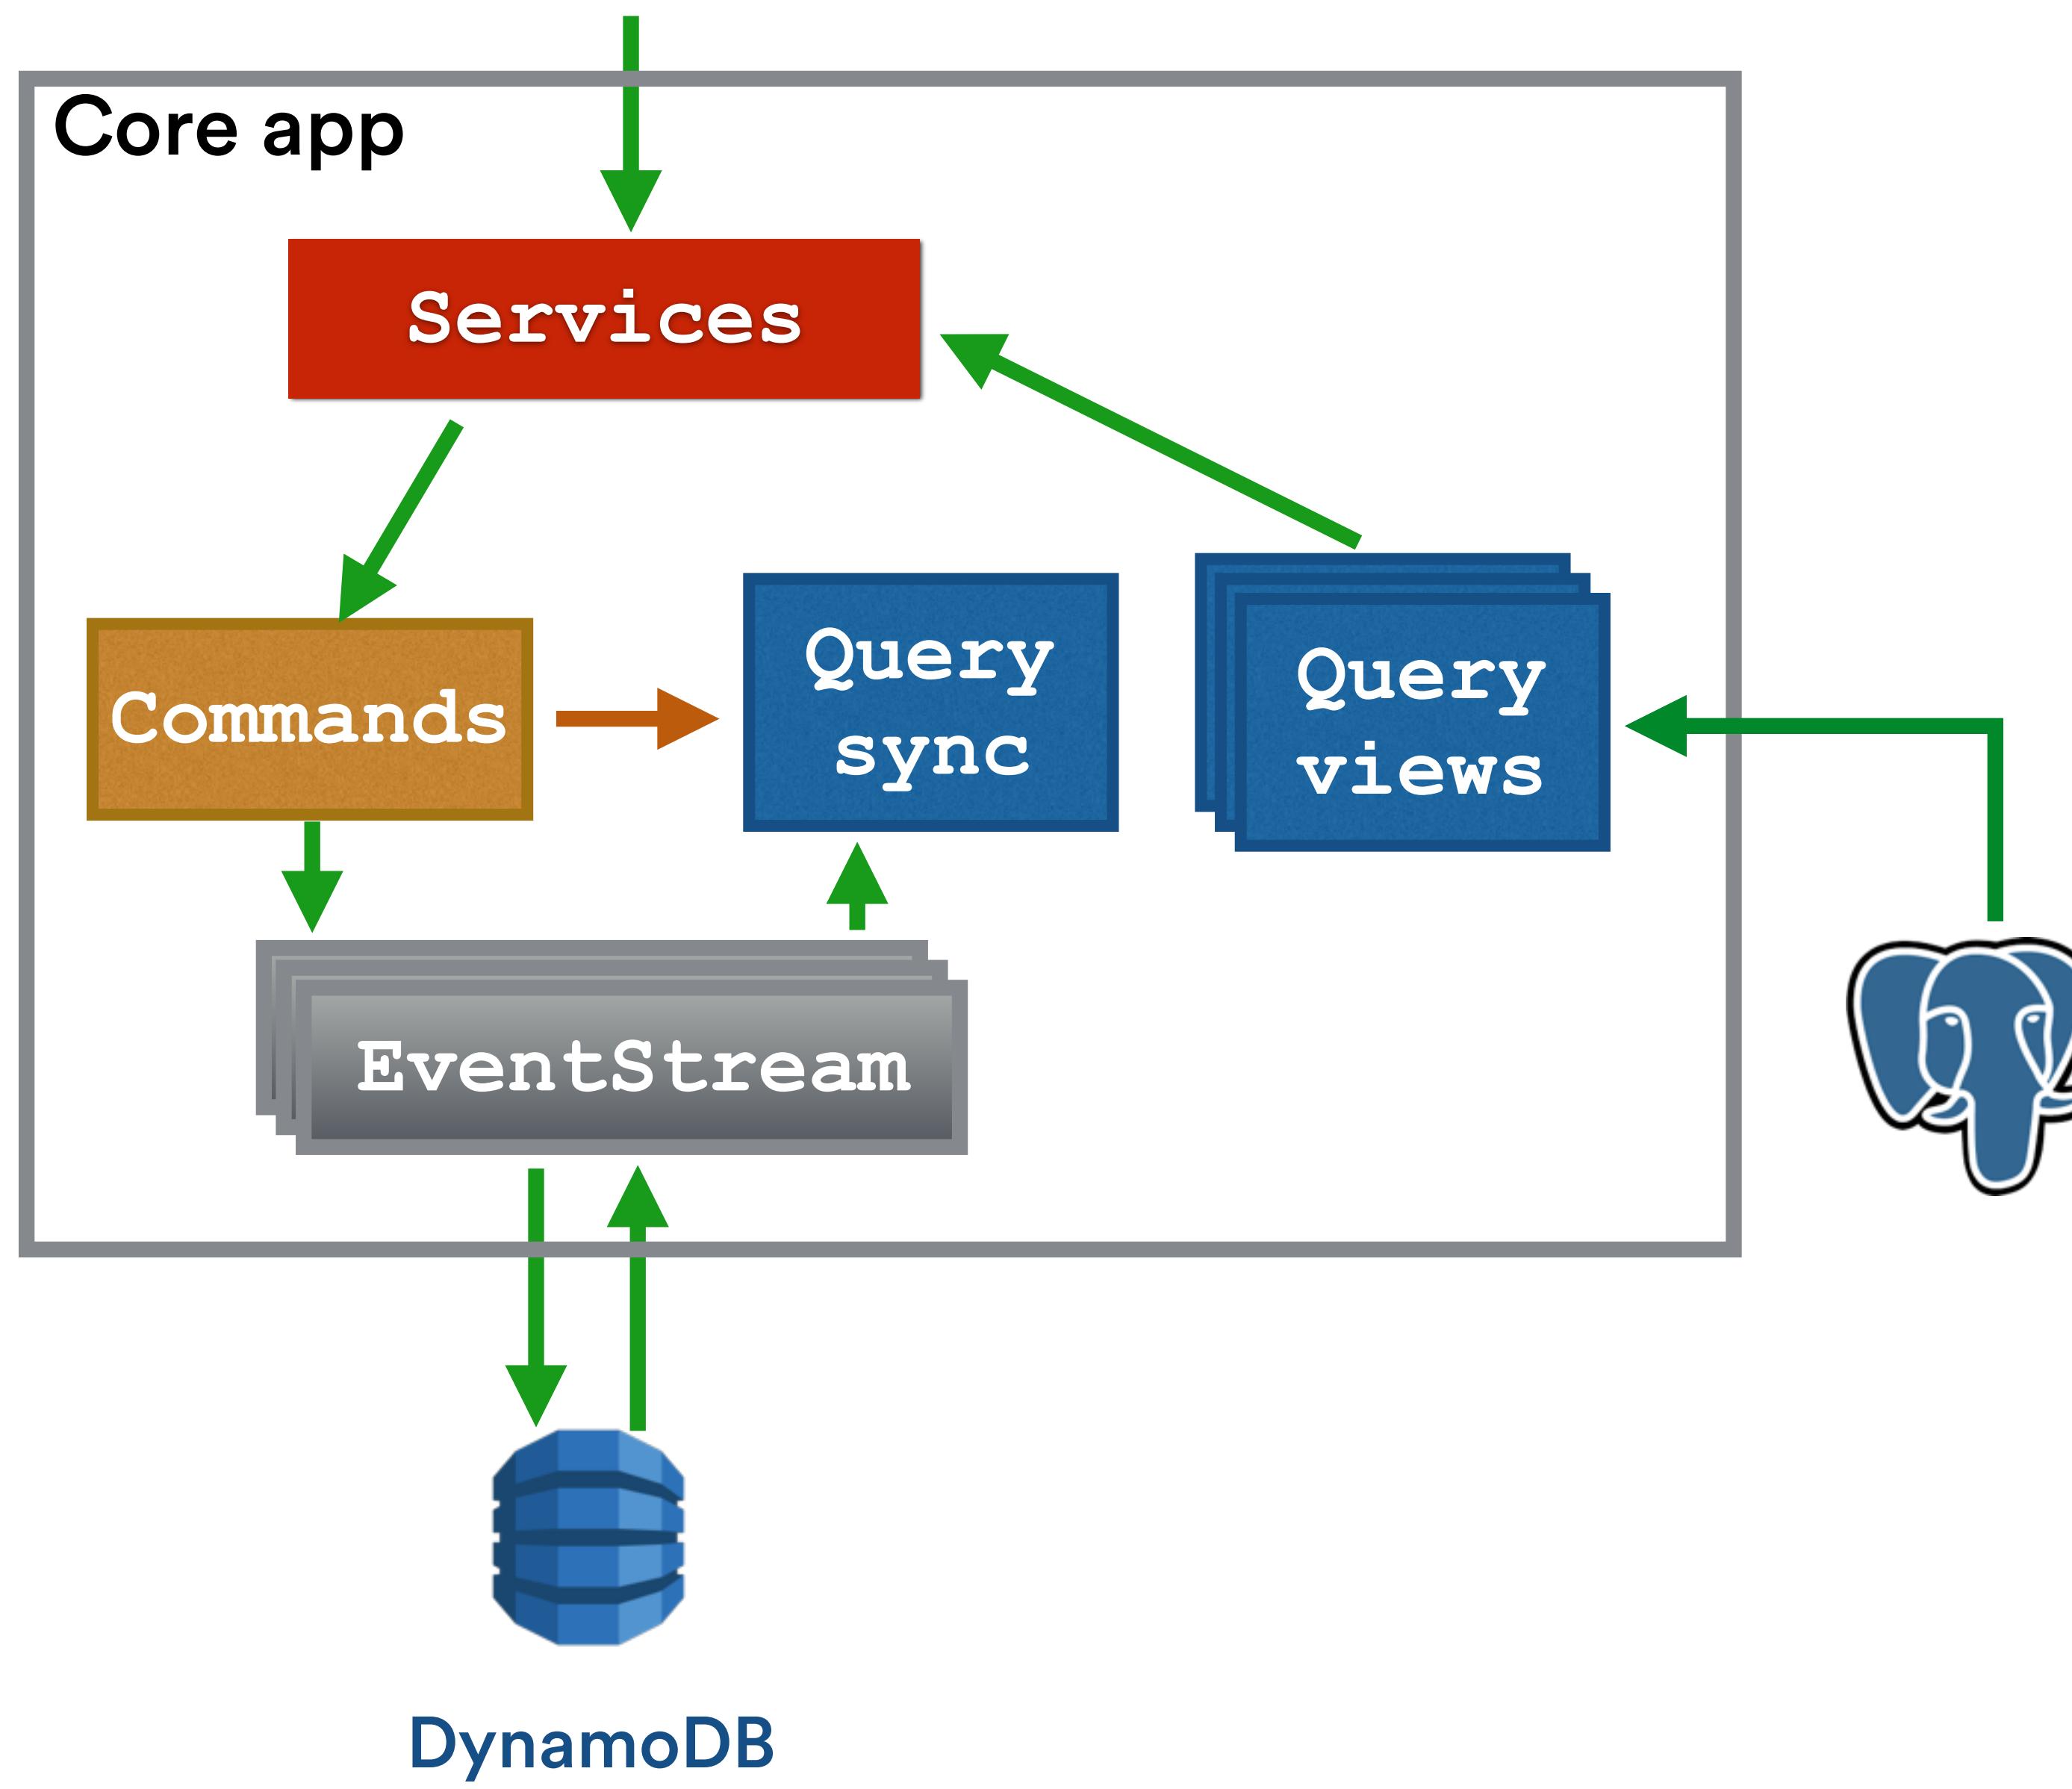
e.g. Add User



**REST calls**  
e.g. Add User

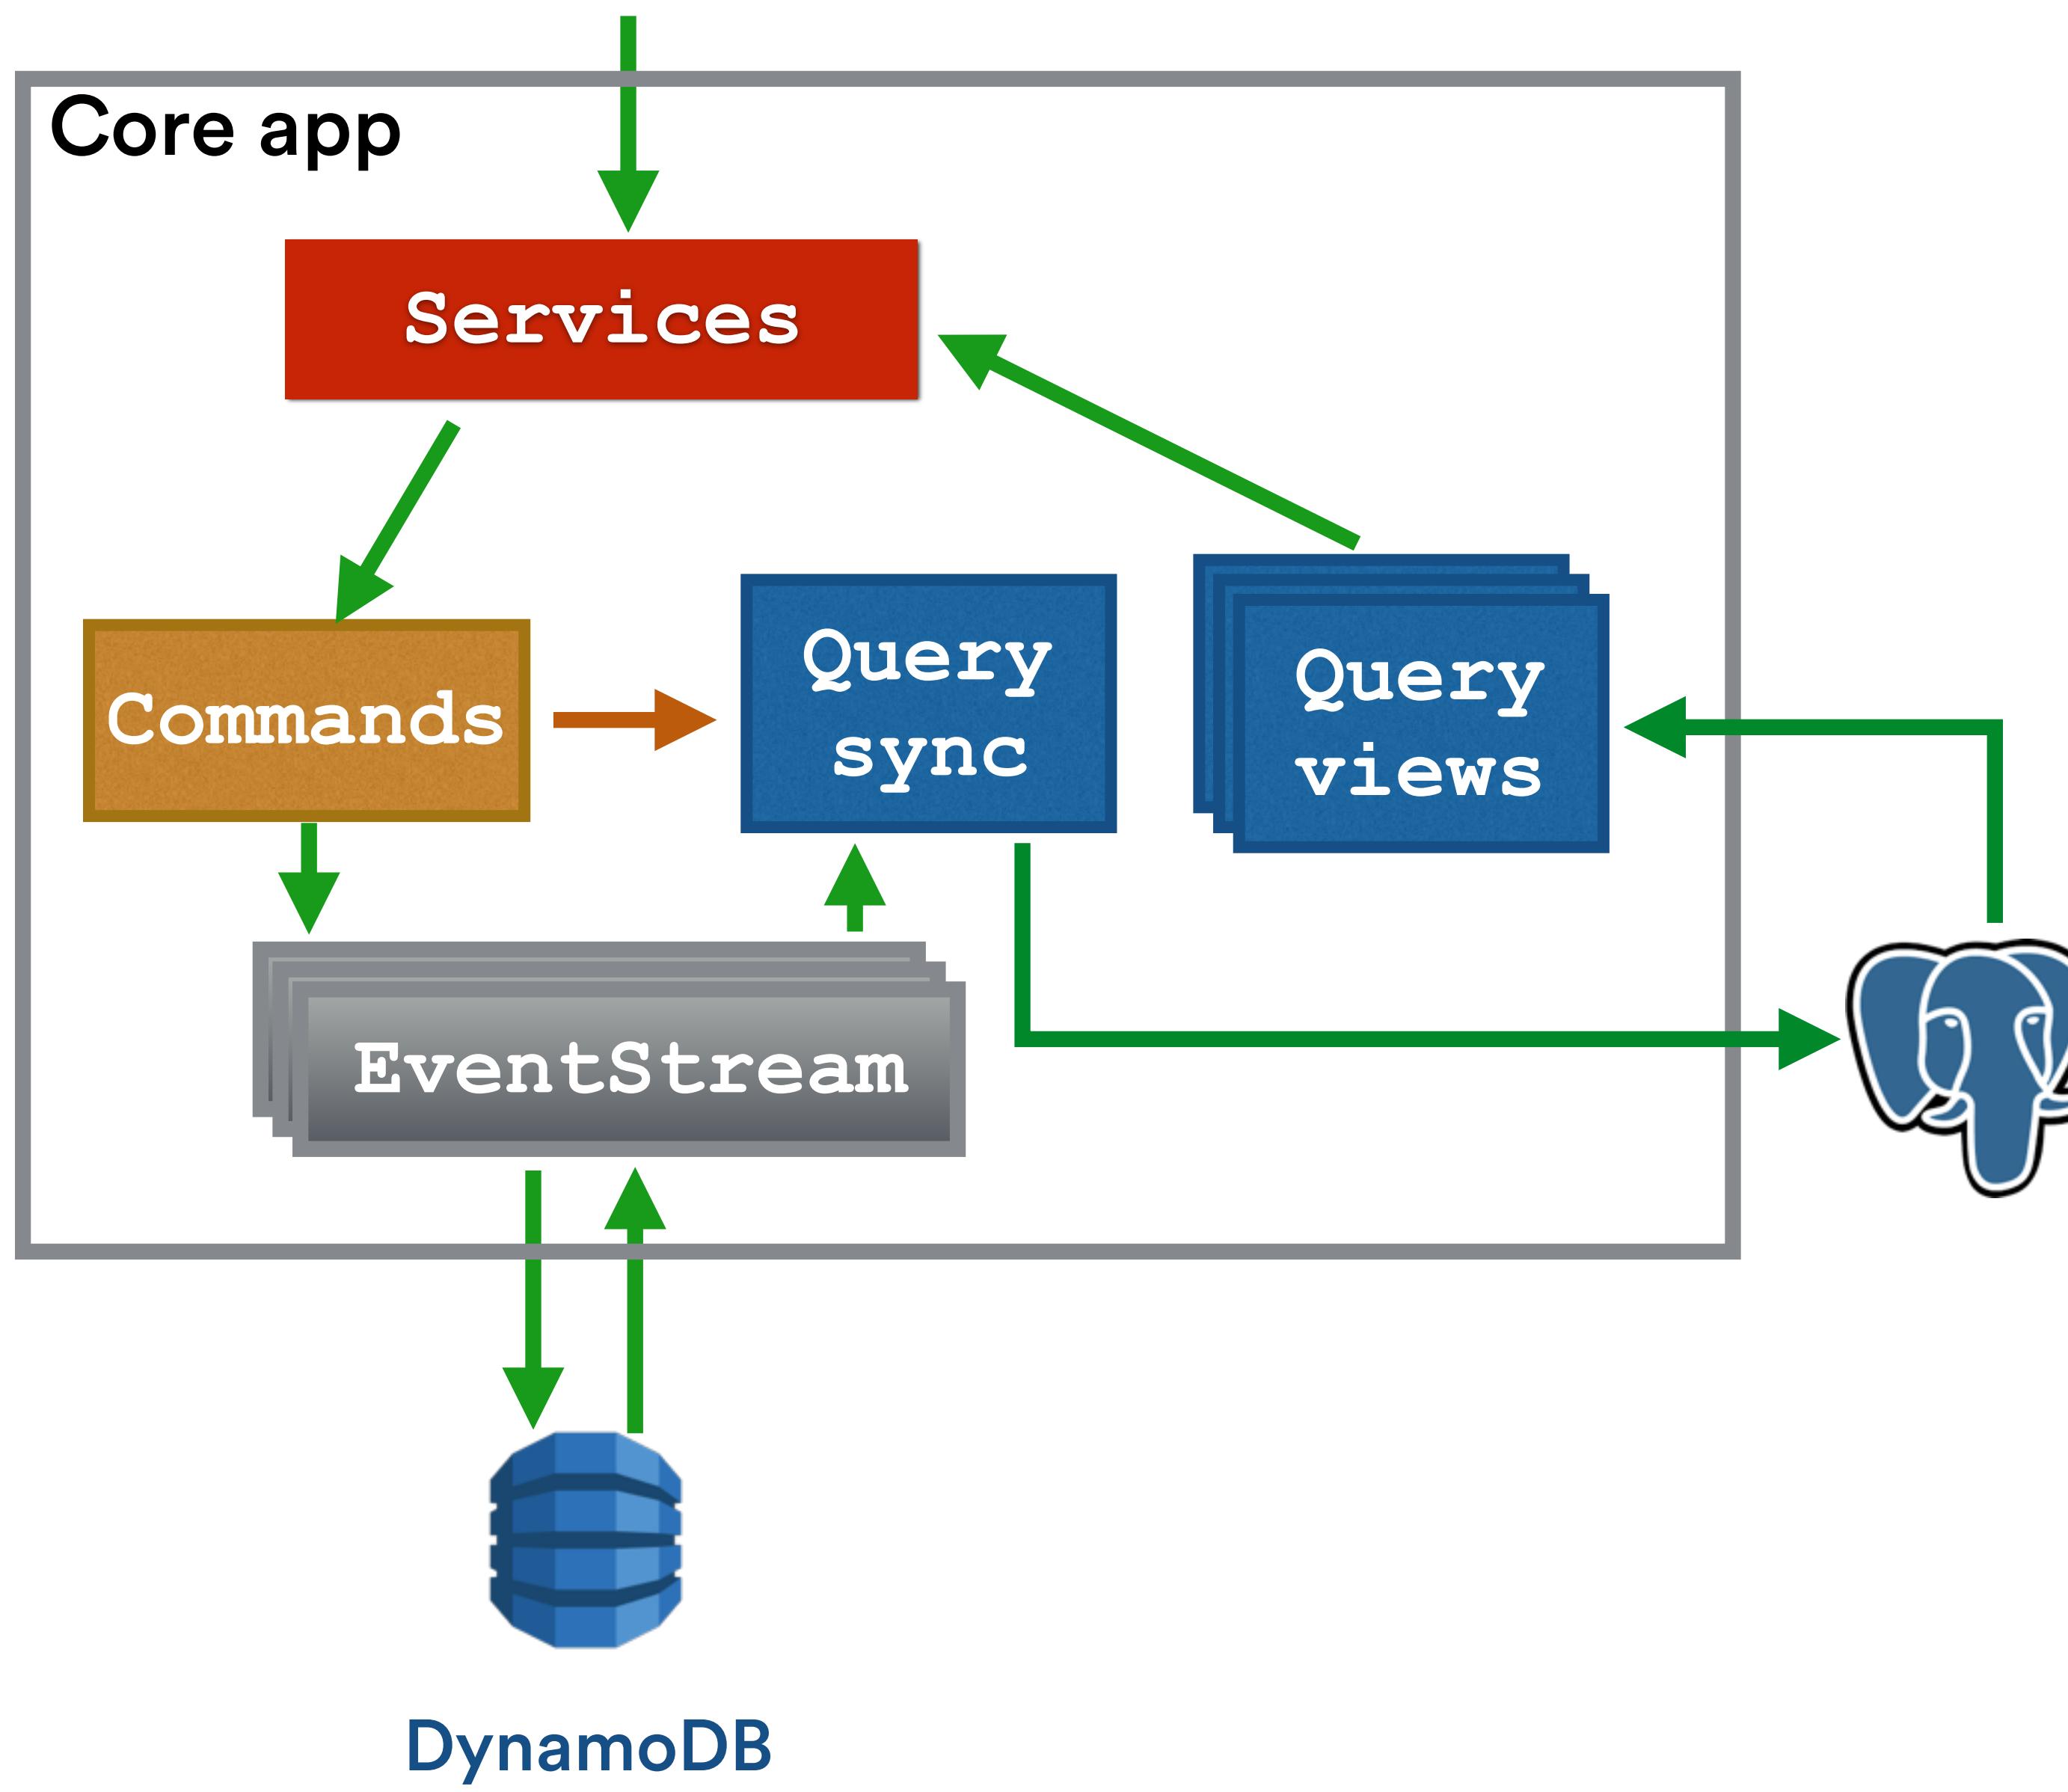


**REST calls**  
e.g. Add User

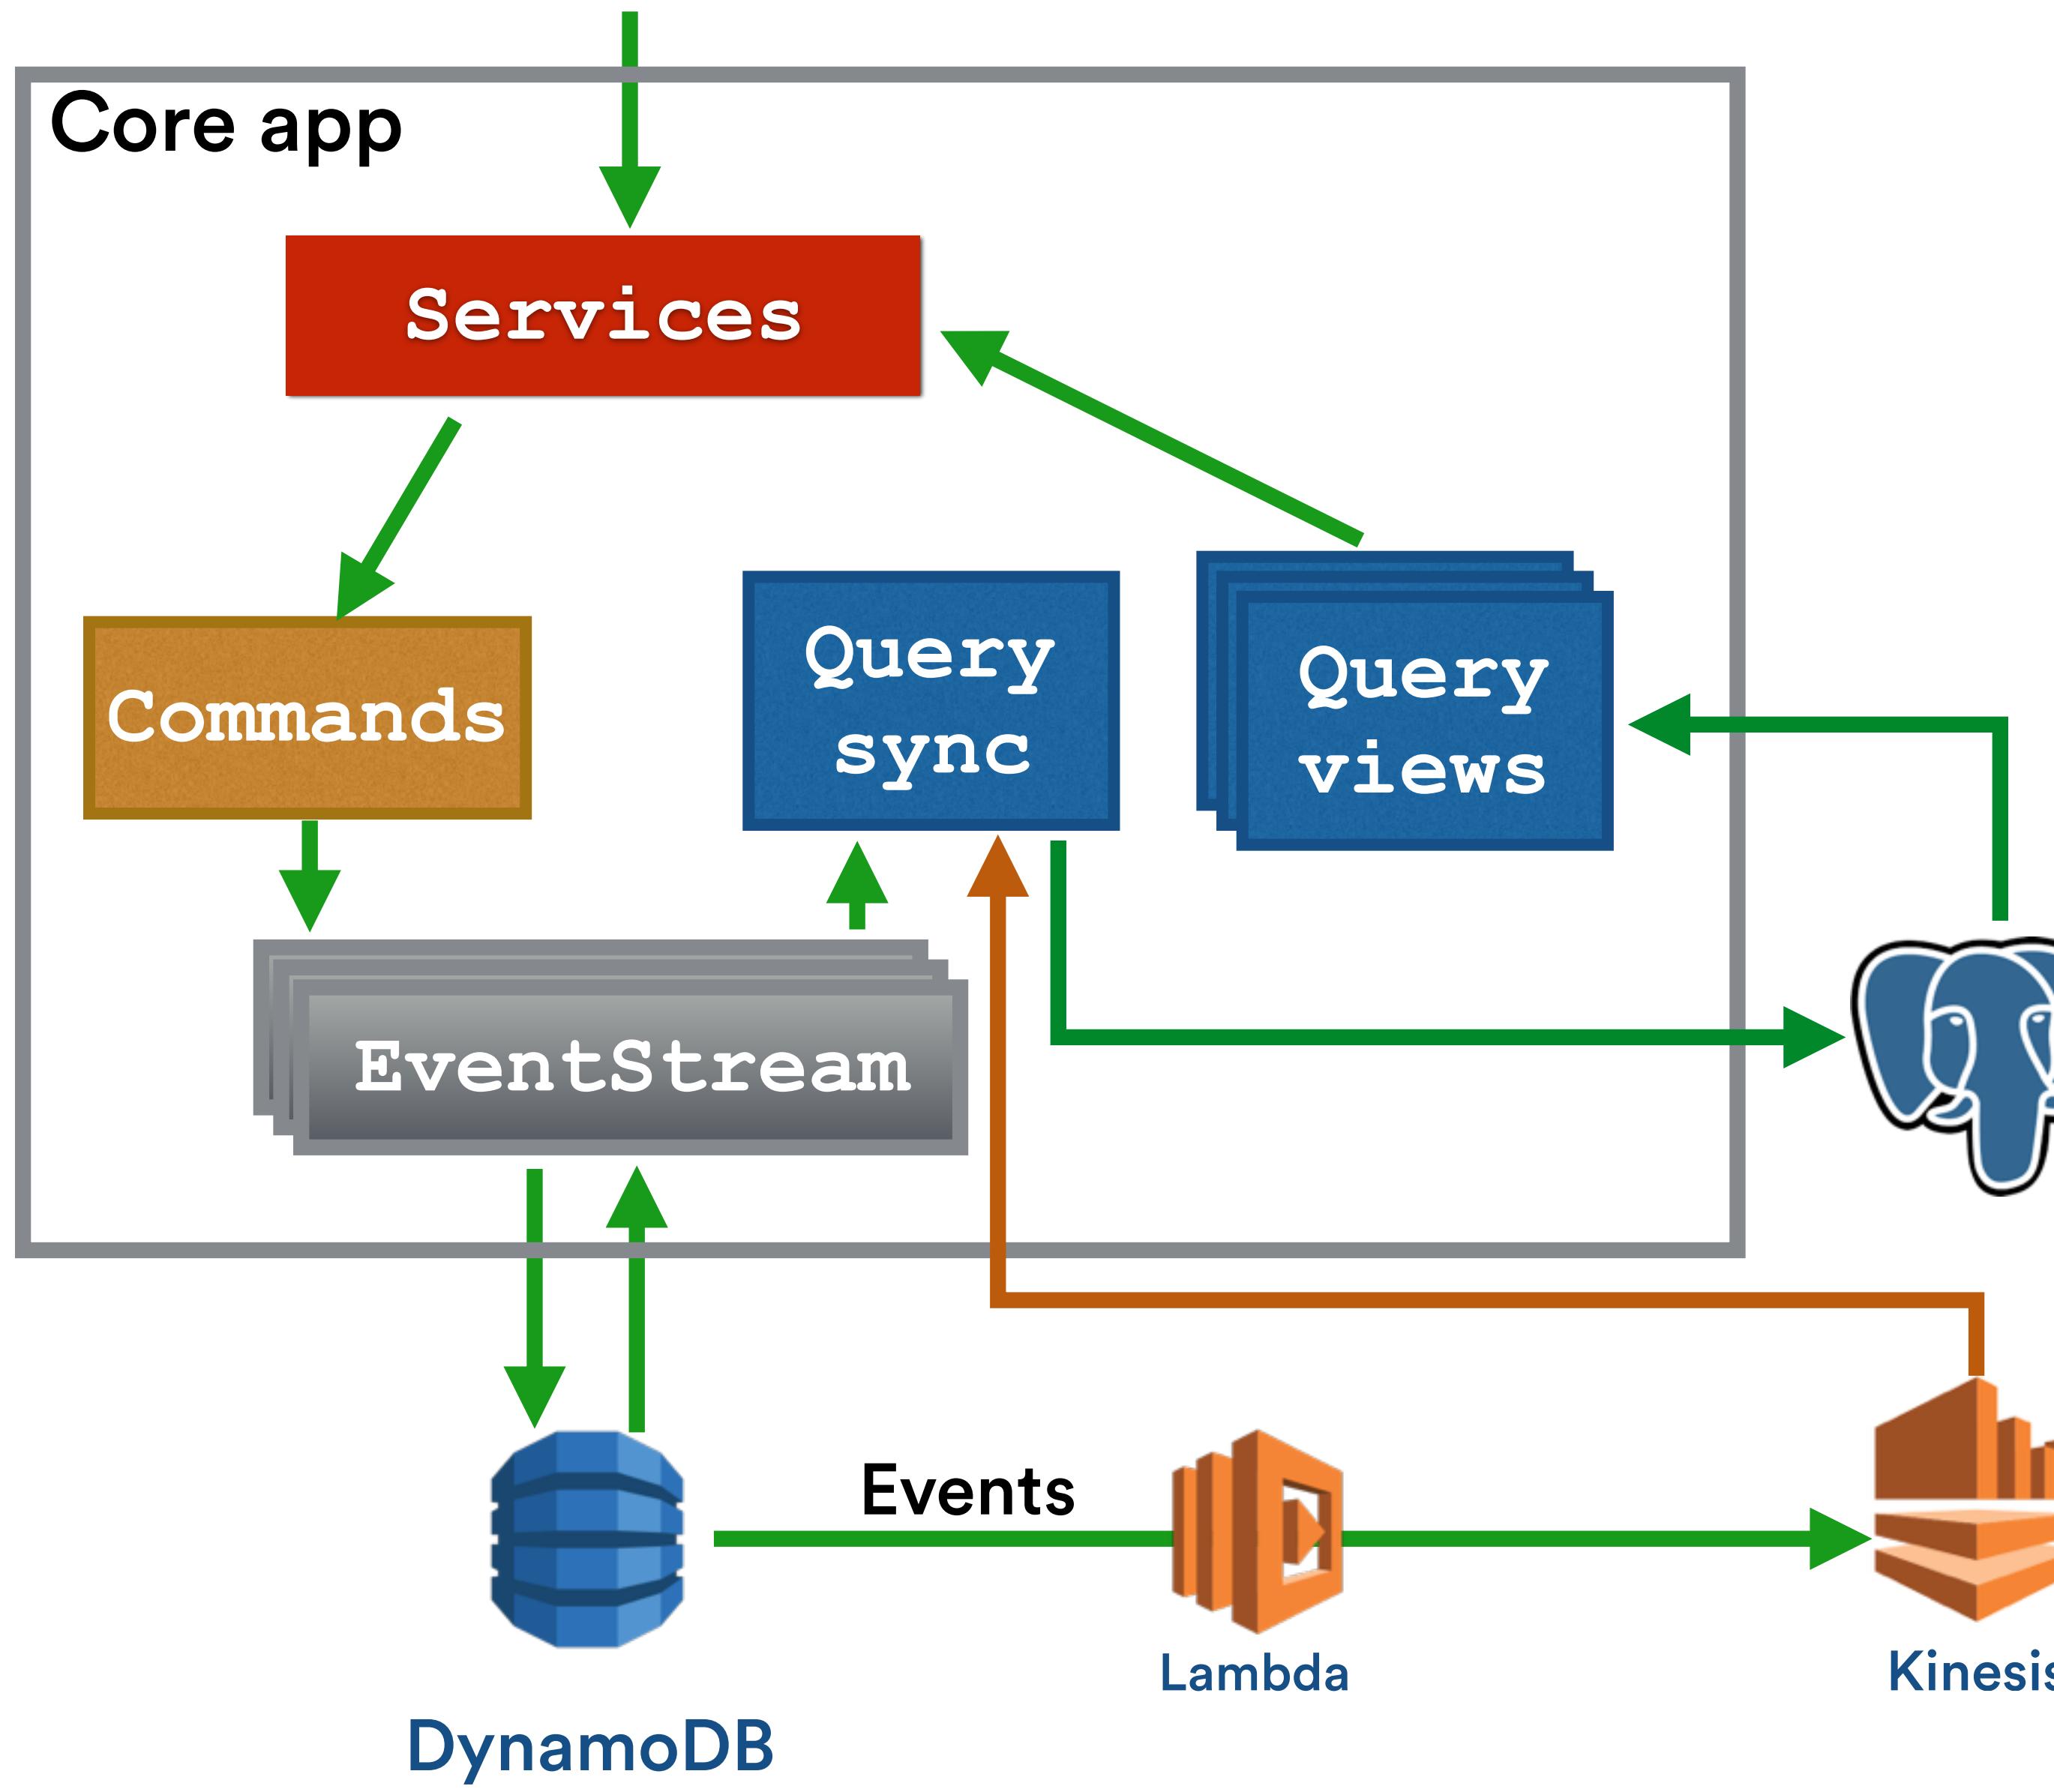


# REST calls

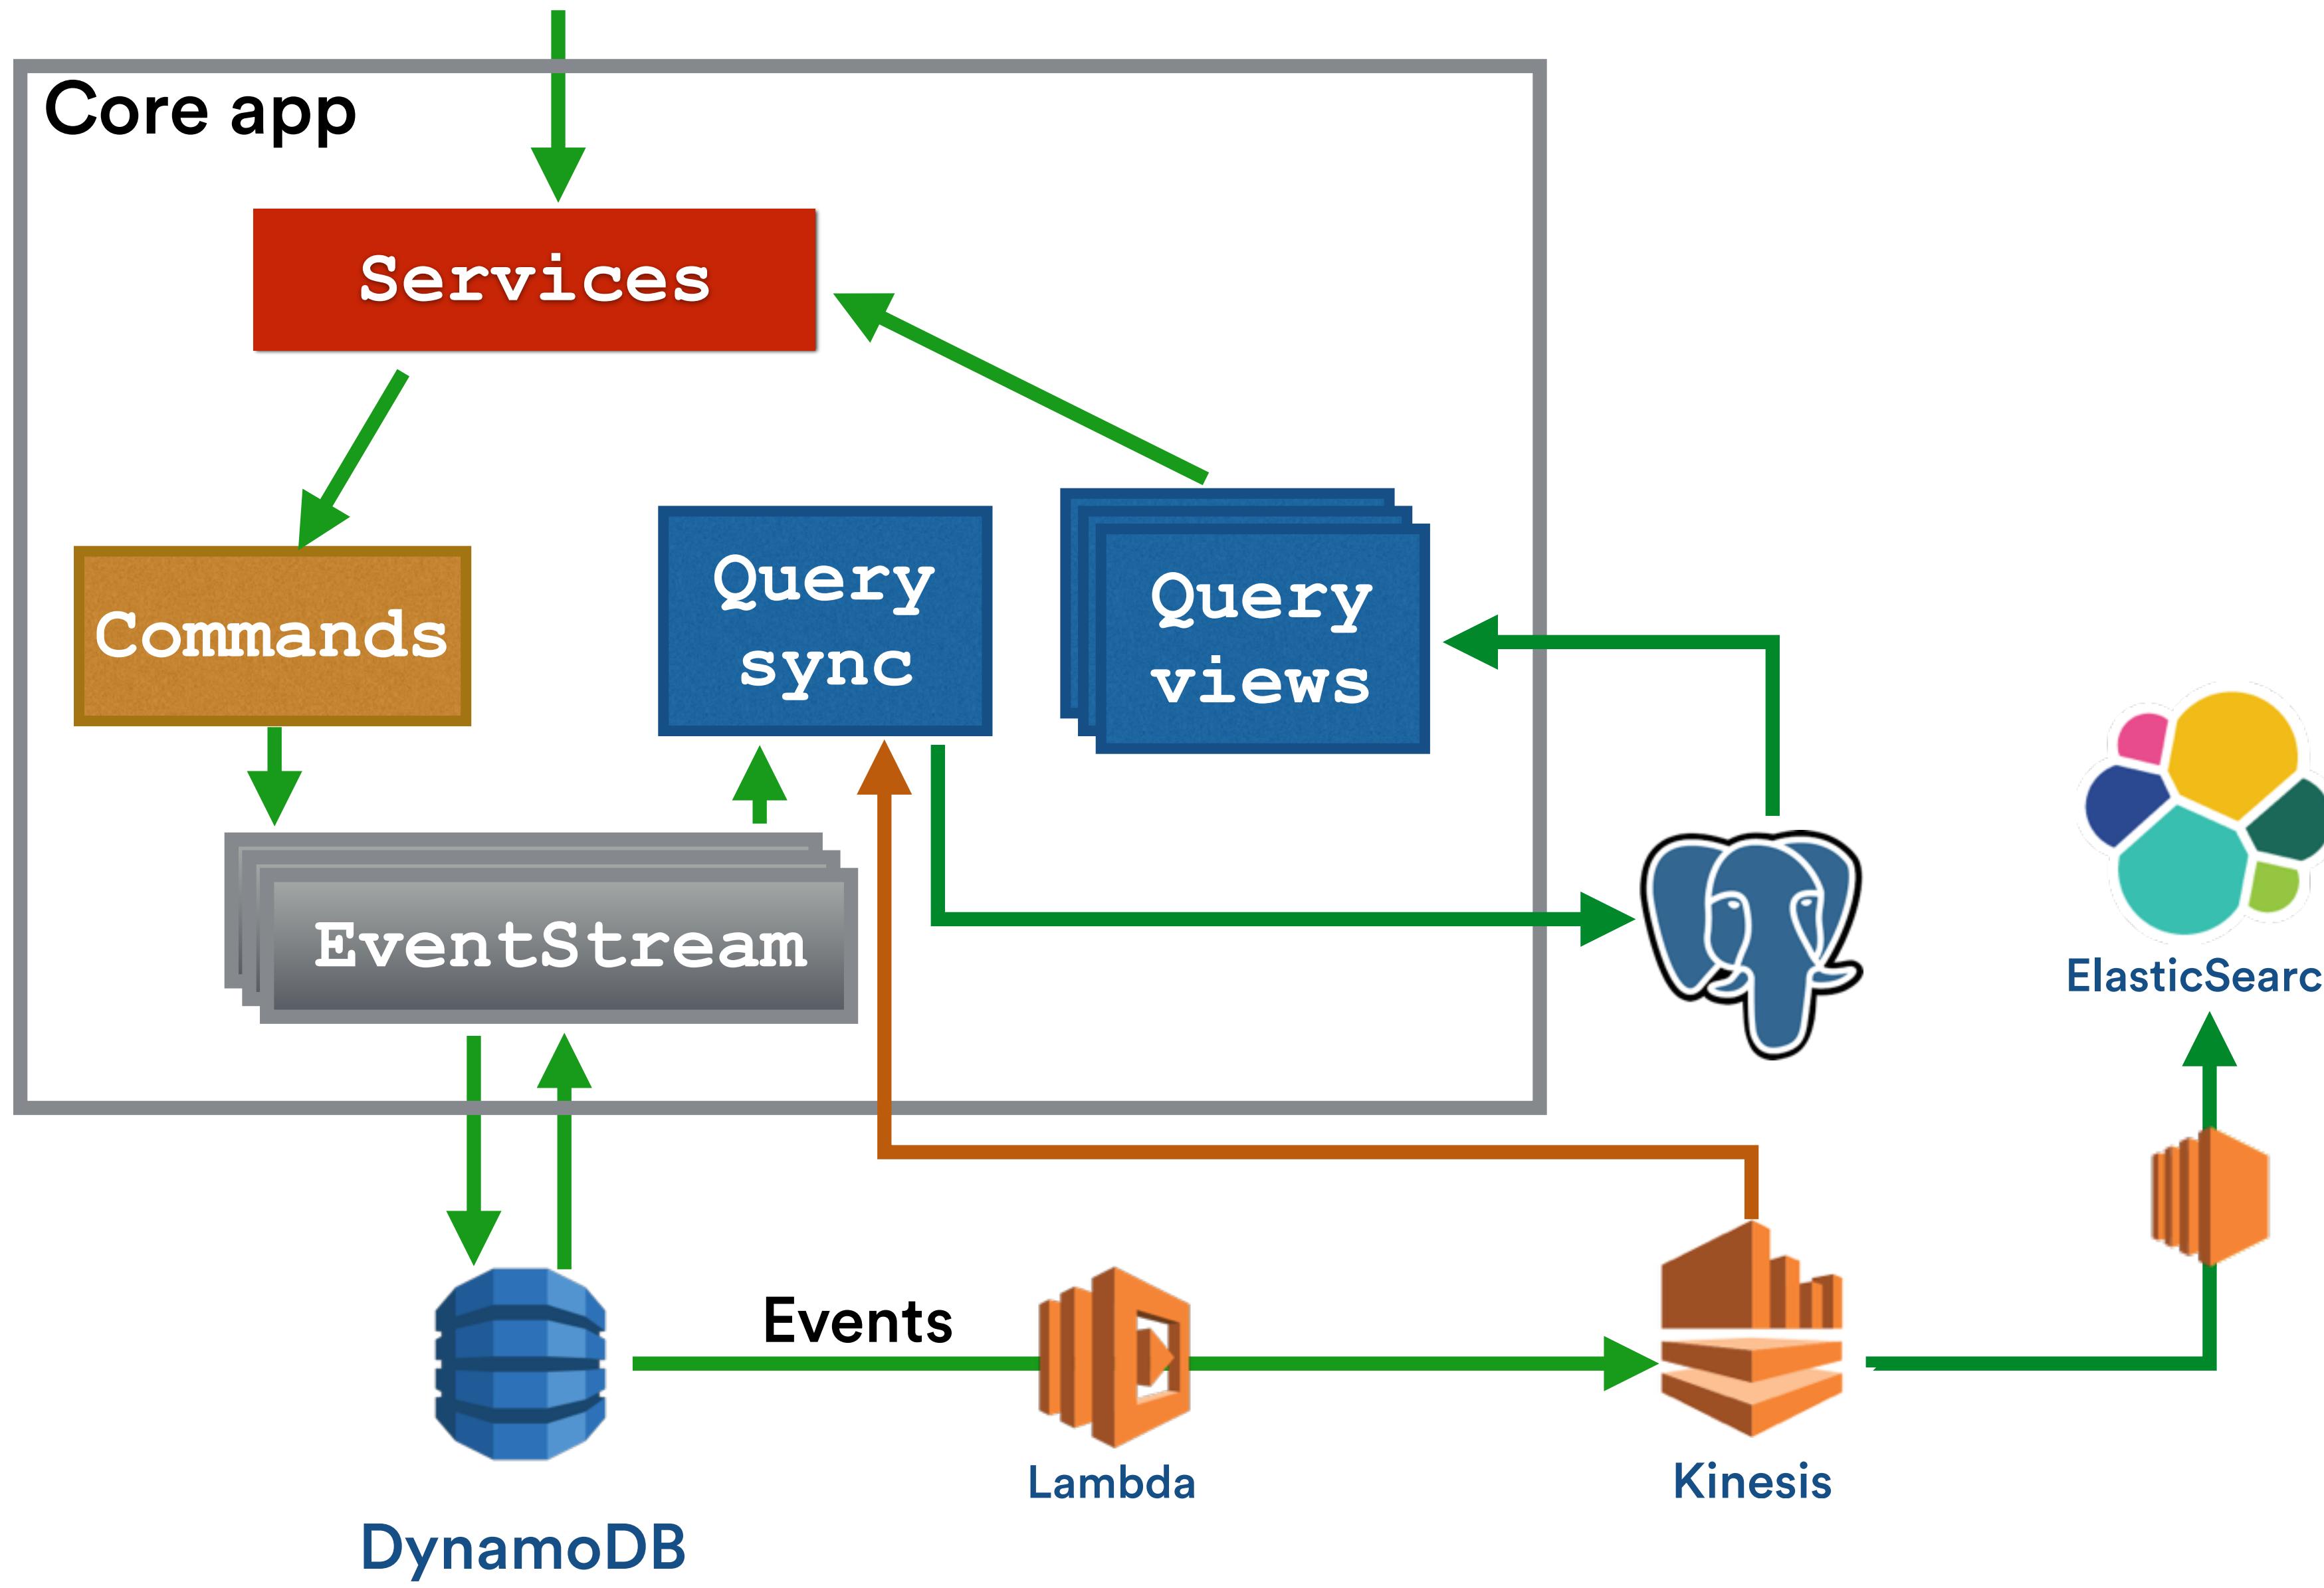
e.g. Add User



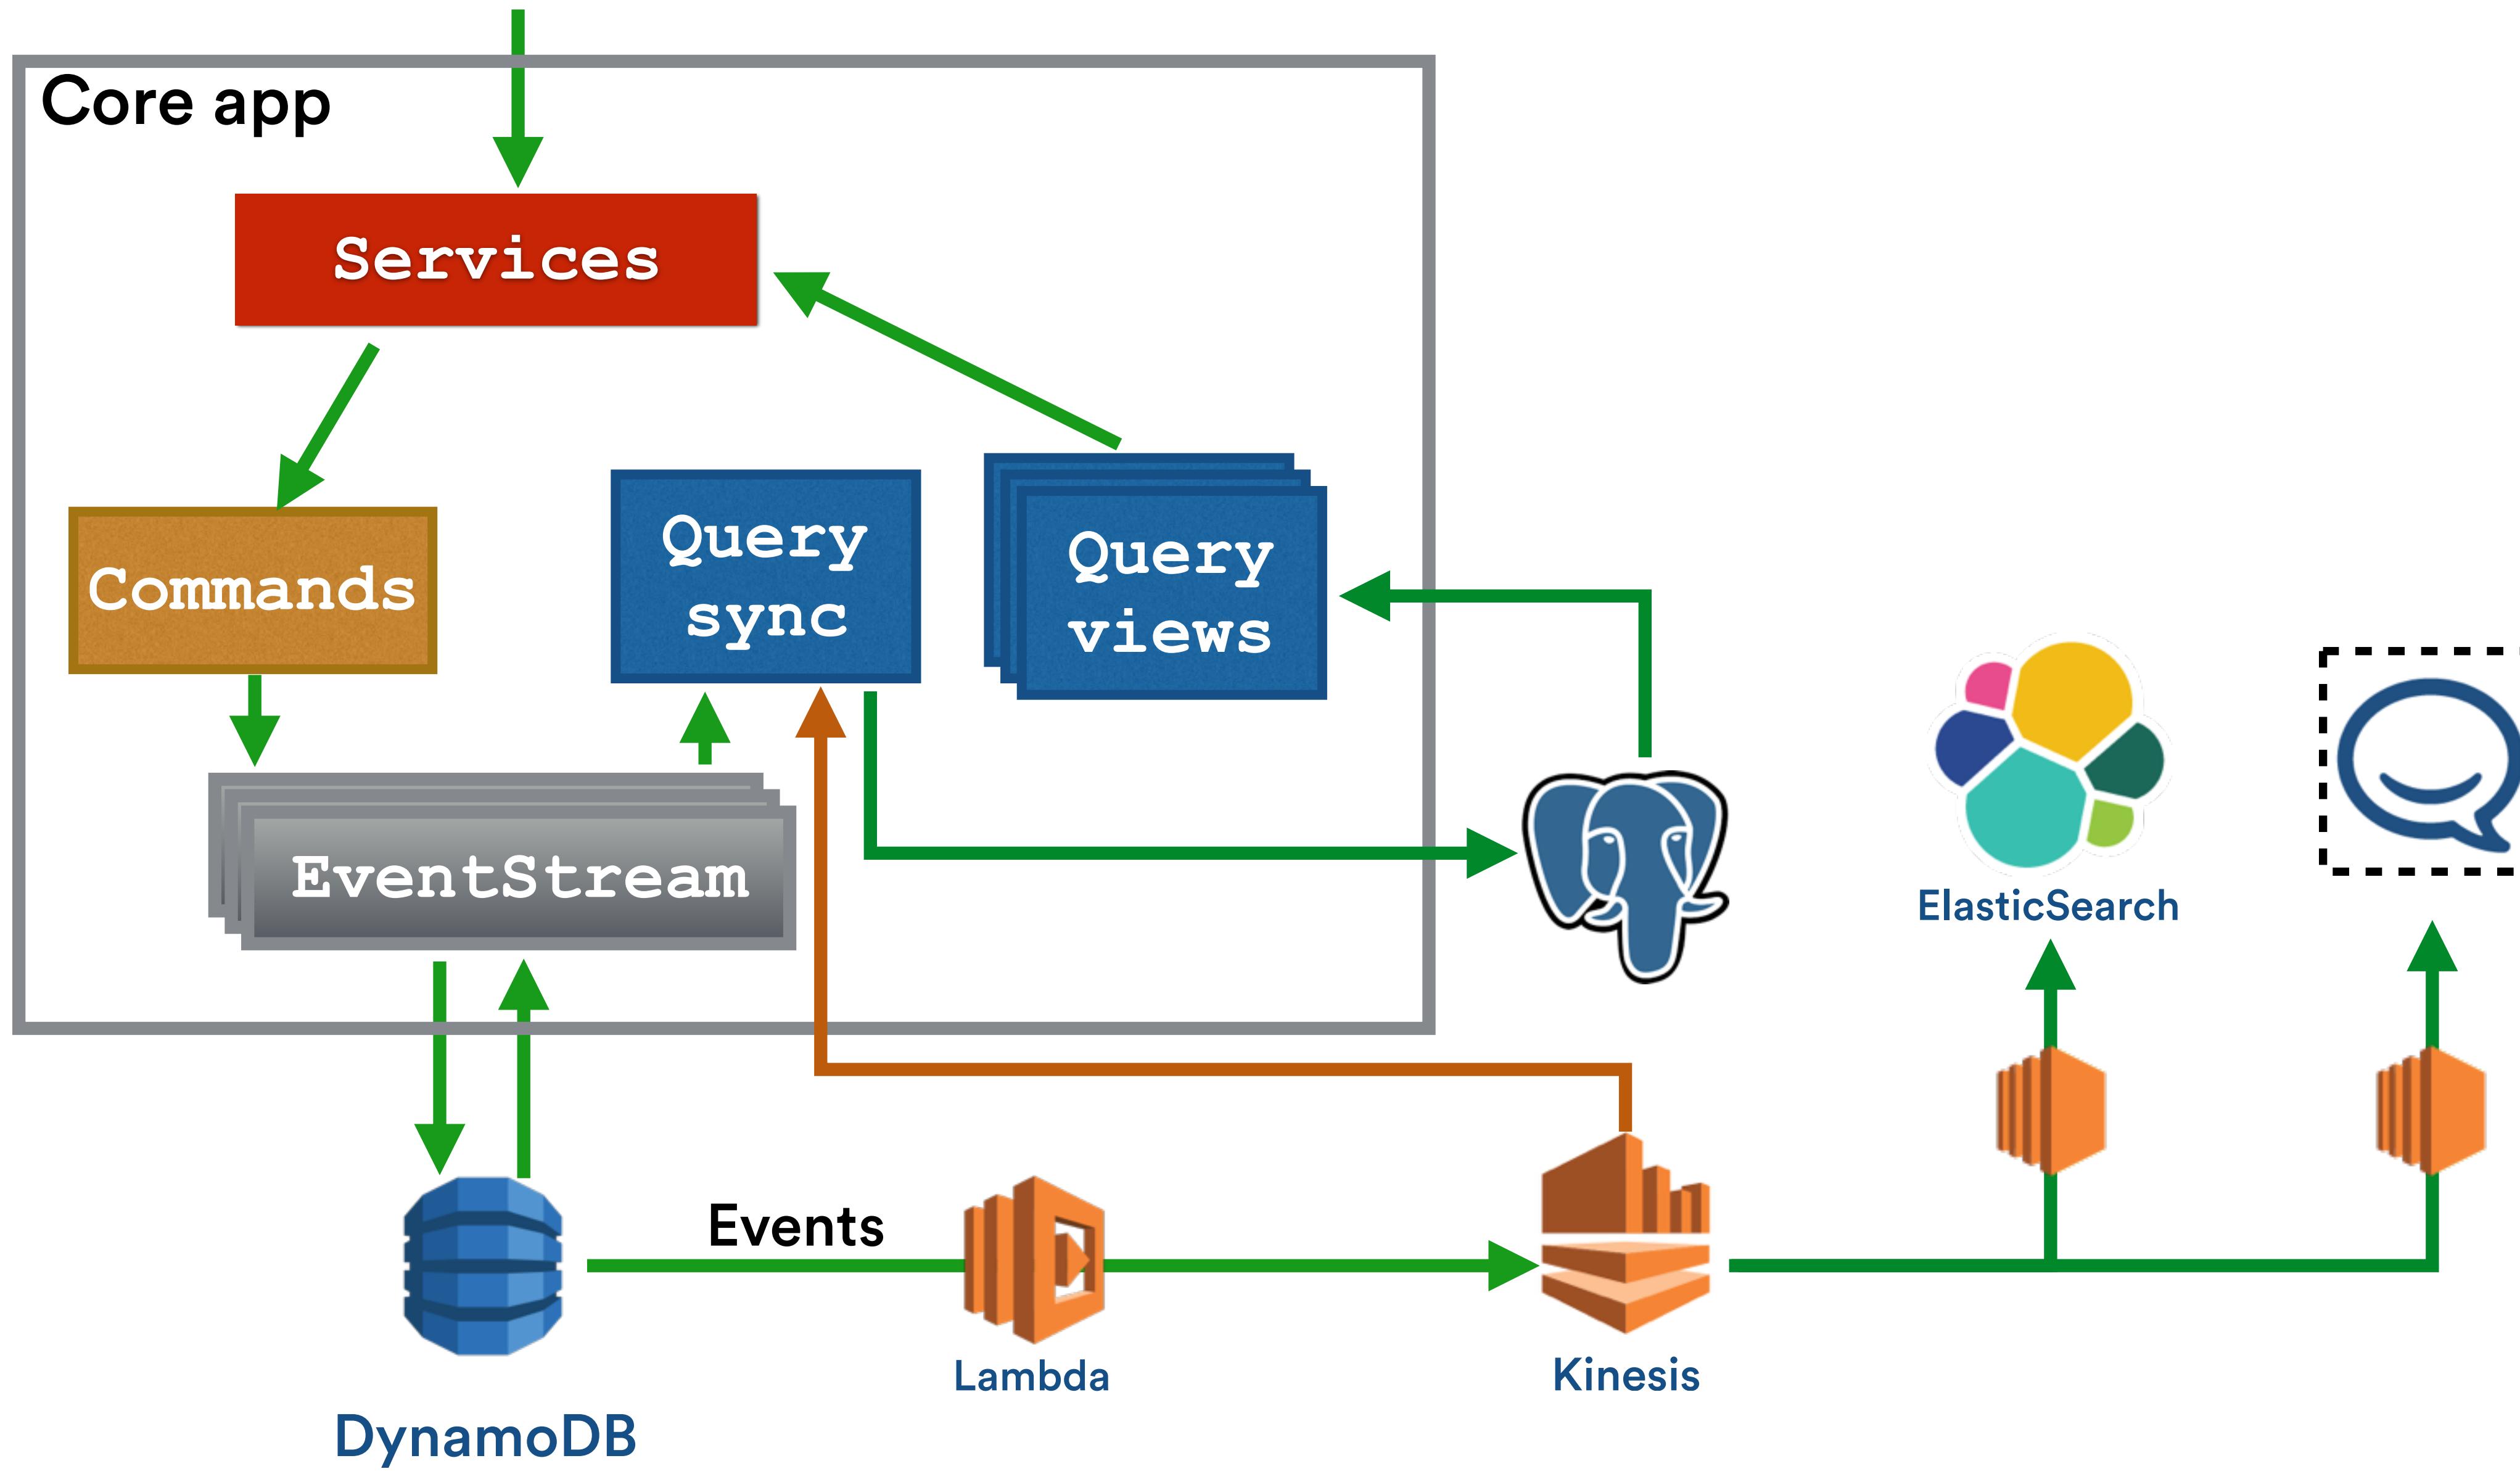
**REST calls**  
e.g. Add User



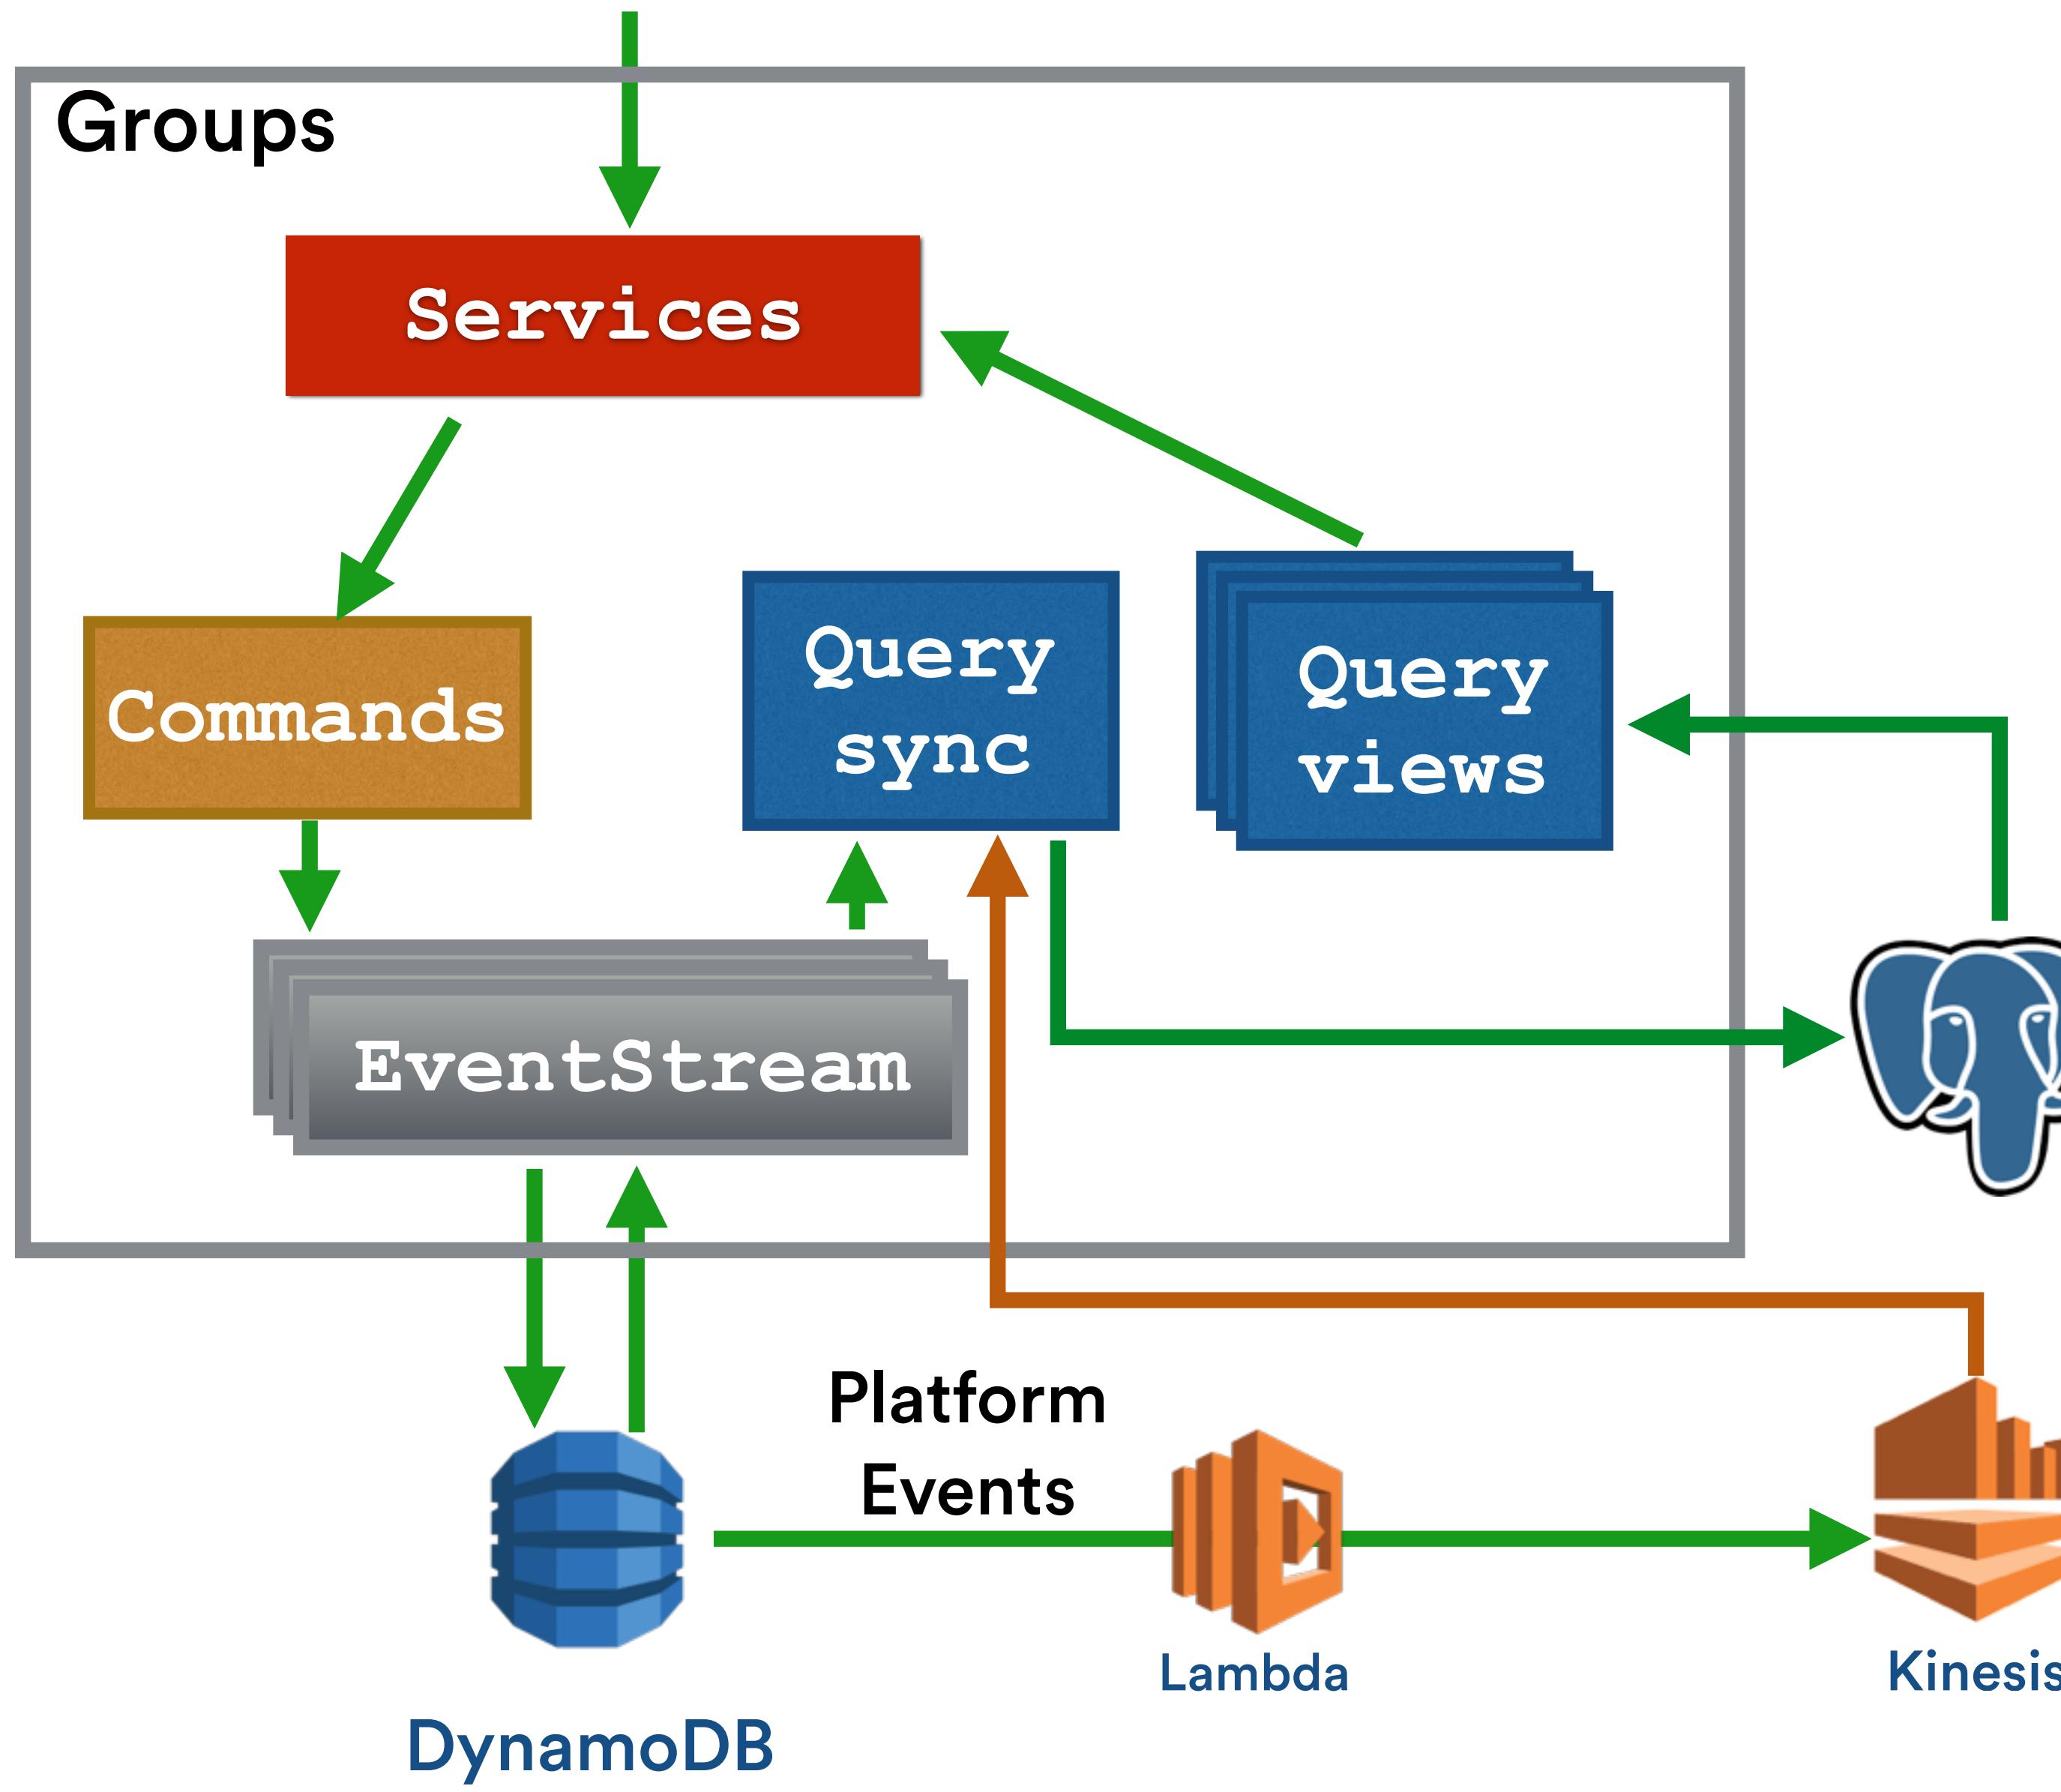
**REST calls**  
e.g. Add User



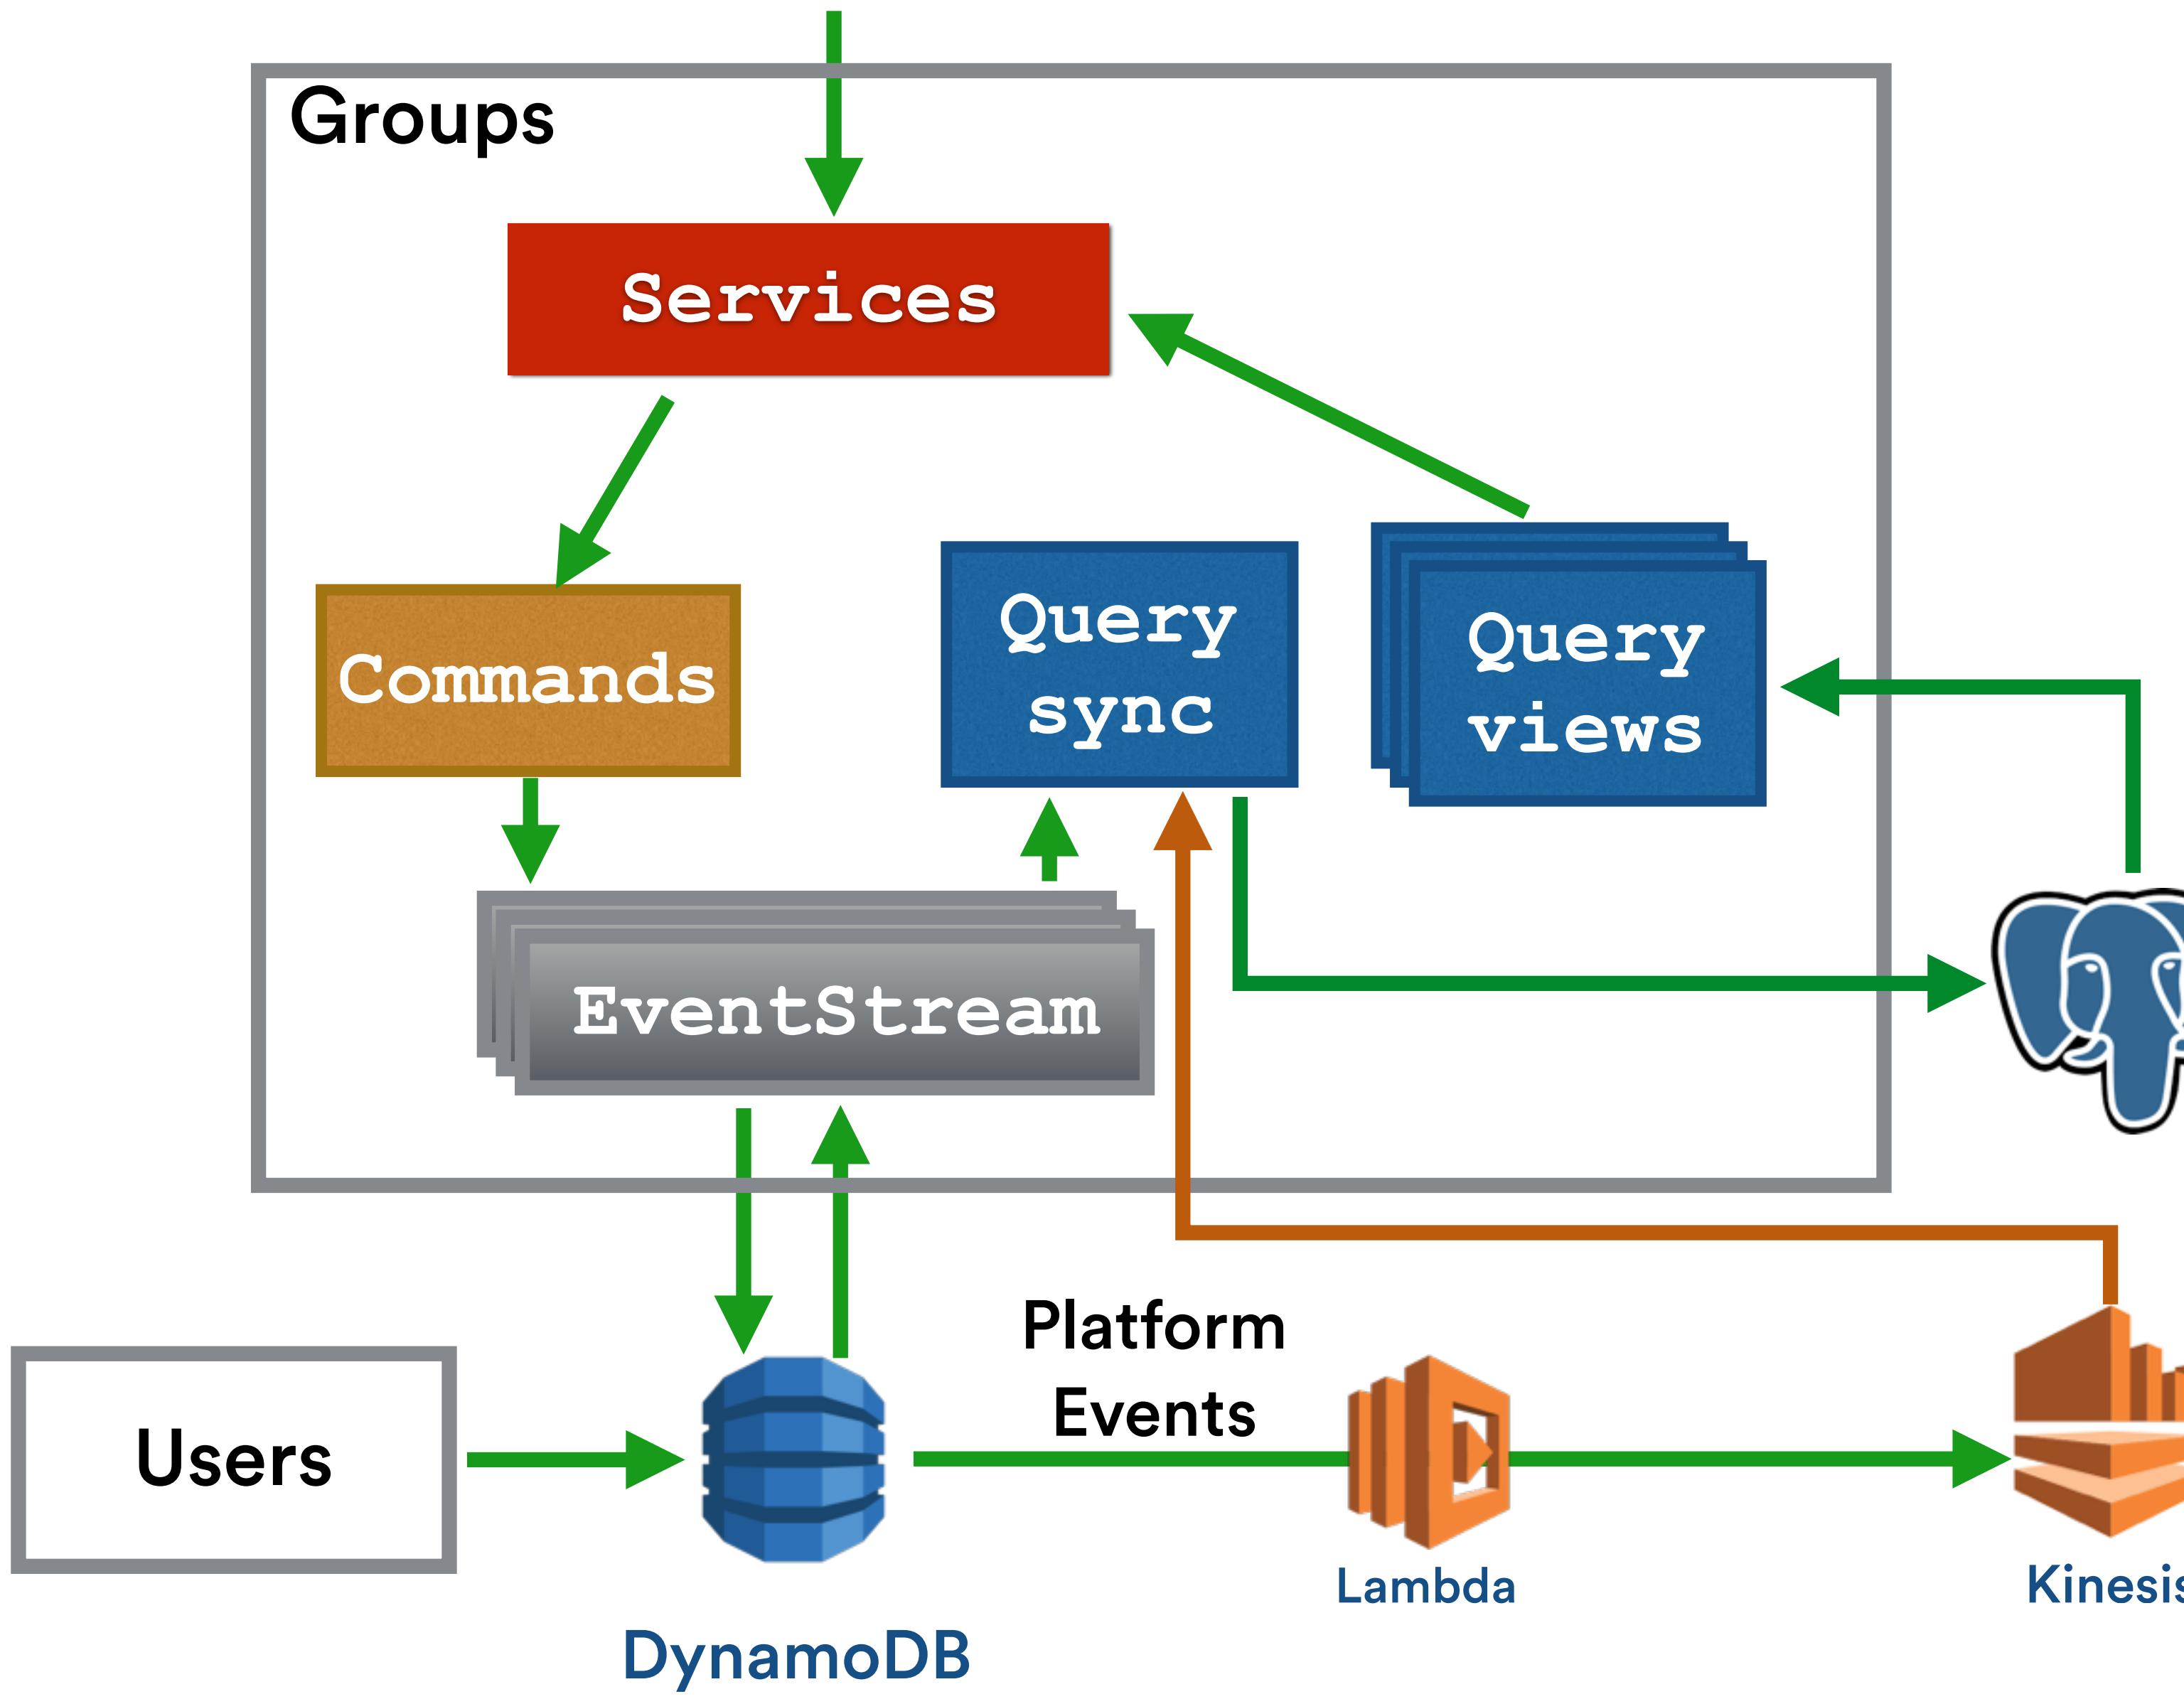
**REST calls**  
e.g. Add User



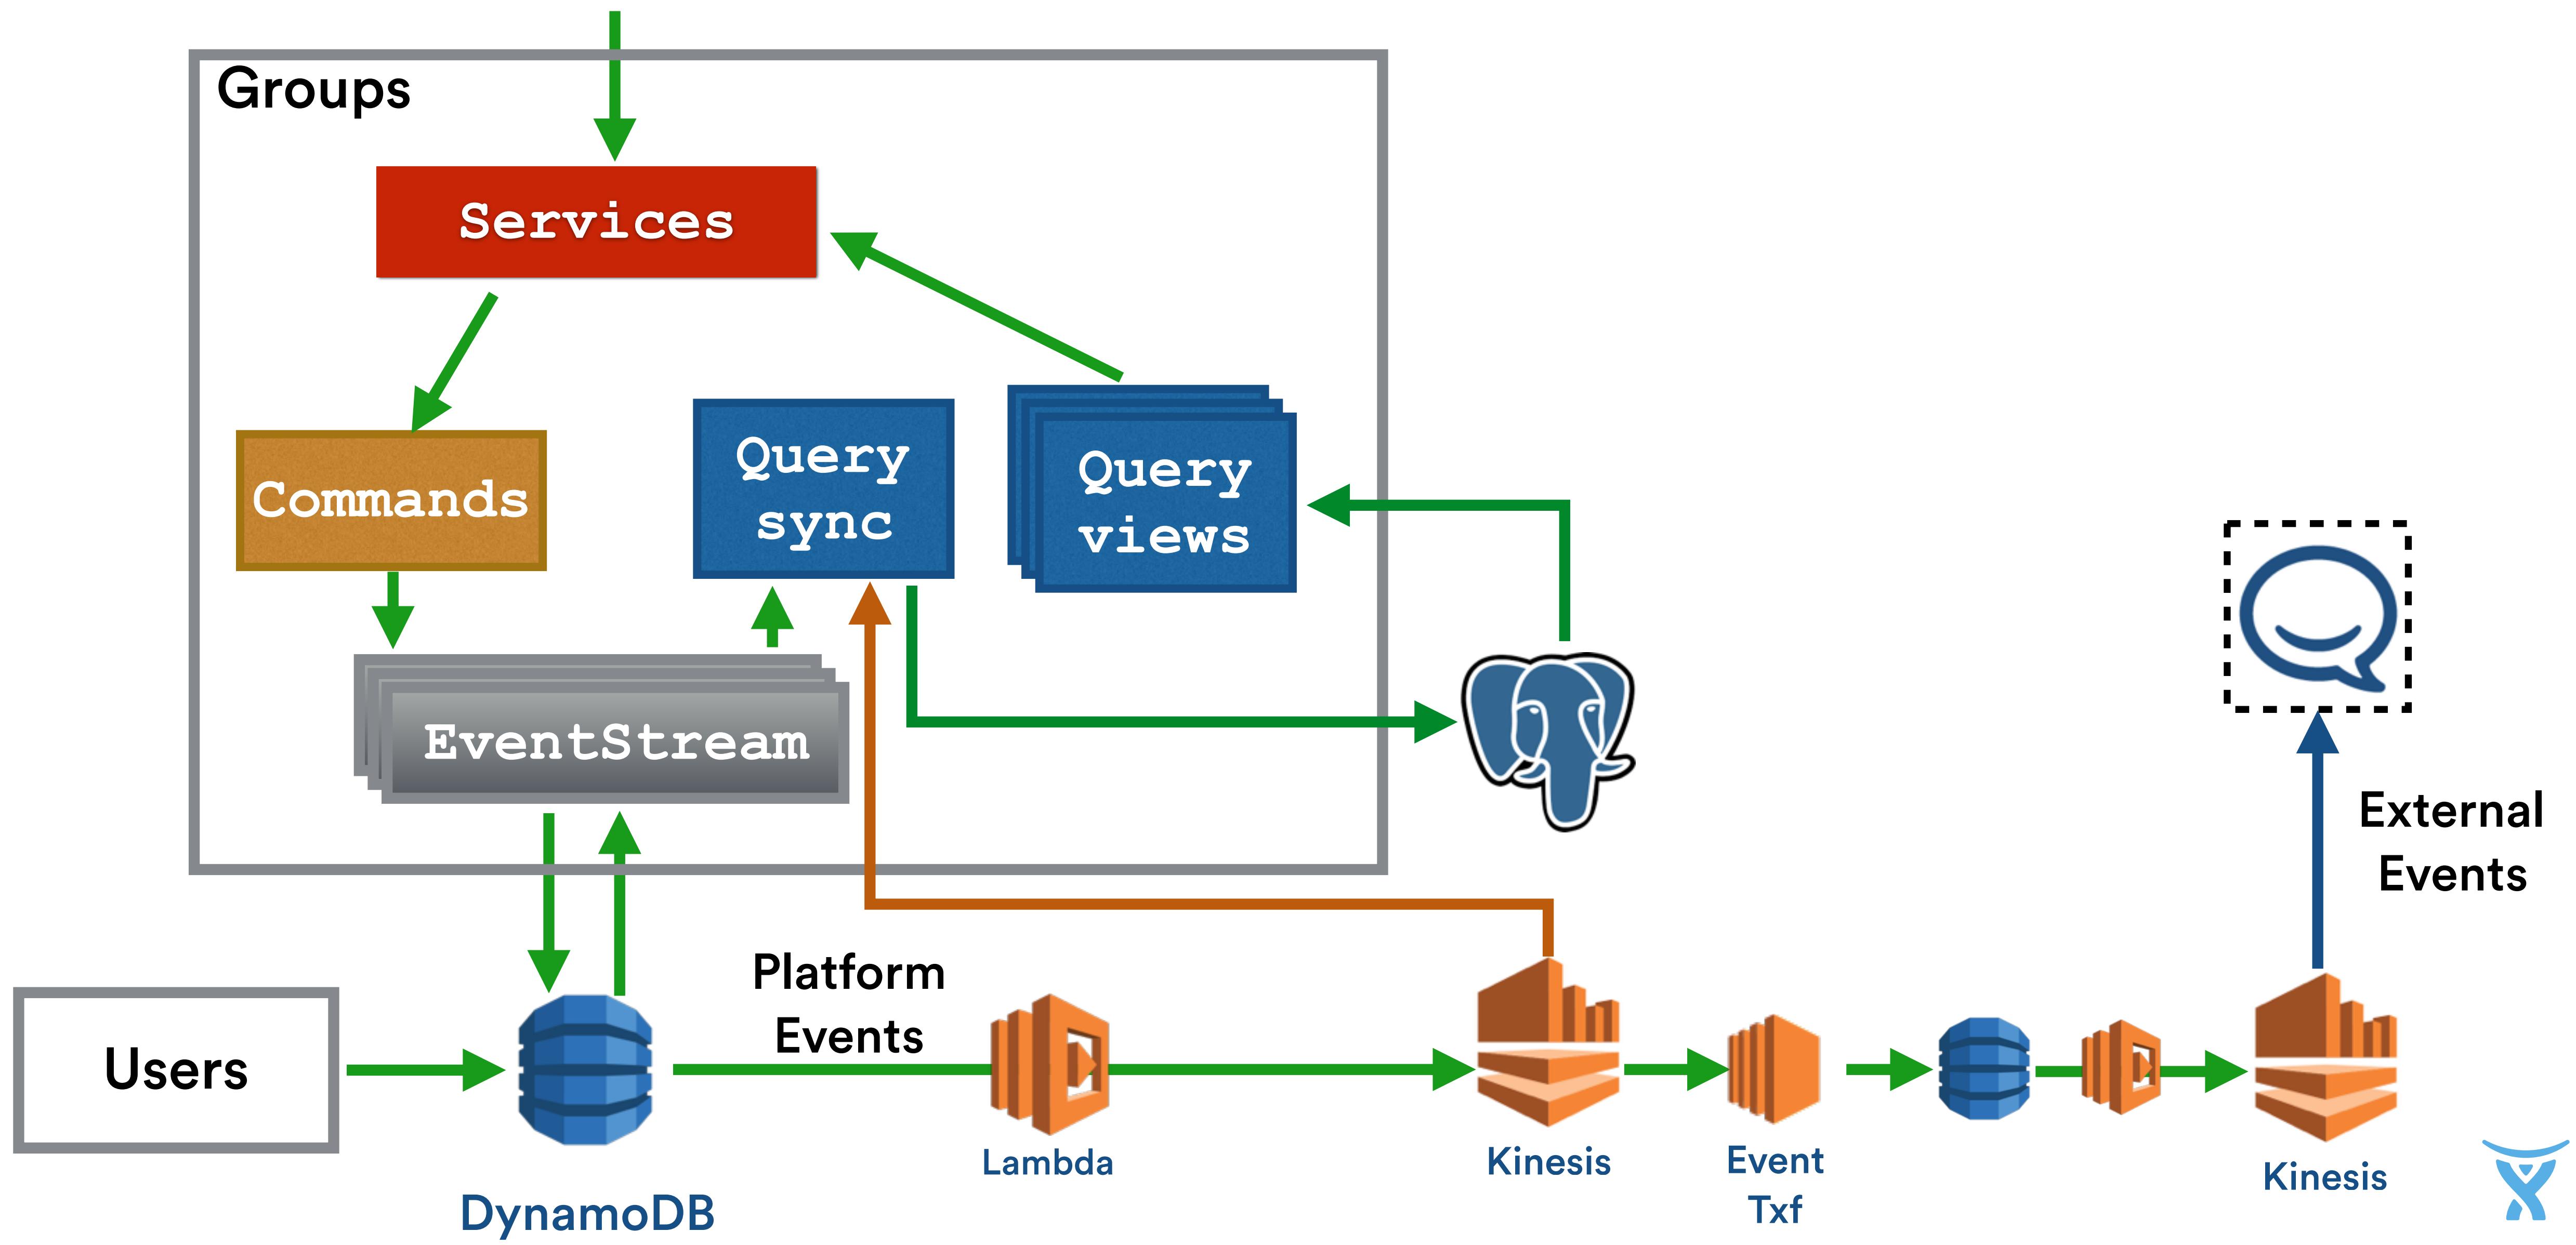
**REST calls**  
e.g. Add User



**REST calls**  
e.g. Add User



**REST calls**  
e.g. Add User



# Delving into code

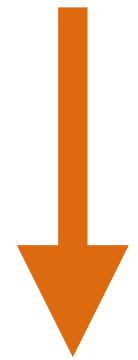
# Modeling Events

Key	Seq	Event	Time
3	123	<b>SetUsername(Maggie)</b>	0
3	124	<b>SetName(Lisa Jr)</b>	10
1	100	<b>SetUsername(homers)</b>	15



# Modeling Events

Ordered ‘Sequence’ (key K + seq S)

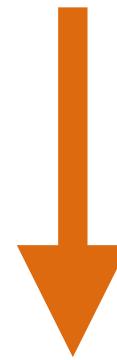


Key	Seq	Event	Time
3	123	<b>SetUsername(Maggie)</b>	0
3	124	<b>SetName(Lisa Jr)</b>	10
1	100	<b>SetUsername(homers)</b>	15



# Modeling Events

Ordered ‘Sequence’ (key K + seq S)



Key	Seq	Event	Time
3	123	<b>SetUsername(Maggie)</b>	0
3	124	<b>SetName(Lisa Jr)</b>	10
1	100	<b>SetUsername(homers)</b>	15



Flexible Value (E)



# Event storage API

```
trait EventStorage[F[_], K, S, E] {  
}  
}
```



# Event storage API

```
trait EventStorage[F[_], K, S, E] {  
  
  def put(event: Event[K, S, E]):  
    F[Error \\  
      Event[K, S, E]]  
  
}
```



# Event storage API

```
trait EventStorage[F[_], K, S, E] {  
  
  def put(event: Event[K, S, E]):  
    F[Error \\  
      Event[K, S, E]]  
  
  def latest(key: K): OptionT[F, Event[K, S, E]]  
}  
}
```



# Event storage API

```
trait EventStorage[F[_], K, S, E] {  
  
  def put(event: Event[K, S, E]):  
    F[Error \\  
      Event[K, S, E]]  
  
  def latest(key: K): OptionT[F, Event[K, S, E]]  
  
  def get(key: K): Stream[Event[K, S, E]]?  
}
```



# Event storage API

```
import scalaz.stream.Process

trait EventStorage[F[_], K, S, E] {

  def put(event: Event[K, S, E]): F[Error \\/ Event[K, S, E]]

  def latest(key: K): OptionT[F, Event[K, S, E]]

  def get(key: K): Process[F, Event[K, S, E]]

}
```



# Saving events



# Saving events

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
  def save(k: K, e: E): F[SaveResult[S]] =  
}  
}
```



# Saving events

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
  def save(k: K, e: E): F[SaveResult[S]] =  
    for {  
      old <- store.latest(k).run  
    } yield ???  
}
```



# Saving events

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
    def save(k: K, e: E): F[SaveResult[S]] =  
        for {  
            old <- store.latest(k).run  
            putResult <- eventStore.put(k, Event.next(old, e))  
        } yield ???  
}
```



# Saving events

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
    def save(k: K, e: E): F[SaveResult[S]] =  
        for {  
            old <- store.latest(k).run  
            putResult <- eventStore.put(k, Event.next(old, e))  
            saveResult <- putResult match {  
                case \/- (ev) =>  
                    SaveResult.success(ev.id.seq)  
            }  
        } yield saveResult  
}
```



# Saving events

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
  def save(k: K, e: E): F[SaveResult[S]] =  
    for {  
      old <- store.latest(k).run  
      putResult <- eventStore.put(k, Event.next(old, e))  
      saveResult <- putResult match {  
        case \/- (ev) =>  
          SaveResult.success(ev.id.seq)  
        case -\/(Error.DuplicateEventId) =>  
          save(k, e) // With back off + jitter  
      }  
    } yield saveResult  
}
```



# Saving events

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
    def save(k: K, e: E): F[SaveResult[S]] =  
        for {  
            old <- store.latest(k).run  
            putResult <- eventStore.put(k, Event.next(old, e))  
            saveResult <- putResult match {  
                case \/- (ev) =>  
                    SaveResult.success(ev.id.seq)  
                case -\/(Error.DuplicateEventId) =>  
                    save(k, e) // With back off + jitter  
                case -\/(Error.Rejected(reasons)) =>  
                    SaveResult.reject(reasons)  
            }  
        } yield saveResult  
}
```



# What about data constraints?



# Operation: Constraint as a type

```
case class Operation[S, E] (apply: Option[S] => OpResult[E])
```



# Operation: Constraint as a type

```
case class Operation[S, E] (apply: Option[S] => OpResult[E])
```

```
sealed trait OpResult[E]
```

```
case class Success[E] (e: E) extends OpResult[E]
```

```
case class Reject[E] (reasons: List[Reason]) extends OpResult[E]
```



# Operation: Constraint as a type

```
object Operation {  
    def ifSeq(seq: Option[S], e: E): Operation[S, E] =  
        Operation { s =>  
            if (s === seq) Success(e)  
            else Reject(List(Reason("Sequence mismatch")))  
        }  
    ...  
}
```



# Save without Constraints

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
  def save(k: K, e: E): F[SaveResult[S]] =  
    for {  
      old <- store.latest(k).run  
  
      putResult <- eventStore.put(k, Event.next(old, e))  
  
      saveResult <- putResult match {  
        case \/- (ev) =>  
          SaveResult.success(query.acc(old, ev).value)  
        case -\/(Error.DuplicateEventId) =>  
          save(k, e)  
        case -\/(Error.Rejected(reasons)) =>  
          SaveResult.reject(reasons)  
      }  
    } yield saveResult  
}
```



# Save with Constraints

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
  def save(k: K, op: Operation[S, E]): F[SaveResult[S]] =  
    for {  
      old <- store.latest(k).run  
  
      saveResult <- putResult match {  
        case \/- (ev) =>  
          SaveResult.success(query.acc(old, ev).value)  
        case -\/(Error.DuplicateEventId) =>  
          save(k, e)  
        case -\/(Error.Rejected(reasons)) =>  
          SaveResult.reject(reasons)  
      }  
    } yield saveResult  
}
```



# Save with Constraints

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
  def save(k: K, op: Operation[S, E]): F[SaveResult[S]] =  
    for {  
      old <- store.latest(k).run  
      opResult = op.apply(old)  
  
      saveResult <- putResult match {  
        case \/- (ev) =>  
          SaveResult.success(query.acc(old, ev).value)  
        case -\/(Error.DuplicateEventId) =>  
          save(k, e)  
        case -\/(Error.Rejected(reasons)) =>  
          SaveResult.reject(reasons)  
      }  
    } yield saveResult  
}
```



# Save with Constraints

```
case class SaveAPI[K, S, E](store: EventStorage...) {  
    def save(k: K, op: Operation[S, E]): F[SaveResult[S]] =  
        for {  
            old <- store.latest(k).run  
            opResult = op.apply(old)  
            putResult <- opResult match {  
                case Success(e) => eventStore.put(k, Event.next(old.seq, e))  
                case Reject(rs) => SaveResult.reject(rs)  
            }  
            saveResult <- putResult match {  
                case \/- (ev) =>  
                    SaveResult.success(query.acc(old, ev).value)  
                case -\/(Error.DuplicateEventId) =>  
                    save(k, e)  
                case -\/(Error.Rejected(reasons)) =>  
                    SaveResult.reject(reasons)  
            }  
        } yield saveResult  
}
```



# Querying by pulling



# Querying event streams

```
def get(k: Key): F[Option[Val]] =  
  eventStore.get(???)
```



# Querying event streams

```
type Ev = Event[K, S, E]

def get(k: Key): F[Option[Val]] =
  streamFold(acc) {
    eventStore.get(?)
  }

def streamFold(
  f: (Option[Val], Ev) => Option[Val]
)(stream: Process[F, Ev]): F[Option[S, Val]] = ???
```



# Querying event streams

```
import scalaz.stream.process1

def get(k: Key): F[Option[Val]] =
  streamFold(acc) {
    eventStore.get(?)
  }

def streamFold(
  f: (Option[Val], Ev) => Option[Val]
)(stream: Process[F, Ev]): F[Option[Val]] =
  stream.pipe {
    process1.fold(None)(f)
  }.runLastOr(None)
```



# Querying event streams

```
import scalaz.stream.process1

def get(toStreamKey: Key => K) (k: Key): F[Option[Val]] =
  streamFold(acc) {
    eventStore.get(toStreamKey(k))
  }

def streamFold(
  f: (Option[Val], Ev) => Option[Val]
) (stream: Process[F, Ev]): F[Option[Val]] =
  stream.pipe {
    process1.fold(None) (f)
  }.runLastOr(None)
```



# Querying event streams

```
case class QueryAPI[F[_], K, S, E, Key, Val] (
    toStreamKey: Key => K,
    acc: (k: Key) => (s: Snapshot[S, Val], e: Ev)
        => Snapshot[S, Val],
    store: EventStorage[F, K, S, E]) {
    def get(k: Key): F[Option[Val]] =
        ...
}
```



# Querying by pushing



# Query synchroniser is a scalaz-stream Channel...

```
class QuerySynchroniser(db: DB) {  
  
  def synchronise(s: Process[Task, StreamEvent]) =  
  
}  
}
```



# Query synchroniser is a scalaz-stream Channel...

```
class QuerySynchroniser(db: DB) {  
  
  def synchronise(s: Process[Task, StreamEvent]) =  
    (s through writeEvent).run  
  
  private val writeEvent: Channel[Task, StreamEvent, Unit] =  
    Process.constant { e: StreamEvent =>  
  
    }  
}
```



# Query synchroniser is a scalaz-stream Channel...

```
class QuerySynchroniser(db: DB) {  
  
  def synchronise(s: Process[Task, StreamEvent]) =  
    (s through writeEvent).run  
  
  private val writeEvent: Channel[Task, StreamEvent, Unit] =  
    Process.constant { e: StreamEvent =>  
      for {  
        r1 <- StreamState.transitionSeq(db, e.id.key, e.id.seq)  
      } yield ()  
    }  
}
```



# Query synchroniser is a scalaz-stream Channel...

```
class QuerySynchroniser(db: DB) {  
  
  def synchronise(s: Process[Task, StreamEvent]) =  
    (s through writeEvent).run  
  
  private val writeEvent: Channel[Task, StreamEvent, Unit] =  
    Process.constant { e: StreamEvent =>  
      for {  
        r1 <- StreamState.transitionSeq(db, e.id.key, e.id.seq)  
        _ <- if (r1.success) db.process(e) else db.rollback  
      } yield ()  
    }  
}
```



# ...that you pass it an event stream

```
class KinesisProcessor(qs: QuerySynchroniser,  
                      es: EventStorage[...]) extends IRecordProcessor {  
  
    override def processRecords(rs: ProcessRecordInput) = {  
  
    }  
}
```



# ...that you pass it an event stream

```
class KinesisProcessor(qs: QuerySynchroniser,  
                      es: EventStorage[...]) extends IRecordProcessor {  
  
    override def processRecords(rs: ProcessRecordInput) = {  
        // Either emit events from Kinesis, or read events from Dynamo  
        val stream = es.get((rs.getRecords...))  
  
    }  
}
```



# ...that you pass it an event stream

```
class KinesisProcessor(qs: QuerySynchroniser,  
                      es: EventStorage[...]) extends IRecordProcessor {  
  
    override def processRecords(rs: ProcessRecordInput) = {  
        // Either emit events from Kinesis, or read events from Dynamo  
        val stream = es.get((rs.getRecords...))  
  
        qs.synchronise(stream).attemptRunFor(...)  
    }  
  
}
```





## 3.5 Design Principles



# Principle 1: Events as an API

# Insert / Update Delta      vs      ‘Set’ events



# Insert / Update Delta      vs      ‘Set’ events

UserAdded(id, name, email1)



# Insert / Update Delta

vs

# ‘Set’ events

UserAdded(id, name, email1)

UserUpdated(id, email = *Some(email)*)



# Insert / Update Delta

vs

# ‘Set’ events

UserAdded(id, name, email1)

UserUpdated(id, email = *Some(email)*)

UserNameSet(id, name)

UserEmailSet(id, email1)



# Insert / Update Delta

vs

# ‘Set’ events

UserAdded(id, name, email1)

UserUpdated(id, email = *Some*(email))

UserNameSet(id, name)

UserEmailSet(id, email1)

UserEmailSet(id, email2)



# Insert / Update Delta

vs

# ‘Set’ events

UserAdded(id, name, email1)

UserUpdated(id, email = *Some*(email))

UserNameSet(id, name)

UserEmailSet(id, email1)

UserEmailSet(id, email2)

Fits nicely with CRUD + PATCH



# Insert / Update Delta

vs

# ‘Set’ events

UserAdded(id, name, email1)

UserUpdated(id, email = *Some*(email))

UserNameSet(id, name)

UserEmailSet(id, email1)

UserEmailSet(id, email2)

Fits nicely with CRUD + PATCH

Assume insert before update



# Insert / Update Delta

vs

# ‘Set’ events

UserAdded(id, name, email1)

UserUpdated(id, email = *Some*(email))

UserNameSet(id, name)

UserEmailSet(id, email1)

UserEmailSet(id, email2)

Fits nicely with CRUD + PATCH

Encourages idempotent processing

Assume insert before update



# Insert / Update Delta

vs

# ‘Set’ events

UserAdded(id, name, email1)

UserUpdated(id, email = *Some*(email))

UserNameSet(id, name)

UserEmailSet(id, email1)

UserEmailSet(id, email2)

Fits nicely with CRUD + PATCH

Encourages idempotent processing

Assume insert before update

Single code path for query sync



# Insert / Update Delta

vs

# ‘Set’ events

UserAdded(id, name, email1)

UserUpdated(id, email = *Some*(email))

UserNameSet(id, name)

UserEmailSet(id, email1)

UserEmailSet(id, email2)

Fits nicely with CRUD + PATCH

Encourages idempotent processing

Assume insert before update

Single code path for query sync

Minimally sized events to avoid conflict



# Insert / Update Delta

vs

UserAdded(id, name, email1)  
UserUpdated(id, email = *Some*(email))

Fits nicely with CRUD + PATCH

Assume insert before update

# ‘Set’ events

UserNameSet(id, name)  
UserEmailSet(id, email1)  
UserEmailSet(id, email2)

Encourages idempotent processing

Single code path for query sync

Minimally sized events to avoid conflict





## Principle 2: Split event streams

# Single stream

# Multiple streams

# Single stream

Transactions and  
consistent data  
resolution

# Multiple streams

# Single stream

Transactions and  
consistent data  
resolution

# Multiple streams

Sharding for  
throughput

# Single stream

Transactions and  
consistent data  
resolution

# Multiple streams

Sharding for  
throughput

No ordering  
between streams

# Single stream

~~Transactions and  
consistent data  
resolution~~

# Multiple streams

Sharding for  
throughput

No ordering  
between streams

Better latency  
vs consistency  
compromise

# Rules for splitting streams

1. Place independent events on different streams



# Rules for splitting streams

1. Place independent events on different streams
2. Split streams by event type and unique Id



# Rules for splitting streams

1. Place independent events on different streams
2. Split streams by event type and unique Id
3. Identify the ‘transactions’ you really need



# Rules for splitting streams

1. Place independent events on different streams
2. Split streams by event type and unique Id
3. Identify the ‘transactions’ you really need
4. Use hierarchical streams to maximise number of streams



# Rules for splitting streams

1. Place independent events on different streams
2. Split streams by event type and unique Id
3. Identify the ‘transactions’ you really need
4. Use hierarchical streams to maximise number of streams
5. Splitting and joining streams later is possible



But... no  
guaranteed order  
between streams



# Query views get populated eventually



# Query views get populated eventually

- A field should only be updated by a single event stream



# Query views get populated eventually

- A field should only be updated by a single event stream
- No foreign key constraints



# Query views get populated eventually

- A field should only be updated by a single event stream
- No foreign key constraints
- In general, unique or data constraints ‘enforced’ on write





Principle 3:  
Let go of  
transactions  
and consistency

# Why do we need transactions?



# Why do we need transactions?

- Enforce business constraints e.g. uniqueness



# Why do we need transactions?

- Enforce business constraints e.g. uniqueness
- Guaranteed to see what I just wrote



# Write and Read Consistency



But CAP theorem...



# CAP or PACELC?



# CAP or PASELC?

During a network **Partition**, choose between  
**Availability versus Consistency**



# CAP or PACELC?

During a network **Partition**, choose between  
**Availability versus Consistency**

**Else** choose between  
**Latency versus Consistency**



There is a middle  
ground...



# Check-and-Set writes

# Optional forced reads



# Check-and-Set writes

Potentially conflicting events on  
same stream

# Optional forced reads



# Check-and-Set writes

Potentially conflicting events on  
same stream

1. Read at seq X
2. Run business rule
3. Stream must be at X to write

# Optional forced reads



# Check-and-Set writes

Potentially conflicting events on  
same stream

1. Read at seq X
2. Run business rule
3. Stream must be at X to write

Potential false positives

# Optional forced reads



# Check-and-Set writes

Potentially conflicting events on  
same stream

1. Read at seq X
2. Run business rule
3. Stream must be at X to write

Potential false positives

Smaller streams alleviate  
problems

# Optional forced reads



# Check-and-Set writes

Potentially conflicting events on same stream

1. Read at seq X
2. Run business rule
3. Stream must be at X to write

Potential false positives

Smaller streams alleviate problems

# Optional forced reads

Query view must be at stream seq X



# Check-and-Set writes

Potentially conflicting events on same stream

1. Read at seq X
2. Run business rule
3. Stream must be at X to write

Potential false positives

Smaller streams alleviate problems

# Optional forced reads

Query view must be at stream seq X

Potential increased latency



# Check-and-Set writes

Potentially conflicting events on same stream

1. Read at seq X
2. Run business rule
3. Stream must be at X to write

Potential false positives

Smaller streams alleviate problems

# Optional forced reads

Query view must be at stream seq X

Potential increased latency

Do not use as default



# Check-and-Set writes

Potentially conflicting events on same stream

1. Read at seq X
2. Run business rule
3. Stream must be at X to write

Potential false positives

Smaller streams alleviate problems

# Optional forced reads

Query view must be at stream seq X

Potential increased latency

Do not use as default  
Enforce timed waits



# Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



# Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag



# Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag
- Pass as request header for:



# Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag
- Pass as request header for:
  - Condition write ('transaction')



# Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

User 4: Seq 23



- Returned on read and write via ETag
- Pass as request header for:
  - Condition write ('transaction')
  - Force query view update ('consistency')



# Tokens to emulate transactions and consistency

User: homer (id 4)

All Users: Seq 100

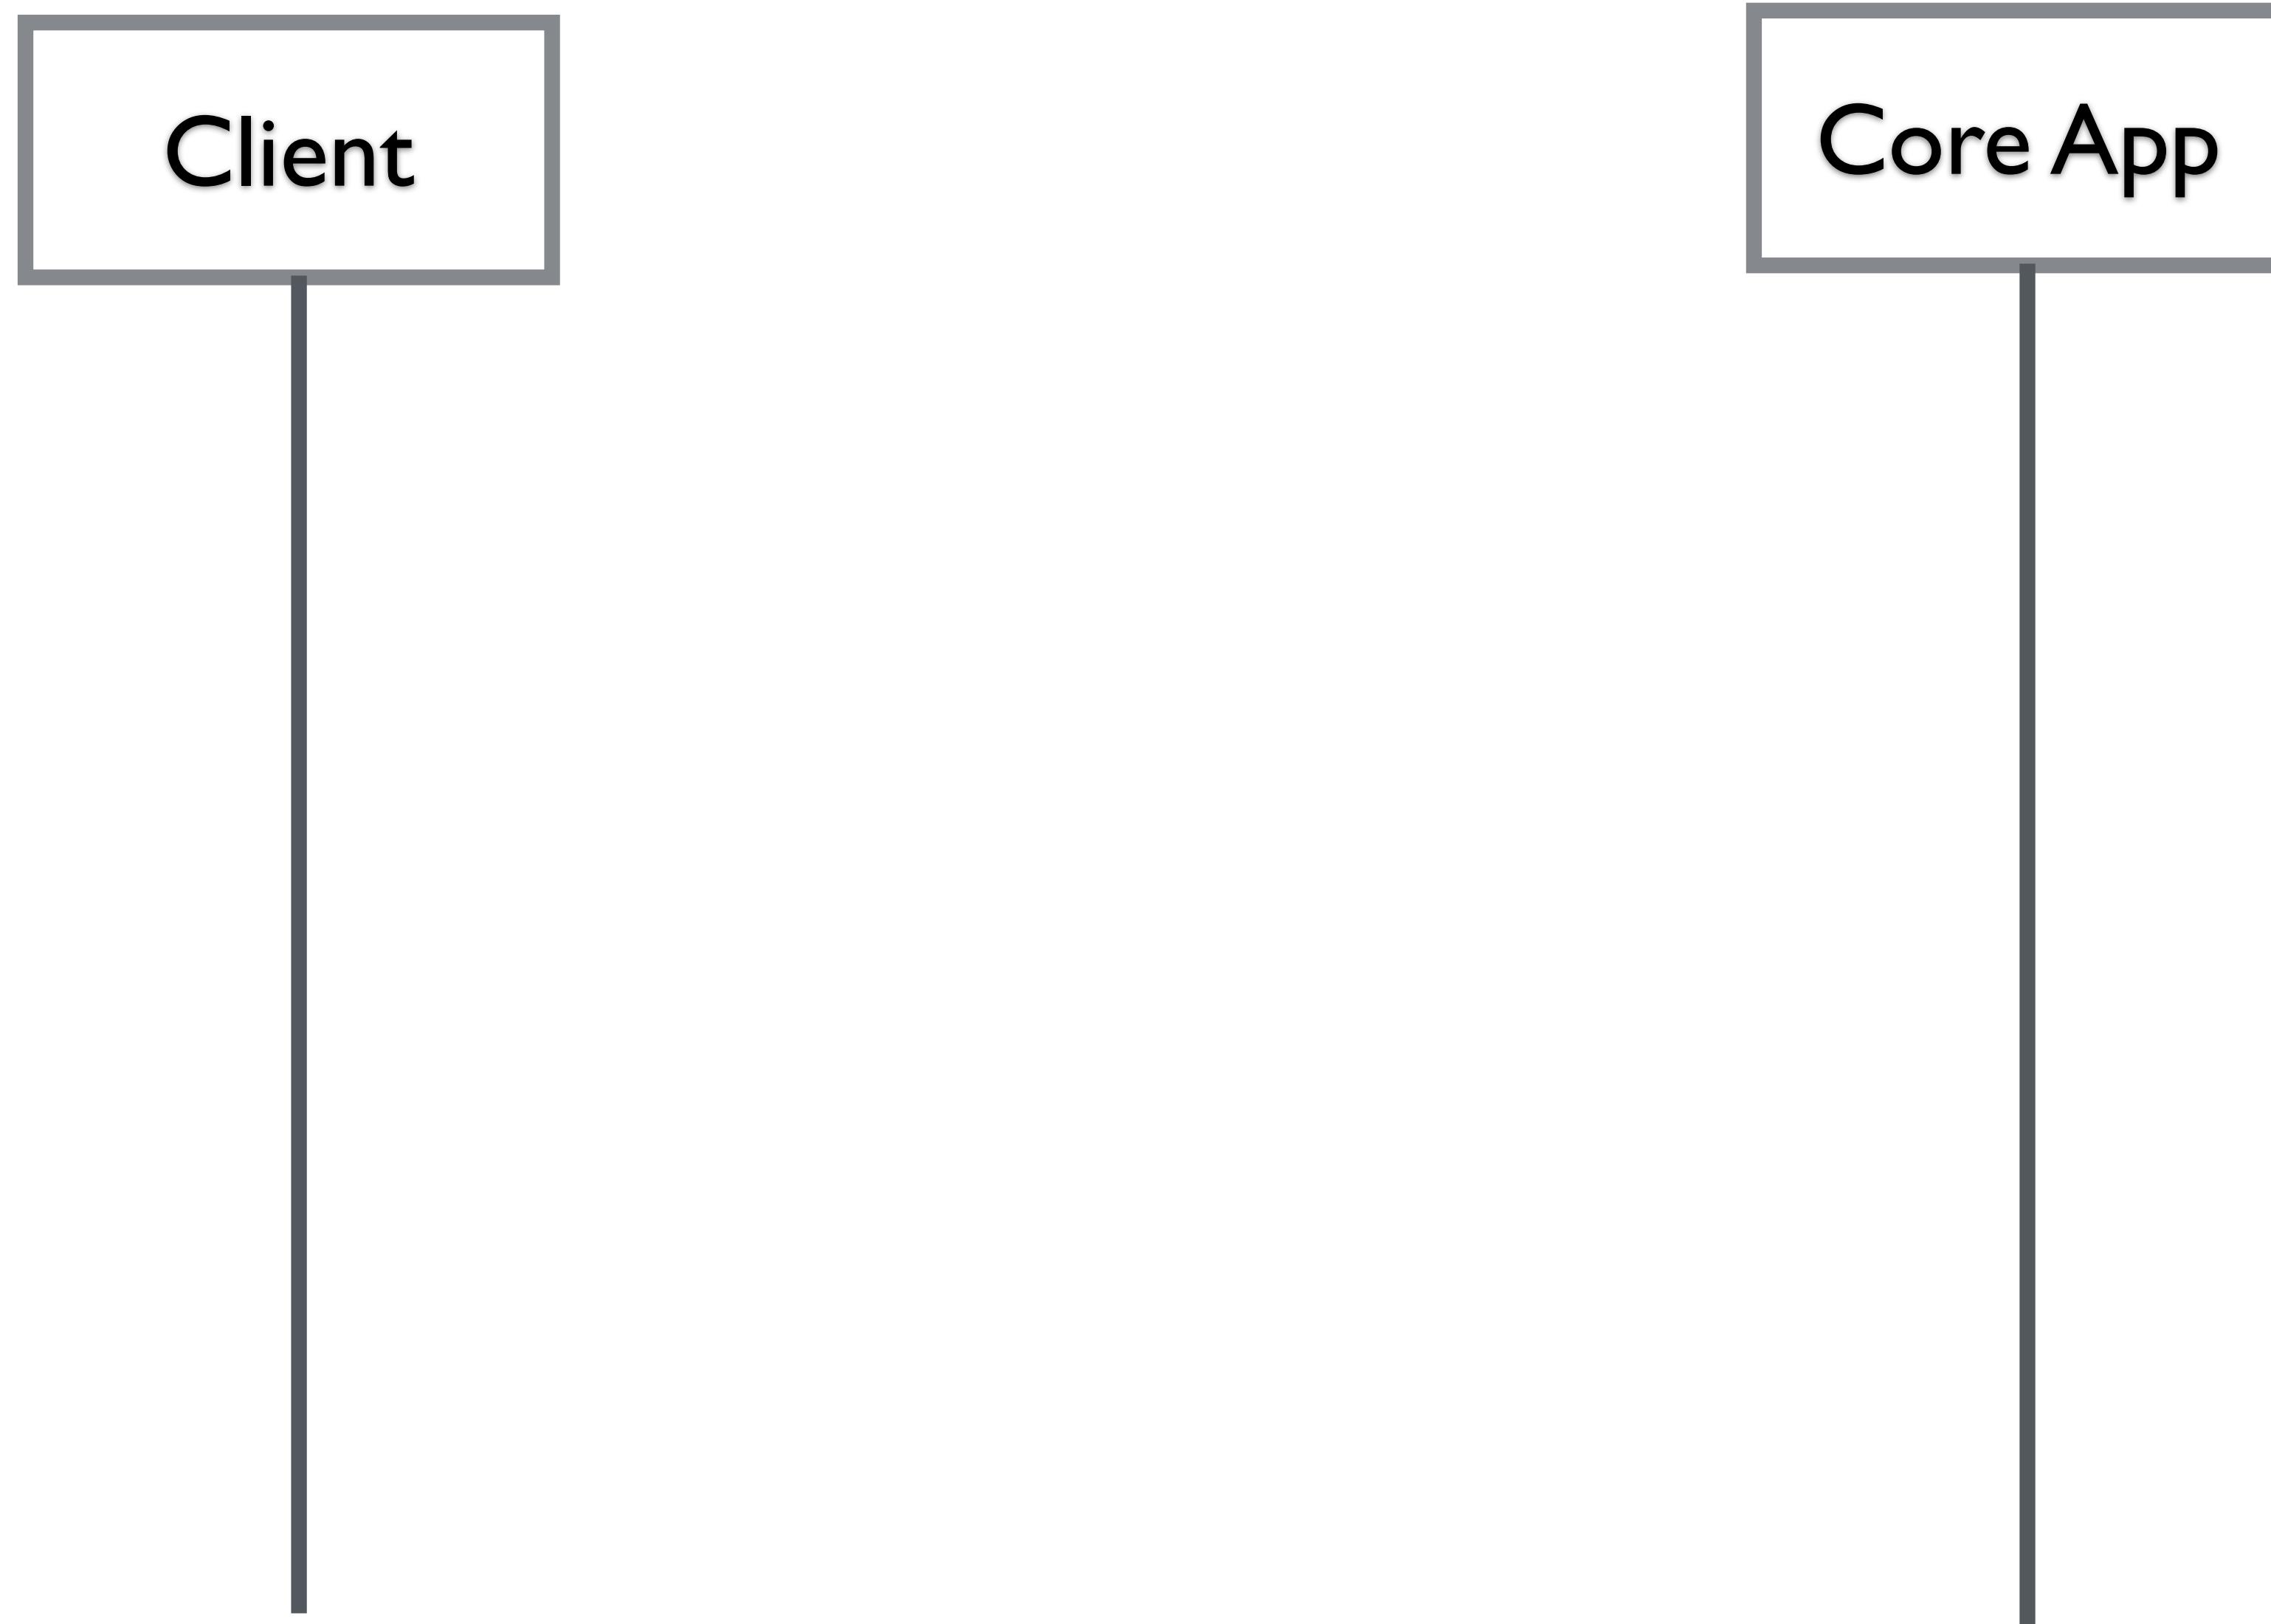
User 4: Seq 23



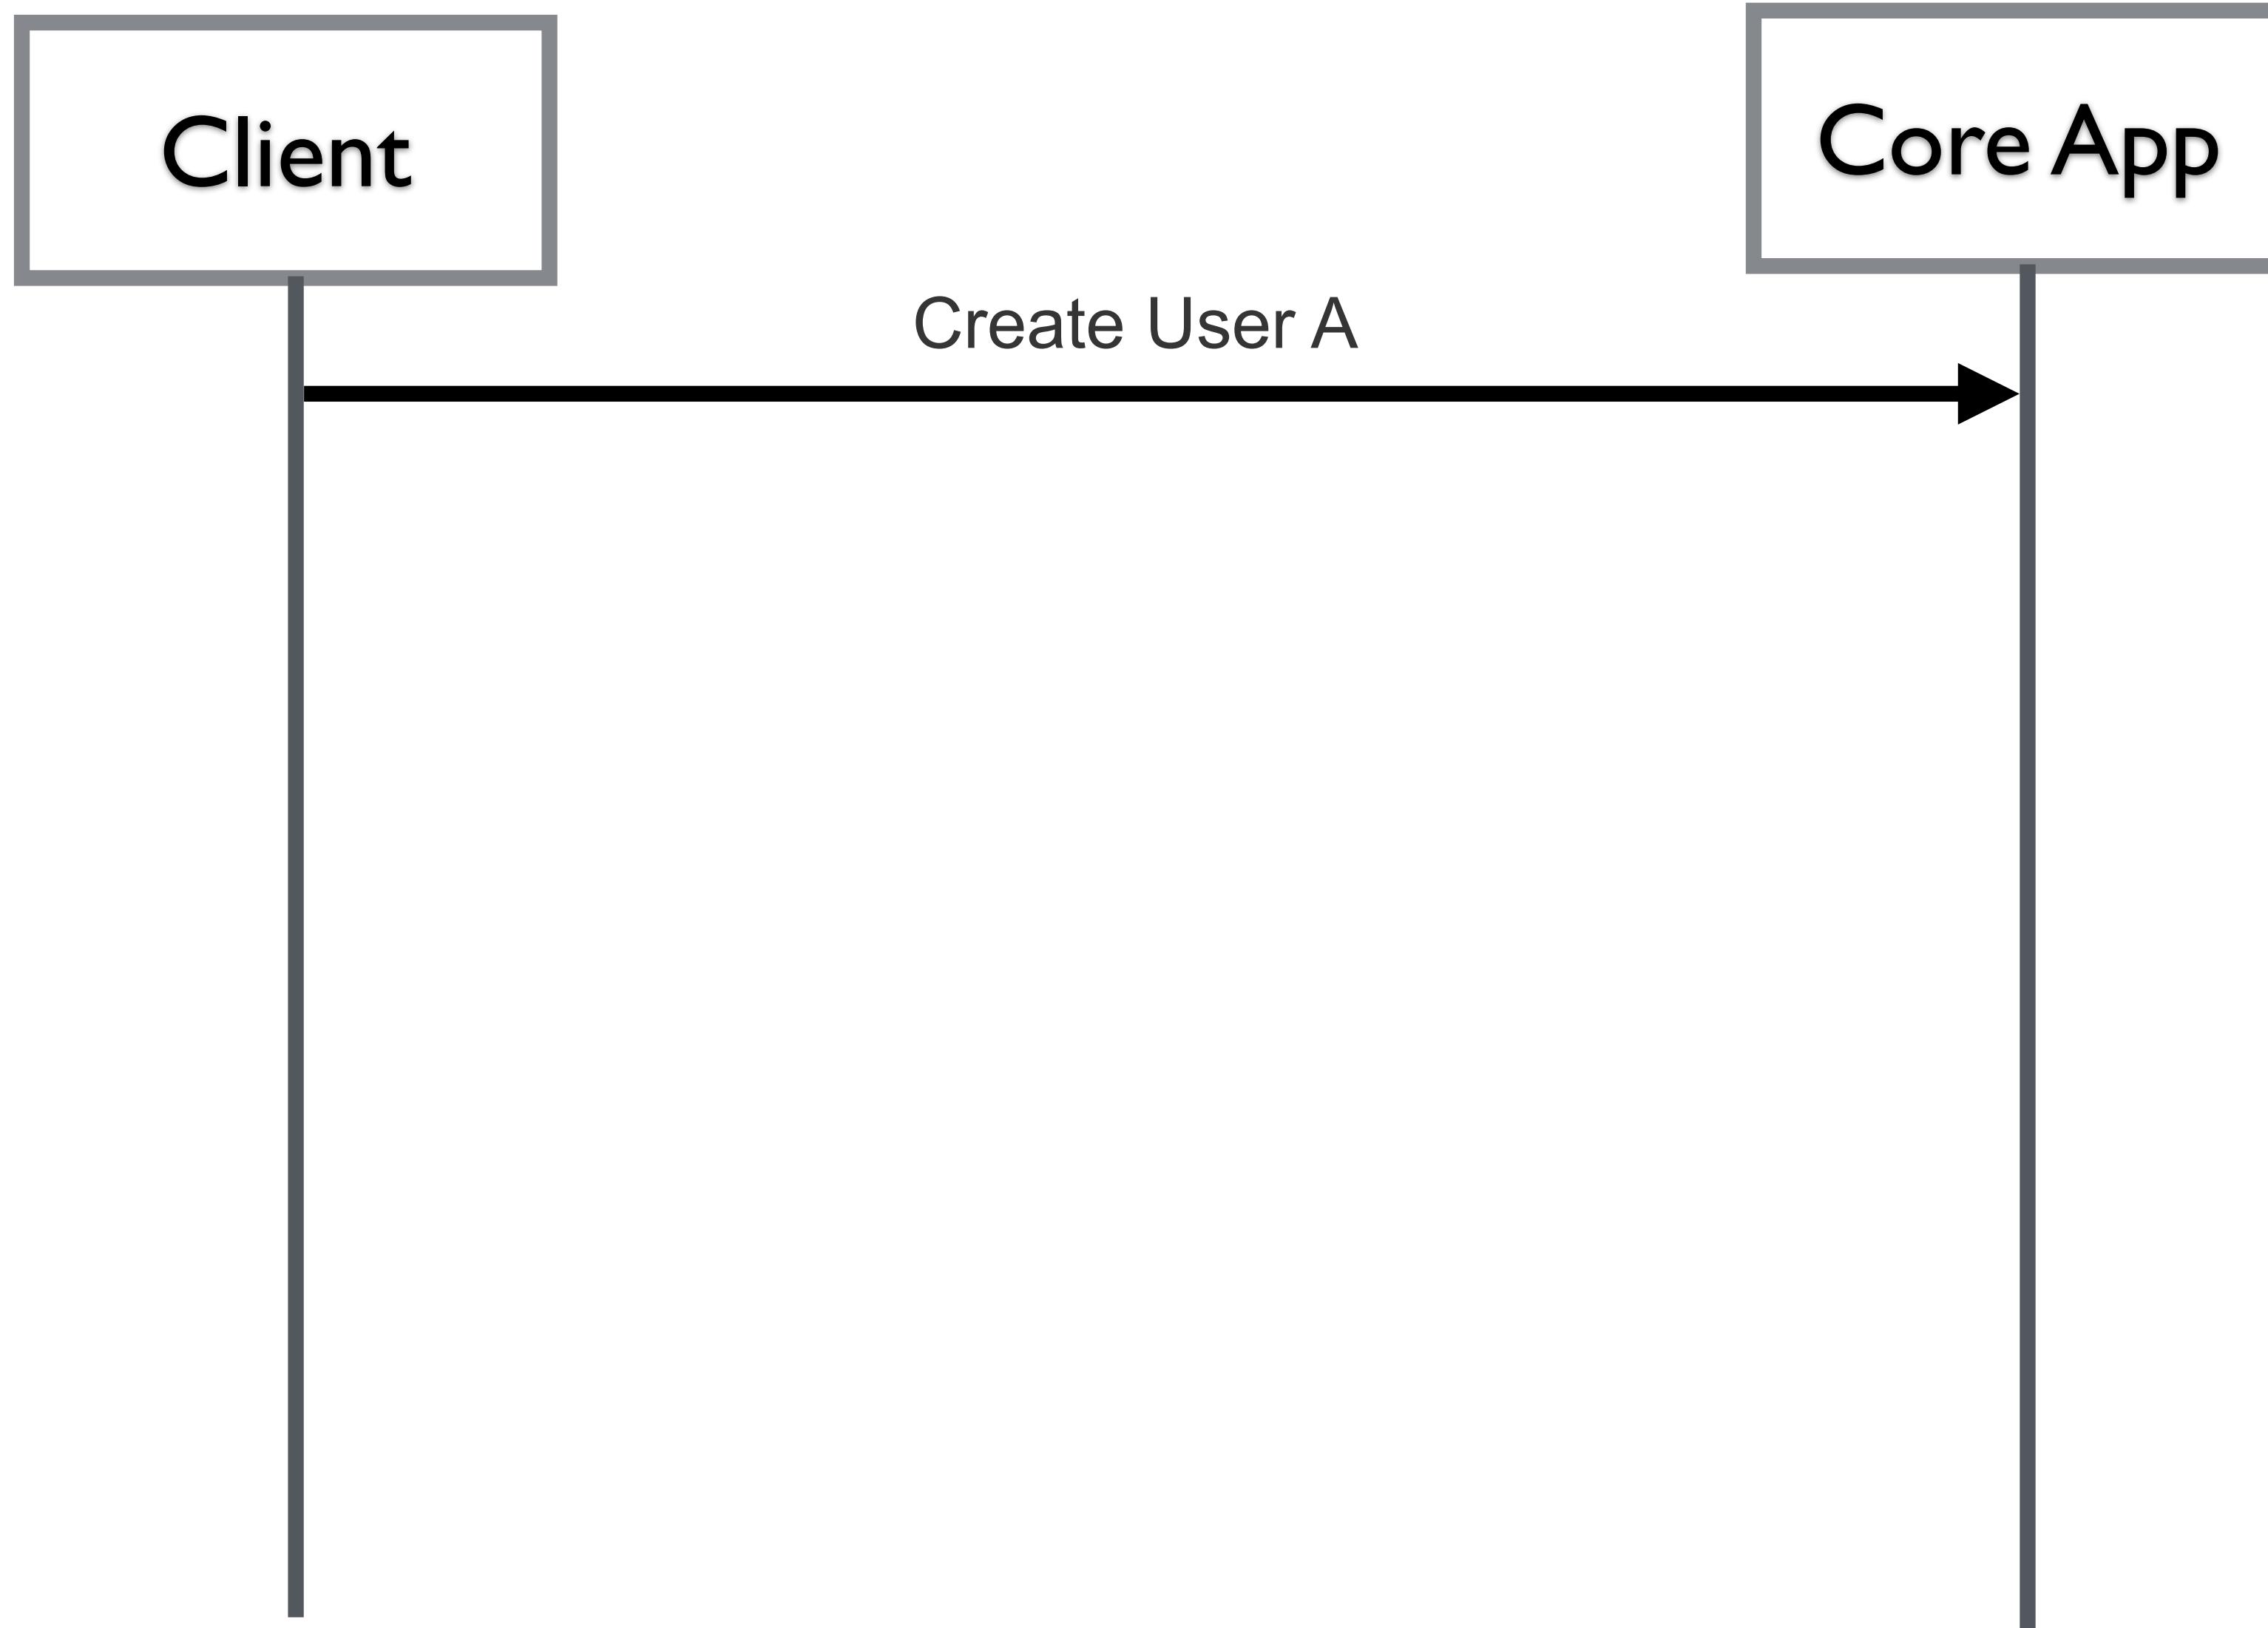
- Returned on read and write via ETag
- Pass as request header for:
  - Condition write ('transaction')
  - Force query view update ('consistency')
  - Caching



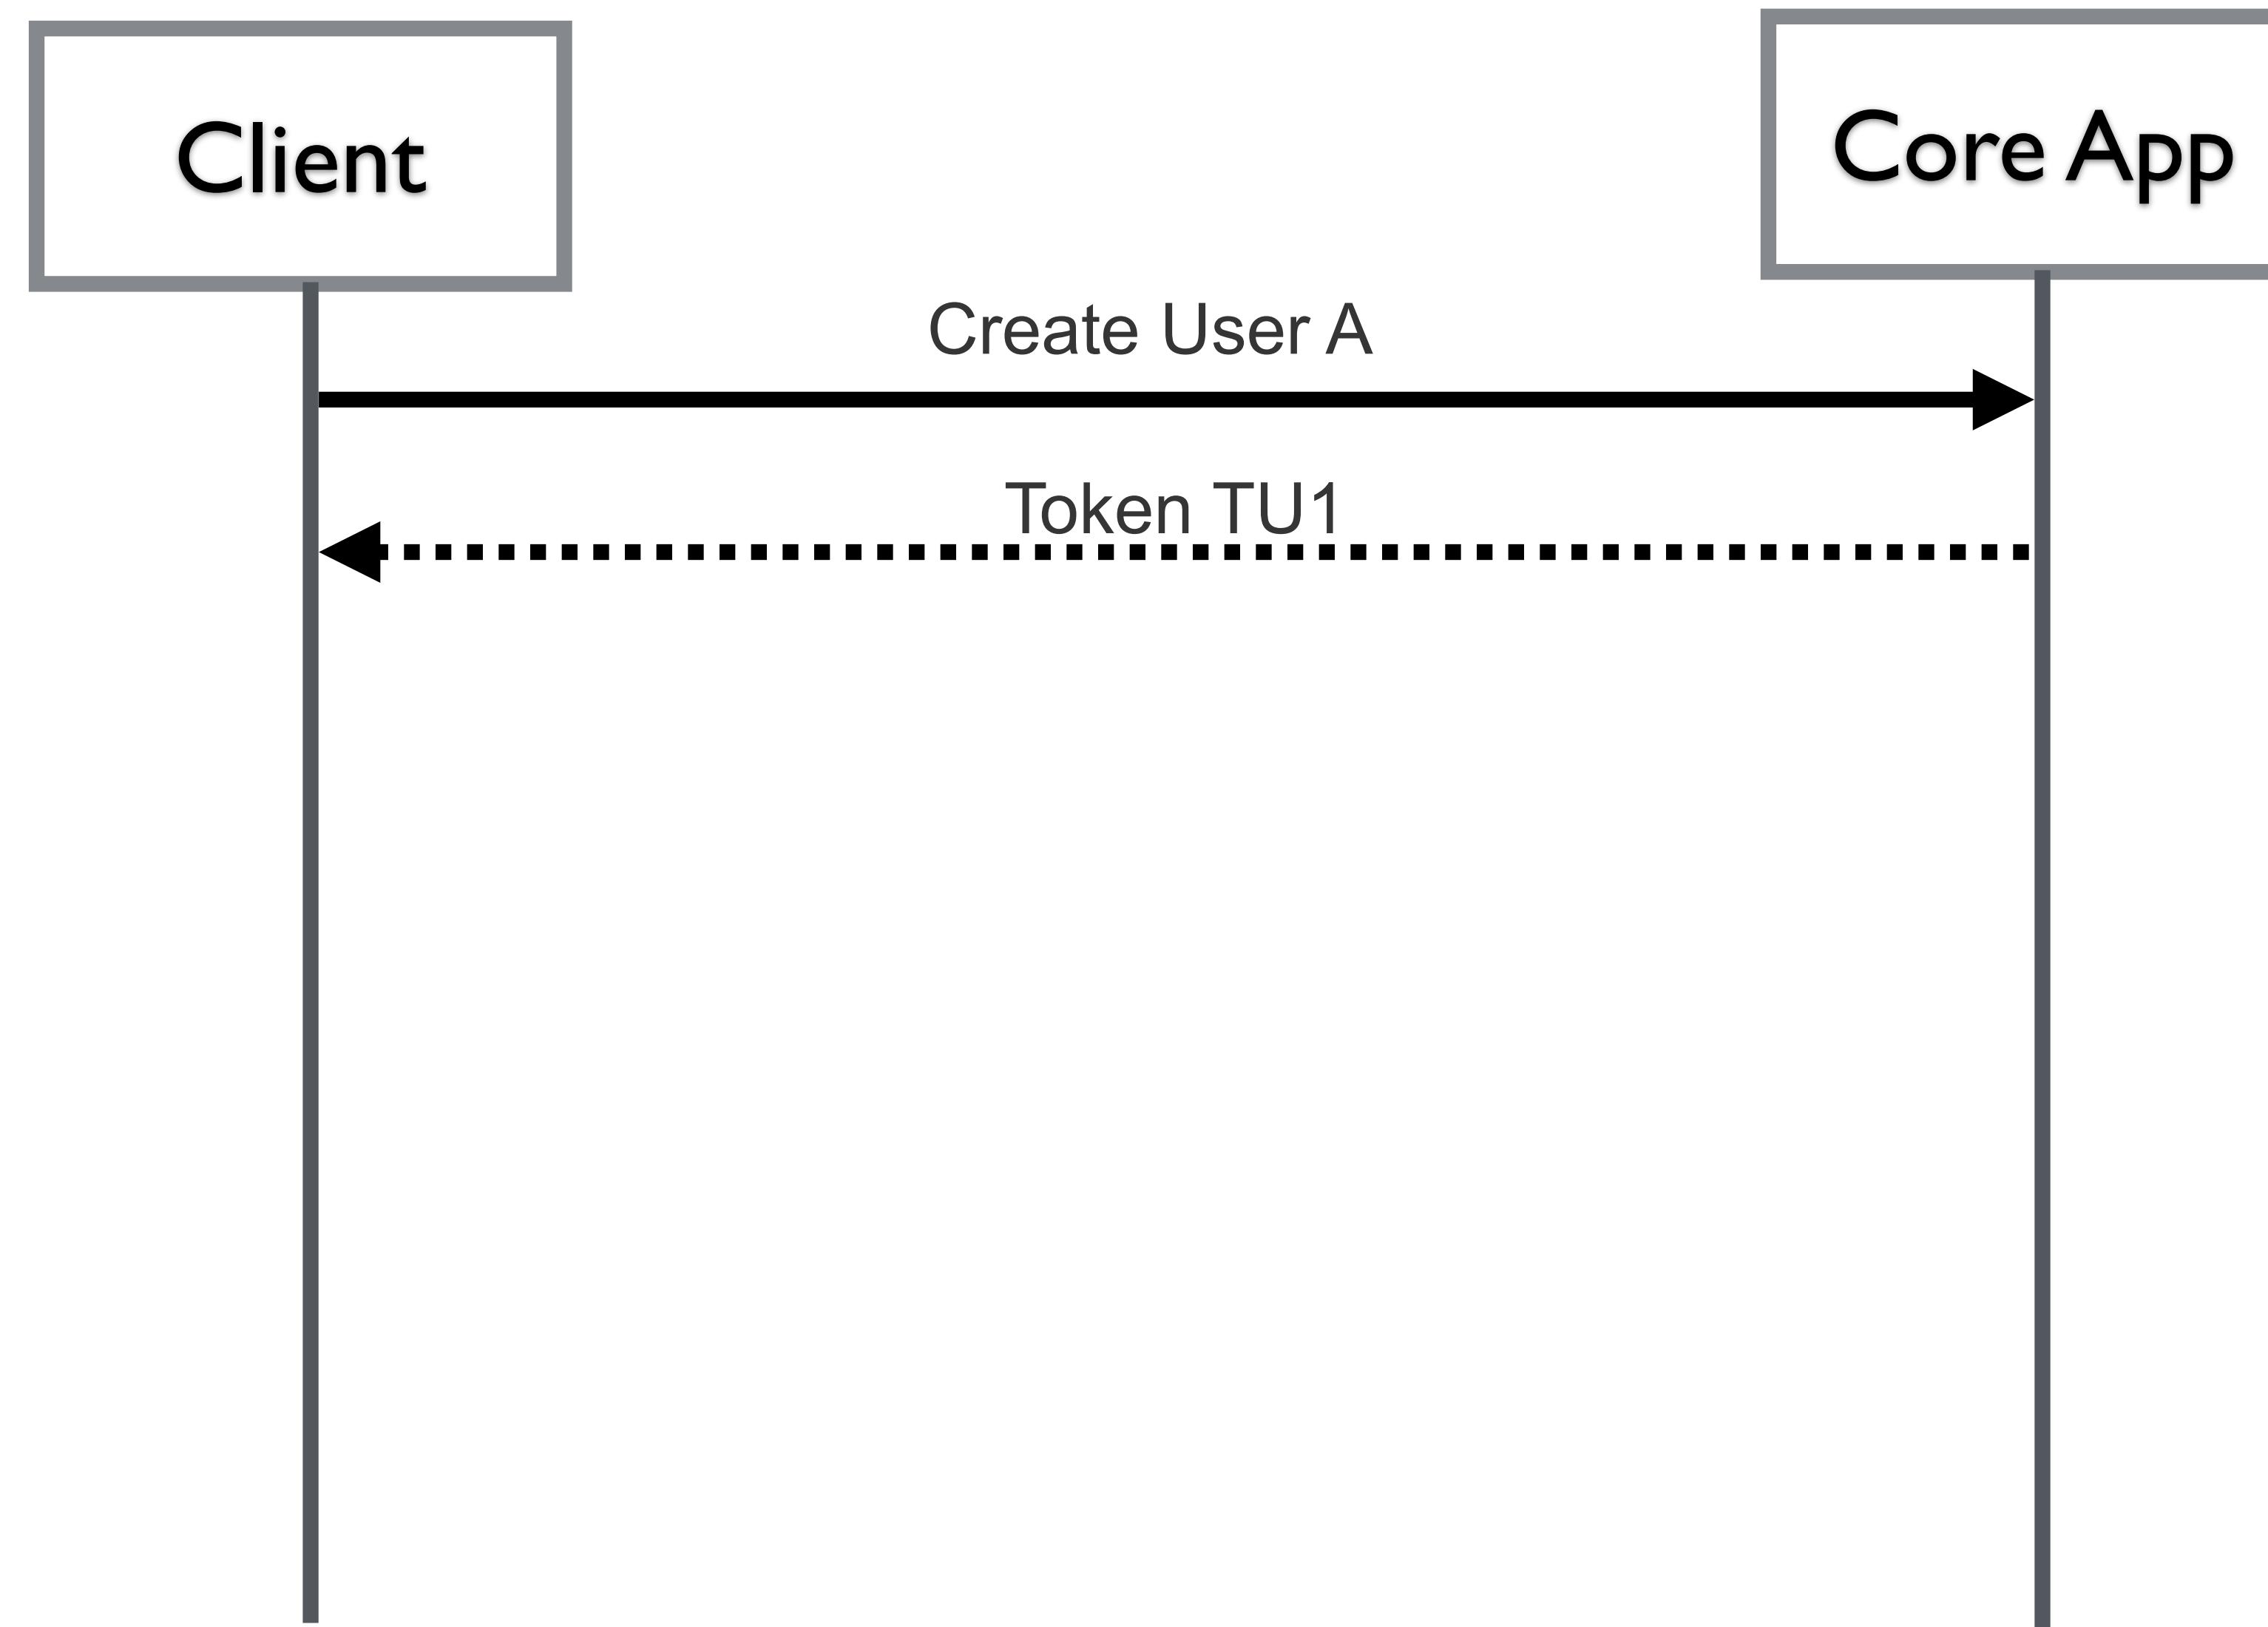
# Using tokens to enforce state



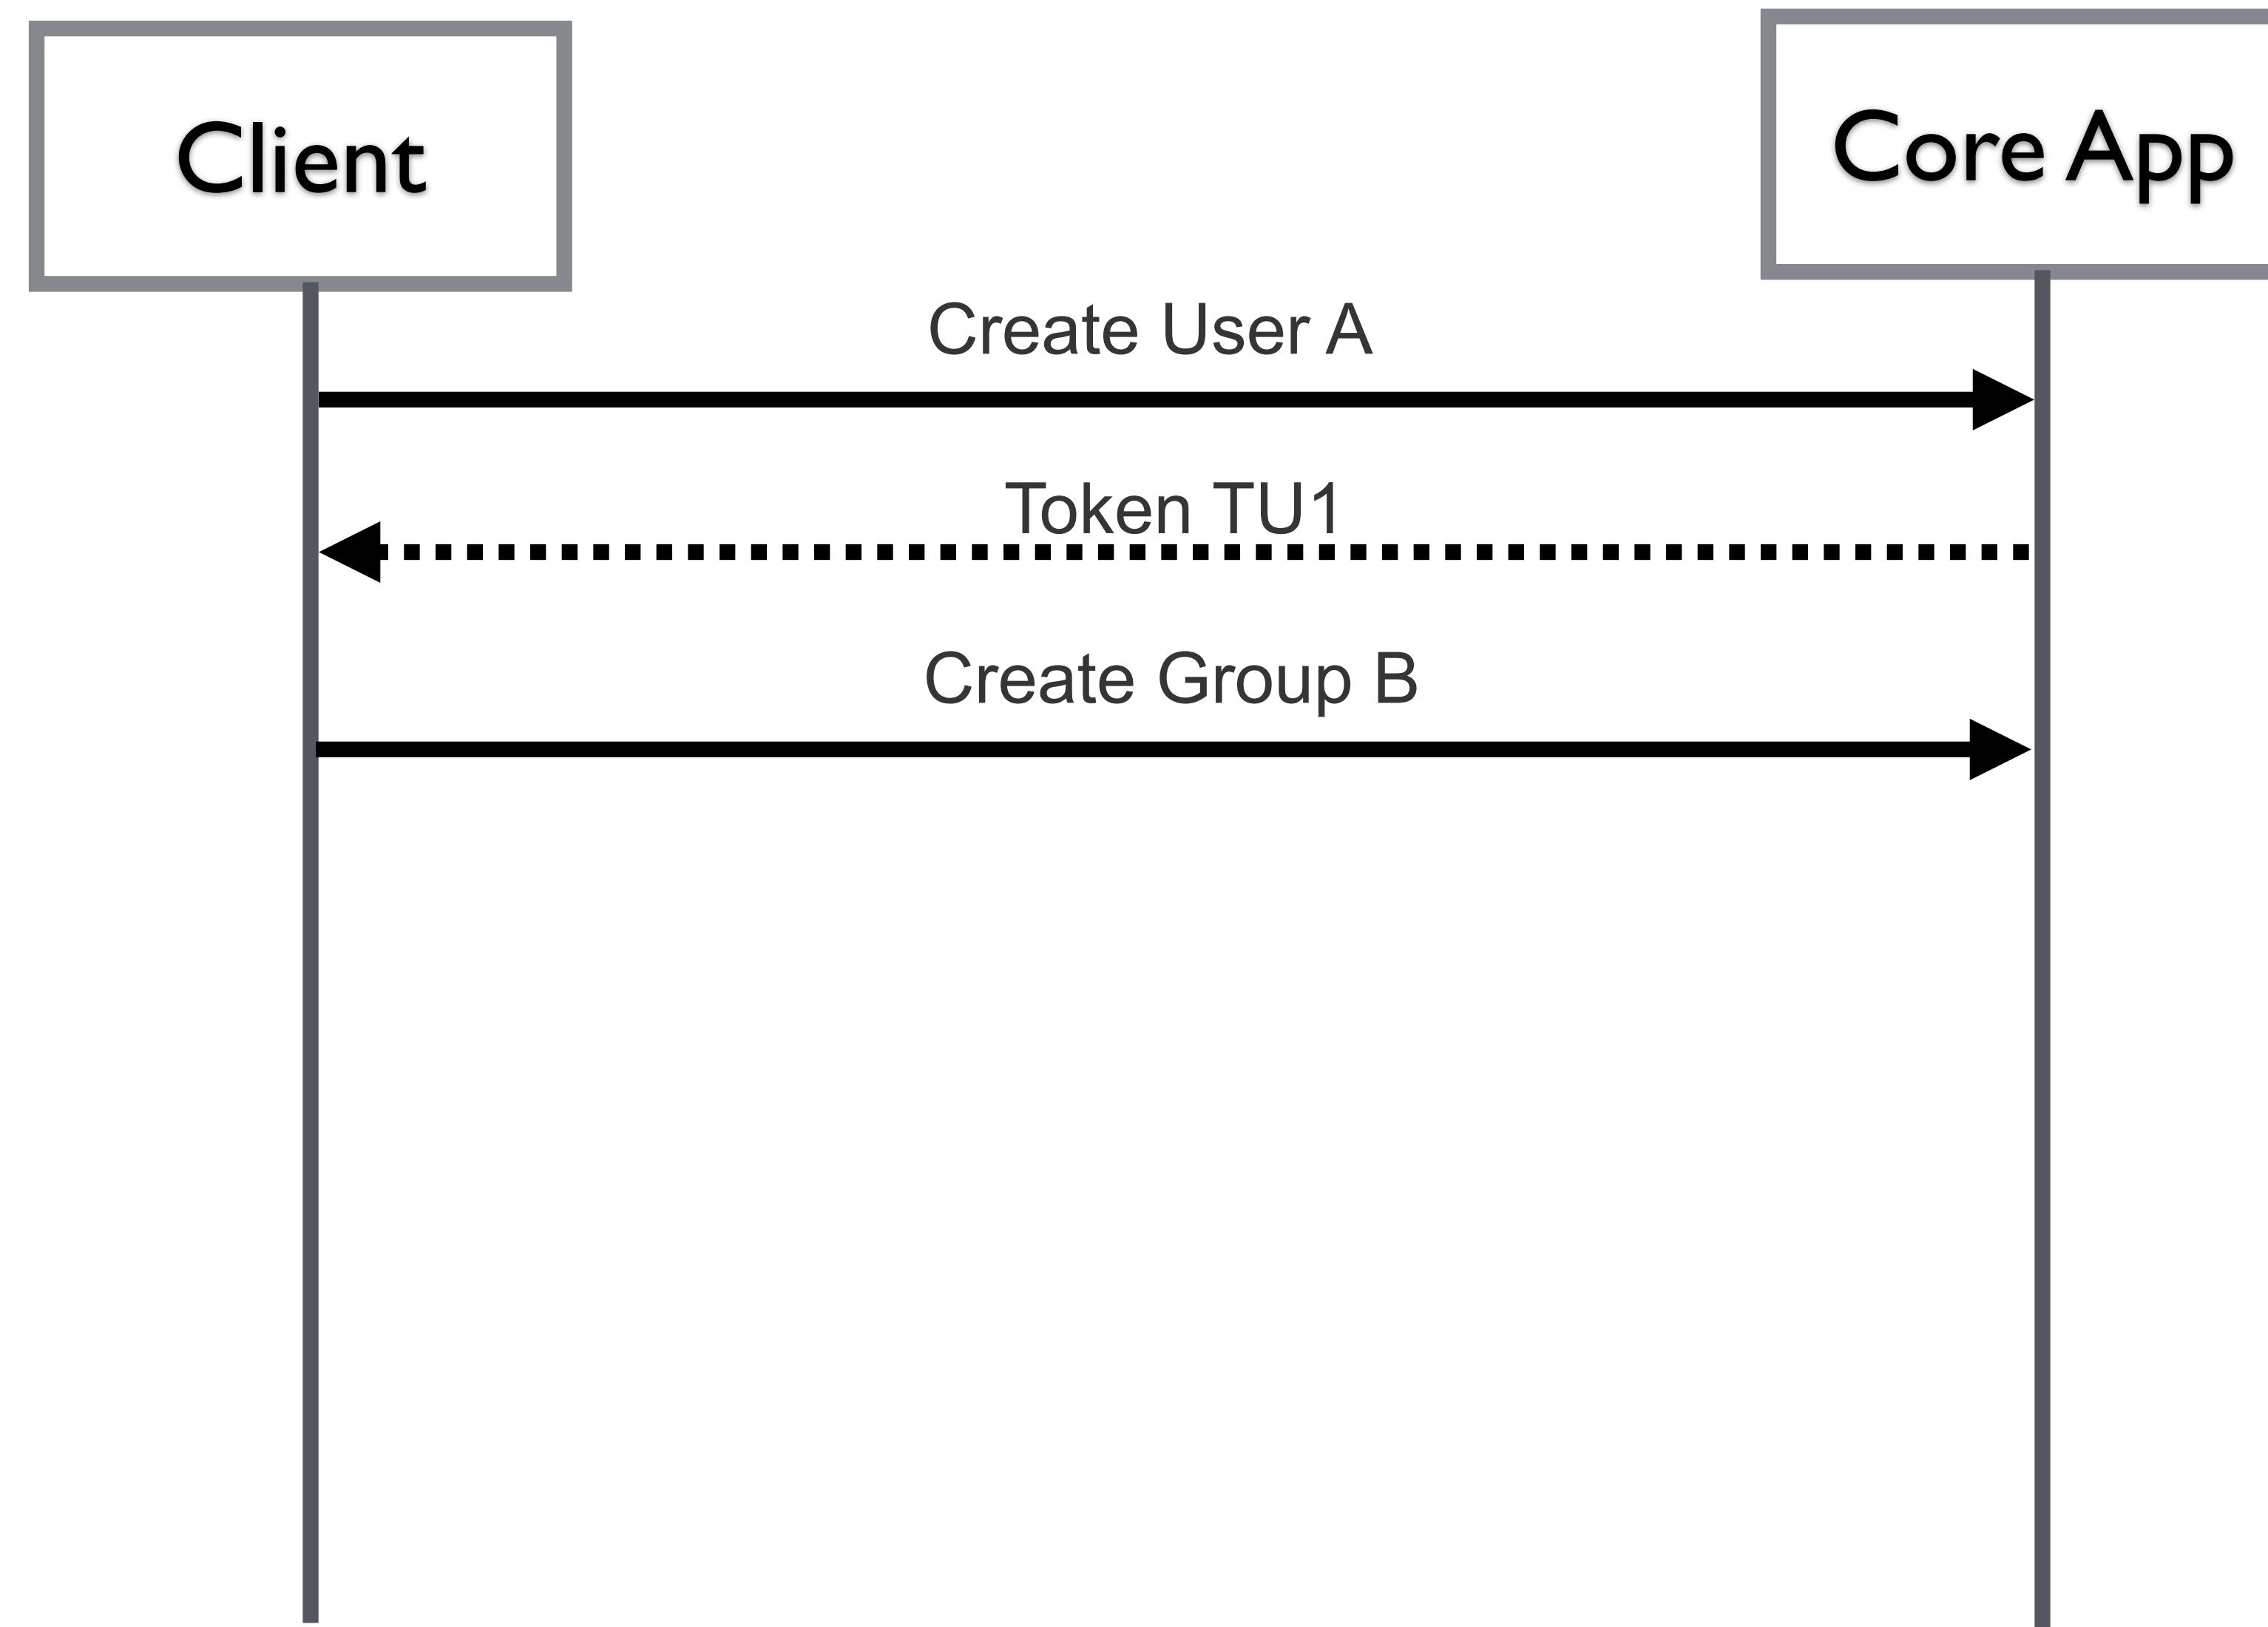
# Using tokens to enforce state



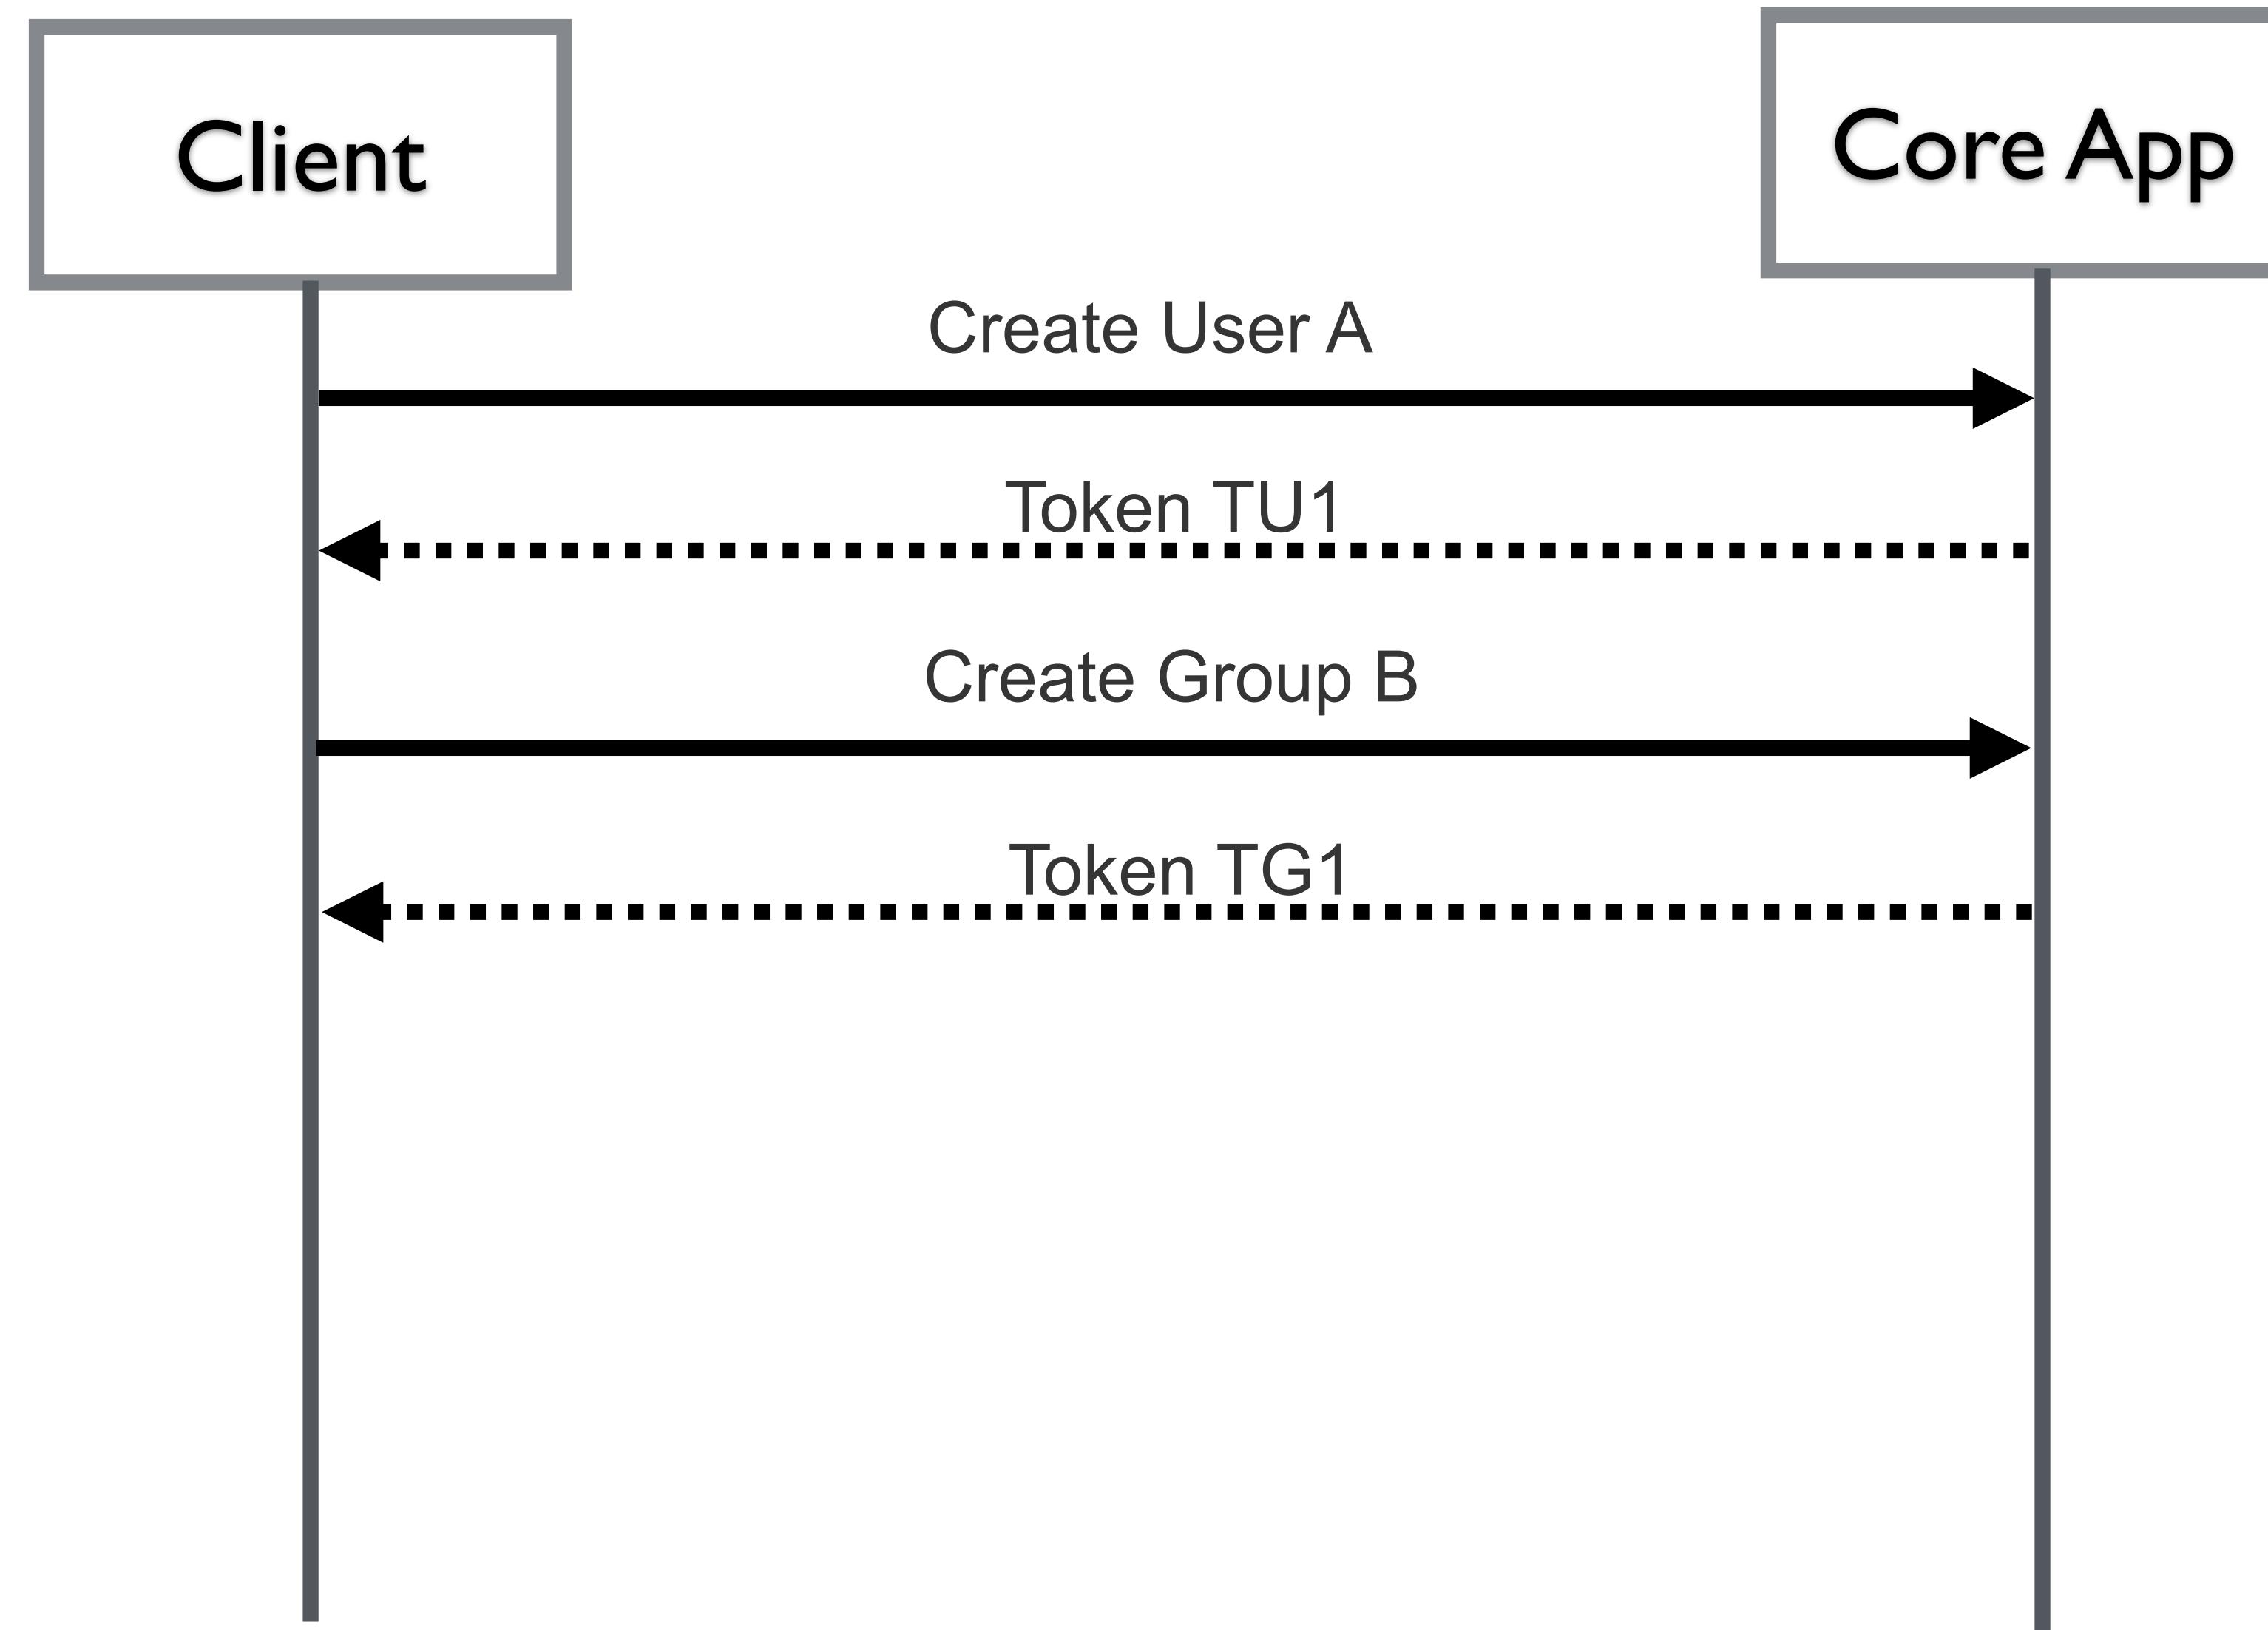
# Using tokens to enforce state



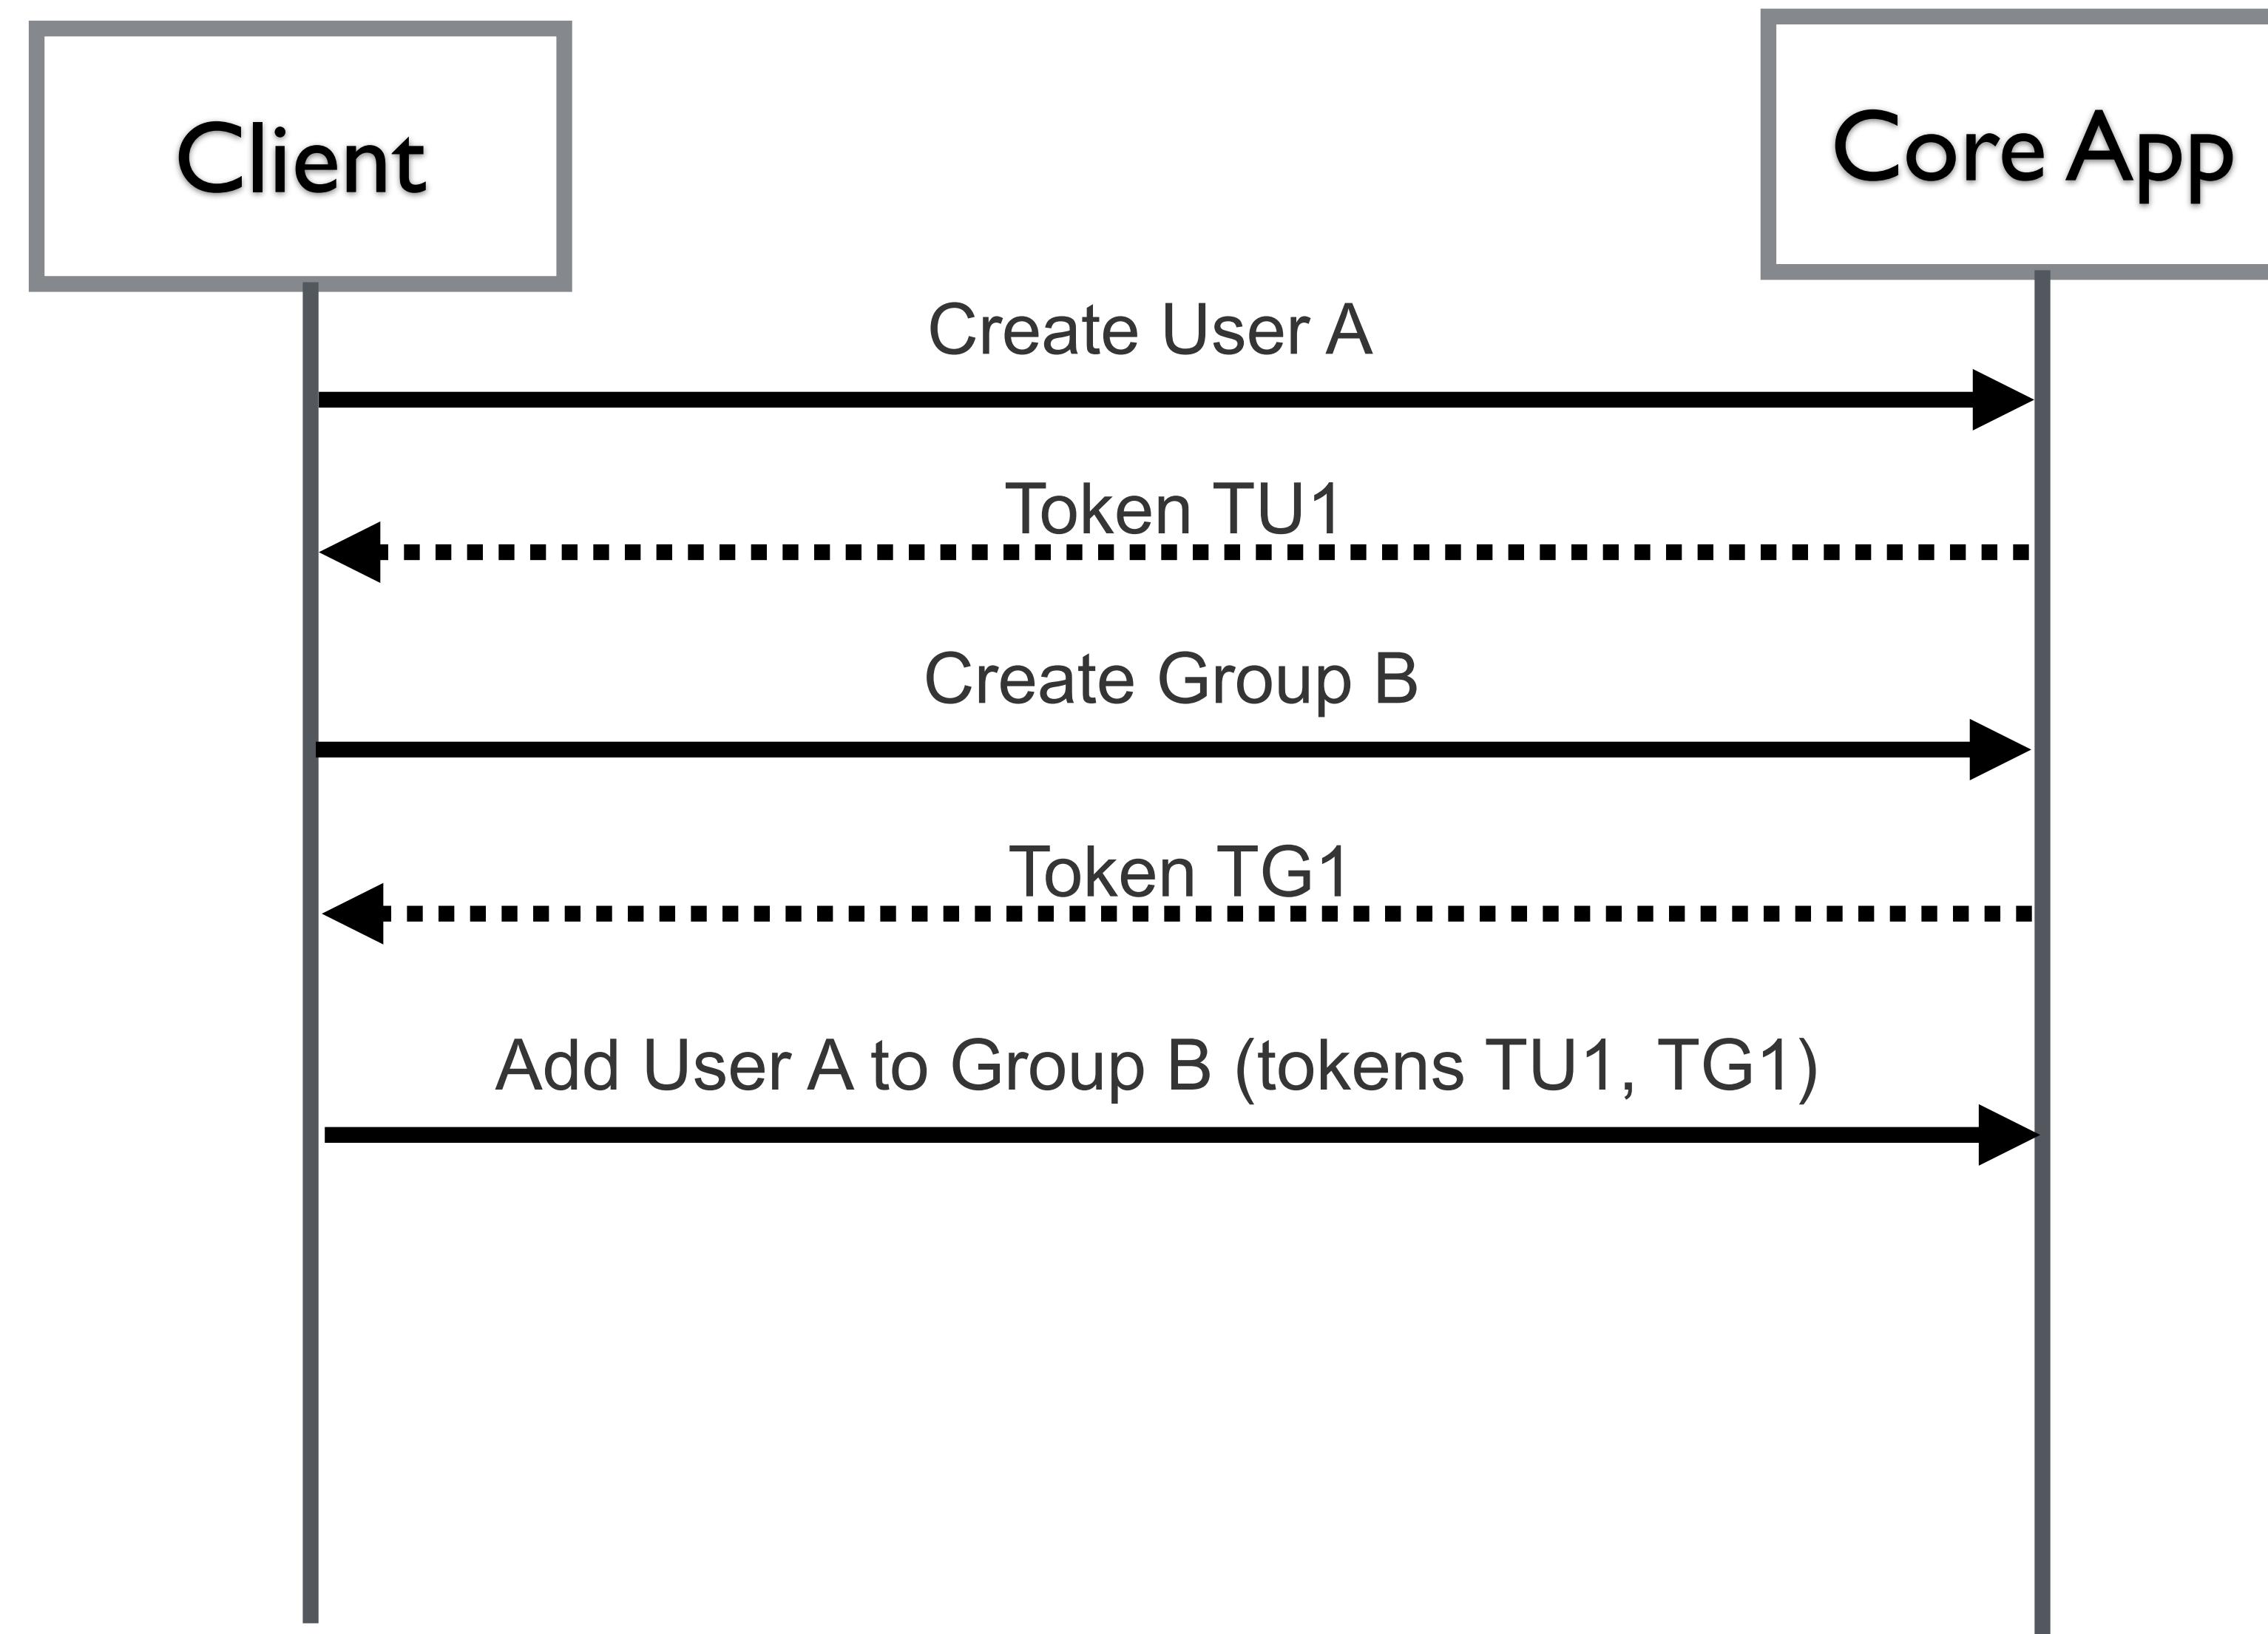
# Using tokens to enforce state



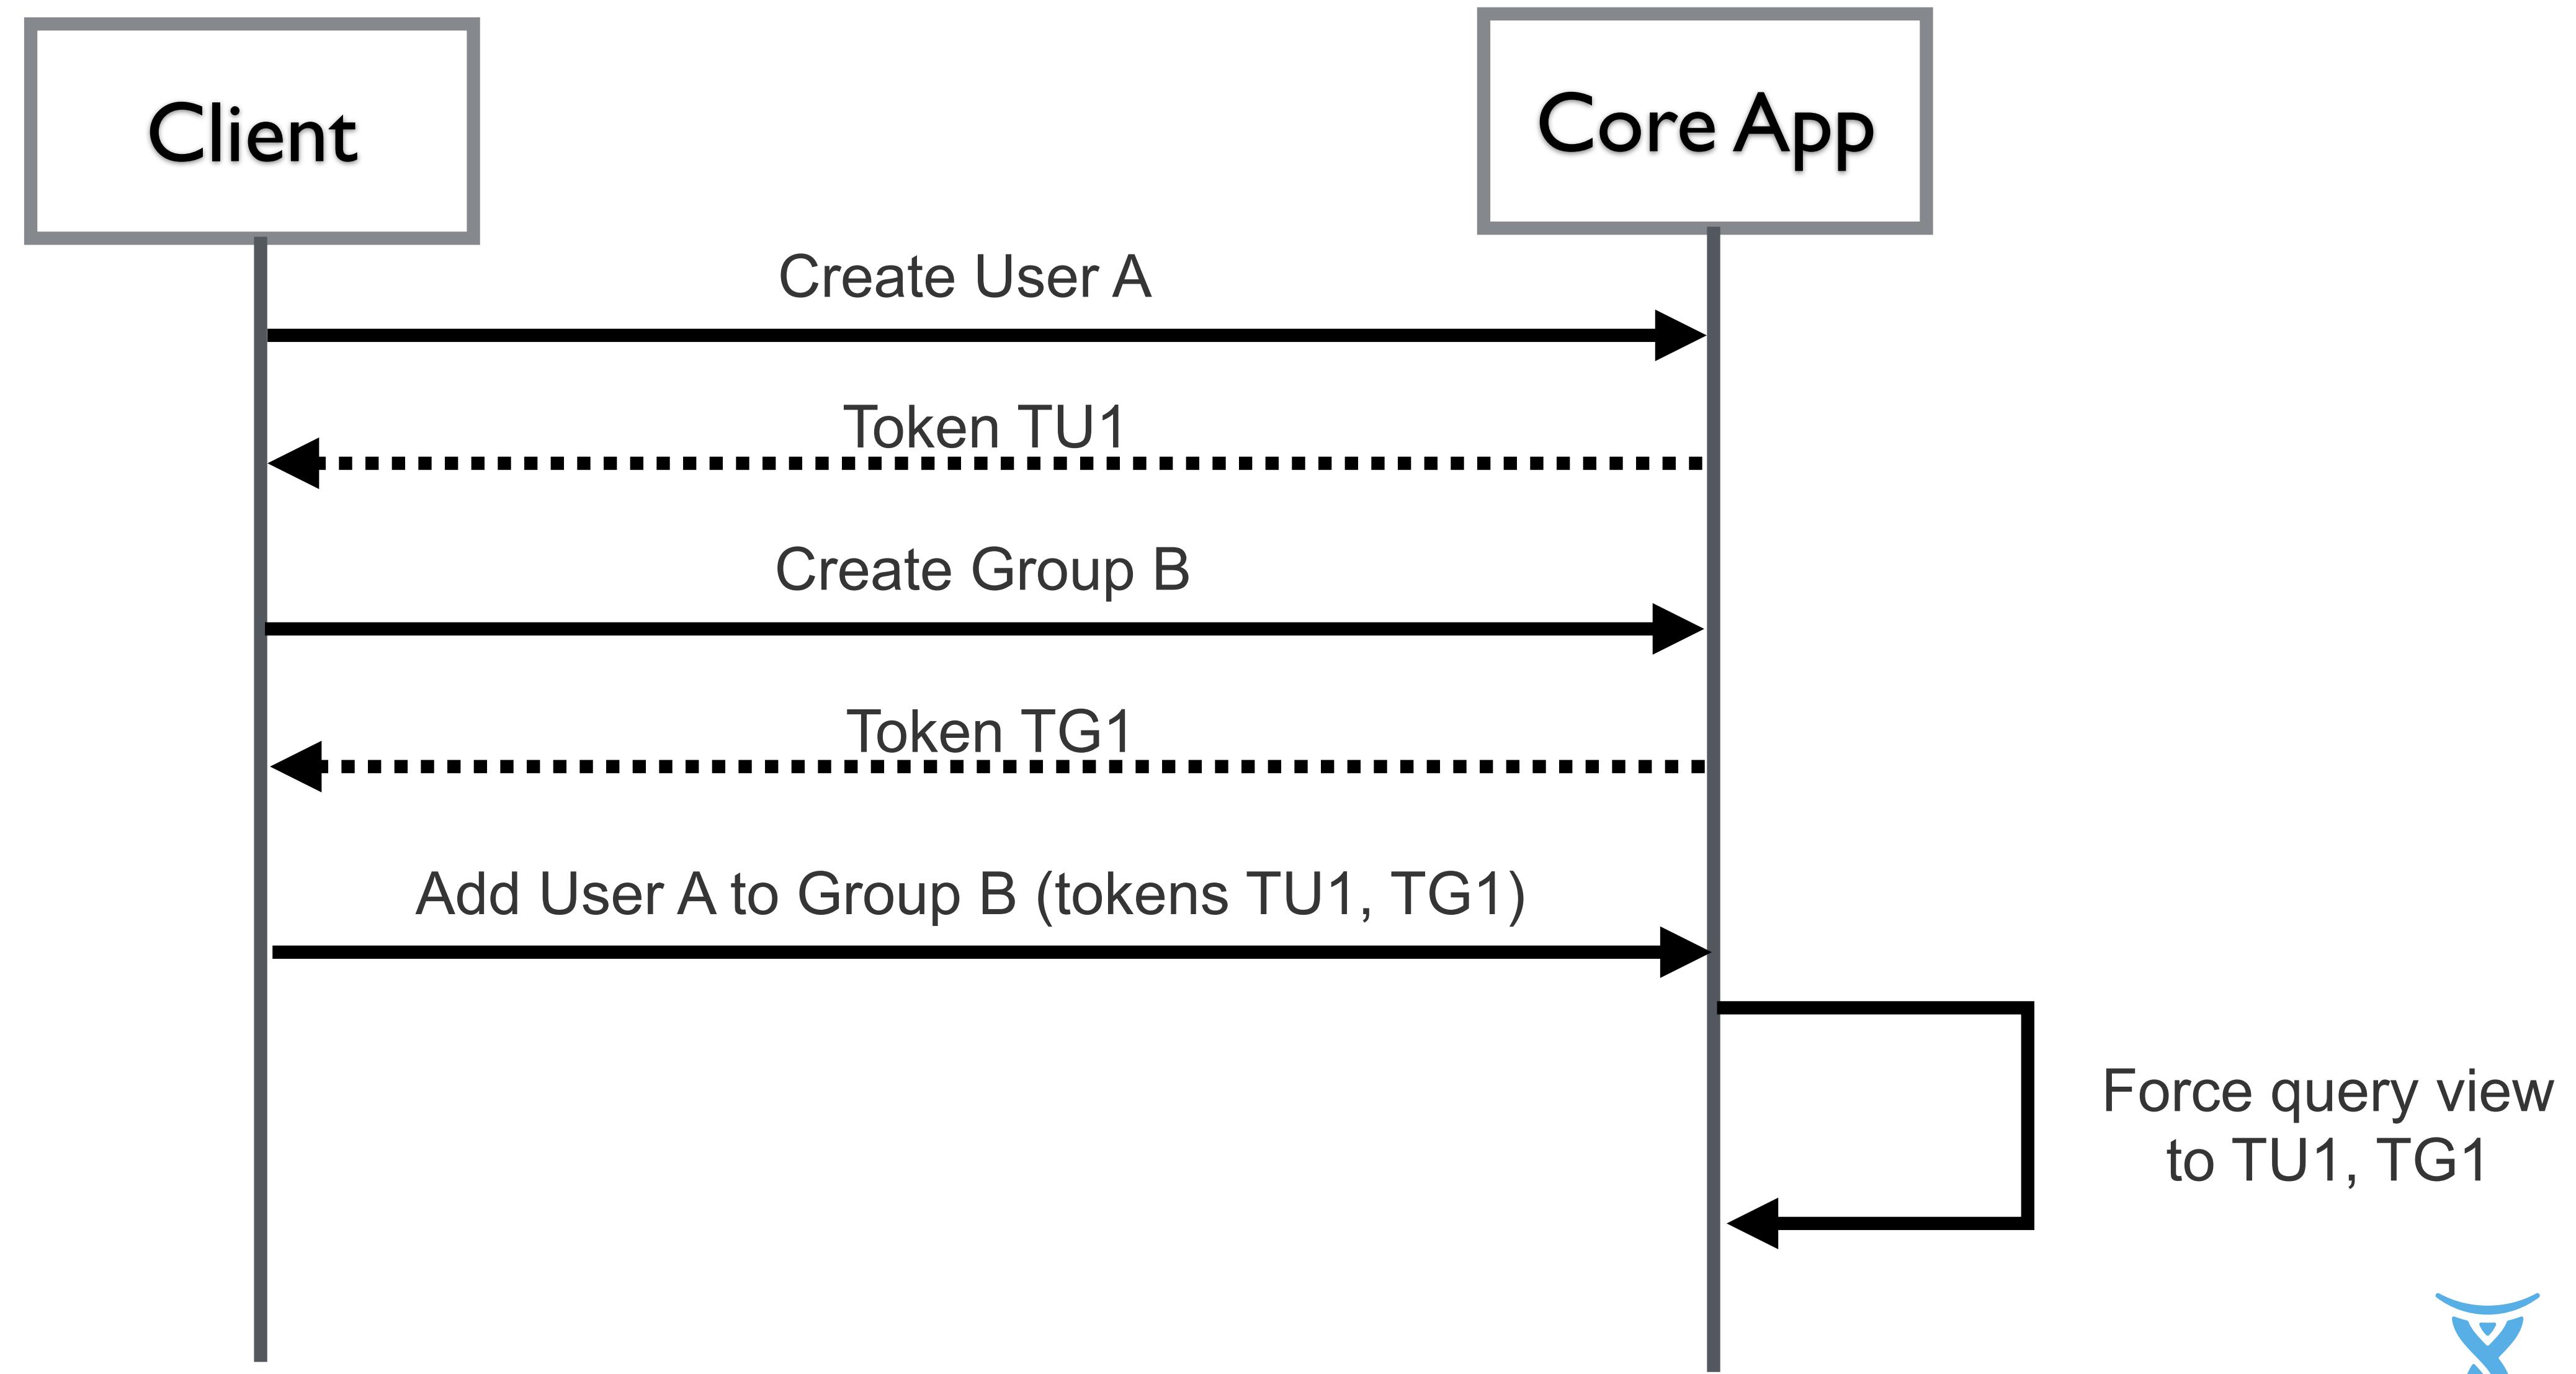
# Using tokens to enforce state



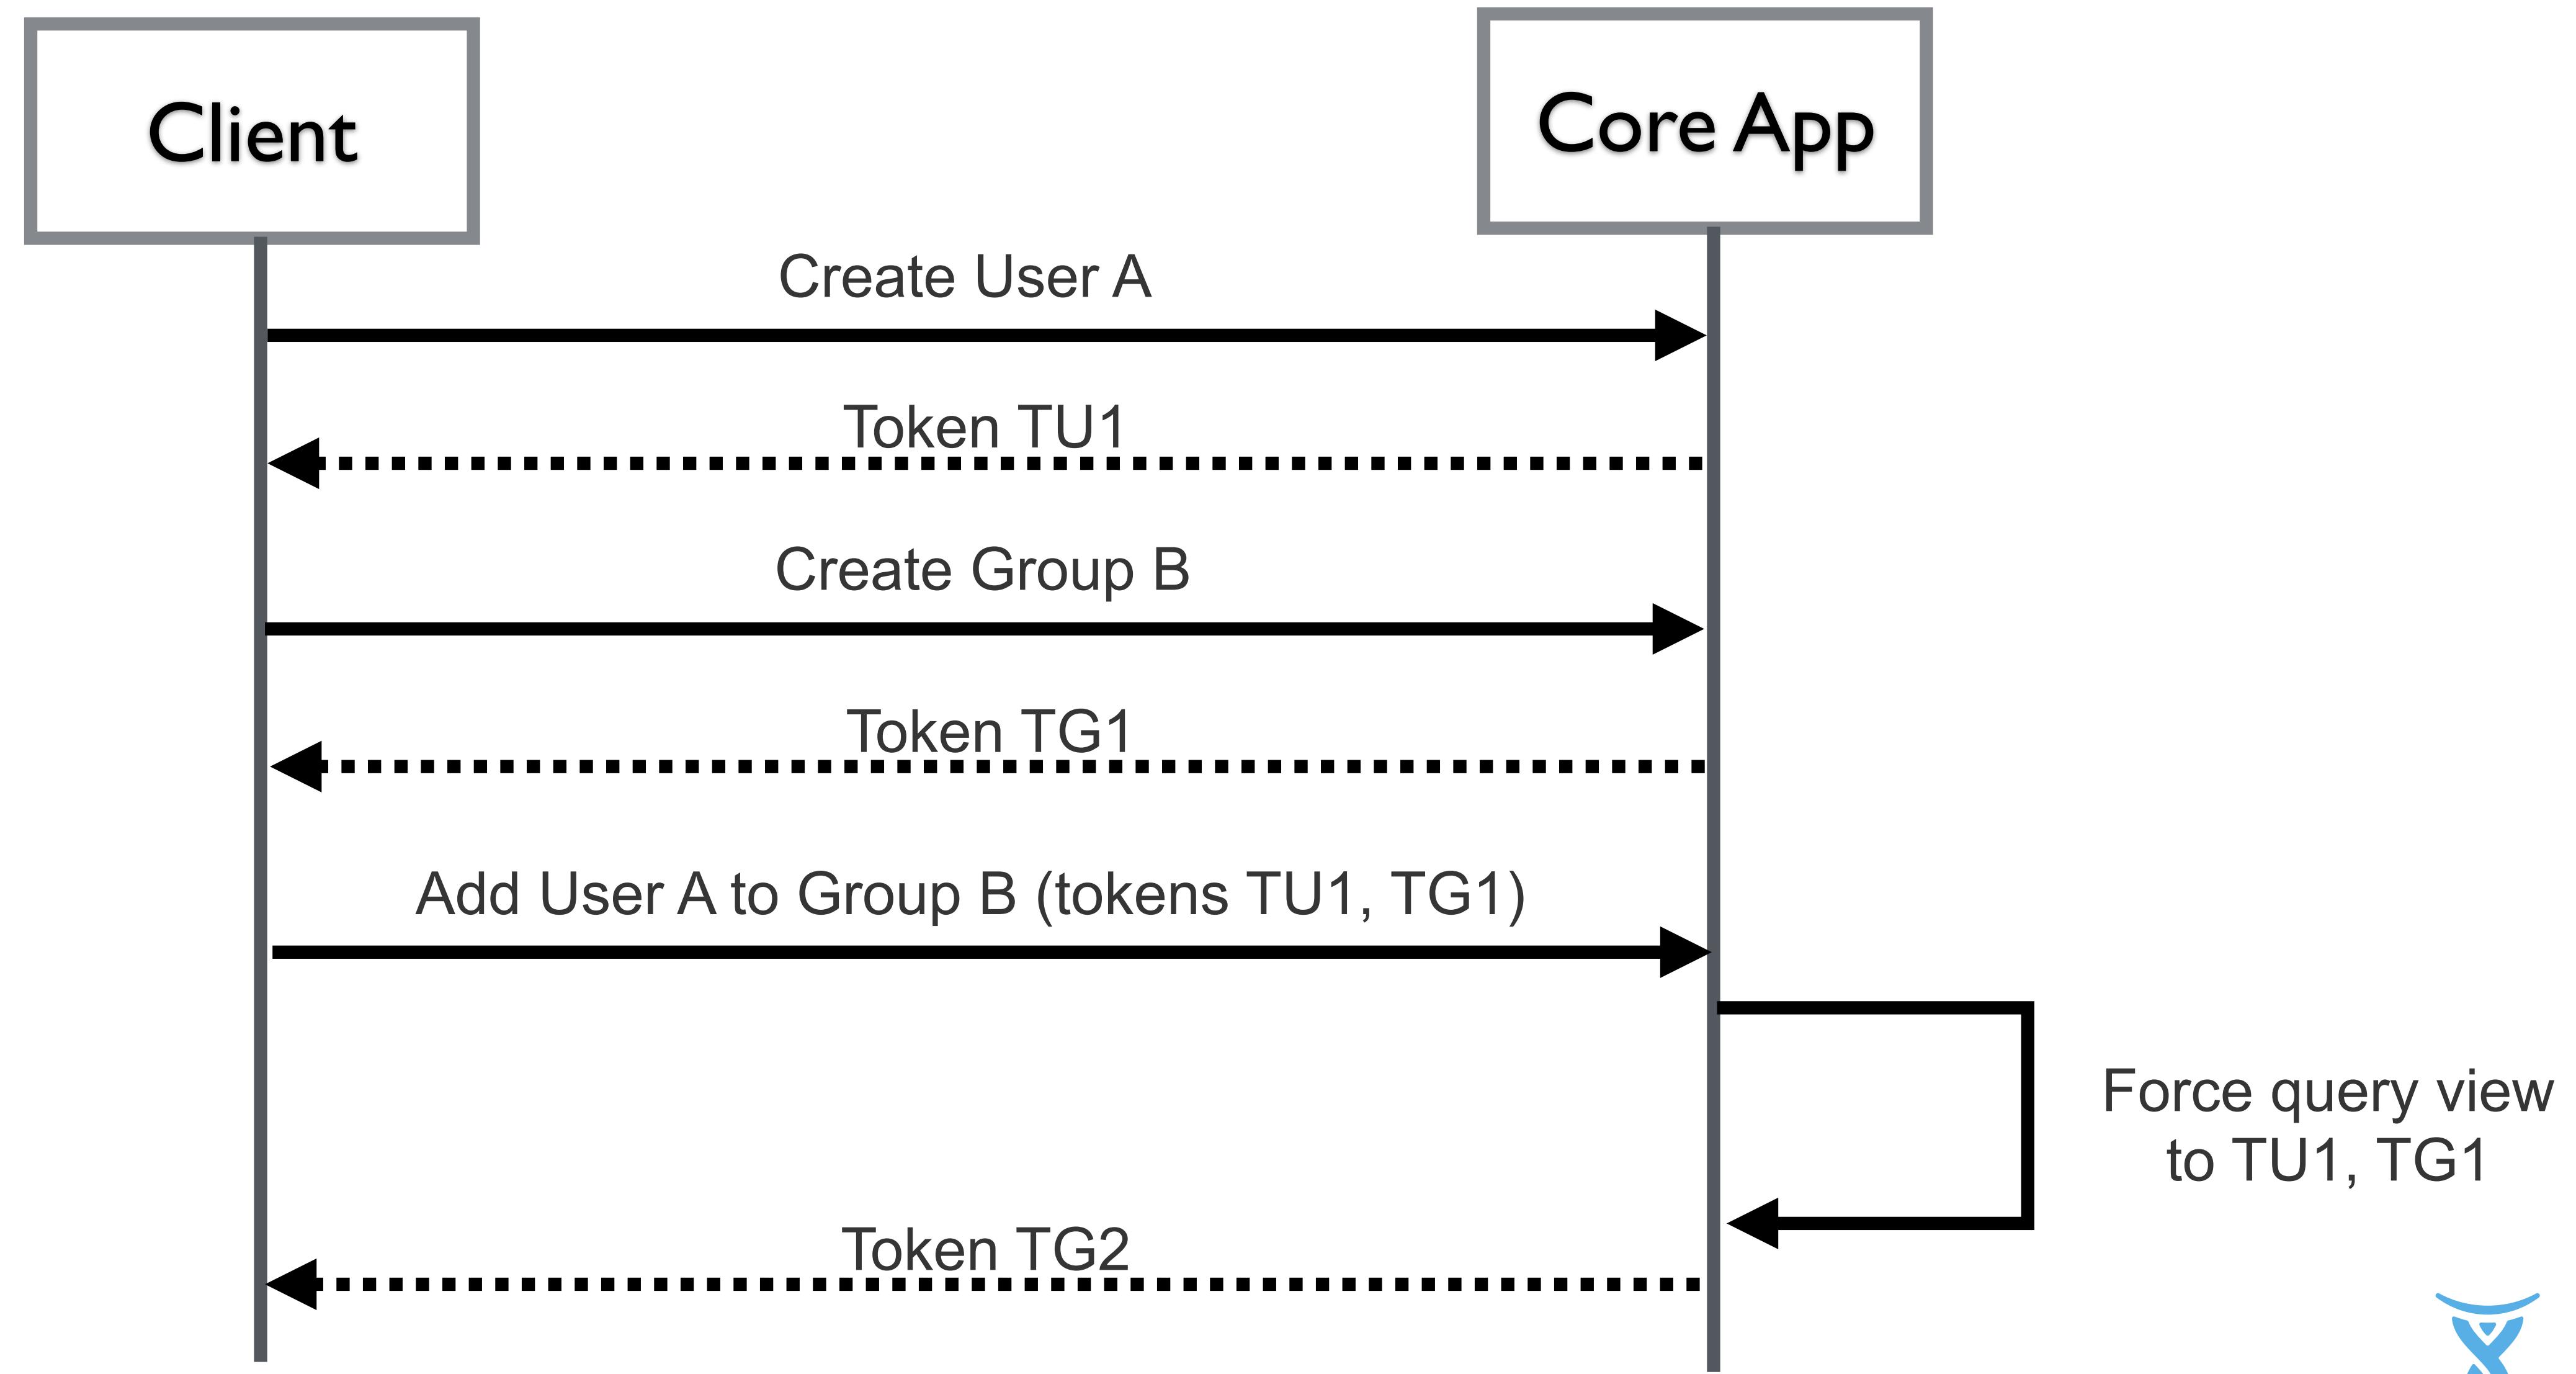
# Using tokens to enforce state



# Using tokens to enforce state



# Using tokens to enforce state





Principle 3.5:  
Conflict resolution  
instead of  
transactions

# Resolve conflicts on query



# Resolve conflicts on query

- Conflicting events on same stream



# Resolve conflicts on query

- Conflicting events on same stream
- Defined resolution algorithm on replay
  - e.g. Last Write Wins
  - Convergent/Commutative Replicated Data Types (CRDTs)



# Resolve conflicts on query

- Conflicting events on same stream No falsely failed transactions
- Defined resolution algorithm on replay
  - e.g. Last Write Wins
  - Convergent/Commutative Replicated Data Types (CRDTs)



# Resolve conflicts on query

- Conflicting events on same stream
  - No falsely failed transactions
- Defined resolution algorithm on replay
  - e.g. Last Write Wins
  - Convergent/Commutative Replicated Data Types (CRDTs)
  - More resilient query views



# Resolve conflicts on query

- Conflicting events on same stream
  - No falsely failed transactions
- Defined resolution algorithm on replay
  - e.g. Last Write Wins
    - More resilient query views
  - Convergent/Commutative Replicated Data Types (CRDTs)
    - Handles multi-region writes



# Resolve conflicts on query

- Conflicting events on same stream
  - No falsely failed transactions
- Defined resolution algorithm on replay
  - e.g. Last Write Wins
    - More resilient query views
  - Convergent/Commutative Replicated Data Types (CRDTs)
    - Handles multi-region writes
    - Potential temporary glitches



# Resolve conflicts on query

- Conflicting events on same stream
- Defined resolution algorithm on replay
  - e.g. Last Write Wins
  - Convergent/Commutative Replicated Data Types (CRDTs)

No falsely failed transactions

More resilient query views

Handles multi-region writes

Potential temporary glitches

Needs to be implemented on all query nodes





# Summary

# Key takeaways

- Start small and challenge everything!



# Key takeaways

- Start small and challenge everything!
- Incremental architecture for incremental demos



# Key takeaways

- Start small and challenge everything!
- Incremental architecture for incremental demos
- Think “Events as an API” - small idempotent events



# Key takeaways

- Start small and challenge everything!
- Incremental architecture for incremental demos
- Think “Events as an API” - small idempotent events
- Split streams for reasonable latency vs consistency



# Key takeaways

- Start small and challenge everything!
- Incremental architecture for incremental demos
- Think “Events as an API” - small idempotent events
- Split streams for reasonable latency vs consistency
- Accept weaker transactions and eventual consistency



*“We should using  
event sourcing more  
than we do”*

Martin Fowler (very loosely paraphrased)



event sourcing lib:  
[bitbucket.org/atlassianlabs/eventsr](http://bitbucket.org/atlassianlabs/eventsr)

aws-scala lib:  
<http://bitbucket.org/atlassian/aws-scala>





*Please*

**Remember to  
rate this session**

*Thank you!*



Join the conversation #scaladays



Join the conversation #scaladays